



Sequential CPU Implementation



Y86 Instruction Set

P259

Byte	0	1	2	3	4	5
nop	0	0				
halt	1	0				
rrmovl rA, rB	2	0	rA	rB		
irmovl V, rB	3	0	8	rB	V	
rmmovl rA, D(rB)	4	0	rA	rB	D	
mrmmovl D(rB), rA	5	0	rA	rB	D	
Op1 rA, rB	6	fn	rA	rB		
jXX Dest	7	fn	Dest			
call Dest	8	0	Dest			
ret	9	0				
pushl rA	A	0	rA	8		
popl rA	B	0	rA	8		


addl	6	0				
subl	6	1				
andl	6	2				
xorl	6	3				
jmp	7	0				
jle	7	1				
j1	7	2				
je	7	3				
jne	7	4				
jge	7	5				
jg	7	6				



Building Blocks P278, P279, P280

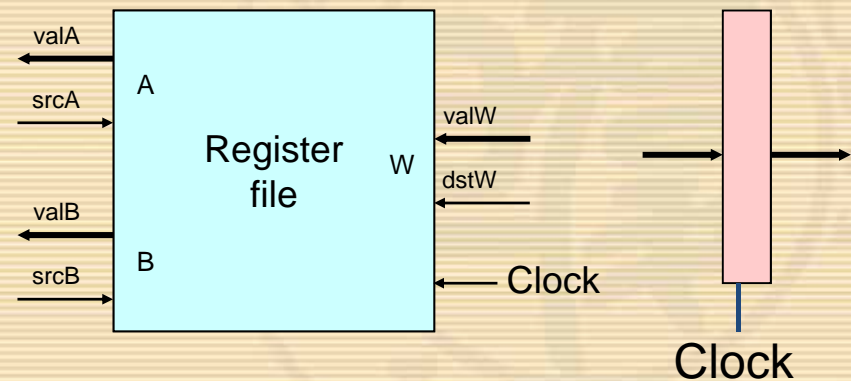
- Combinational Logic

- Compute Boolean functions of inputs
- Continuously respond to input changes
- Operate on data and implement control



- Storage Elements

- Store bits
- Addressable memories
- Non-addressable registers
- Loaded only as clock rises






How function works?

- Operations
- Movs
- Push/Pop
- Jump
- Call




Instruction Execution Stages

- Fetch
 - Read instruction from instruction memory
- Decode
 - Read program registers
- Execute
 - Compute value or address
- Memory
 - Read or write data
- Write Back
 - Write program registers
- PC
 - Update program counter





Instruction Decoding



- Instruction Format

- | | |
|--------------------------|------------|
| – Instruction byte | icode:ifun |
| – Optional register byte | rA:rB |
| – Optional constant word | valC |



Executing Arith./Logical Operation

OP1 rA, rB



- Fetch
 - Read 2 bytes
- Decode
 - Read operand registers
- Execute
 - Perform operation
 - Set condition codes
- Memory
 - Do nothing
- Write back
 - Update register
- PC Update
 - Increment PC by 2



Stage Computation: Arith/Log. Ops

	OPI rA, rB	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC}+2$	Read instruction byte Read register byte Compute next PC
Decode	$\text{valA} \leftarrow R[\text{rA}]$ $\text{valB} \leftarrow R[\text{rB}]$	Read operand A Read operand B
Execute	$\text{valE} \leftarrow \text{valB OP valA}$ Set CC	Perform ALU operation Set condition code register
Memory		
Write back	$R[\text{rB}] \leftarrow \text{valE}$	Write back result
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

- Formulate instruction execution as sequence of simple steps
- Use same general form for all instructions



Executing `rrmovl`

`rrmovl rA, rB`

2	0	rA	rB
---	---	----	----

- Fetch
 - Read 2 bytes
- Decode
 - Read operand register rA
- Execute
 - Do nothing
- Memory
 - Do nothing
- Write back
 - Update register
- PC Update
 - Increment PC by 2



Stage Computation: `rrmovl`

	<code>rrmovl rA, rB</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC}+2$	Read instruction byte Read register byte Compute next PC
Decode	$\text{valA} \leftarrow R[\text{rA}]$	Read operand A
Execute	$\text{valE} \leftarrow 0 + \text{valA}$	Perform ALU operation
Memory		*valE
Write back	$R[\text{rB}] \leftarrow \text{valE}$	Write back result
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

- Formulate instruction execution as sequence of simple steps
- Use same general form for all instructions



Executing `irmovl`

`irmovl V, rB`



- Fetch
 - Read 6 bytes
- Decode
 - Do nothing
- Execute
 - Do nothing
- Memory
 - Do nothing
- Write back
 - Update register
- PC Update
 - Increment PC by 6



Stage Computation: `irmovl`

	<code>irmovl V, rB</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valC} \leftarrow M_4[\text{PC}+2]$ $\text{valP} \leftarrow \text{PC}+6$	Read instruction byte Read register byte Read constant value Compute next PC
Decode		
Execute	$\text{valE} \leftarrow 0 + \text{valC}$	Perform ALU operation
Memory		
Write back	$R[\text{rB}] \leftarrow \text{valE}$	Write back result
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

- Formulate instruction execution as sequence of simple steps
- Use same general form for all instructions



Executing `rmmovl`

`rmmovl rA, D(rB)`

4	0	rA	rB	D
---	---	----	----	---

- Fetch
 - Read 6 bytes
- Decode
 - Read operand registers
- Execute
 - Compute effective address
- Memory
 - Write to memory
- Write back
 - Do nothing
- PC Update
 - Increment PC by 6



Stage Computation: `rmmovl`

	<code>rmmovl rA, D(rB)</code>	
Fetch	<code>icode:ifun $\leftarrow M_1[PC]$</code> <code>rA:rB $\leftarrow M_1[PC+1]$</code> <code>valC $\leftarrow M_4[PC+2]$</code> <code>valP $\leftarrow PC+6$</code>	Read instruction byte Read register byte Read displacement D Compute next PC
Decode	<code>valA $\leftarrow R[rA]$</code> <code>valB $\leftarrow R[rB]$</code>	Read operand A Read operand B
Execute	<code>valE $\leftarrow valB + valC$</code>	Compute effective address (sum of the displacement and the base register value)
Memory	<code>$M_4[valE] \leftarrow valA$</code>	Write value to memory
Write back		
PC update	<code>PC $\leftarrow valP$</code>	Update PC

– Use ALU for address computation



Executing `mrmovl`

`mrmovl D(rB),rA`

5	0	rA	rB	D
---	---	----	----	---

- Fetch
 - Read 6 bytes
- Decode
 - Read operand register rB
- Execute
 - Compute effective address
- Memory
 - Read from memory
- Write back
 - Update register rA
- PC Update
 - Increment PC by 6



Stage Computation: `mrmovl`

	<code>mrmovl D(rB) , rA</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valC} \leftarrow M_4[\text{PC}+2]$ $\text{valP} \leftarrow \text{PC}+6$	Read instruction byte Read register byte Read displacement D Compute next PC
Decode	$\text{valB} \leftarrow R[\text{rB}]$	Read operand B
Execute	$\text{valE} \leftarrow \text{valB} + \text{valC}$	Compute effective address
Memory	$\text{valM} \leftarrow M_4[\text{valE}]$	Read data from memory
Write back	$R[\text{rA}] \leftarrow \text{valM}$	Update register rA
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

– Use ALU for address computation



Executing pushl

pushl rA

a	0	rA	8
---	---	----	---

- Fetch
 - Read 2 bytes
- Decode
 - Read stack pointer and rA
- Execute
 - Decrement stack pointer by 4
- Memory
 - Store valA at the address of new stack pointer
- Write back
 - Update stack pointer
- PC Update
 - Increment PC by 2



Stage Computation: `pushl`

	<code>pushl rA</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC}+2$	Read instruction byte Read register byte Compute next PC
Decode	$\text{valA} \leftarrow R[\text{rA}]$ $\text{valB} \leftarrow R[\text{\%esp}]$	Read valA Read stack pointer
Execute	$\text{valE} \leftarrow \text{valB} + (-4)$	Decrement stack pointer
Memory	$M_4[\text{valE}] \leftarrow \text{valA}$	Store to stack
Write back	$R[\text{\%esp}] \leftarrow \text{valE}$	Update stack pointer *在write back之前实际上写入的元素在堆栈外。
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

– Use ALU to Decrement stack pointer



Executing popl

popl rA



- Fetch
 - Read 2 bytes
- Decode
 - Read stack pointer
- Execute
 - Increment stack pointer by 4
- Memory
 - Read from old stack pointer
- Write back
 - Update stack pointer
 - Write result to register
- PC Update
 - Increment PC by 2




Stage Computation: `popl`

	<code>popl rA</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC}+2$	Read instruction byte Read register byte Compute next PC
Decode	$\text{valA} \leftarrow R[\%esp]$ $\text{valB} \leftarrow R[\%esp]$	Read stack pointer Read stack pointer
Execute	$\text{valE} \leftarrow \text{valB} + 4$	Increment stack pointer
Memory	$\text{valM} \leftarrow M_4[\text{valA}]$	Read from stack
Write back	$R[\%esp] \leftarrow \text{valE}$ $R[\text{rA}] \leftarrow \text{valM}$	Update stack pointer Write back result
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

- Use ALU to increment stack pointer
- Must update two registers
 - Popped value
 - New stack pointer



Executing Jumps



- Fetch
 - Read 5 bytes
 - Increment PC by 5
- Decode
 - Do nothing
- Execute
 - Determine whether to take branch based on jump condition and condition codes
- Memory
 - Do nothing
- Write back
 - Do nothing
- PC Update
 - Set PC to Dest if branch taken or to incremented PC if not branch



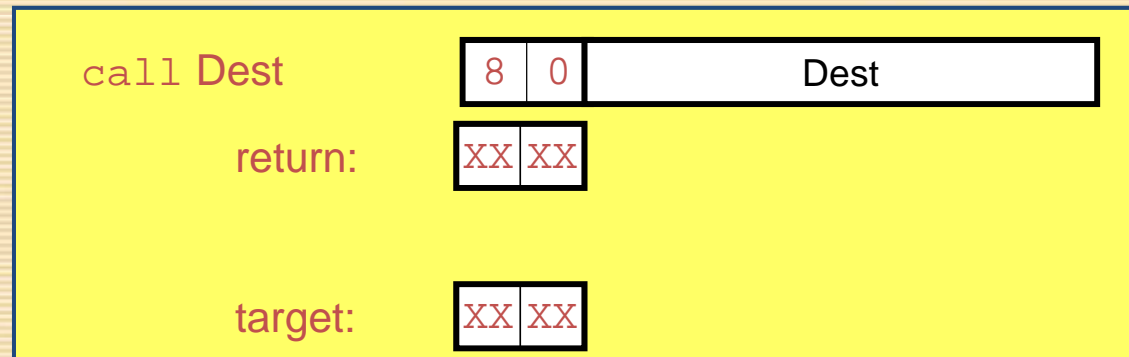
Stage Computation: Jumps

	jXX Dest	
Fetch	icode:ifun $\leftarrow M_1[PC]$	Read instruction byte
	valC $\leftarrow M_4[PC+1]$	Read destination address
	valP $\leftarrow PC+5$	Fall through address
Decode		
Execute	Bch $\leftarrow \text{Cond}(CC, \text{ifun})$	Take branch?
Memory		
Write back		
PC update	PC $\leftarrow \text{Bch} ? \text{valC} : \text{valP}$	Update PC

- Compute both addresses
- Choose based on setting of condition codes and branch condition



Executing call



- Fetch
 - Read 5 bytes
 - Increment PC by 5
- Decode
 - Read stack pointer
- Execute
 - Decrement stack pointer by 4
- Memory
 - Write incremented PC to new value of stack pointer
- Write back
 - Update stack pointer
- PC Update
 - Set PC to Dest



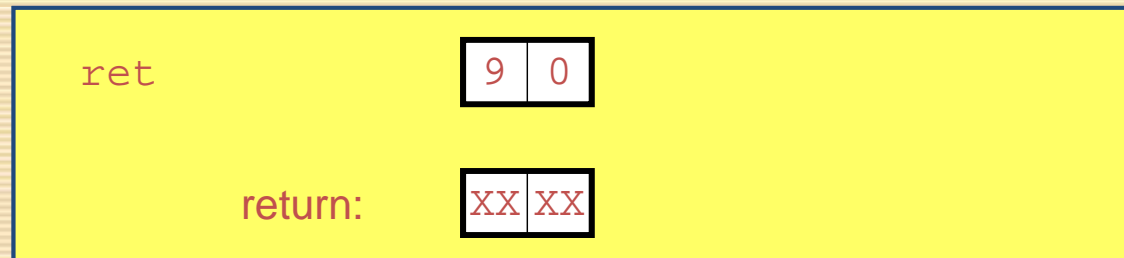
Stage Computation: `call`

	<code>call Dest</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$	Read instruction byte
	$\text{valC} \leftarrow M_4[\text{PC}+1]$	Read destination address
	$\text{valP} \leftarrow \text{PC}+5$	Compute return point
Decode	$\text{valB} \leftarrow R[\%esp]$	Read stack pointer
Execute	$\text{valE} \leftarrow \text{valB} + -4$	Decrement stack pointer
Memory	$M_4[\text{valE}] \leftarrow \text{valP}$	Write return value on stack
Write back	$R[\%esp] \leftarrow \text{valE}$	Update stack pointer
PC update	$\text{PC} \leftarrow \text{valC}$	Set PC to destination

- Use ALU to decrement stack pointer
- Store incremented PC



Executing `ret`



- Fetch
 - Read 1 byte
- Decode
 - Read stack pointer
- Execute
 - Increment stack pointer by 4
- Memory
 - Read return address from old stack pointer
- Write back
 - Update stack pointer
- PC Update
 - Set PC to return address



Stage Computation: `ret`

	<code>ret</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$	Read instruction byte
Decode	$\text{valA} \leftarrow R[\%esp]$ $\text{valB} \leftarrow R[\%esp]$	Read operand stack pointer Read operand stack pointer
Execute	$\text{valE} \leftarrow \text{valB} + 4$	Increment stack pointer
Memory	$\text{valM} \leftarrow M_4[\text{valA}]$	Read return address
Write back	$R[\%esp] \leftarrow \text{valE}$	Update stack pointer
PC update	$\text{PC} \leftarrow \text{valM}$	Set PC to return address

- Use ALU to increment stack pointer
- Read return address from memory



Computation Steps

		OPI rA, rB	
Fetch	icode,ifun	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$	Read instruction byte
	rA,rB	$\text{rA:rB} \leftarrow M_1[\text{PC}+1]$	Read register byte
	valC		[Read constant word]
	valP	$\text{valP} \leftarrow \text{PC}+2$	Compute next PC
Decode	valA, srcA	$\text{valA} \leftarrow R[\text{rA}]$	Read operand A
	valB, srcB	$\text{valB} \leftarrow R[\text{rB}]$	Read operand B
Execute	valE	$\text{valE} \leftarrow \text{valB OP valA}$	Perform ALU operation
	Cond code	Set CC	Set condition code register
Memory	valM		[Memory read/write]
Write	dstE	$R[\text{rB}] \leftarrow \text{valE}$	Write back ALU result
back	dstM		[Write back memory result]
PC update	PC	$\text{PC} \leftarrow \text{valP}$	Update PC

- All instructions follow same general pattern
- Differ in what gets computed on each step



Computation Steps

		call Dest	
Fetch	icode,ifun	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$	Read instruction byte
	rA,rB		[Read register byte]
	valC	$\text{valC} \leftarrow M_4[\text{PC}+1]$	Read constant word
	valP	$\text{valP} \leftarrow \text{PC}+5$	Compute next PC
Decode	valA, srcA	$\text{valB} \leftarrow R[\%esp]$	[Read operand A]
	valB, srcB		Read operand B
Execute	valE	$\text{valE} \leftarrow \text{valB} + -4$	Perform ALU operation
	Cond code		[Set condition code reg.]
Memory	valM	$M_4[\text{valE}] \leftarrow \text{valP}$	Memory read/write
Write	dstE	$R[\%esp] \leftarrow \text{valE}$	Write back ALU result
back	dstM		[Write back memory result]
PC update	PC	$\text{PC} \leftarrow \text{valC}$	Update PC

- All instructions follow same general pattern
- Differ in what gets computed on each step



Computed Values

•Fetch

icode	Instruction code
ifun	Instruction function
rA	Instr. Register A
rB	Instr. Register B
valC	Instruction constant
valP	Incremented PC

•Decode § Write back

valA	Register value A
valB	Register value B

•Execute

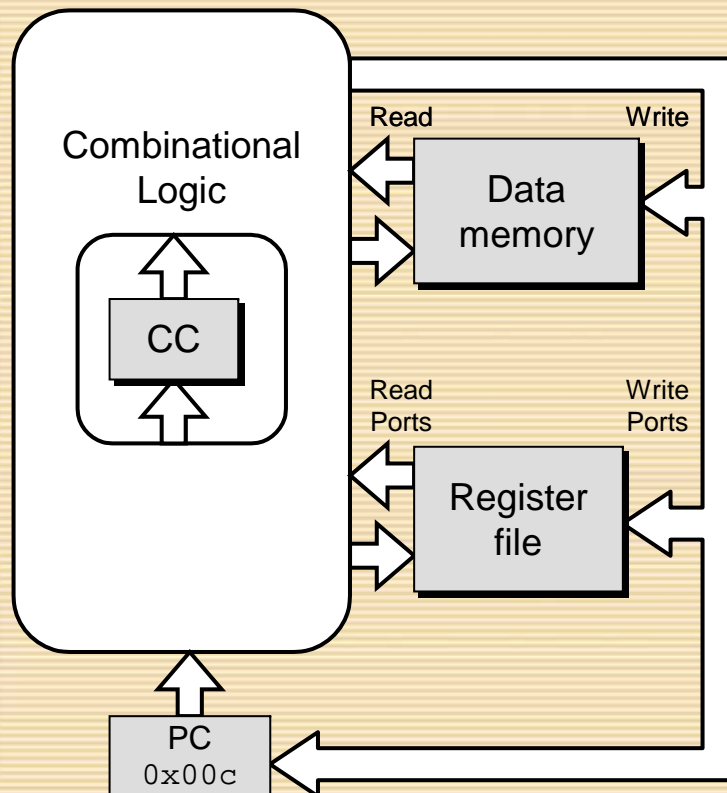
– valE	ALU result
– Bch	Branch flag

•Memory

– valM	Value from memory
--------	-------------------



SEQ Operation P297



- State

- Program counter register (PC)
- Condition code register (CC)
- Register File
- Memories
 - Access same memory space
 - Data: for reading/writing program data
 - Instruction: for reading instructions

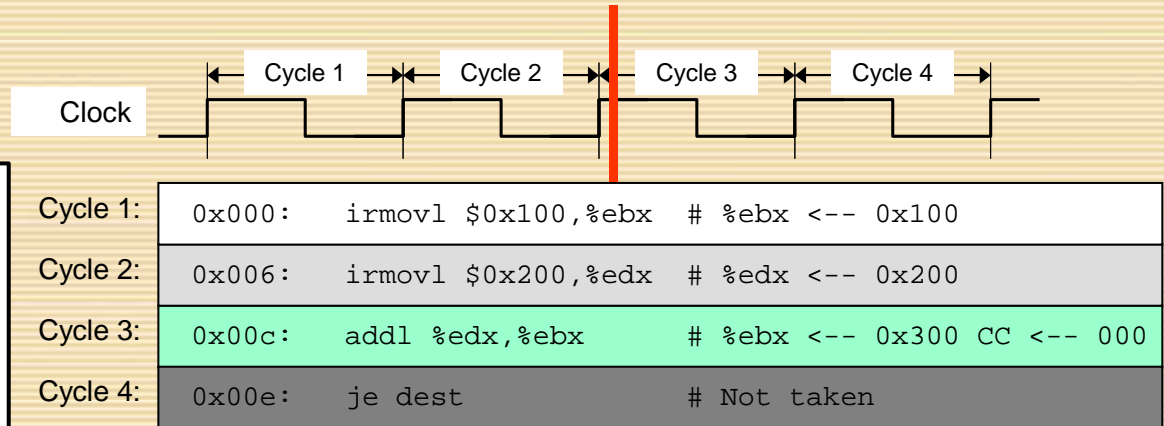
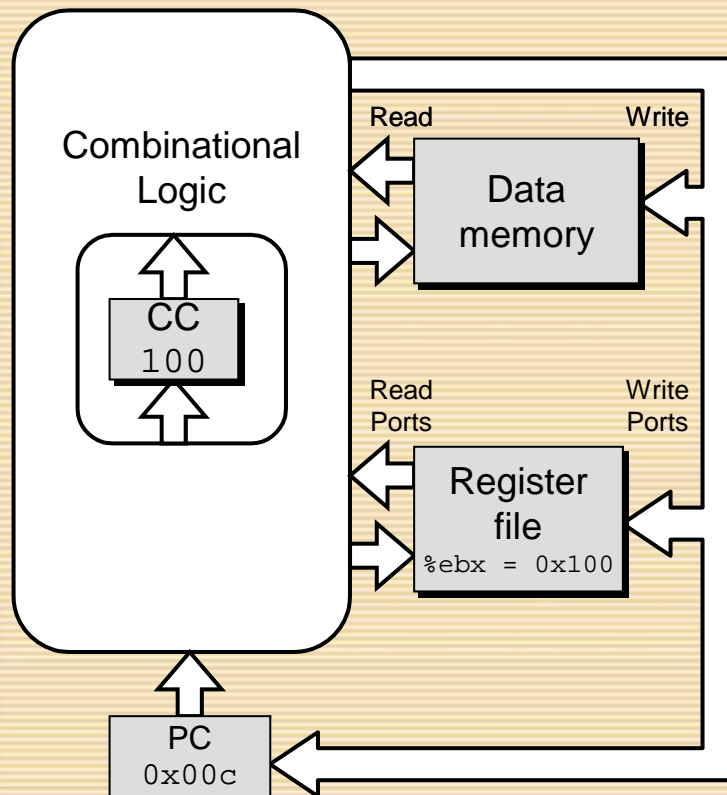
All updated as clock rises

- Combinational Logic

- ALU
- Control logic
- Memory reads
 - Instruction memory
 - Register file
 - Data memory




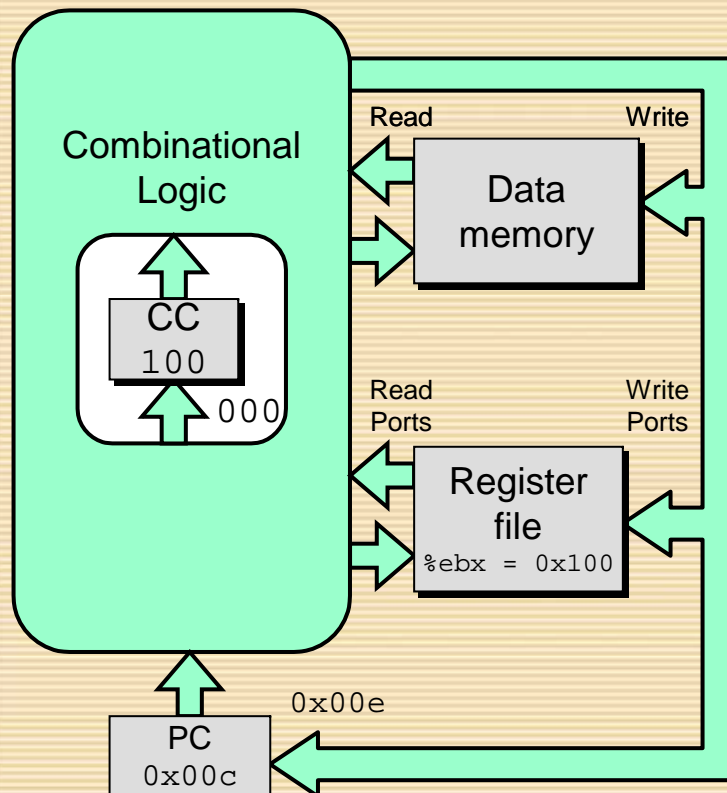
SEQ Operation



- state set according to second `irmovl` instruction
- combinational logic starting to react to state changes



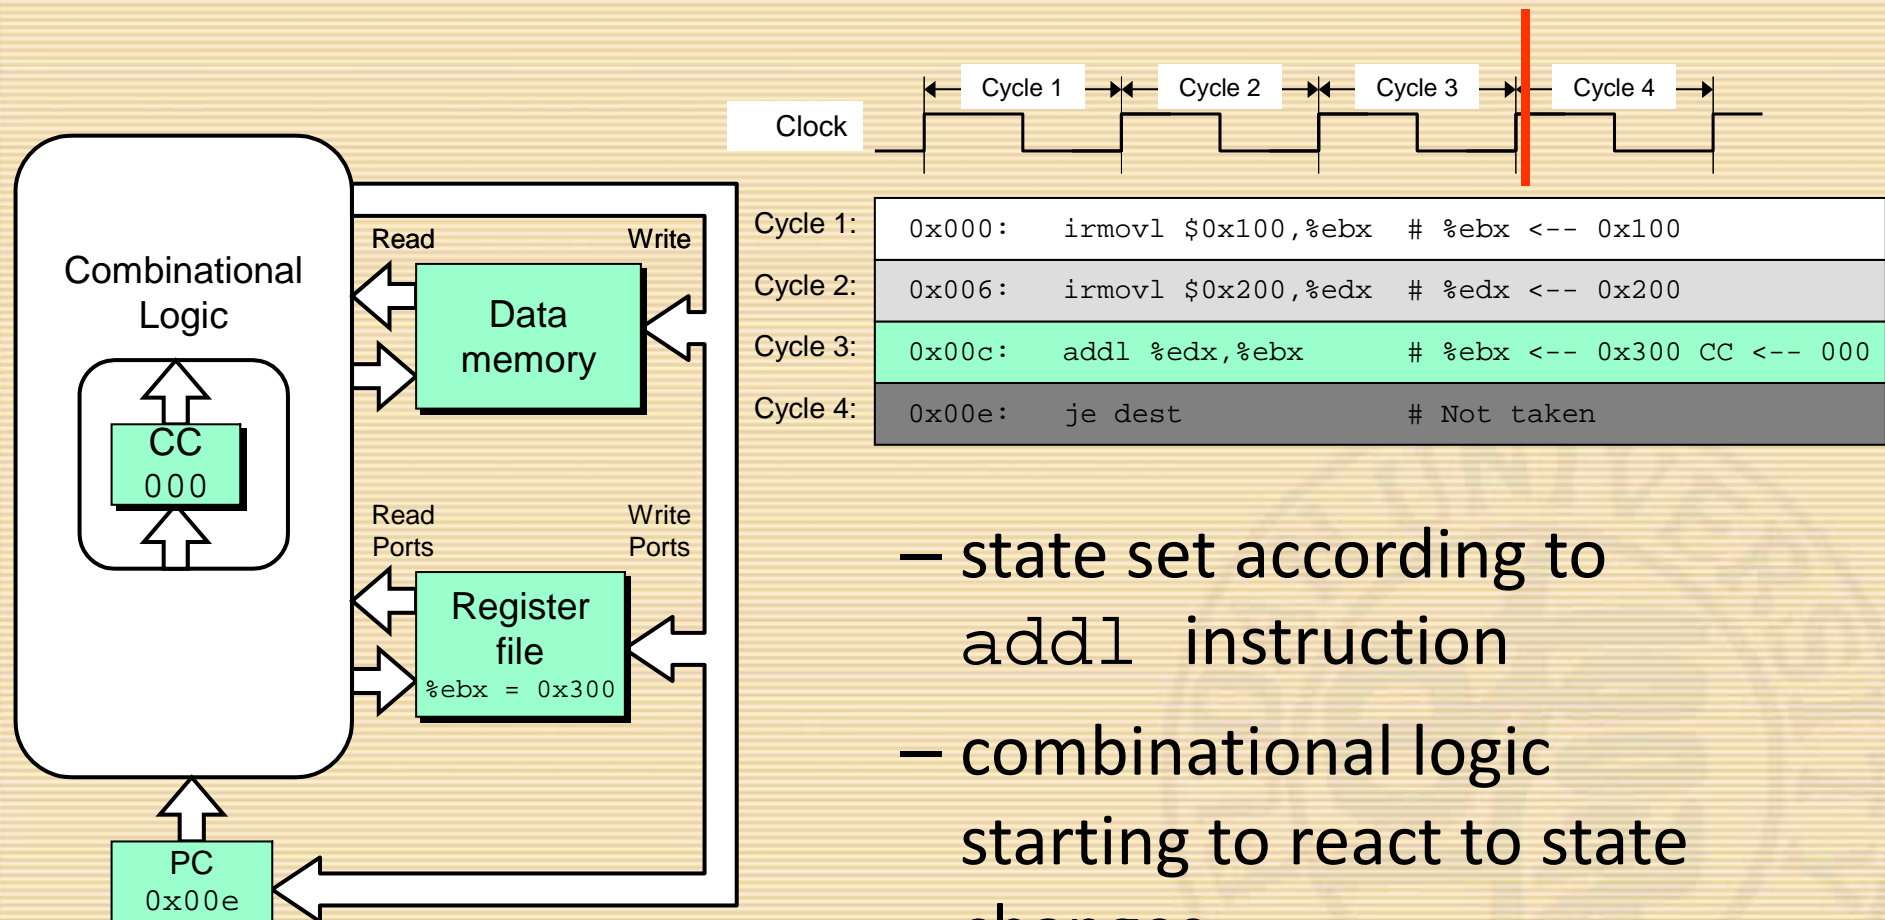
SEQ Operation



- state set according to second `irmovl` instruction
- combinational logic generates results for `addl` instruction



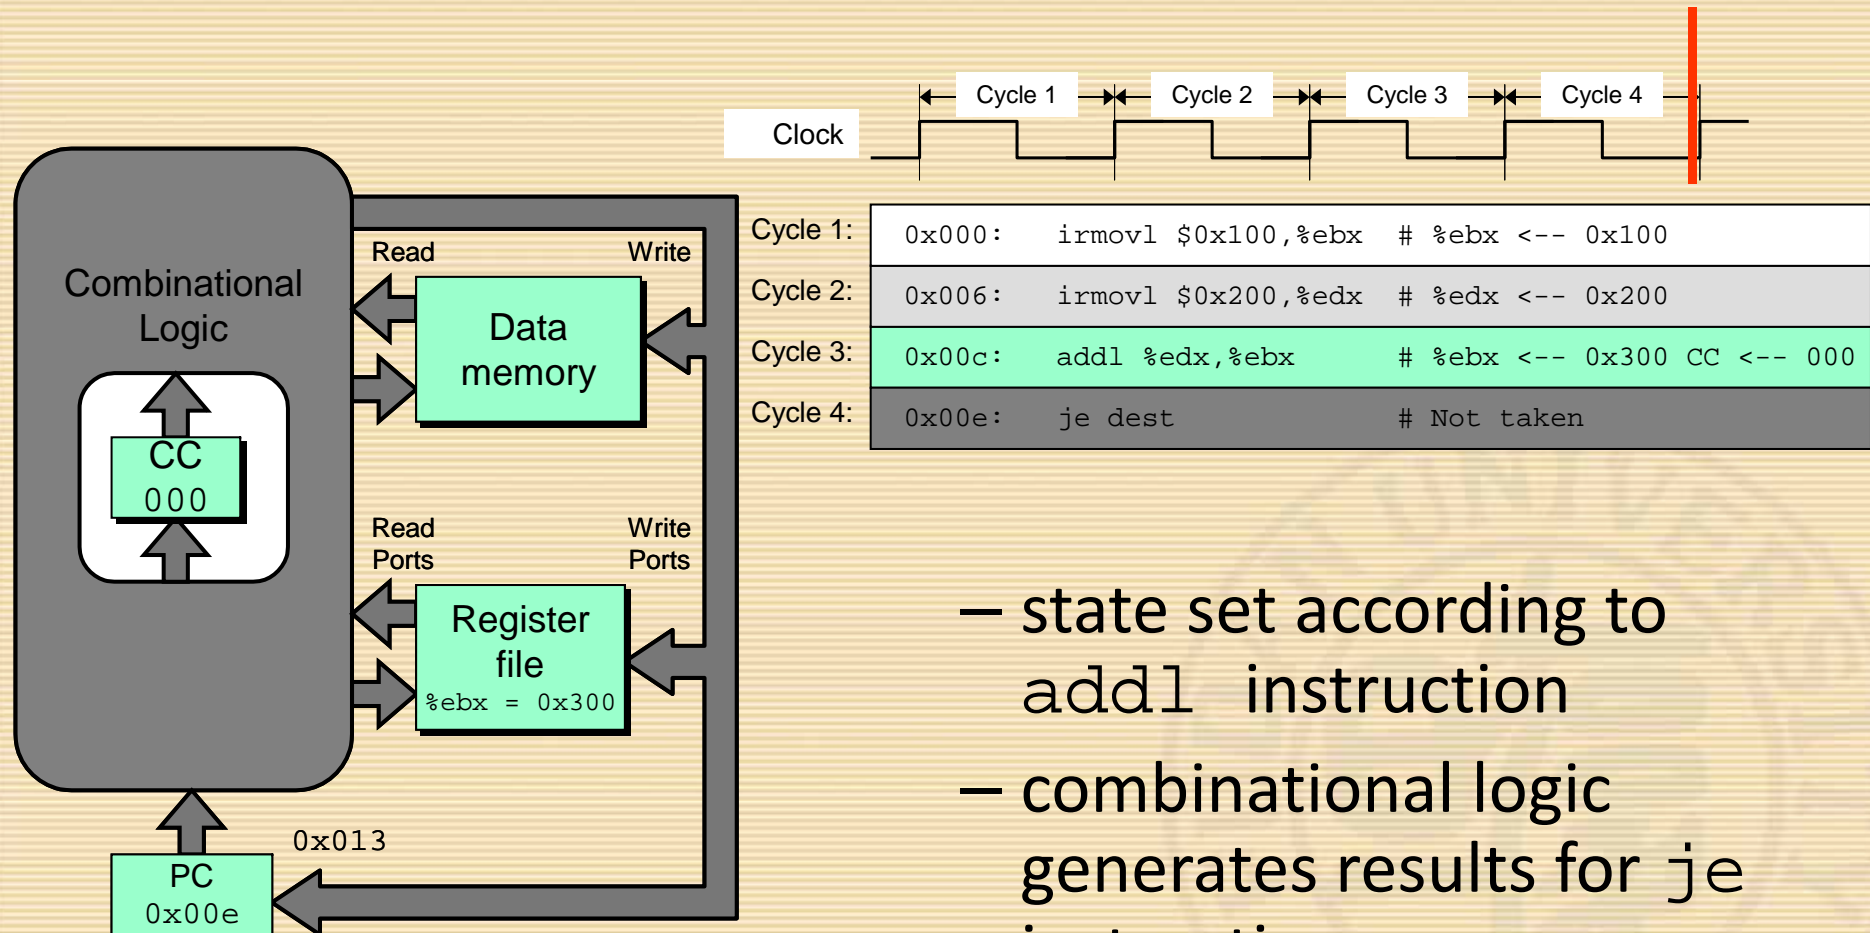
SEQ Operation



- state set according to addl instruction
- combinational logic starting to react to state changes



SEQ Operation



- state set according to `addl` instruction
- combinational logic generates results for `je` instruction



SEQ Semantics

- Achieve the same effect as a sequential execution of the assignment shown in the tables of Figures 4.16 to 4.19
 - Though all of the state updates occur simultaneously at the clock rises to the next cycle.
 - A problem: `popl %esp` need to sequentially write two registers. So the register file control logic must process it.
- Principle: never needs to read back the state updated by an instruction in order to complete the processing of this instruction (P295)
 - If so, the update must happen in the instruction cycle
 - E.g. `pushl` semantics
 - E.g. no one instruction need to both set and then read the condition codes.



SEQ CPU Implementation



Main Topics

- The implementation of a sequential CPU ---- SEQ
 - Every Instruction finished in one cycle.
 - Instruction executes in sequential
 - No two instructions execute in parallel or overlap
- An revised version of SEQ ---- SEQ+
 - Modify the PC Update stage of SEQ
 - to show the difference between ISA and implementation



Some Macros P299

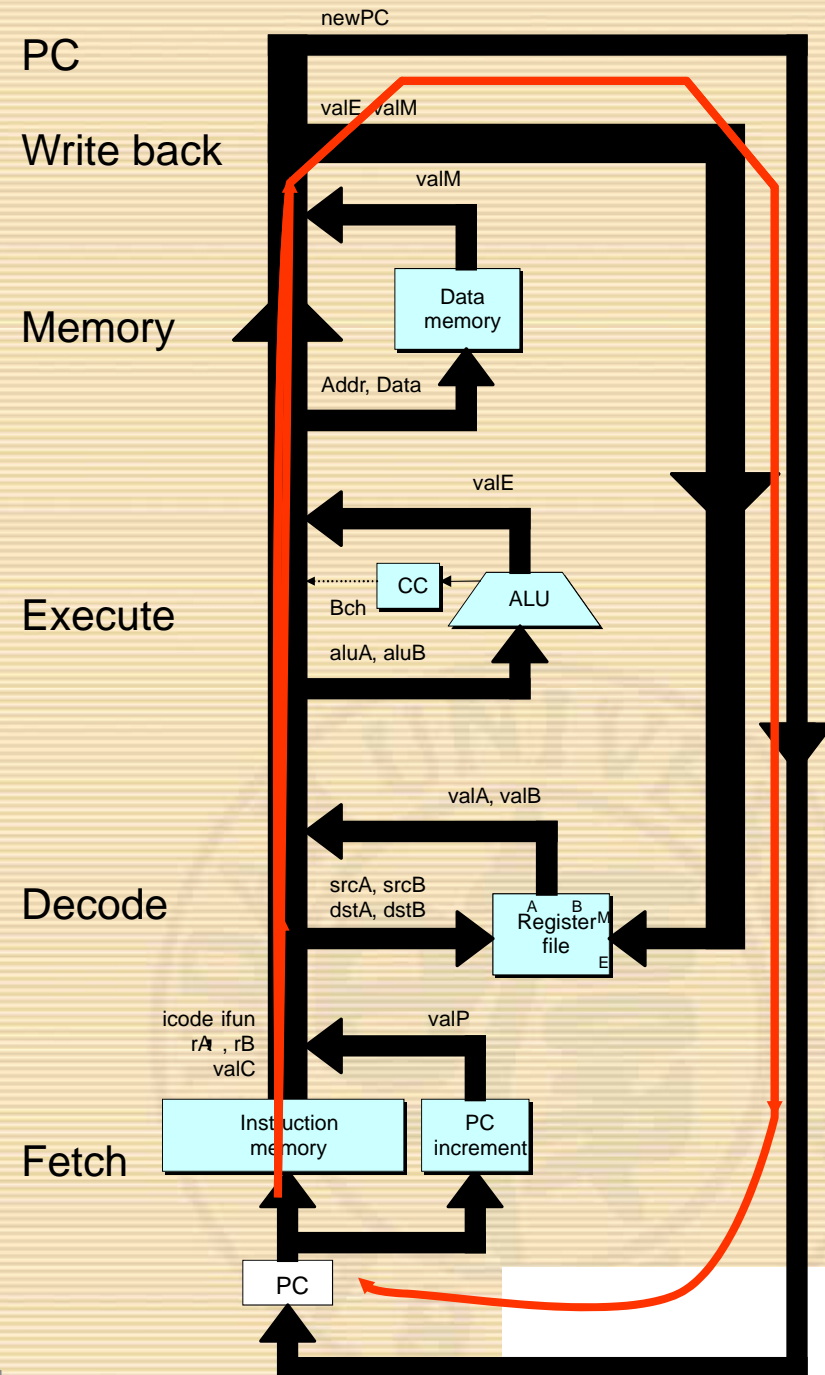
Name	Value	Meaning
INOP	0	Code for <code>nop</code> instruction
IHALT	1	Code for <code>halt</code> instruction
IRRMOVL	2	Code for <code>rrmovl</code> instruction
IIRMOVL	3	Code for <code>irmovl</code> instruction
IRMMOVL	4	Code for <code>rmmovl</code> instruction
IMRMOVL	5	Code for <code>mrmmovl</code> instruction
IOPL	6	Code for integer op instructions
IJXX	7	Code for jump instructions
.....
IPOPL	B	Code for <code>popl</code> instruction
RESP	6	Register ID for <code>%esp</code>
RENONE	8	Indicates no register file access
ALUADD	0	Function for addition operation



SEQ Hardware Structure

P292

- Stages
 - Fetch
 - Read instruction from memory
 - Decode
 - Read program registers
 - Execute
 - Compute value or address
 - Memory
 - Read or write data
 - Write Back
 - Write program registers
 - PC
 - Update program counter
- Instruction Flow
 - Read instruction at address specified by PC
 - Process through stages
 - Update program counter





Difference between semantics and implementation

- ISA
 - Every stage may update some states, these updates occur sequentially
- SEQ
 - All the state update operations occur simultaneously at clock rising (except memory and CC)



SEQ Hardware P293

- Key

- Blue boxes: predesigned hardware blocks
 - E.g., memories, ALU
- Gray boxes: control logic
 - Describe in HCL
- White ovals: labels for signals
- Thick lines: 32-bit word values
- Thin lines: 4-8 bit values
- Dotted lines: 1-bit values

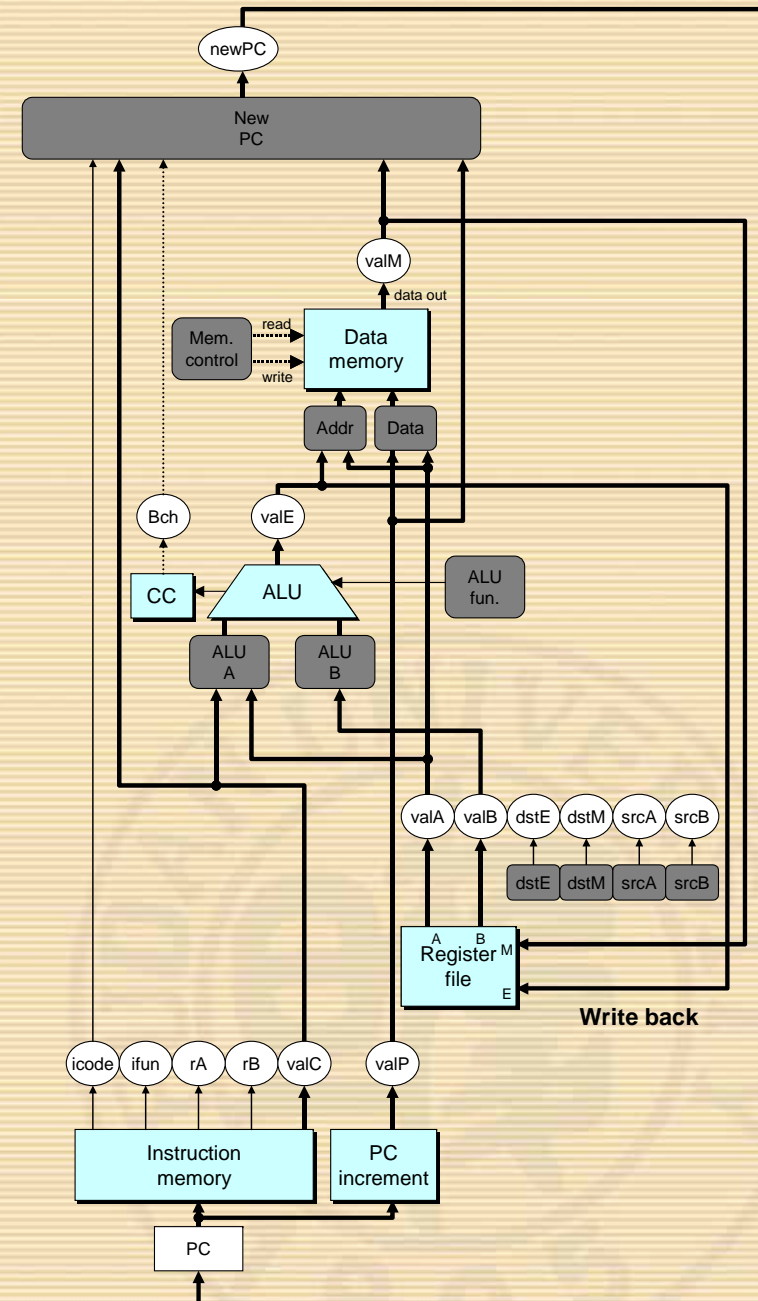
PC

Memory

Execute

Decode

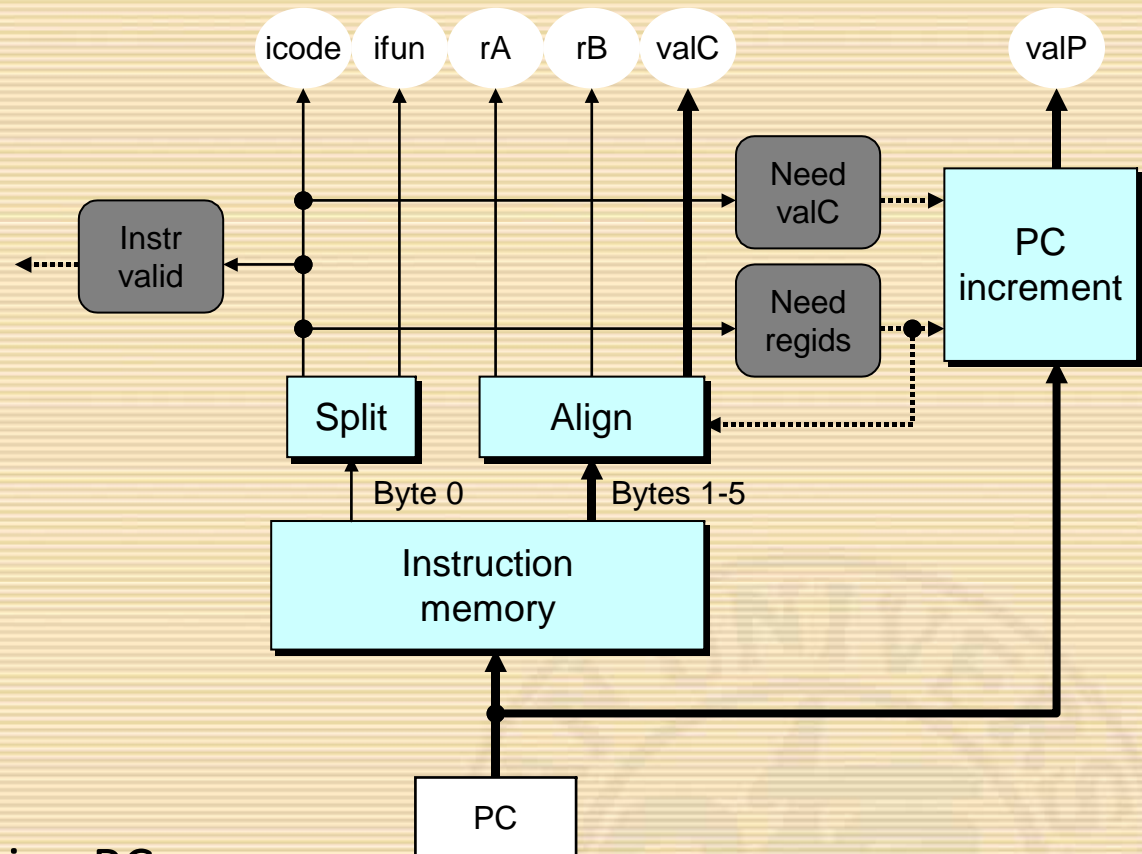
Fetch





Fetch Logic

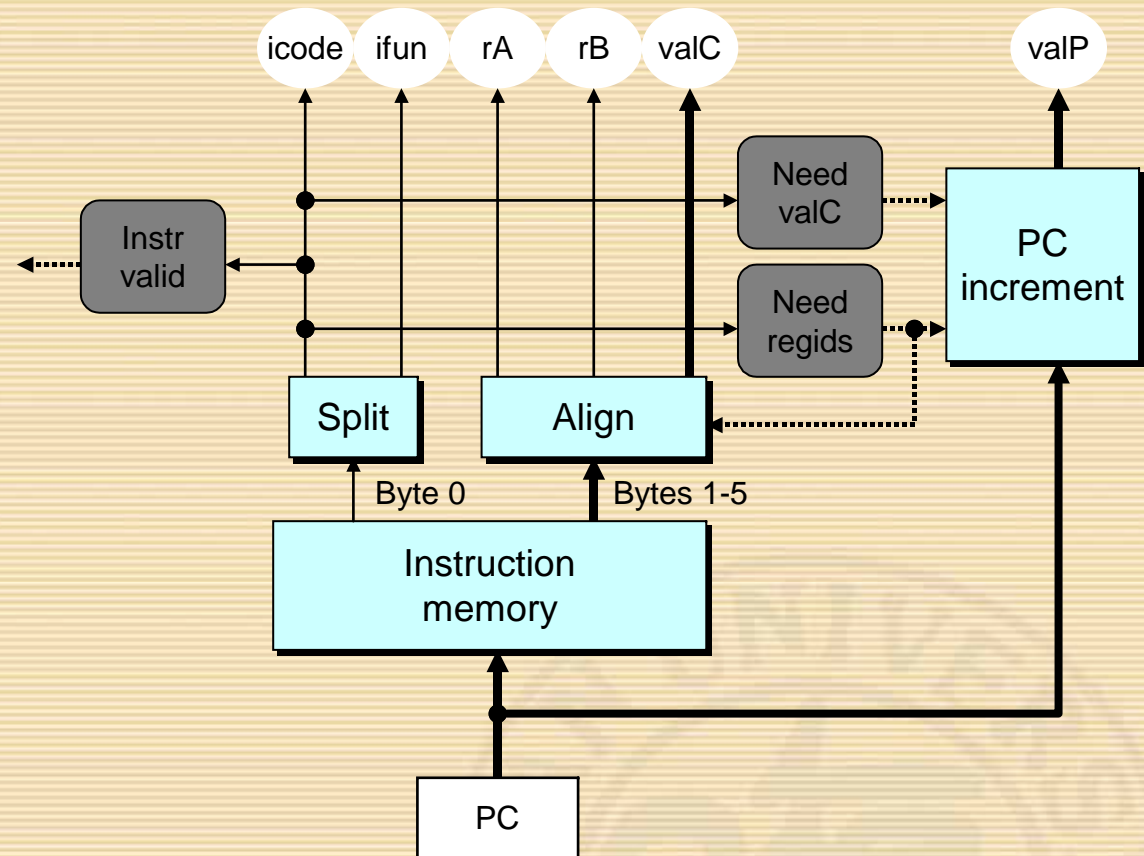
Figure 4.25 P299



- Predefined Blocks
 - PC: Register containing PC
 - Instruction memory: Read 6 bytes (PC to PC+5)
 - Split: Divide instruction byte into icode and ifun
 - Align: Get fields for rA, rB, and valC



Fetch Logic

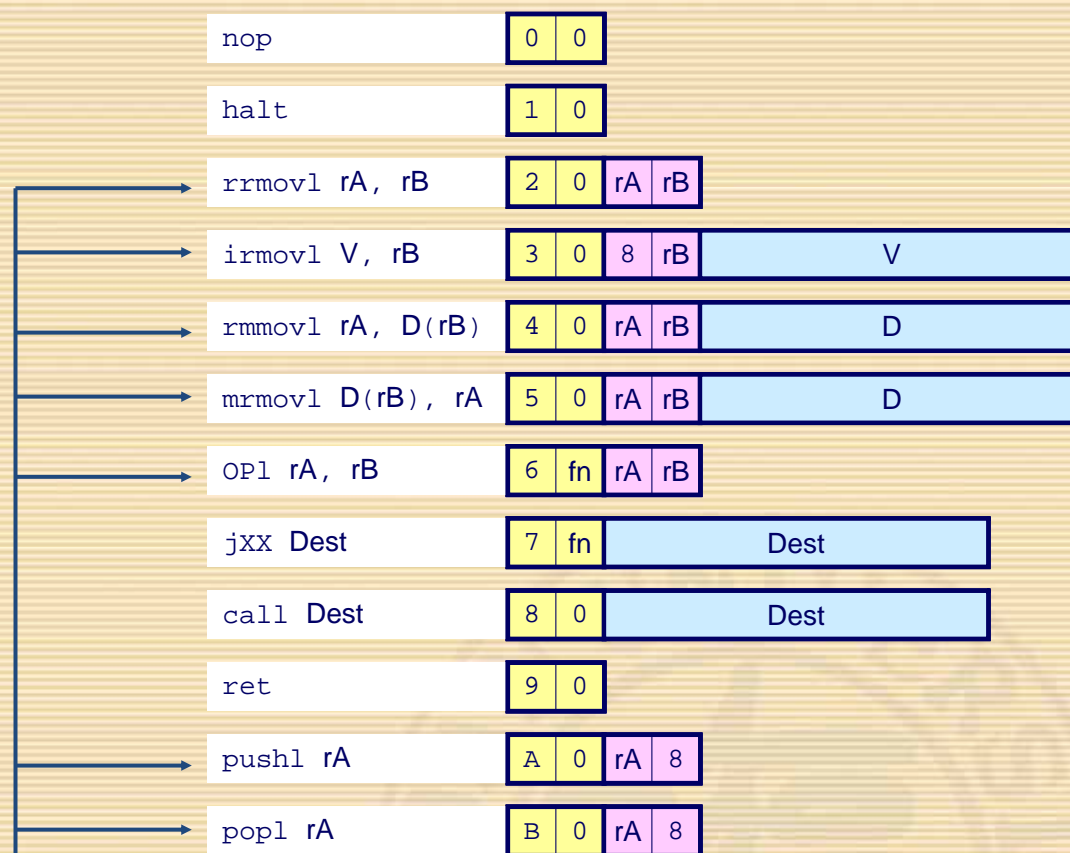


- Control Logic

- Instr. Valid: Is this instruction valid?
- Need regids: Does this instruction have a register bytes?
- Need valC: Does this instruction have a constant word?



Fetch Control Logic P300



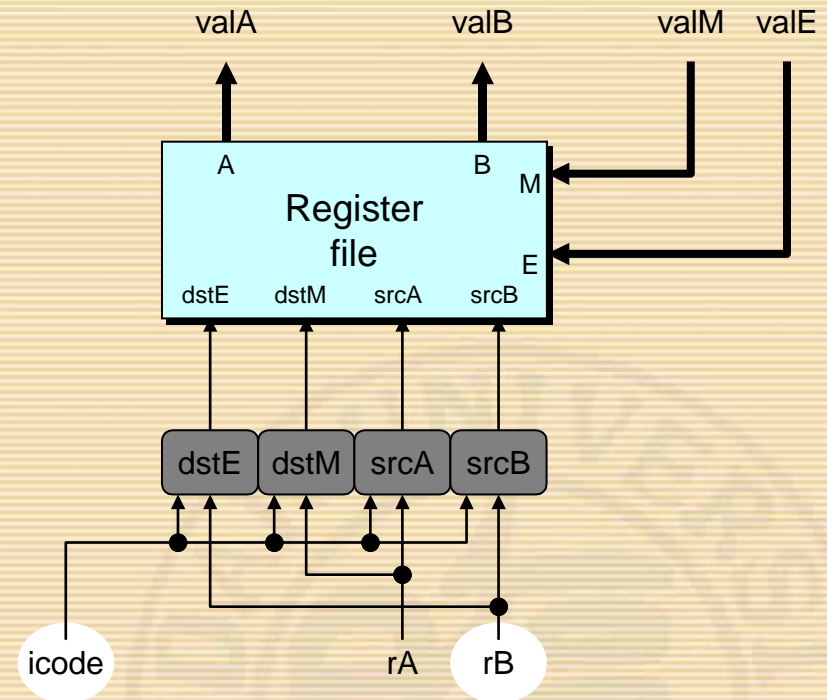
```
bool need_regids =  
    icode in { IRRMOVL, IOPL, IPUSHL, IPOPL,  
               IIRMOVL, IRMMOVL, IMRMOVL };  
bool instr_valid = icode in  
    { INOP, IHALT, IRRMOVL, IIRMOVL, IRMMOVL, IMRMOVL,  
      IOPL, IJXX, ICALL, IRET, IPUSHL, IPOPL };
```



Decode & Write-Back Logic

Figure 4.26 P300

- Register File
 - Read ports A, B
 - Write ports E, M
 - Addresses are register IDs or 8 (no access)
- Control Logic
 - srcA, srcB: read port addresses
 - dstE, dstM: write port addresses





Y86 Instruction Set

P259

Byte 0 1 2 3 4 5

nop

0 0

halt

1 0

rrmovl rA, rB

2 0 rA rB

irmovl V, rB

3 0 8 rB V

rmmovl rA, D(rB)

4 0 rA rB D

mrmovl D(rB), rA

5 0 rA rB D

Op1 rA, rB

6 fn rA rB

jXX Dest

7 fn Dest

call Dest

8 0 Dest

ret

9 0

pushl rA

A 0 rA 8

popl rA

B 0 rA 8

addl

6 0

subl

6 1

andl

6 2

xorl

6 3

jmp

7 0

jle

7 1

j1

7 2

je

7 3

jne

7 4

jge

7 5

jg

7 6



A Source

	OPl rA, rB	
Decode	valA \leftarrow R[rA]	Read operand A
	rmmovl rA, D(rB)	
Decode	valA \leftarrow R[rA]	Read operand A
	popl rA	
Decode	valA \leftarrow R[%esp]	Read stack pointer
	jXX Dest	
Decode		No operand
	call Dest	
Decode		No operand
	ret	
Decode	valA \leftarrow R[%esp]	Read stack pointer

```
int srcA = [  
    icode in { IRRMOVL, IRMMOVL, IOPL, IPUSHL } : rA;  
    icode in { IPOPL, IRET } : RESP;  
    1 : RNONE; # Don't need register  
];
```



E Destination

	OPl rA, rB	
Write-back	$R[rB] \leftarrow \text{valE}$	Write back result
	rmmovl rA, D(rB)	
Write-back		None
	popl rA	
Write-back	$R[\%esp] \leftarrow \text{valE}$	Update stack pointer
	jXX Dest	
Write-back		None
	call Dest	
Write-back	$R[\%esp] \leftarrow \text{valE}$	Update stack pointer
	ret	
Write-back	$R[\%esp] \leftarrow \text{valE}$	Update stack pointer

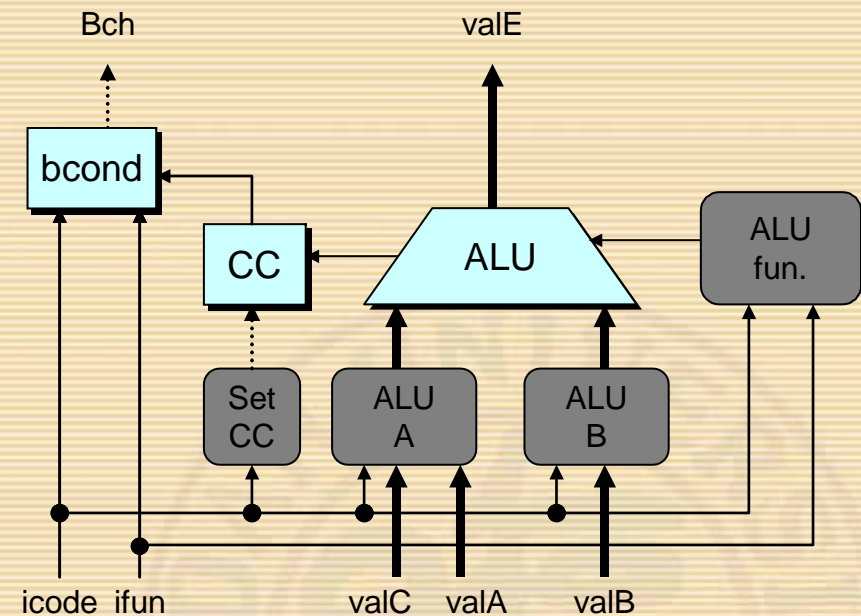
```
int dstE = [  
    icode in { IRRMOVL, IIRMOVL, IOPL } : rB;  
    icode in { IPUSHL, IPOPL, ICALL, IRET } : RESP;  
    1 : RNONE; # Don't need register  
];
```




Execute Logic

Figure 4.27 P302

- Units
 - ALU
 - Implements 4 required functions
 - Generates condition code values
 - CC
 - Register with 3 condition code bits
 - bcond
 - Computes branch flag
- Control Logic
 - Set CC: Should condition code register be loaded?
 - ALU A: Input A to ALU
 - ALU B: Input B to ALU
 - ALU fun: What function should ALU compute?





Y86 Instruction Set

P259

Byte	0	1	2	3	4	5
nop	0	0				
halt	1	0				
rrmovl rA, rB	2	0	rA	rB		
irmovl V, rB	3	0	8	rB	V	
rmmovl rA, D(rB)	4	0	rA	rB	D	
mrmovl D(rB), rA	5	0	rA	rB	D	
OpI rA, rB	6	fn	rA	rB		
jXX Dest	7	fn	Dest			
call Dest	8	0	Dest			
ret	9	0				
pushl rA	A	0	rA	8		
popl rA	B	0	rA	8		

addl	6	0				
subl	6	1				
andl	6	2				
xorl	6	3				
jmp	7	0				
jle	7	1				
j1	7	2				
je	7	3				
jne	7	4				
jge	7	5				
jg	7	6				



ALU A Input

	OPl rA, rB	
Execute	$valE \leftarrow valB \text{ OP } valA$	Perform ALU operation
	rmmovl rA, D(rB)	
Execute	$valE \leftarrow valB + valC$	Compute effective address
	popl rA	
Execute	$valE \leftarrow valB + 4$	Increment stack pointer
	jXX Dest	
Execute		No operation
	call Dest	
Execute	$valE \leftarrow valB + -4$	Decrement stack pointer
	ret	
Execute	$valE \leftarrow valB + 4$	Increment stack pointer

```
int aluA = [  
    icode in { IRRMOVL, IOPL } : valA;  
    icode in { IIRMOVL, IRMMOVL, IMRMOVL } : valC;  
    icode in { ICALL, IPUSHL } : -4;  
    icode in { IRET, IPOPL } : 4;  
    # Other instructions don't need ALU  
];
```

jc@fudan.edu.cn



ALU Operation

	OPl rA, rB	
Execute	$valE \leftarrow valB \text{ OP } valA$	Perform ALU operation
	rmmovl rA, D(rB)	
Execute	$valE \leftarrow valB + valC$	Compute effective address
	popl rA	
Execute	$valE \leftarrow valB + 4$	Increment stack pointer
	jXX Dest	
Execute		No operation
	call Dest	
Execute	$valE \leftarrow valB + -4$	Decrement stack pointer
	ret	
Execute	$valE \leftarrow valB + 4$	Increment stack pointer

```
int alufun = [  
    icode == IOPL : ifun;  
    1 : ALUADD;  
];
```



Condition Set


- `Bool set_cc = icode in { IOPL };`
- We will not discuss the detail of Bcond
 - Though it is also a control unit



Memory Logic

Figure 4.28 P303

- Memory
 - Reads or writes memory word
- Control Logic
 - Mem. read: should word be read?
 - Mem. write: should word be written?
 - Mem. addr.: Select address
 - Mem. data.: Select data





Y86 Instruction Set

P259

Byte 0 1 2 3 4 5

nop 0 0

halt 1 0

rrmovl rA, rB 2 0 rA rB

irmovl V, rB 3 0 8 rB V

rmmovl rA, D(rB) 4 0 rA rB D

mrmovl D(rB), rA 5 0 rA rB D

Op1 rA, rB 6 fn rA rB

jXX Dest 7 fn Dest

call Dest 8 0 Dest

ret 9 0

pushl rA A 0 rA 8

popl rA B 0 rA 8

addl 6 0

subl 6 1

andl 6 2

xorl 6 3

jmp 7 0

jle 7 1

j1 7 2

je 7 3

jne 7 4

jge 7 5

1g 7 6



Memory Address

	OPl rA, rB	
Memory		No operation
	rmmovl rA, D(rB)	
Memory	$M_4[\text{valE}] \leftarrow \text{valA}$	Write value to memory
	popl rA	
Memory	$\text{valM} \leftarrow M_4[\text{valA}]$	Read from stack
	jXX Dest	
Memory		No operation
	call Dest	
Memory	$M_4[\text{valE}] \leftarrow \text{valP}$	Write return value on stack
	ret	
Memory	$\text{valM} \leftarrow M_4[\text{valA}]$	Read return address

```
int mem_addr = [  
    icode in { IRMMOVL, IPUSHL, ICALL, IMRMOVL } : valE;  
    icode in { IPOPL, IRET } : valA;  
    # Other instructions don't need address  
];
```



Memory Read


	OPl rA, rB	
Memory		No operation
	rmmovl rA, D(rB)	
Memory	$M_4[valE] \leftarrow valA$	Write value to memory
	popl rA	
Memory	$valM \leftarrow M_4[valA]$	Read from stack
	jXX Dest	
Memory		No operation
	call Dest	
Memory	$M_4[valE] \leftarrow valP$	Write return value on stack
	ret	
Memory	$valM \leftarrow M_4[valA]$	Read return address



PC Update Logic

Figure 4.29 P304

- New PC
 - Select next value of PC





Y86 Instruction Set

P259

Byte 0 1 2 3 4 5

nop 0 0

halt 1 0

rrmovl rA, rB 2 0 rA rB

irmovl V, rB 3 0 8 rB V

rmmovl rA, D(rB) 4 0 rA rB D

mrmovl D(rB), rA 5 0 rA rB D

Op1 rA, rB 6 fn rA rB

jXX Dest 7 fn Dest

call Dest 8 0 Dest

ret 9 0

pushl rA A 0 rA 8

popl rA B 0 rA 8

addl 6 0

subl 6 1

andl 6 2

xorl 6 3

jmp 7 0

jle 7 1

j1 7 2

je 7 3

jne 7 4

jge 7 5

1g 7 6



PC Update

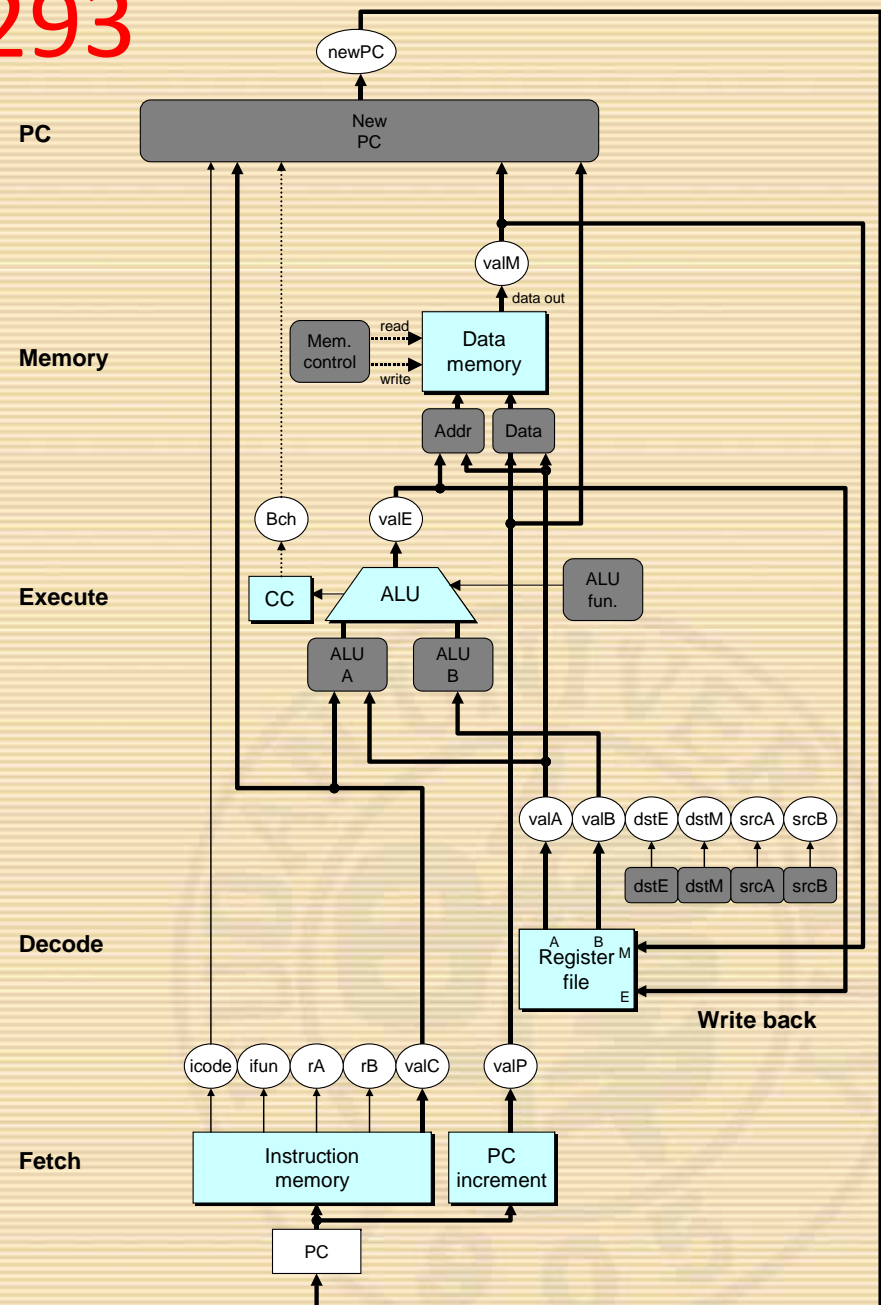
	OPl rA, rB	
PC update	$PC \leftarrow valP$	Update PC
	rmmovl rA, D(rB)	
PC update	$PC \leftarrow valP$	Update PC
	popl rA	
PC update	$PC \leftarrow valP$	Update PC
	jXX Dest	
PC update	$PC \leftarrow Bch ? valC : valP$	Update PC
	call Dest	
PC update	$PC \leftarrow valC$	Set PC to destination
	ret	
PC update	$PC \leftarrow valM$	Set PC to return address

```
int new_pc = [  
    icode == ICALL : valC;  
    icode == IJXX && Bch : valC;  
    icode == IRET : valM;  
    1 : valP;  
];
```




SEQ Hardware P293

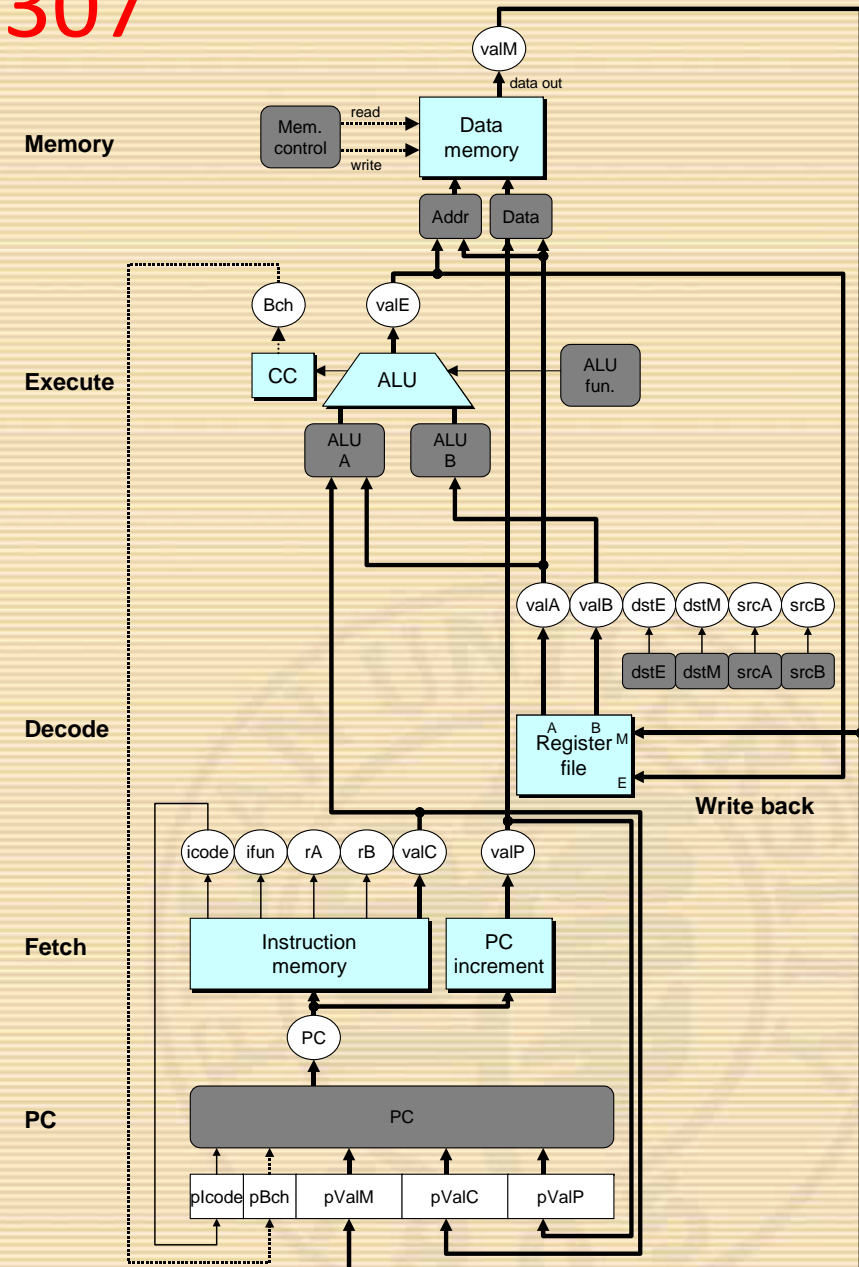
- Stages occur in sequence
- One operation in process at a time





SEQ+ Hardware P307

- Still sequential implementation
- Reorder PC stage to put at beginning
- PC Stage
 - Task is to select PC for current instruction
 - Based on results computed by previous instruction
- Processor State
 - PC is no longer stored in register
 - But, can determine PC based on other stored information





PC Computation

```
Int pc= [  
    plcode == ICALL : pValC;  
    plcode == IJXX && bBch : pValC;  
    Plcode == IRET : pValM;  
    1 : pValP;  
];
```



SEQ Summary

- Implementation
 - Express every instruction as series of simple steps
 - Follow same general flow for each instruction type
 - Assemble registers, memories, predesigned combinational blocks
 - Connect with control logic
- Limitations
 - Too slow to be practical
 - In one cycle, must propagate through instruction memory, register file, ALU, and data memory
 - Would need to run clock very slowly
 - Hardware units only active for fraction of clock cycle



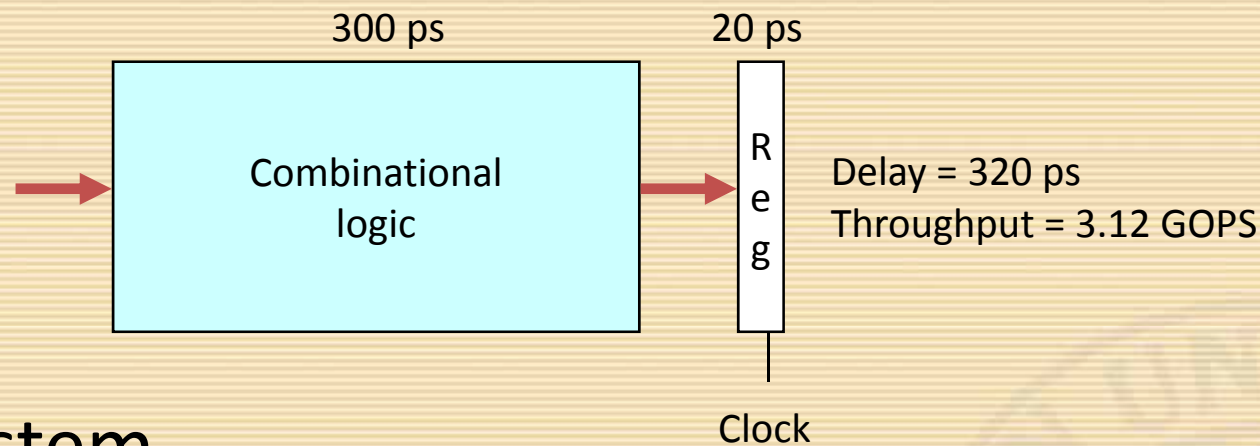
Lecutre



Pipelined Implementation



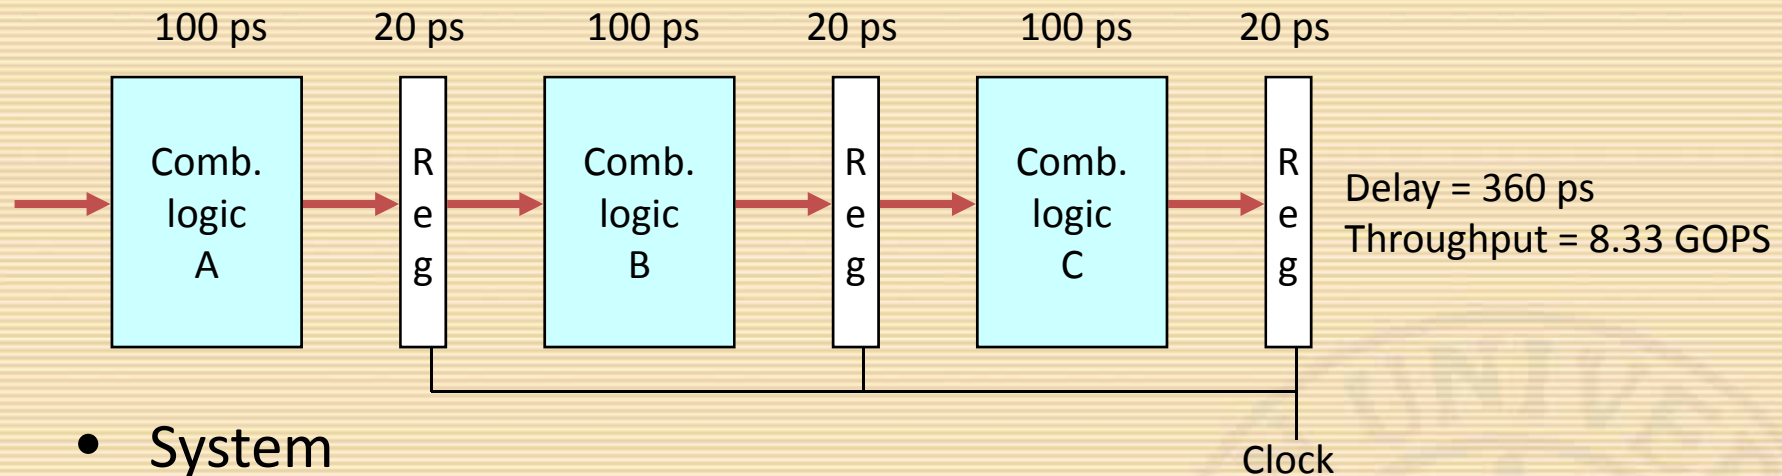
Computational Example



- System
 - Computation requires total of 300 picoseconds
 - Additional 20 picoseconds to save result in register
 - Can must have clock cycle of at least 320 ps



3-Way Pipelined Version




- System

- Divide combinational logic into 3 blocks of 100 ps each
- Can begin new operation as soon as previous one passes through stage A.
 - Begin new operation every 120 ps
- Overall latency increases
 - 360 ps from start to finish



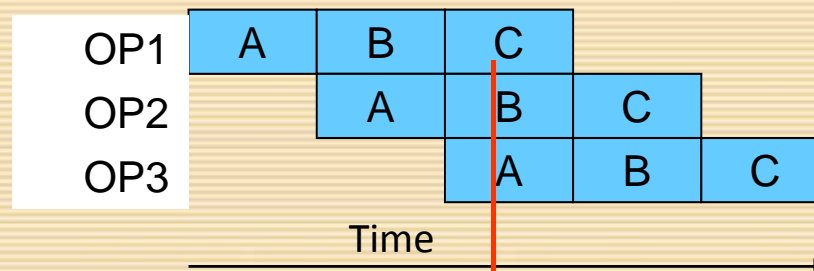
Pipeline Diagrams

- Unpipelined




- Cannot start new operation until previous one completes

- 3-Way Pipelined

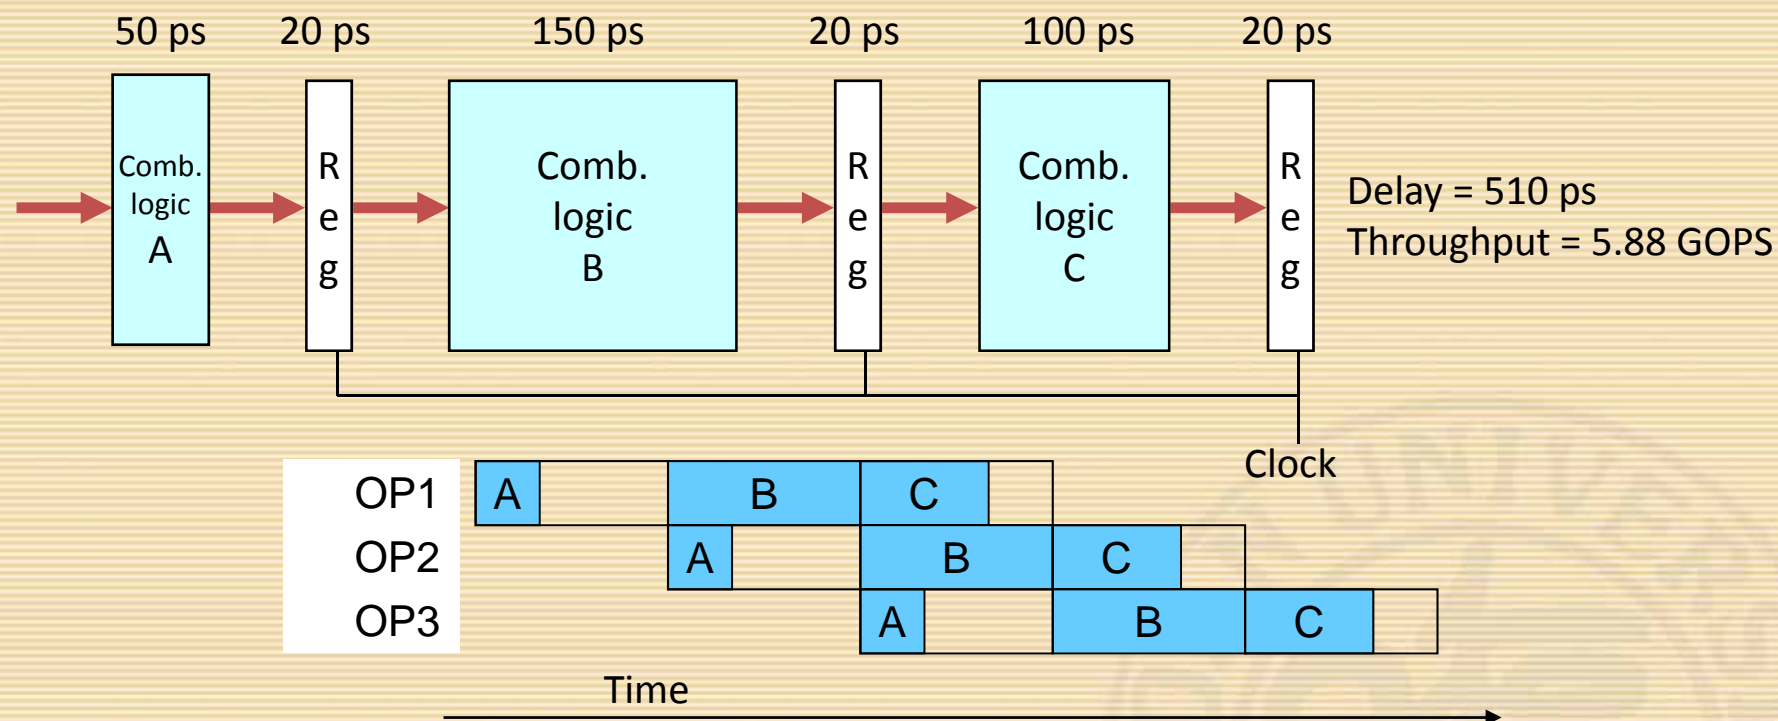


- Up to 3 operations in process simultaneously





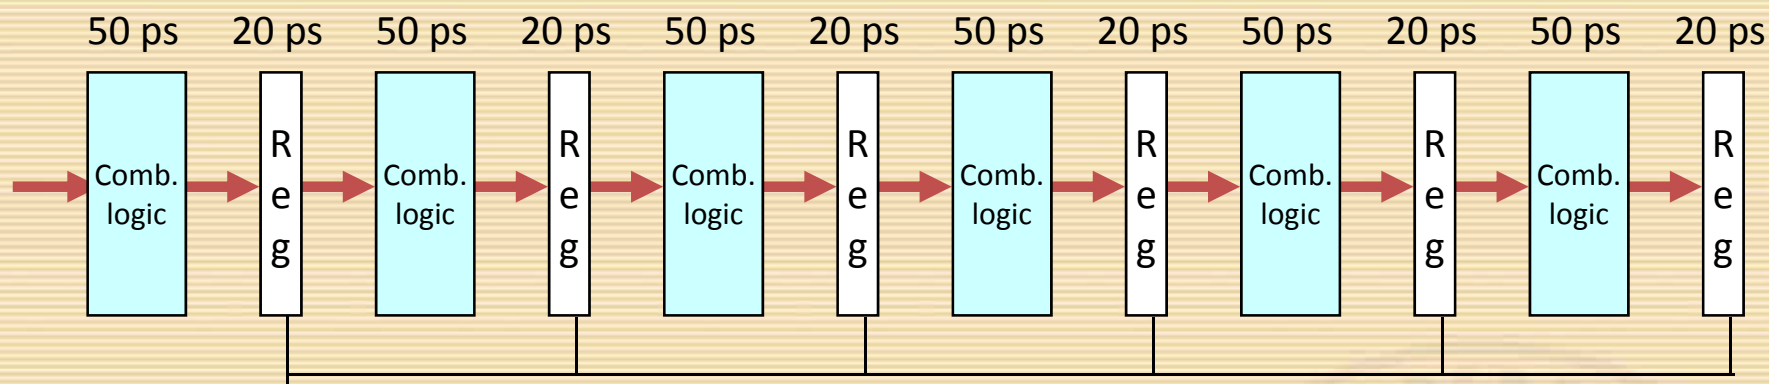
Limitations: Nonuniform Delays



- Throughput limited by slowest stage
- Other stages sit idle for much of the time
- Challenging to partition system into balanced stages



Limitations: Register Overhead



- Clock
- Delay = 420 ps, Throughput = 14.29 GOPS
- As try to deepen pipeline, overhead of loading registers becomes more significant
 - Percentage of clock cycle spent loading register:
 - 1-stage pipeline: 6.25%
 - 3-stage pipeline: 16.67%
 - 6-stage pipeline: 28.57%
 - High speeds of modern processor designs obtained through very deep pipelining

Overhead: 开销



Adding Pipeline Registers

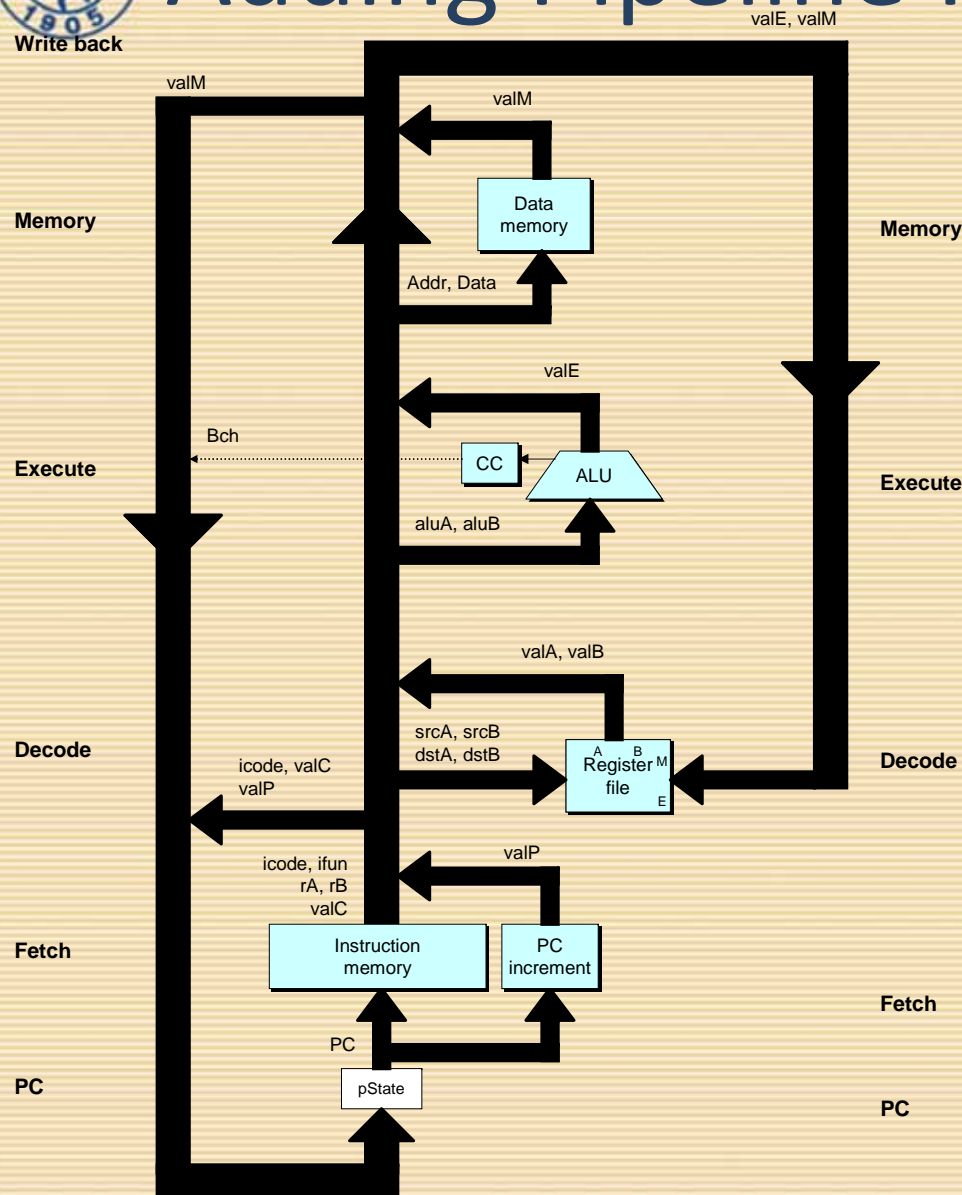


Figure 4.30 P306

jc@fudan.edu.cn

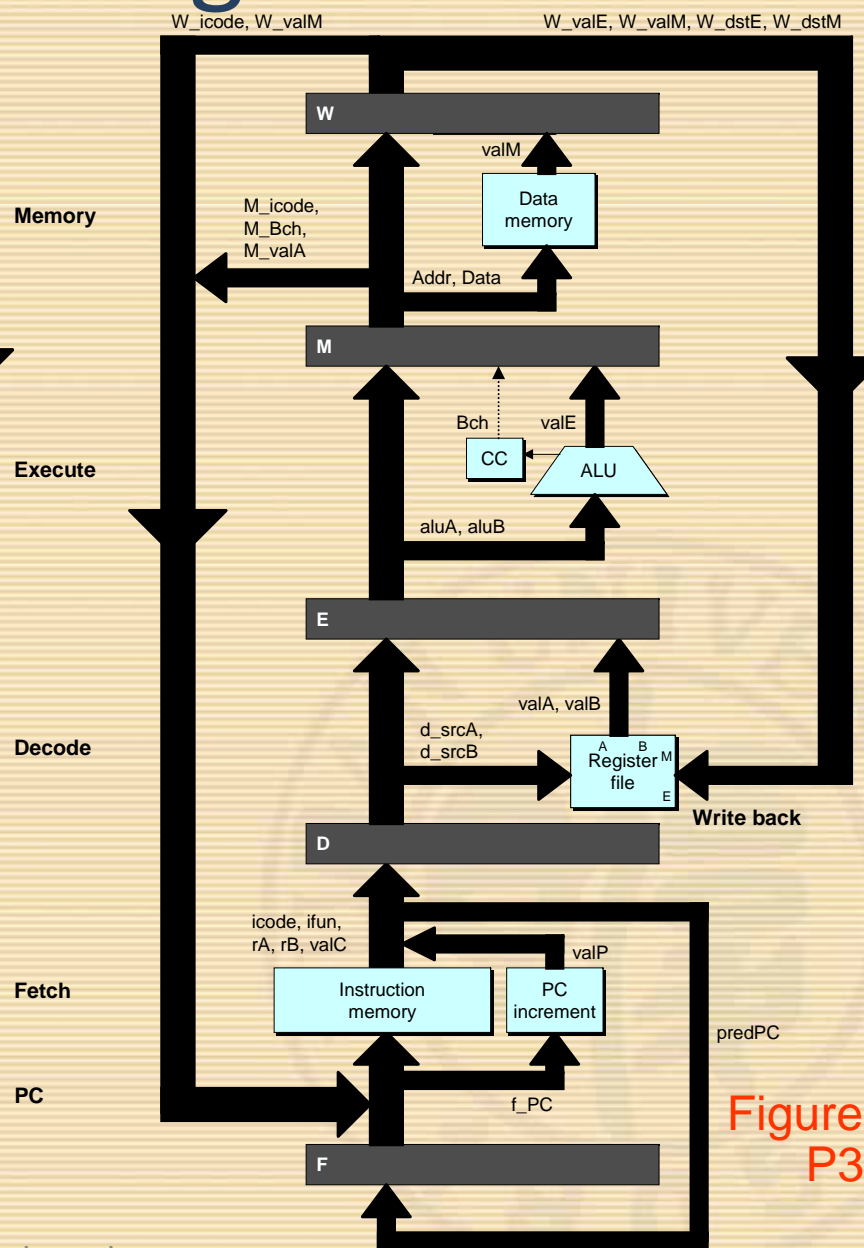
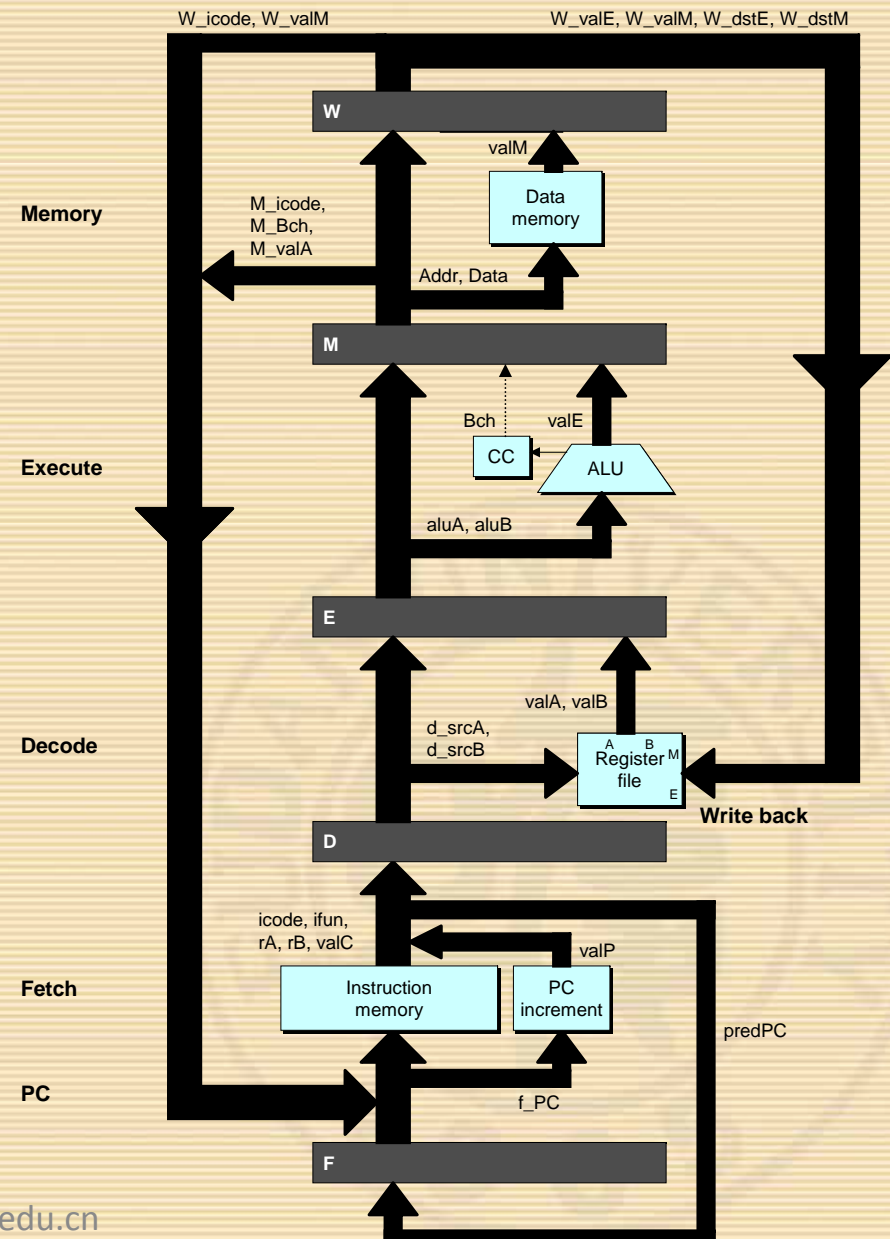


Figure 4.39
P318



Pipeline Stages

- Fetch
 - Select current PC
 - Read instruction
 - Compute incremented PC
- Decode
 - Read program registers
- Execute
 - Operate ALU
- Memory
 - Read or write data memory
- Write Back
 - Update register file






PIPE- Hardware

Figure 4.41 P320

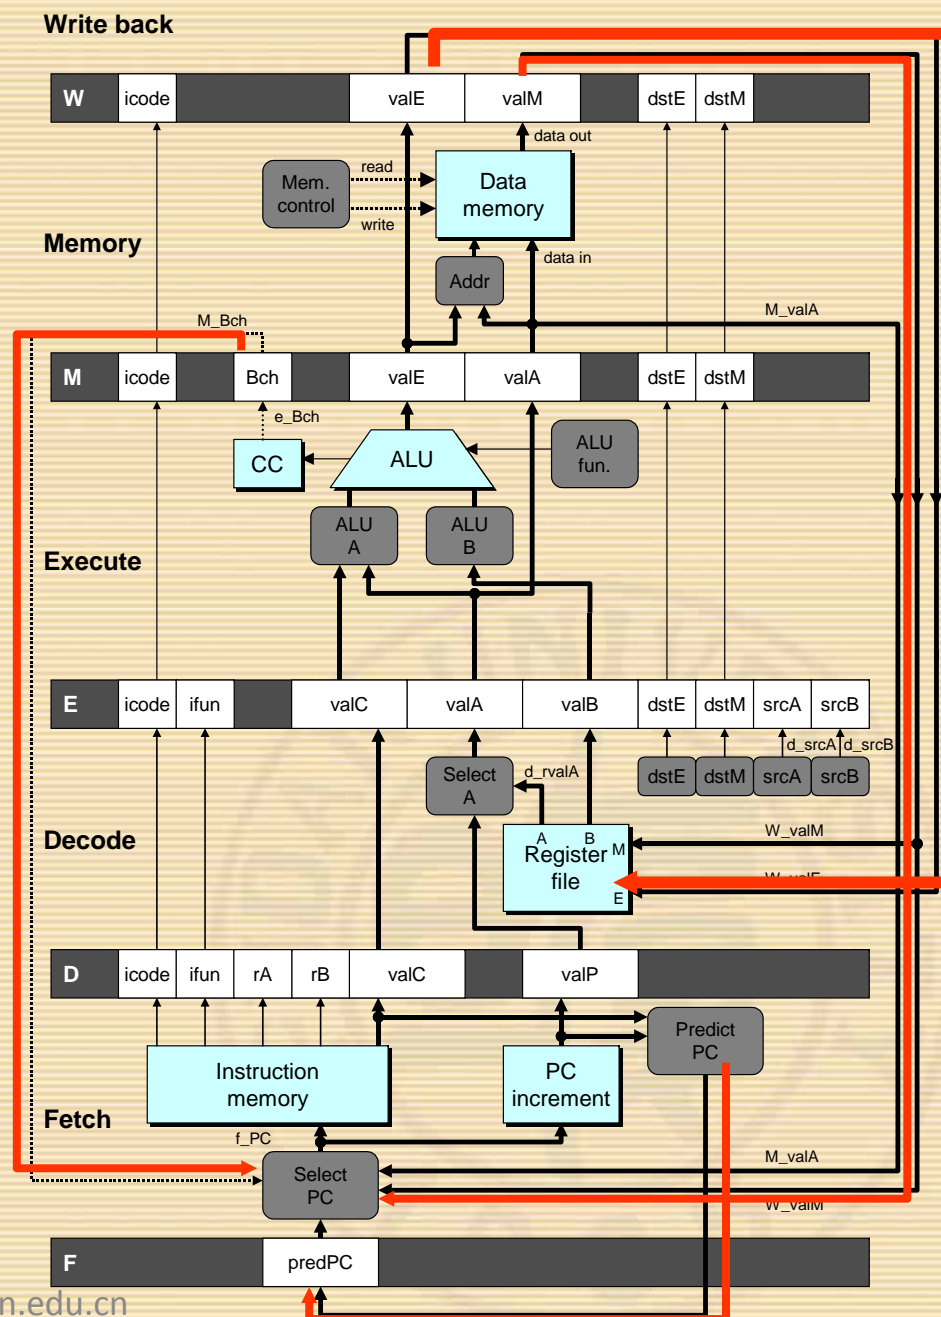
- Pipeline registers hold intermediate values from instruction execution
- Forward (Upward) Paths
 - Values passed from one stage to next
 - Cannot jump past stages
 - e.g., valC passes through decode





Feedback Paths

- Predicted PC
 - Guess value of next PC
- Branch information
 - Jump taken/not-taken
 - Fall-through or target address
- Return point
 - Read from memory
- Register updates
 - To register file write ports





e PC

The diagram illustrates the ePC (enhanced PC) architecture, showing the flow of data and control signals between various components. The components and their interactions are as follows:

- Instruction Memory:** Provides **Byte 0** to the **Split** block and **Bytes 1-5** to the **Align** block.
- Split:** Outputs to the **icode** field of the instruction register.
- Align:** Outputs to the **rA** and **rB** fields of the instruction register.
- Instruction Register:** A register containing fields: **D**, **icode**, **ifun**, **rA**, **rB**, **valC**, and **valP**.
- Instr valid:** A control signal output from the instruction register.
- Need valC:** A control signal output from the instruction register to the **PC increment** block.
- Need regids:** A control signal output from the instruction register to the **PC increment** block.
- PC increment:** A block that increments the PC based on **Need valC** and **Need regids**. It outputs to the **valC** field of the instruction register.
- Predict PC:** A block that predicts the next PC value. It outputs to the **valP** field of the instruction register.
- Select PC:** A block that selects the next PC value based on **valC** and **valP**. It outputs to the **predPC** field of the **F** register.
- F (Future) Register:** A register containing the **predPC** field.
- Control Signals:**
 - M_icode:** A control signal output from the instruction register.
 - M_Bch:** A control signal output from the instruction register.
 - M_valA:** A control signal output from the instruction register.
 - W_icode:** A control signal output from the instruction register.
 - W_valM:** A control signal output from the instruction register.

- jc@fudan.edu.cn



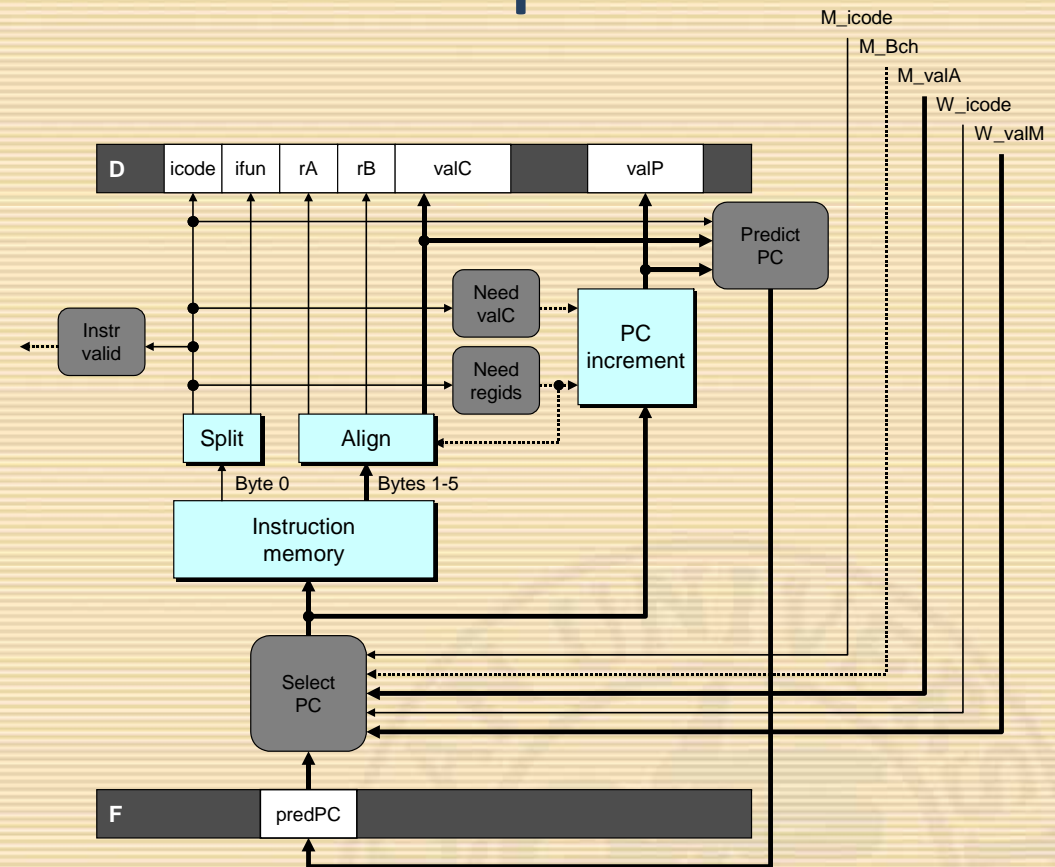
Our Prediction Strategy

- Instructions that Don't Transfer Control
 - Predict next PC to be valP
 - Always reliable
- Call and Unconditional Jumps
 - Predict next PC to be valC (destination)
 - Always reliable
- Conditional Jumps
 - Predict next PC to be valC (destination)
 - Only correct if branch is taken
 - Typically right 60% of time
- Return Instruction
 - Don't try to predict



Recovering from PC Misprediction

Figure 4.56 P338



- Mispredicted Jump
 - Will see branch flag once instruction reaches memory stage
 - Can get fall-through PC from valA
- Return Instruction
 - Will get return PC when `ret` reaches write-back stage



Select PC

P338 and Figure 4.56


- **int f_PC = [**
- **#mispredicted branch. Fetch at incremented PC**
- **M_icode == IJXX && !M_Bch : M_valA;**
- **#completion of RET instruction**
- **W_icode == IRET : W_valM;**
- **#default: Use predicted value of PC**
- **1: F_predPC**
- **];**
- **Int new_F_predPC = [**
- **f_icode in {IJXX, ICALL} : f_valC;**
- **1: f_valP;**
- **];**




Pipeline Demonstration

Figure 4.40 P319

```
irmovl    $1,%eax    #I1
irmovl    $2,%ecx    #I2
irmovl    $3,%edx    #I3
irmovl    $4,%ebx    #I4
halt                      #I5
```

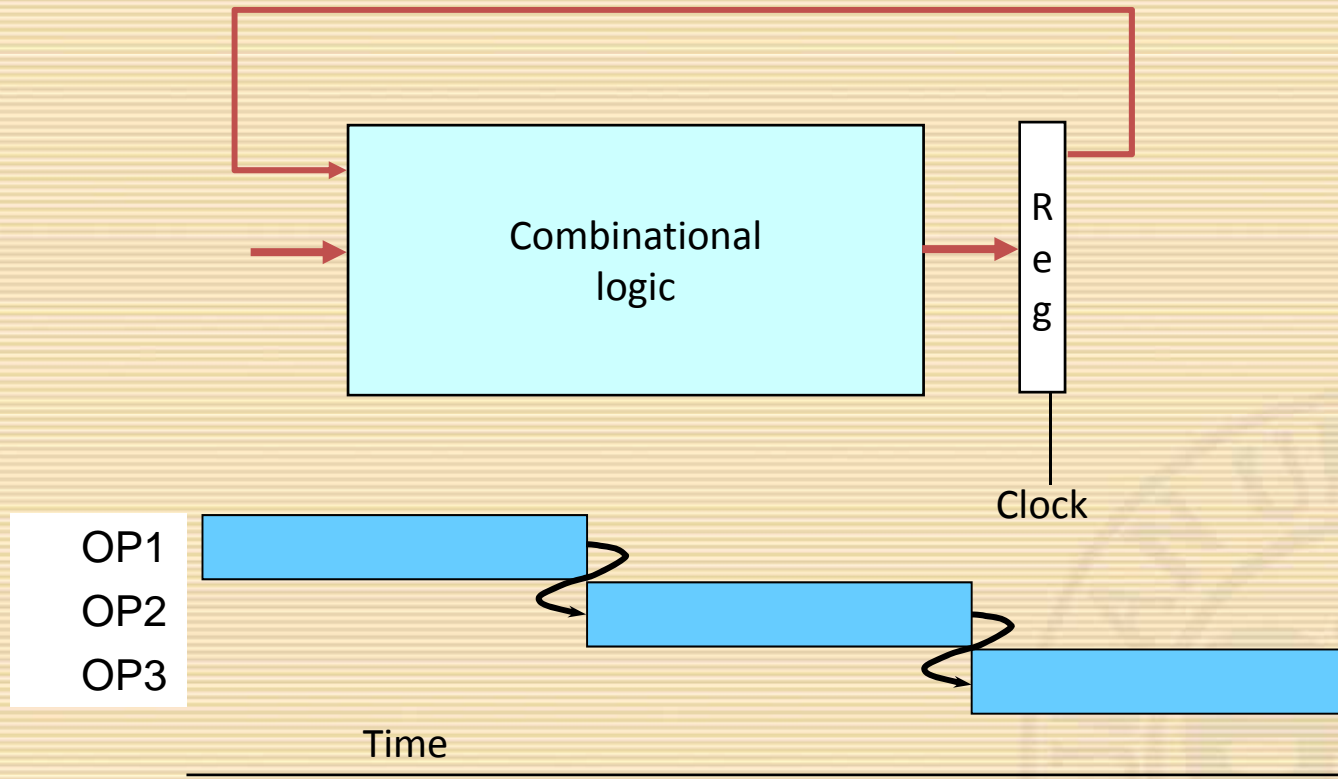


Cycle 5





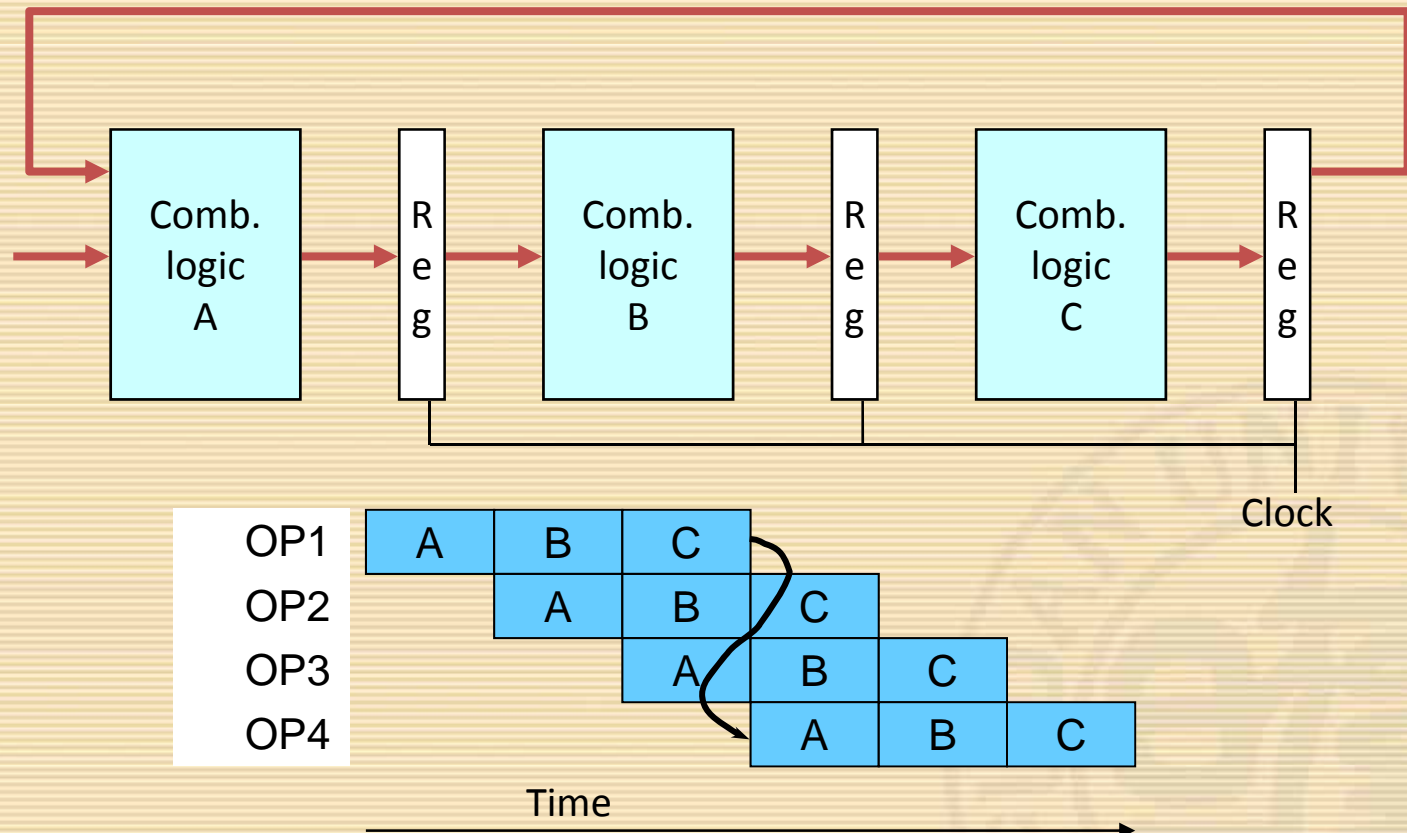
Data Dependencies



- System
 - Each operation depends on result from preceding one



Data Hazards



- Result does not feed back around in time for next operation
- Pipelining has changed behavior of system



Data Dependencies in Processors

```
1    irmovl $50, %eax
2    addl %eax, %ebx
3    mrmovl 100(%ebx), %edx
```

- Result from one instruction used as operand for another
 - Read-after-write (RAW) dependency
- Very common in actual programs
- Must make sure our pipeline handles these properly
 - Get correct results
 - Minimize performance impact



Control Dependence

- Example:
 - Loop:
 - `Subl %edx, %ebx`
 - `Jne targ`
 - `Irmovl $10, %edx`
 - `Jmp loop`
 - Targ:
 - `Halt`
 - The `jne` instruction create a control dependency.
 - Which instruction will be executed?



Data Dependencies

```
# demo-h0.ys  
0x000: irmovl $10,%edx  
0x006: irmovl $3,%eax  
0x00c: addl %edx,%eax  
0x00e: halt
```


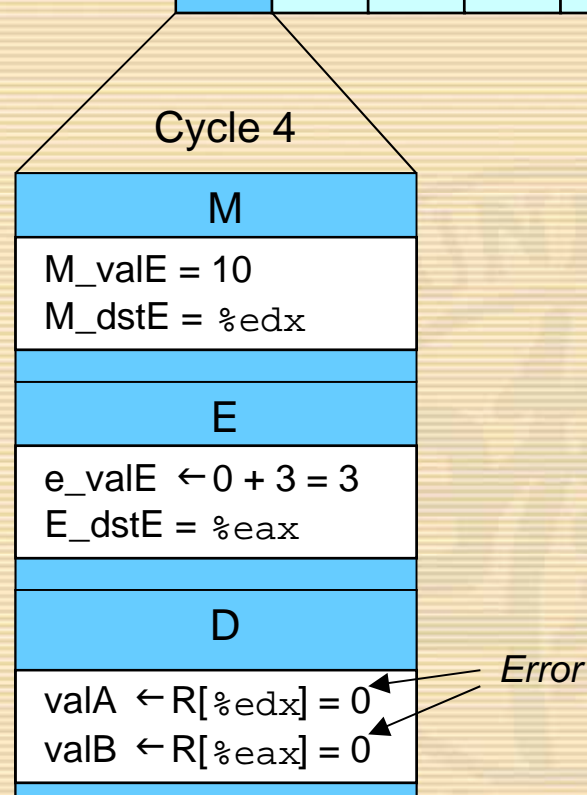


Figure 4.45 P327





Data Dependencies: 1 nop

```
# demo-h1.ys
```

```
0x000: irmovl $10,%edx
```

```
0x006: irmovl $3,%eax
```

```
0x00c: nop
```

```
0x00d: addl %edx,%eax
```

```
0x00f: halt
```


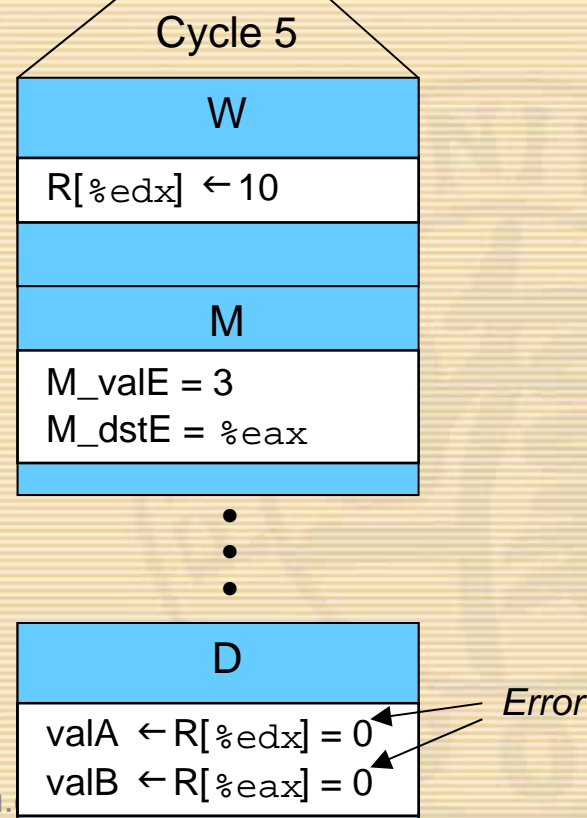


Figure 4.44 P326

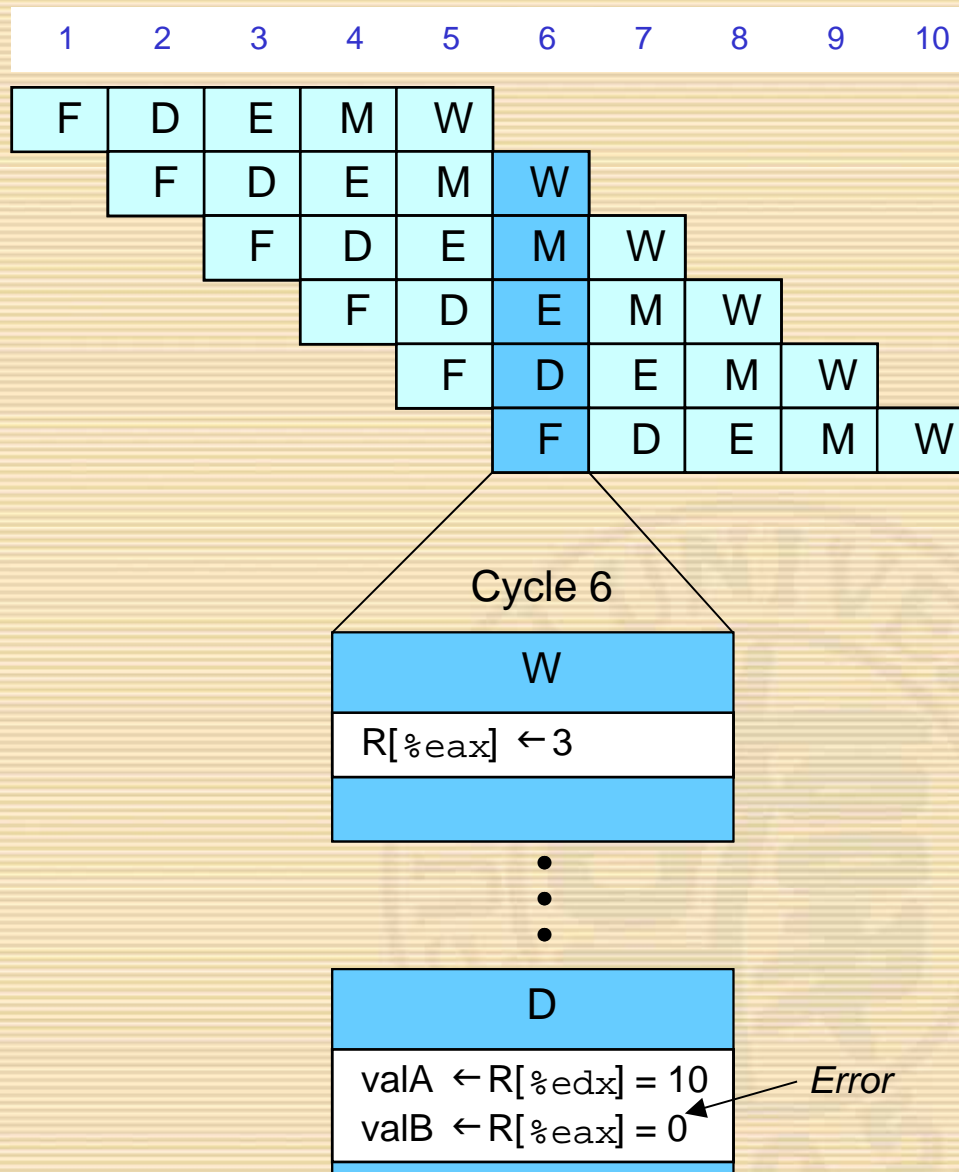




Data Dependencies: 2 nops

```
# demo-h2.js
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: nop
0x00e: addl %edx,%eax
0x010: halt
```

Figure 4.43 P325





Data Dependencies: 3 nops

```
# demo-h3.ys
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: nop
0x00e: nop
0x00f: addl %edx,%eax
0x011: halt
```

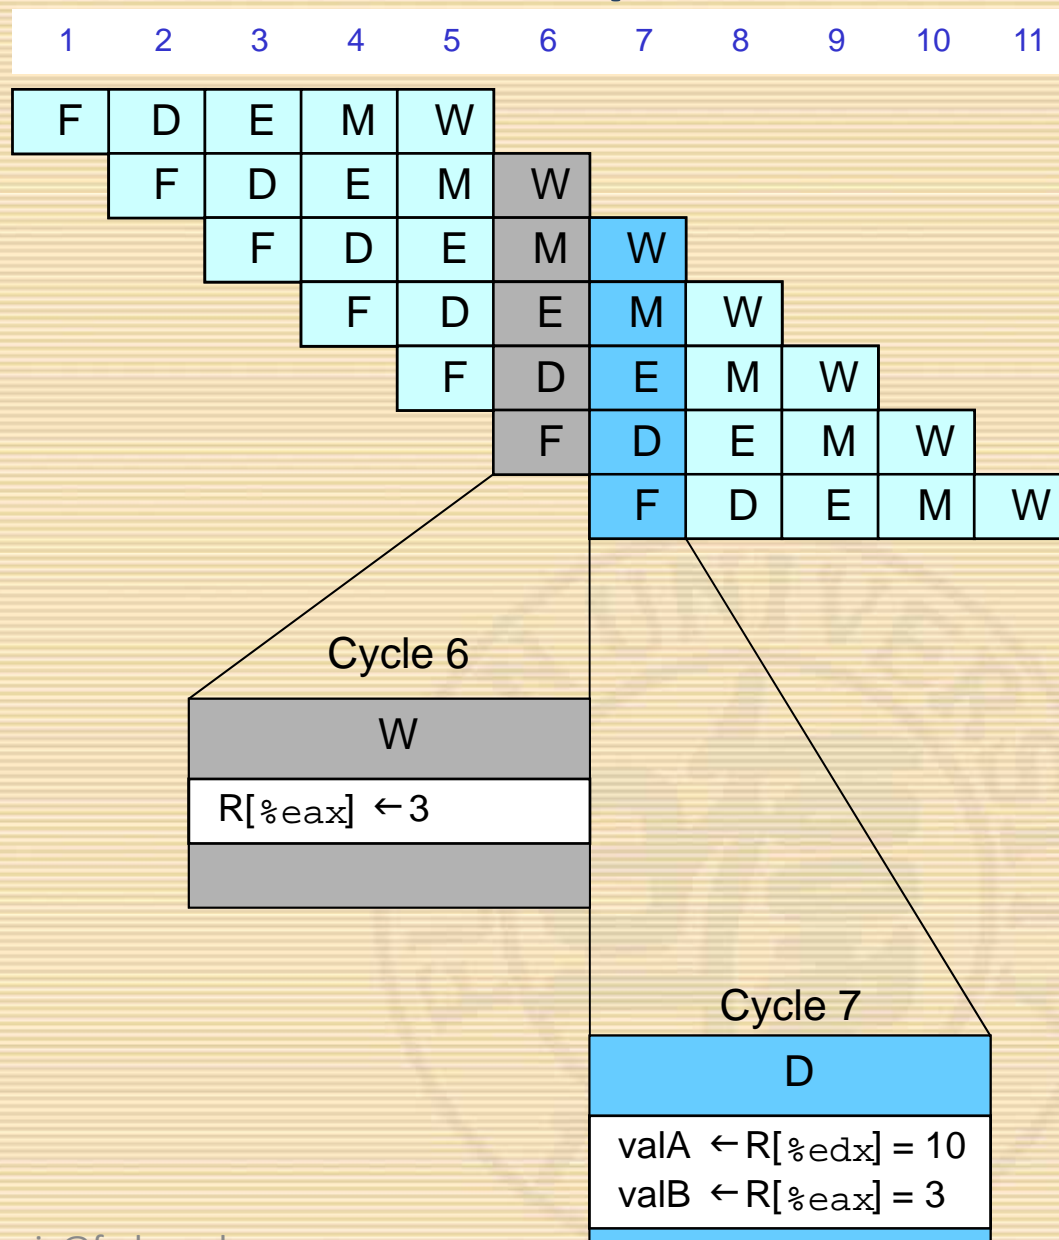


Figure 4.42 P324



Classes of Data Hazards

- Hazards can potentially occur when one instruction updates part of the program state that read by a later instruction
- Program states:
 - Program registers
 - The hazards already identified.
 - Condition codes
 - Both written and read in the execute stage.
 - No hazards can arise
 - Program counter
 - Conflicts between updating and reading PC cause control hazards
 - Memory
 - Both written and read in the memory stage.
 - Without self-modified code, no hazards.



Data Dependencies: 2 nops

```
# demo-h2.ys
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: nop
0x00e: addl %edx,%eax
0x010: halt
```

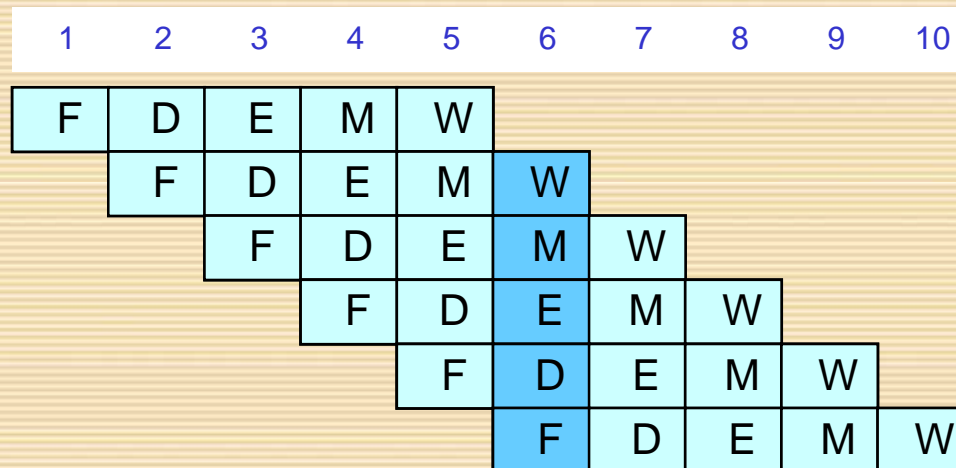
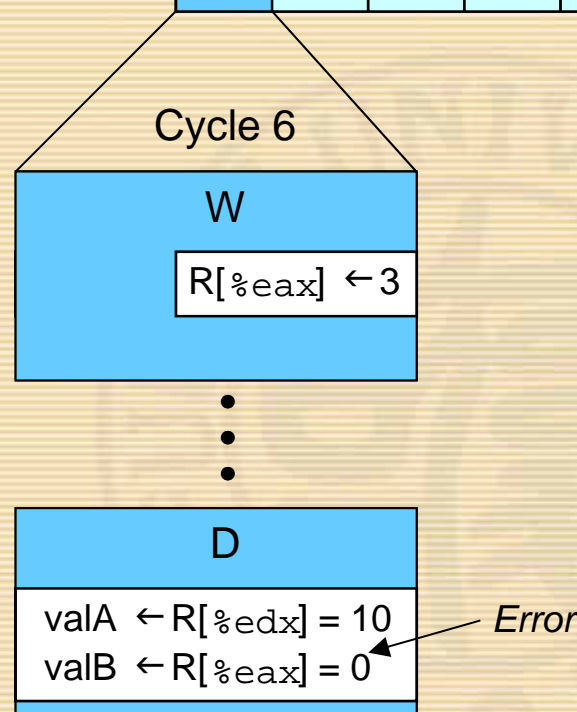


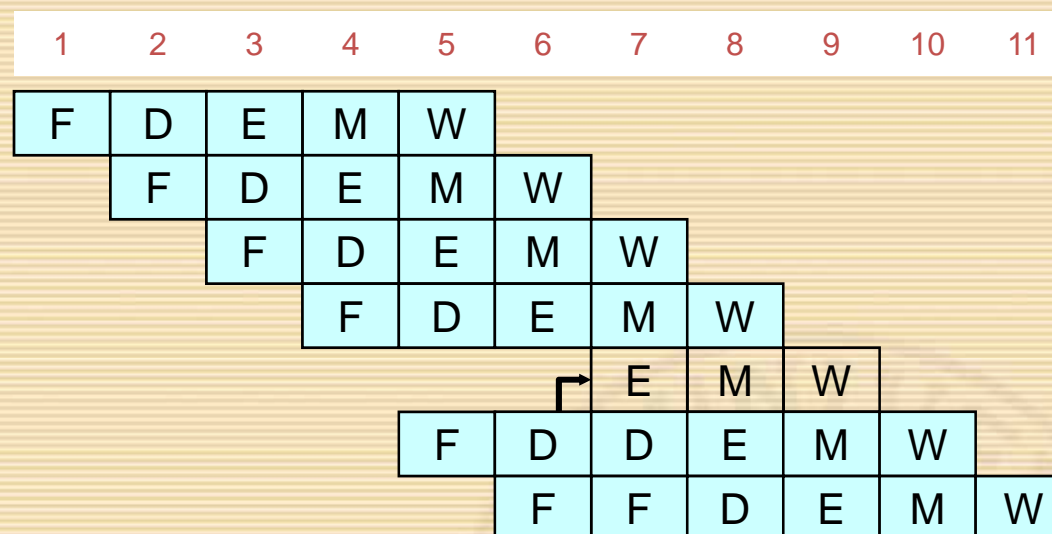
Figure 4.43 P325





Stalling for Data Dependencies

```
# demo-h2.ys
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: nop
        bubble
0x00e: addl %edx,%eax
0x010: halt
```



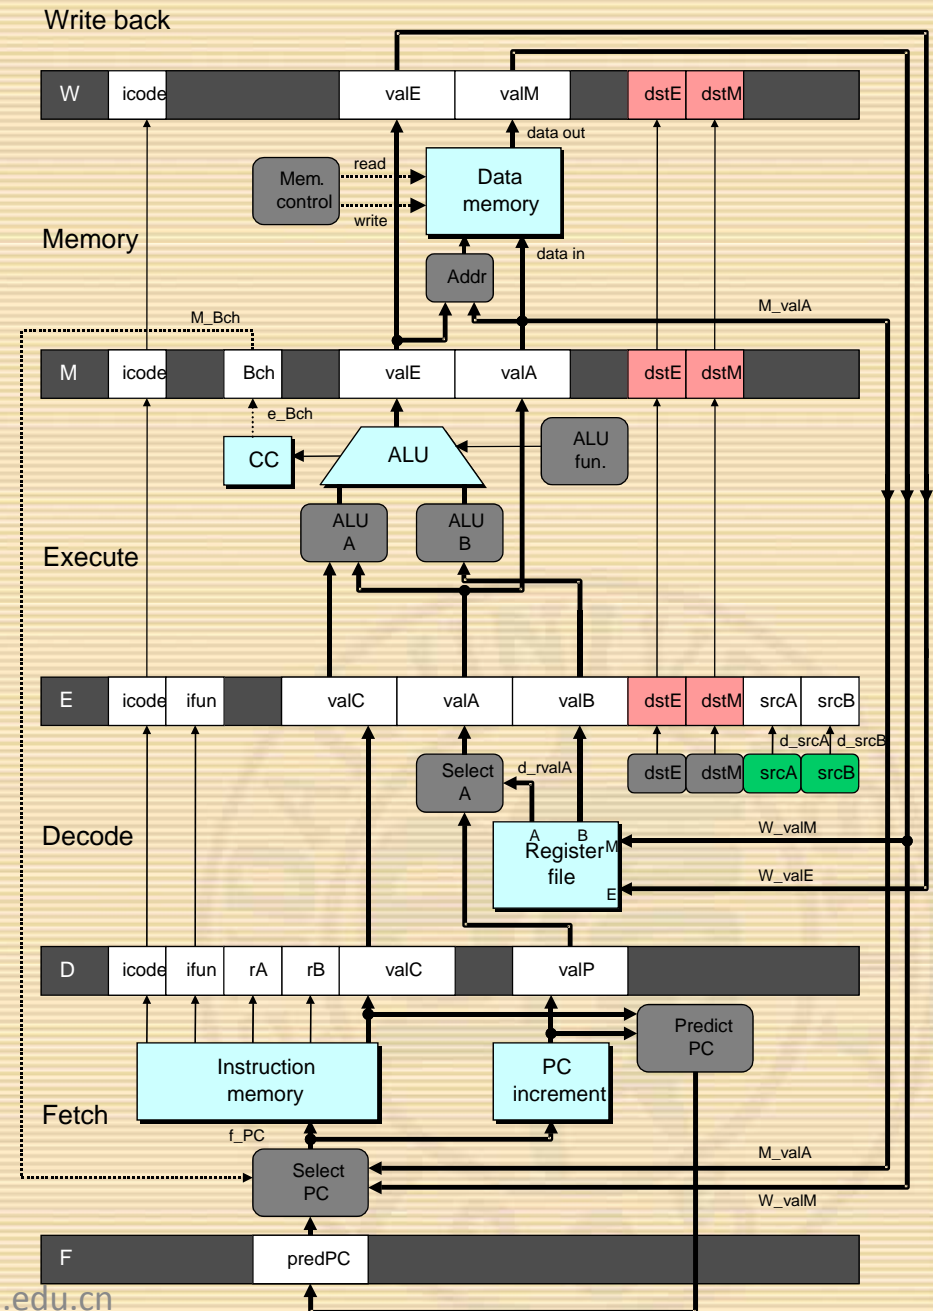
- If instruction follows too closely after one that writes register, slow it down
- Hold instruction in decode
- Dynamically inject nop into execute stage

Stall: 停止, 迟延



Stall Condition

- Source Registers
 - srcA and srcB of current instruction in decode stage
- Destination Registers
 - dstE and dstM fields
 - Instructions in execute, memory, and write-back stages
- Condition
 - $\text{srcA} == \text{dstE}$ or $\text{srcA} == \text{dstM}$
 - $\text{srcB} == \text{dstE}$ or $\text{srcB} == \text{dstM}$
- Special Case
 - Don't stall for register ID 8
 - Indicates absence of register operand





Detecting Stall Condition

```
# demo-h2.ys
```

```
0x000: irmovl $10,%edx
```

```
0x006: irmovl $3,%eax
```

```
0x00c: nop
```

```
0x00d: nop
```

bubble

```
0x00e: addl %edx,%eax
```

```
0x010: halt
```





Figure 4.46 P328

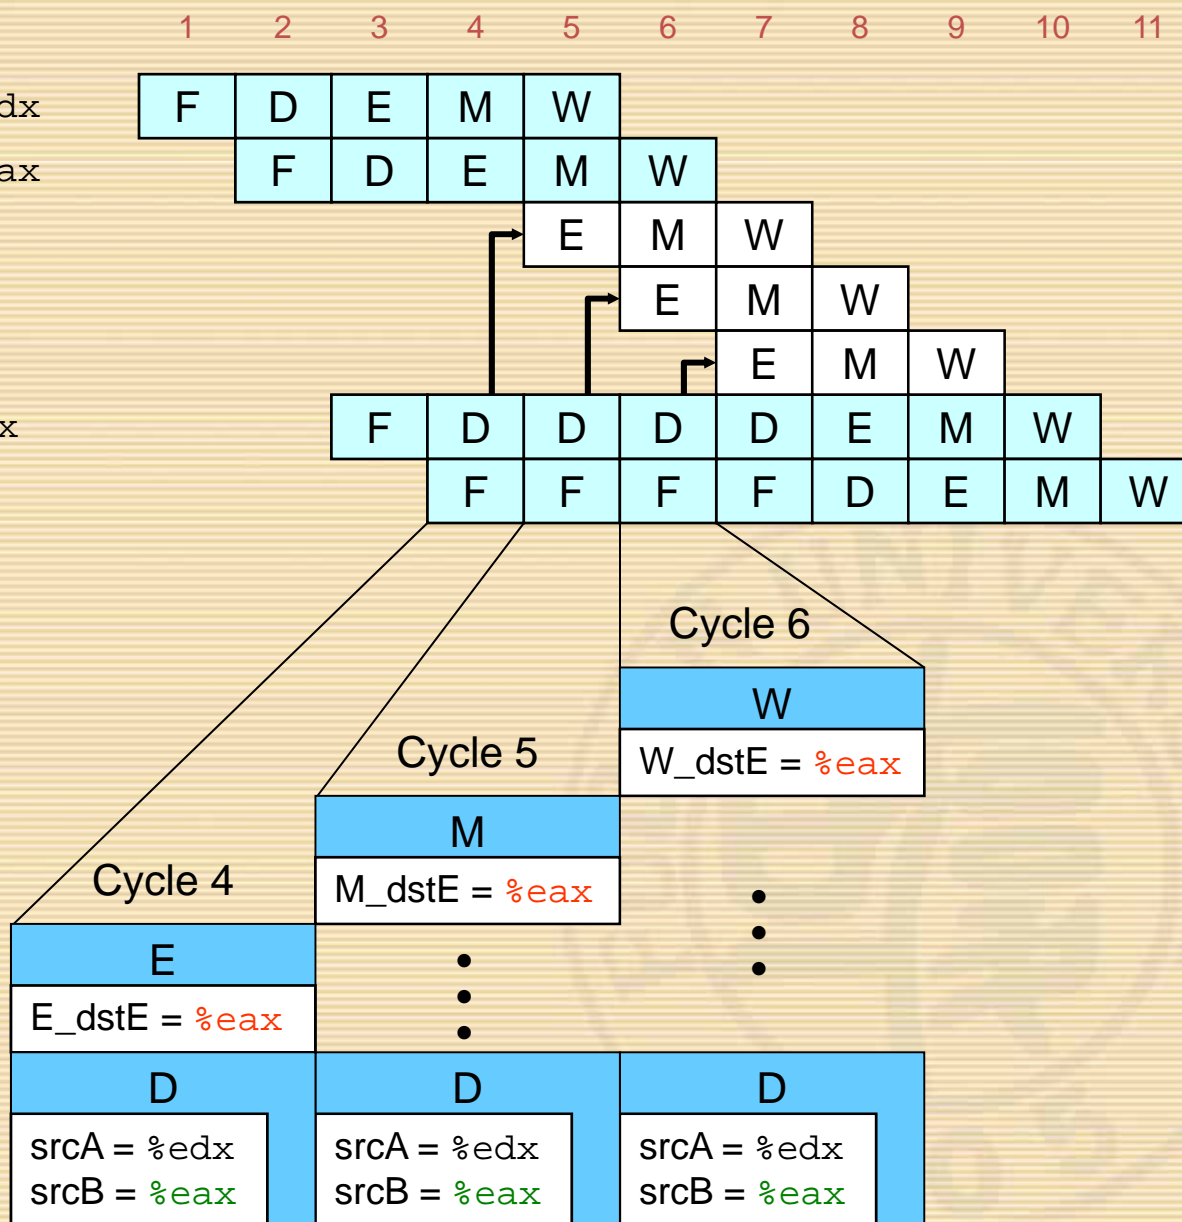




Stalling X 3

```
# demo-h0.ys
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
        bubble
        bubble
        bubble
0x00c: addl %edx,%eax
0x00e: halt
```

Figure 4.48 P329





What Happens When Stalling?

```
# demo-h0.ys
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: addl %edx,%eax
0x00e: halt
```

Figure 4.48 P329

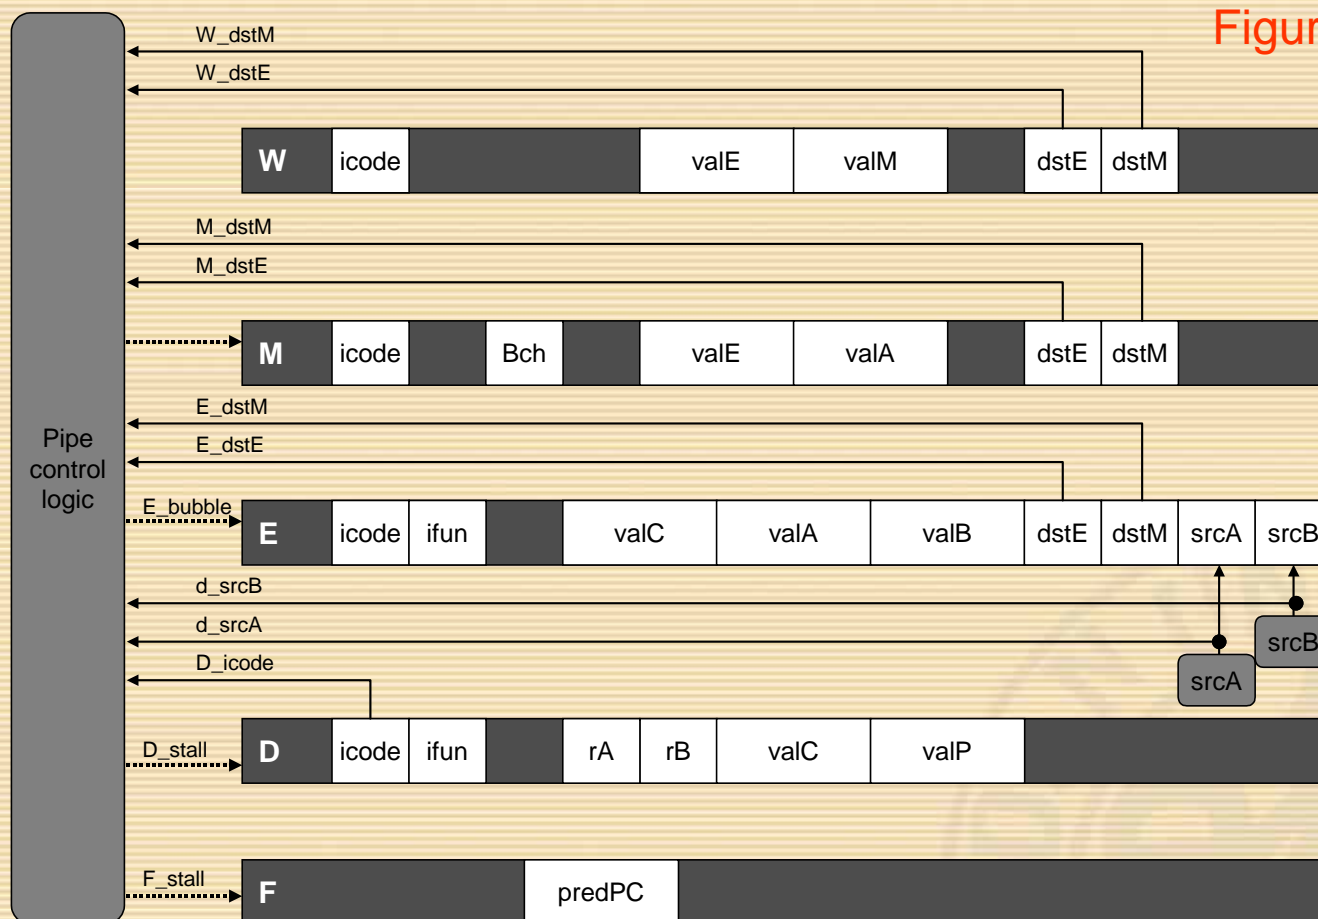
Cycle 8	
Write Back	<i>bubble</i>
Memory	<i>bubble</i>
Execute	0x00c: addl %edx,%eax
Decode	0x00e: halt
Fetch	

- Stalling instruction held back in decode stage
- Following instruction stays in fetch stage
- Bubbles injected into execute stage
 - Like dynamically generated nop's
 - Move through later stages



Implementing Stalling

Figure 4.68 P351

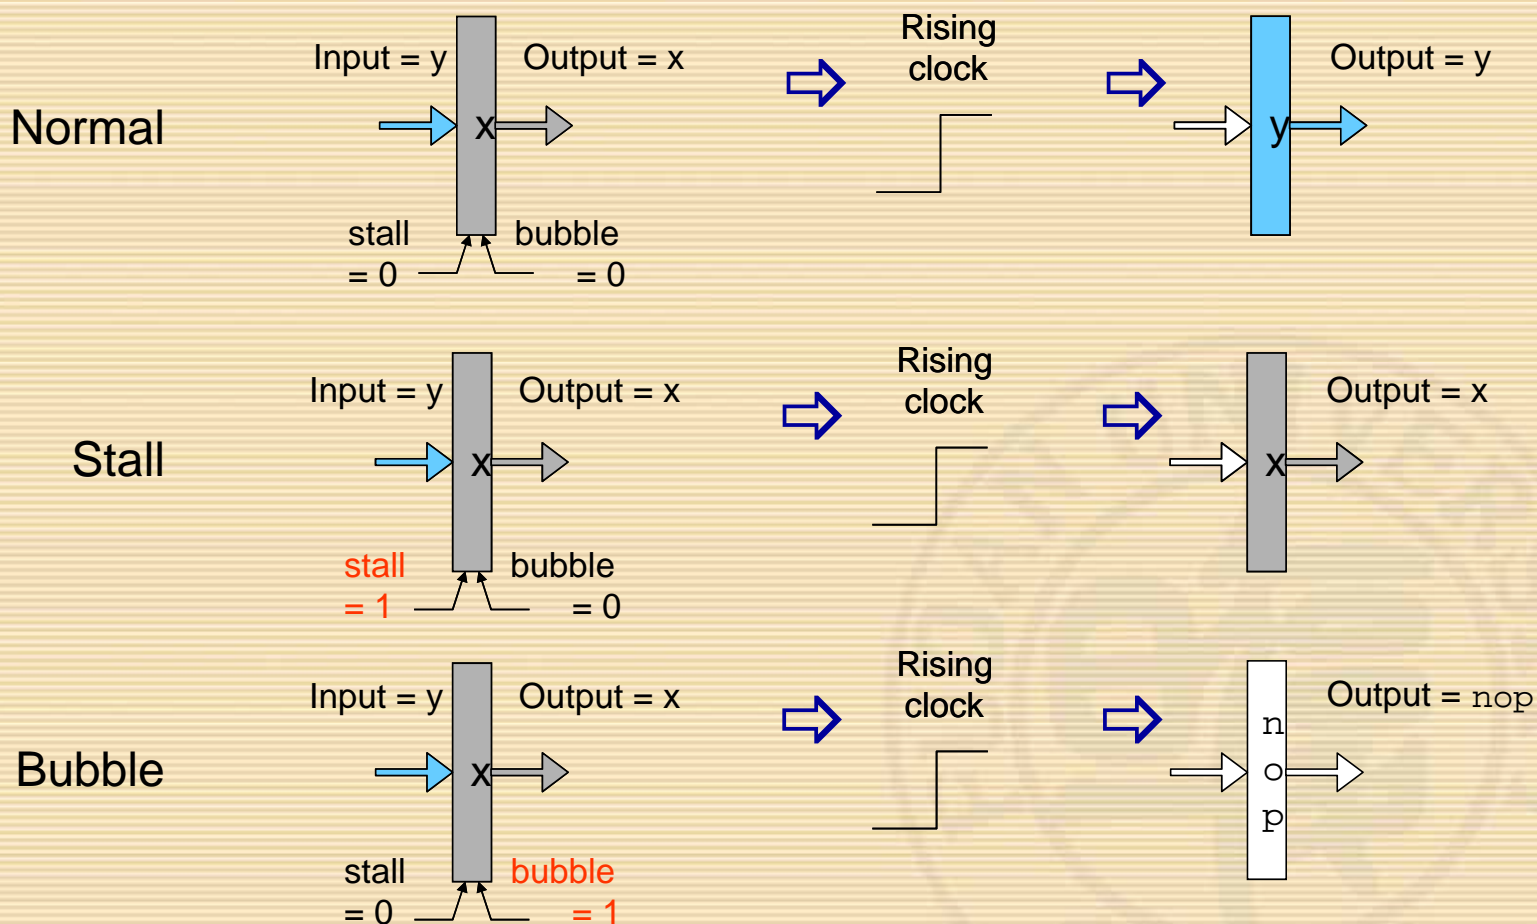


- Pipeline Control
 - Combinational logic detects stall condition
 - Sets mode signals for how pipeline registers should update



Pipeline Register Modes

Figure 4.65 P348





Data Forwarding

- Naive Pipeline
 - Register isn't written until completion of **write-back** stage
 - Source operands read from register file in **decode** stage
 - Needs to be in register file at start of stage
 - Performance is not good
- Observation
 - Value generated in execute or memory stage
- Trick
 - Pass value directly from generating instruction to decode stage
 - Needs to be available at end of decode stage

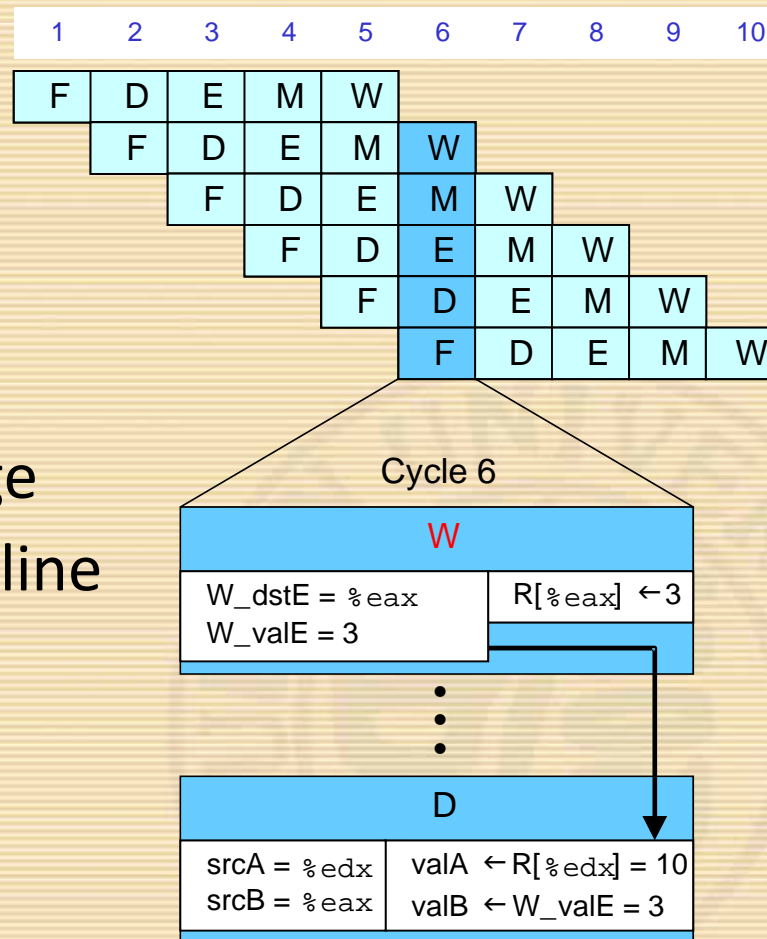


Data Forwarding Example

Figure 4.49 P331

```
# demo-h2.js
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: nop
0x00e: addl %edx,%eax
0x010: halt
```

- `irmovl` in write-back stage
- Destination value in W pipeline register
- Forward as `valB` for decode stage

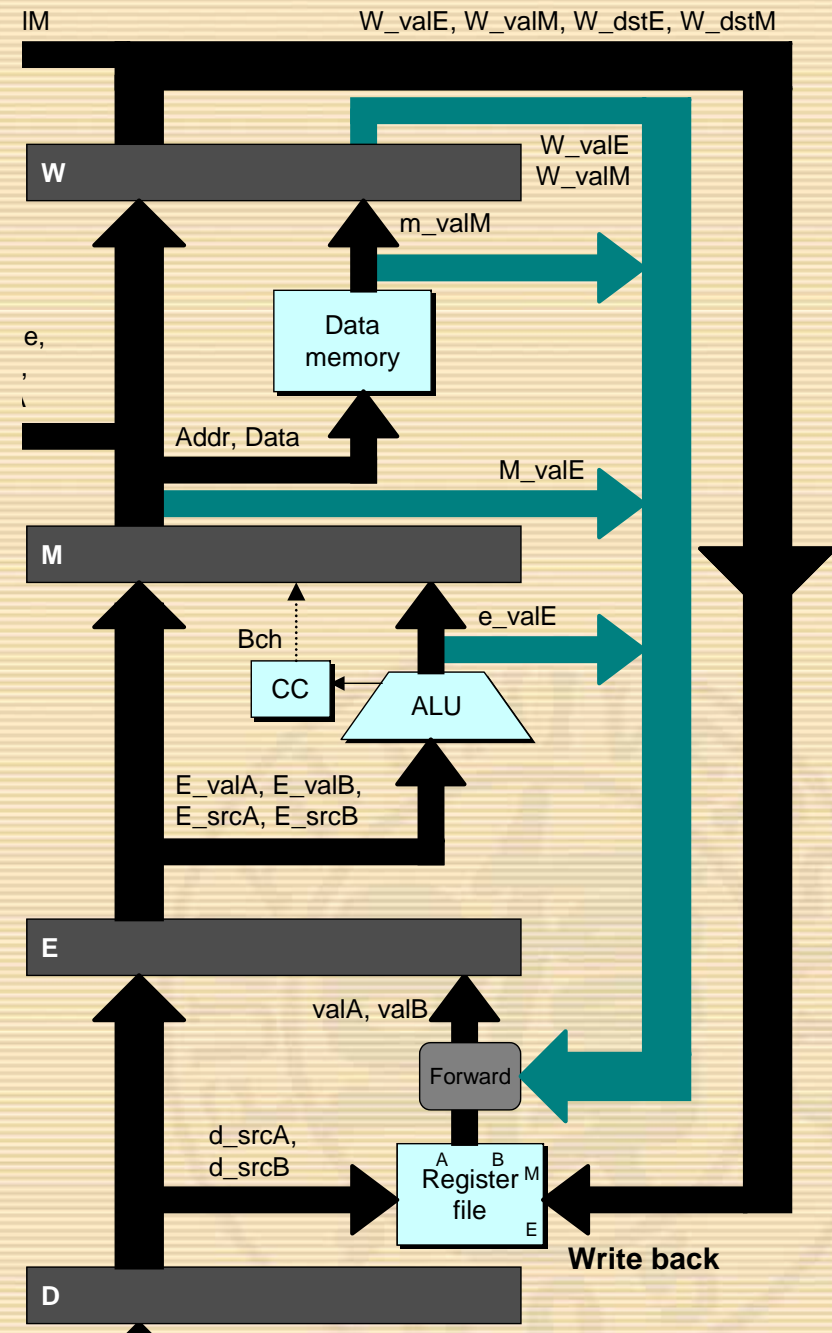




Bypass Paths

P333

- Decode Stage
 - Forwarding logic selects valA and valB
 - Normally from register file
 - Forwarding: get valA or valB from later pipeline stage
- Forwarding Sources
 - Execute: valE
 - Memory: valE, valM
 - Write back: valE, valM



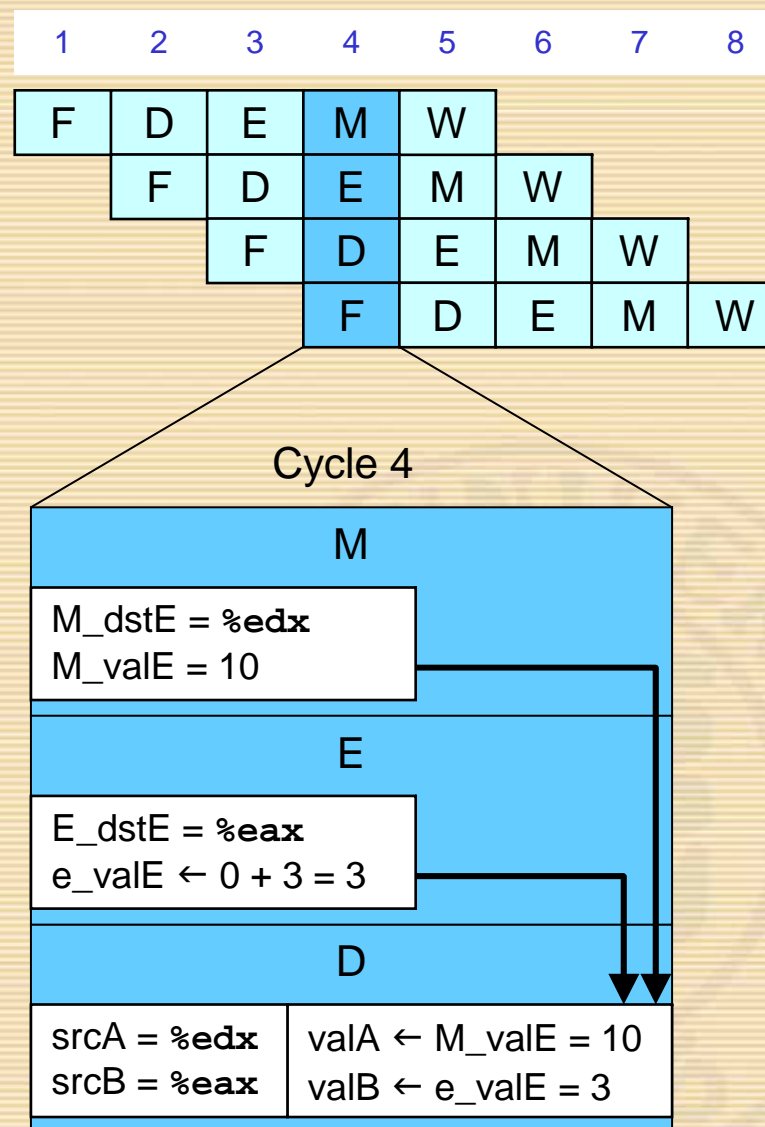


Data Forwarding Example #2

```
# demo-h0.js
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: addl %edx,%eax
0x00e: halt
```

Figure 4.51 P332

- Register `%edx`
 - Generated by ALU during previous cycle
 - Forward from memory stage as `valA`
- Register `%eax`
 - Value just generated by ALU
 - Forward from execute stage as `valB`

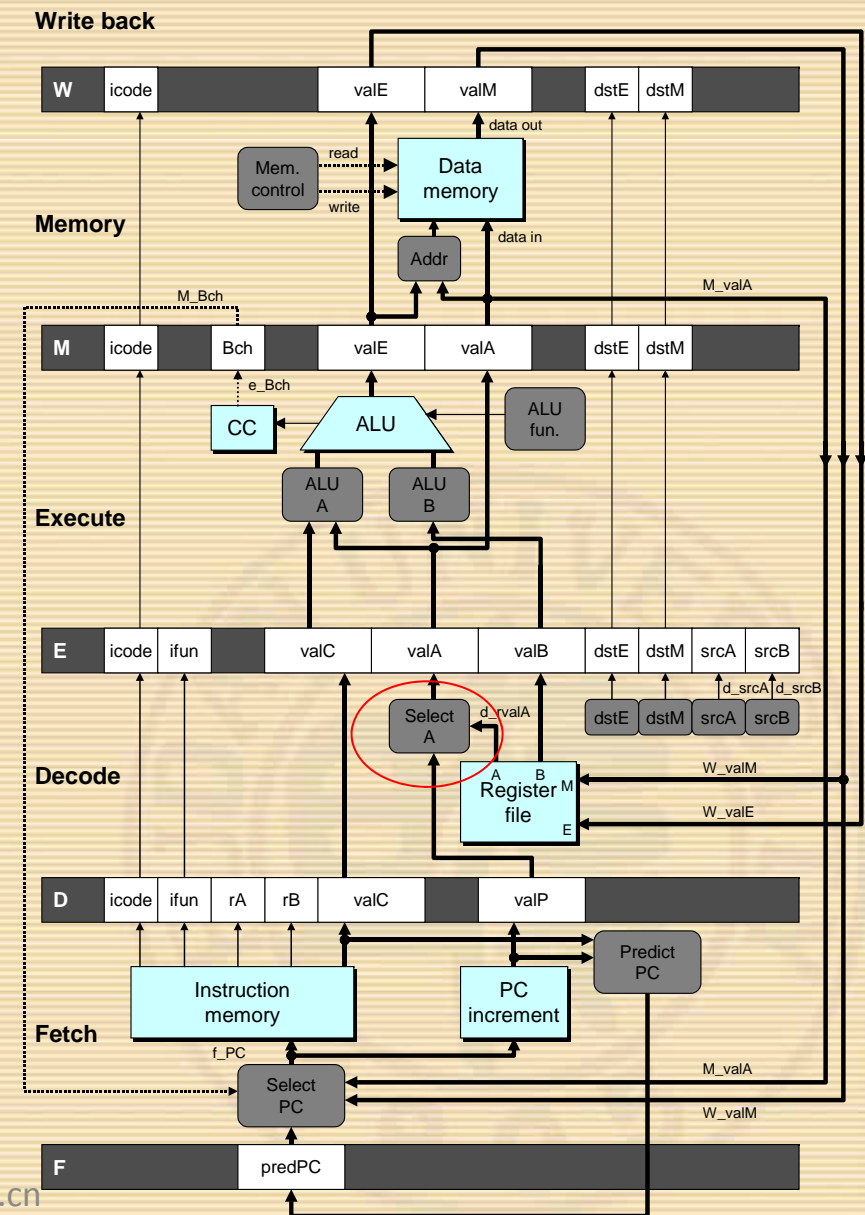




PIPE- Hardware

Figure 4.41 P320

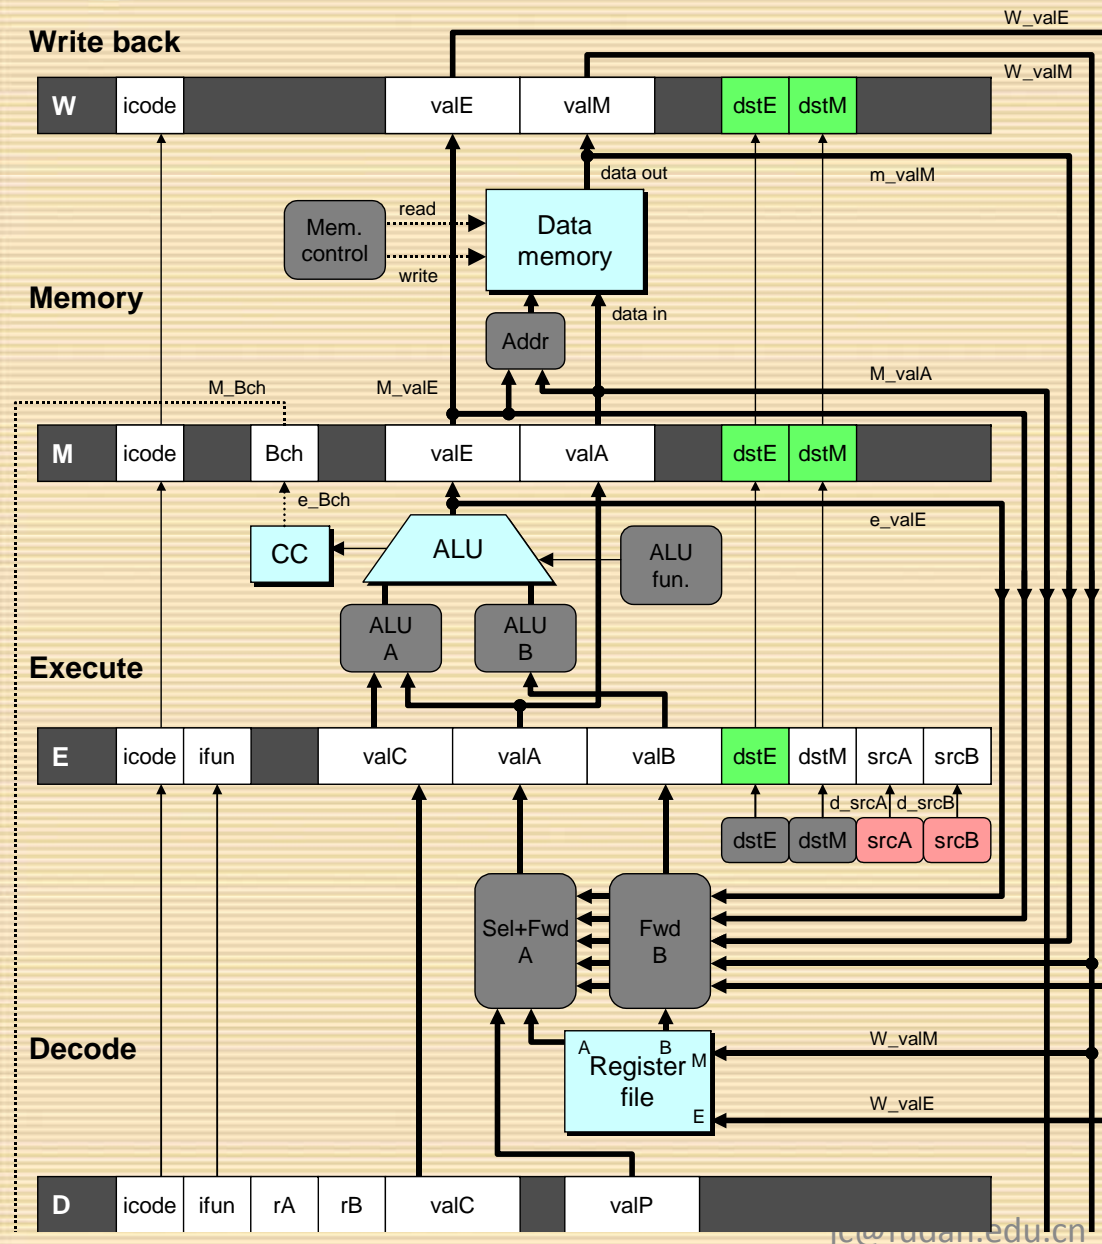
- Pipeline registers hold intermediate values from instruction execution
- Forward (Upward) Paths
 - Values passed from one stage to next
 - Cannot jump past stages
 - e.g., valC passes through decode





Implementing Forwarding

Figure 4.53 P334

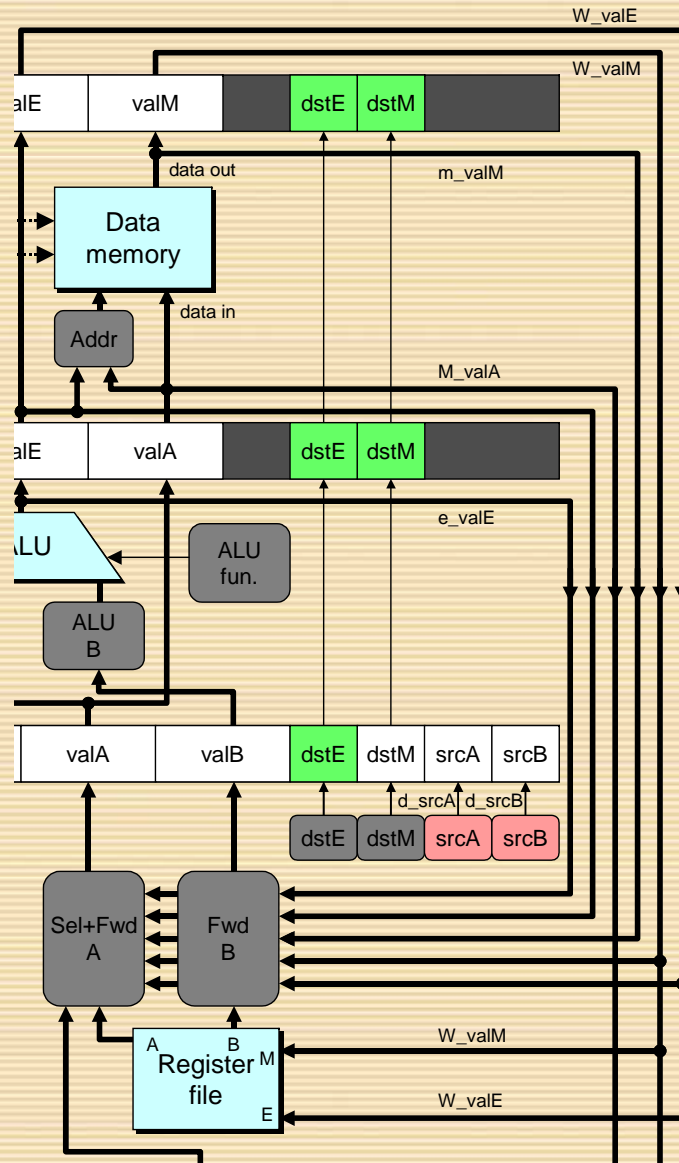


- Add additional feedback paths from E, M, and W pipeline registers into decode stage
- Create logic blocks to select from multiple sources for valA and valB in decode stage



Implementing Forwarding

P340



```
## What should be the A value?
int new_E_valA = [
    # Use incremented PC
    D_icode in { ICALL, IJXX } : D_valP;
    # Forward valE from execute
    d_srcA == E_dstE : e_valE;
    # Forward valM from memory
    d_srcA == M_dstM : m_valM;
    # Forward valE from memory
    d_srcA == M_dstE : M_valE;
    # Forward valM from write back
    d_srcA == W_dstM : W_valM;
    # Forward valE from write back
    d_srcA == W_dstE : W_valE;
    # Use value read from register file
    1 : d_rvalA;
];
```

Figure 4.53 P334



Limitation of Forwarding

demo-luh.y

0x000: irmovl \$128,%edx

0x006: irmovl \$3,%ecx

0x00c: rmmovl %ecx, 0(%edx)


0x012: irmovl \$10,%ebx

0x018: mrmovl 0(%edx),%eax # Load %eax

0x01e: addl %ebx,%eax # Use %eax

0x020: halt

1 2 3 4 5 6 7 8 9 10 11



- Load-use dependency
 - Value needed by end of **decode** stage in cycle 7
 - Value read from memory in **memory** stage of cycle 8


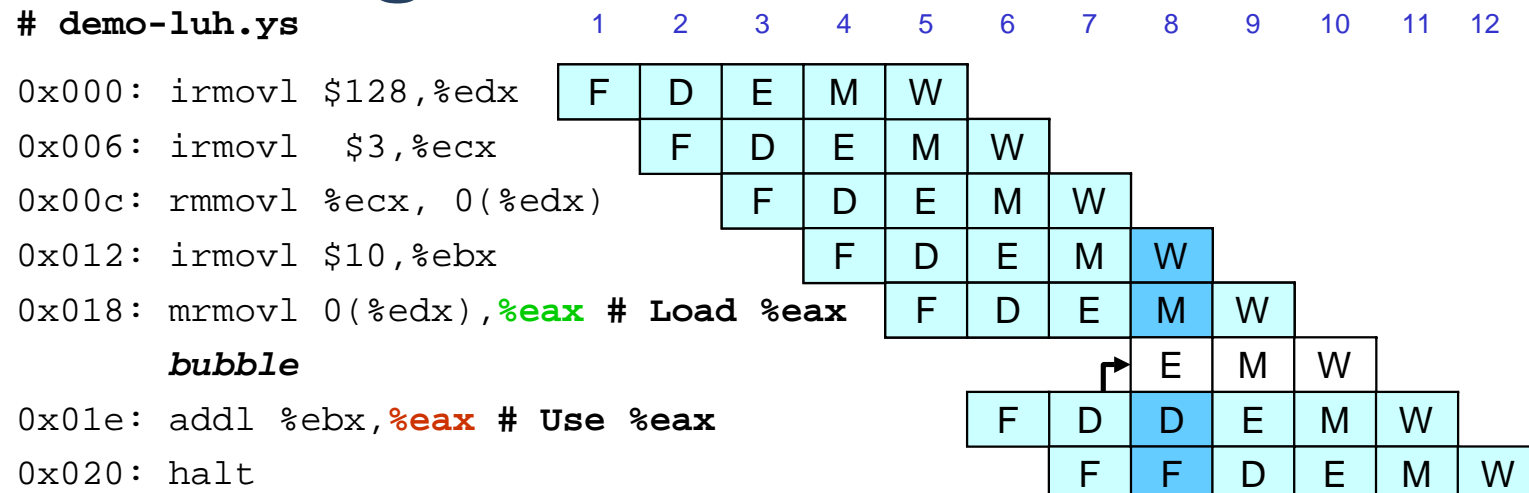


Figure 4.54 P335

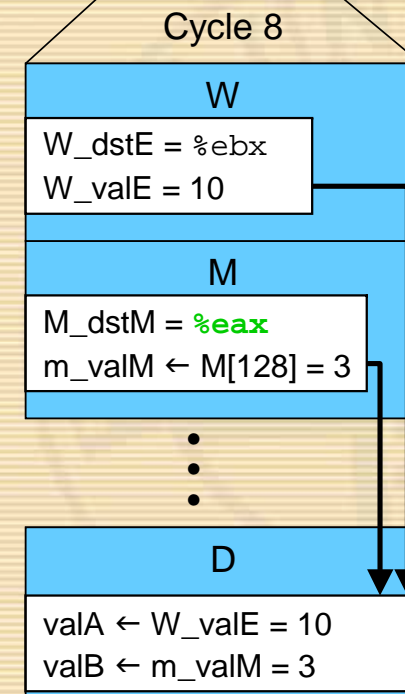


Avoiding Load/Use Hazard



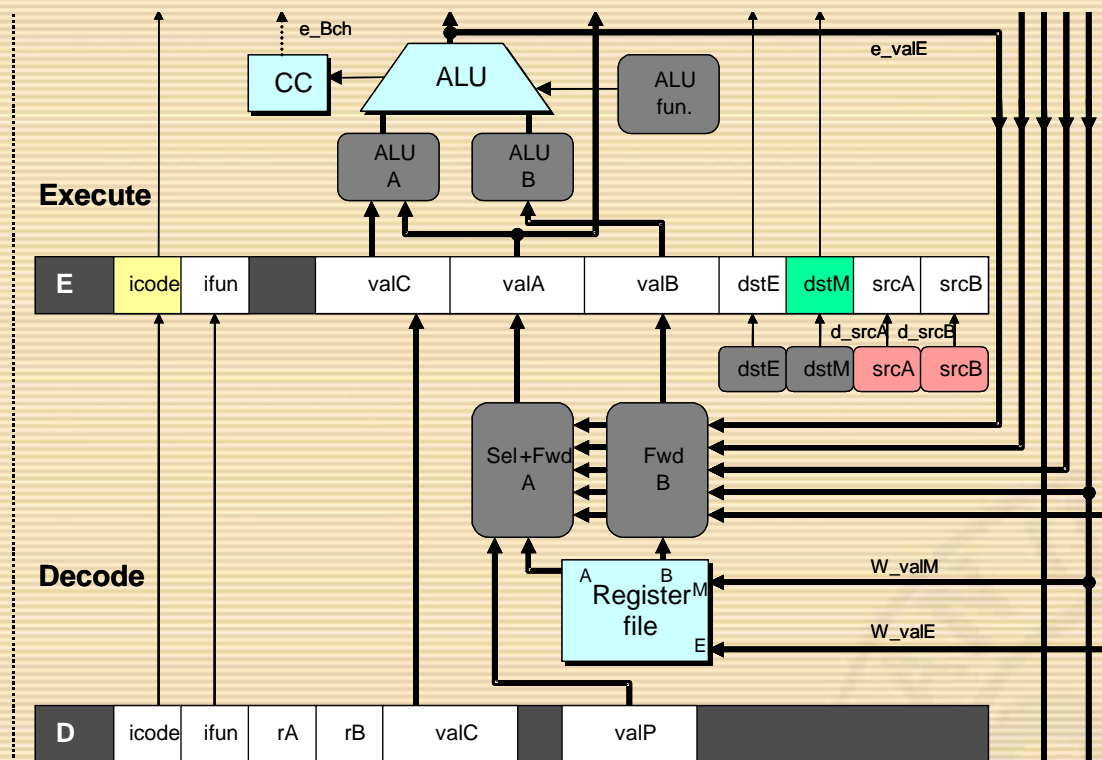
- Stall using instruction for one cycle
- Can then pick up loaded value by forwarding from memory stage

Figure 4.55 P336





Detecting Load/Use Hazard



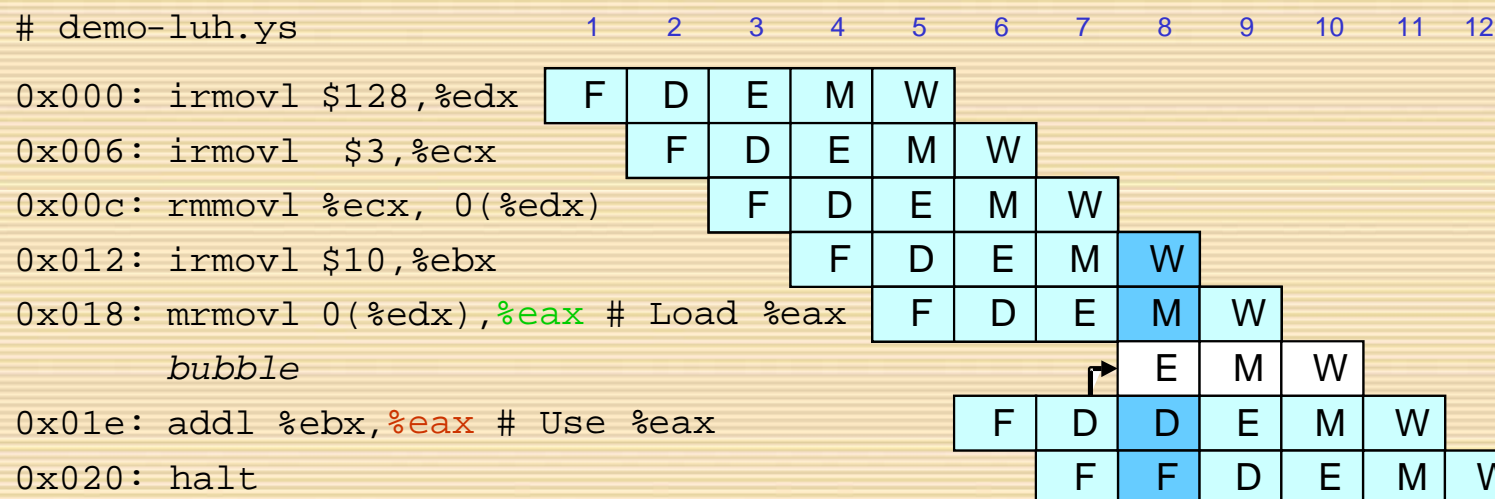
Condition	Trigger
Load/Use Hazard	E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB }

Figure 4.64 P347



Control for Load/Use Hazard

P345



- Stall instructions in fetch and decode stages
- Inject bubble into

Condition	F	D	E	M	W
Load/Use Hazard	stall	stall	bubble	normal	normal

Figure 4.66 P348



Branch Misprediction Example

demo-j.js

```
0x000:    xorl %eax,%eax
0x002:    jne  t                # Not taken
0x007:    irmovl $1, %eax        # Fall through
0x00d:    nop
0x00e:    nop
0x00f:    nop
0x010:    halt
0x011:  t:  irmovl $3, %edx    # Target (Should not execute)
0x017:    irmovl $4, %ecx        # Should not execute
0x01d:    irmovl $5, %edx        # Should not execute
```

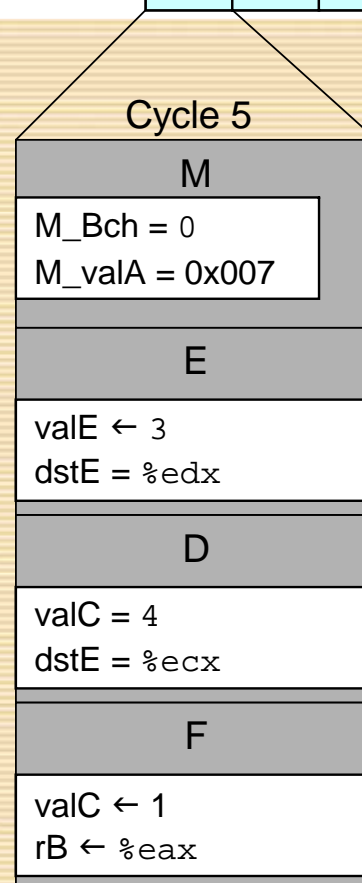
– Should only execute first 7 instructions



Branch Misprediction Trace

# demo-j	1	2	3	4	5	6	7	8	9
0x000: xorl %eax,%eax	F	D	E	M	W				
0x002: jne t # Not taken		F	D	E	M	W			
0x011: t: irmovl \$3, %edx # Target			F	D	E	M	W		
0x017: irmovl \$4, %ecx # Target+1				F	D	E	M	W	
0x007: irmovl \$1, %eax # Fall Through					F	D	E	M	W

- Incorrectly execute two instructions at branch target





Return Example

demo-ret.ys

```
0x000:    irmovl Stack,%esp    # Initialize stack pointer
0x006:    nop                  # Avoid hazard on %esp
0x007:    nop
0x008:    nop
0x009:    call p               # Procedure call
0x00e:    irmovl $5,%esi       # Return point
0x014:    halt
0x020:    .pos 0x20
0x020:    p: nop                # procedure
0x021:    nop
0x022:    nop
0x023:    ret
0x024:    irmovl $1,%eax        # Should not be executed
0x02a:    irmovl $2,%ecx        # Should not be executed
0x030:    irmovl $3,%edx        # Should not be executed
0x036:    irmovl $4,%ebx        # Should not be executed
0x100:    .pos 0x100
0x100:    Stack:                # Stack: Stack pointer
```


– Require lots of nops to avoid data hazards



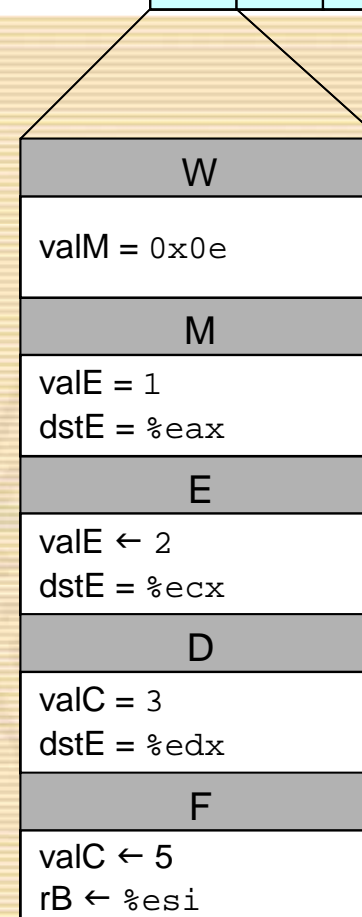
Incorrect Return Example

demo-ret

```
0x023:    ret
0x024:    irmovl $1,%eax # Oops!
0x02a:    irmovl $2,%ecx # Oops!
0x030:    irmovl $3,%edx # Oops!
0x00e:    irmovl $5,%esi # Return
```

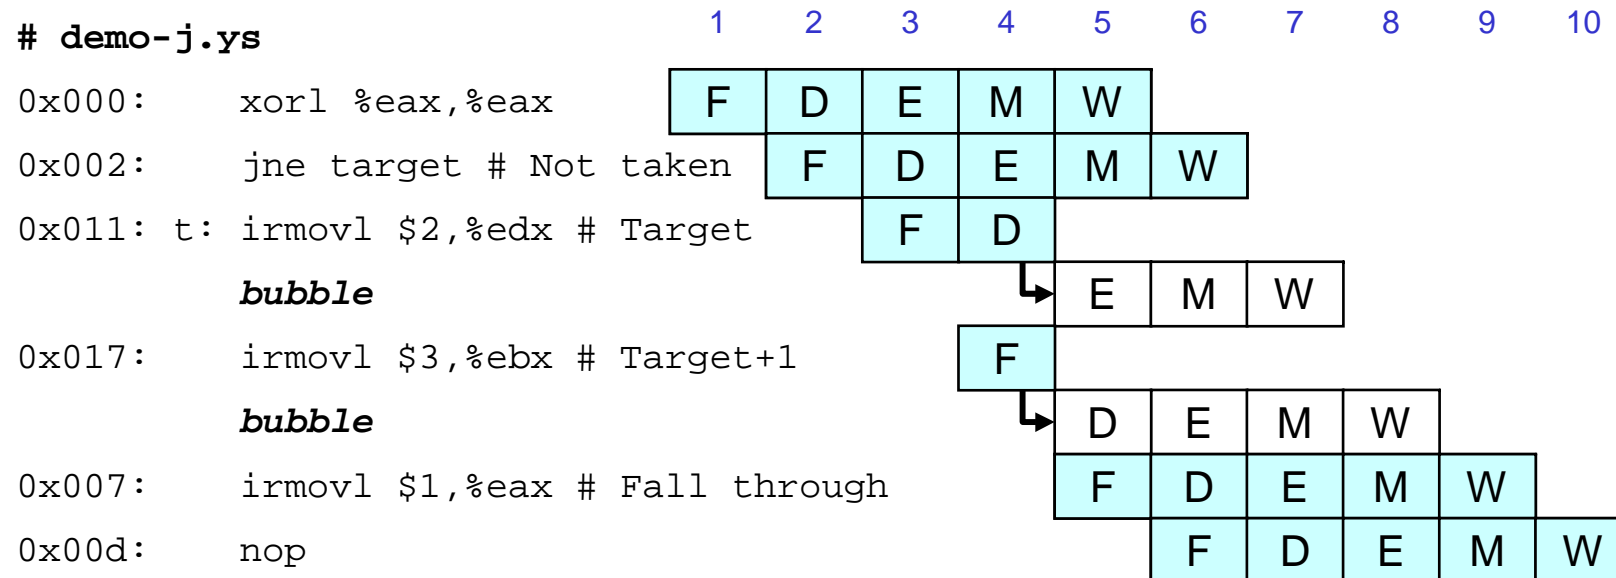


- Incorrectly execute 3 instructions following `ret`





Handling Misprediction



- Predict branch as taken
 - Fetch 2 instructions at target
- Cancel when mispredicted
 - Detect branch not-taken in execute stage
 - On following cycle, replace instructions in execute and decode by bubbles
 - No side effects have occurred yet

Figure 4.63 P346



Detecting Mispredicted Branch

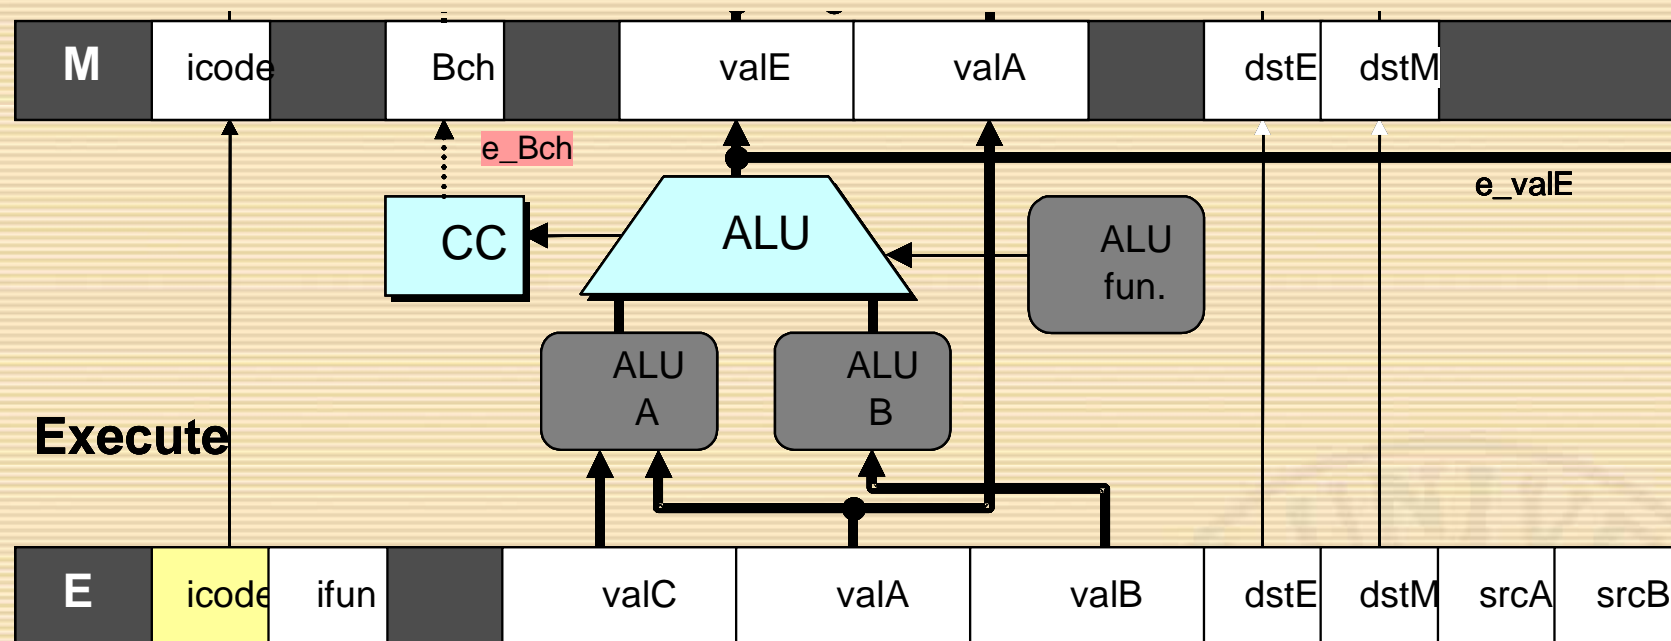


Figure 4.64 P347

Condition	Trigger
Mispredicted Branch	$E_icode = IJXX \ \& \ !e_Bch$



Control for Misprediction

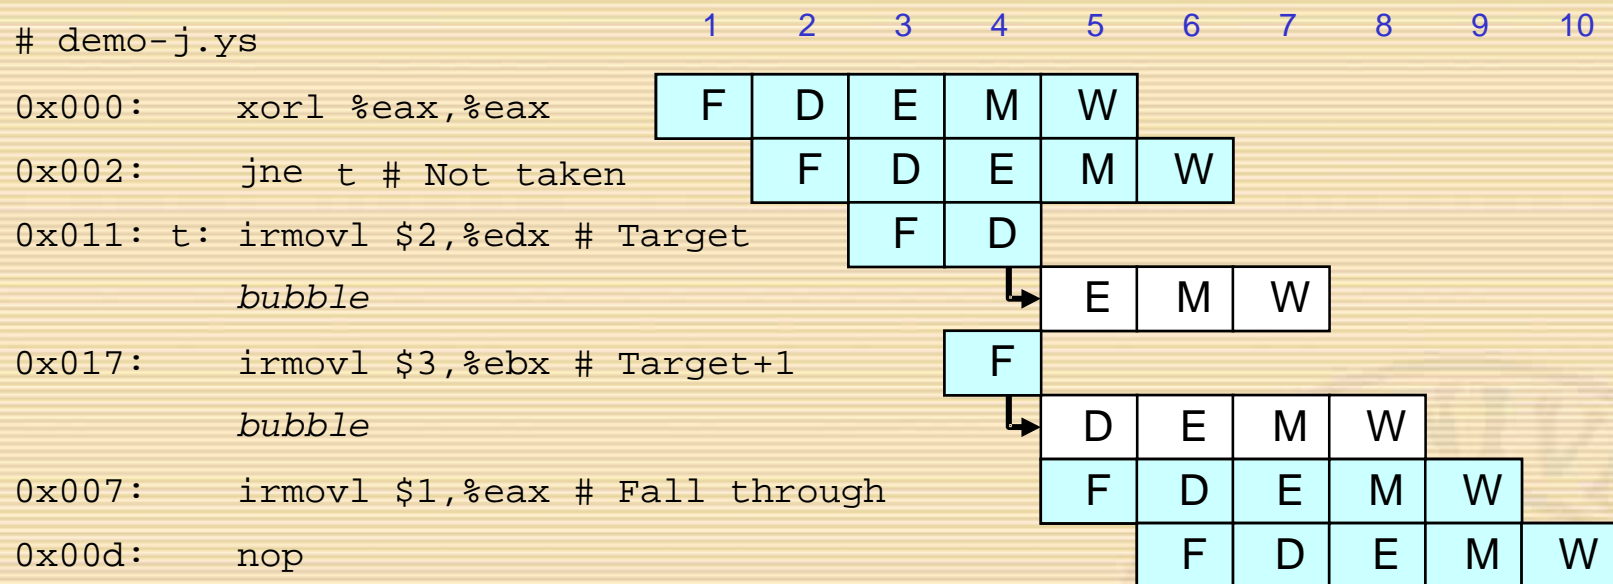


Figure 4.63 P346

Condition	F	D	E	M	W
Mispredicted Branch	normal	bubble	bubble	normal	normal

Figure 4.66 P348



Return Example

demo-retb.ys

```
0x000:    irmovl Stack,%esp    # Initialize stack pointer
0x006:    call p                # Procedure call
0x00b:    irmovl $5,%esi        # Return point
0x011:    halt
0x020:    .pos 0x20
0x020: p:  irmovl $-1,%edi      # procedure
0x026:    ret
0x027:    irmovl $1,%eax         # Should not be executed
0x02d:    irmovl $2,%ecx         # Should not be executed
0x033:    irmovl $3,%edx         # Should not be executed
0x039:    irmovl $4,%ebx         # Should not be executed
0x100:    .pos 0x100
0x100:    Stack:                  # Stack: Stack pointer
```

– Previously executed three additional instructions



Correct Return Example

```
# demo-retb
0x026:    ret
          bubble
          bubble
          bubble
0x00b:    irmovl $5,%esi # Return
```


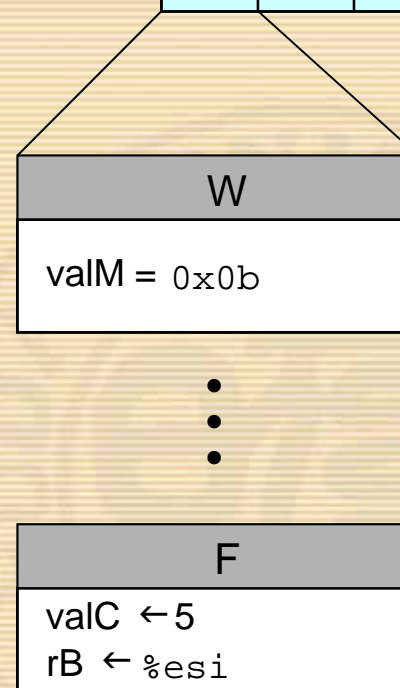



Figure 4.61 P344

- As `ret` passes through pipeline, stall at fetch stage
 - While in decode, execute, and memory stage
 - fetch the same instruction after `ret` 3 times.
- Inject bubble into decode stage
- Release stall when reach write-back stage





jc@fudan.edu.cn



Control for Return

demo-retb


0x026: ret

bubble

bubble

bubble

0x00b: irmovl \$5,%esi # Return



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal

Figure 4.66 P348



Special Control Cases

- Detection Figure 4.64 P347

Condition	Trigger
Processing ret	IRET in { D_icode, E_icode, M_icode }
Load/Use Hazard	E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB }
Mispredicted Branch	E_icode = IJXX & !e_Bch

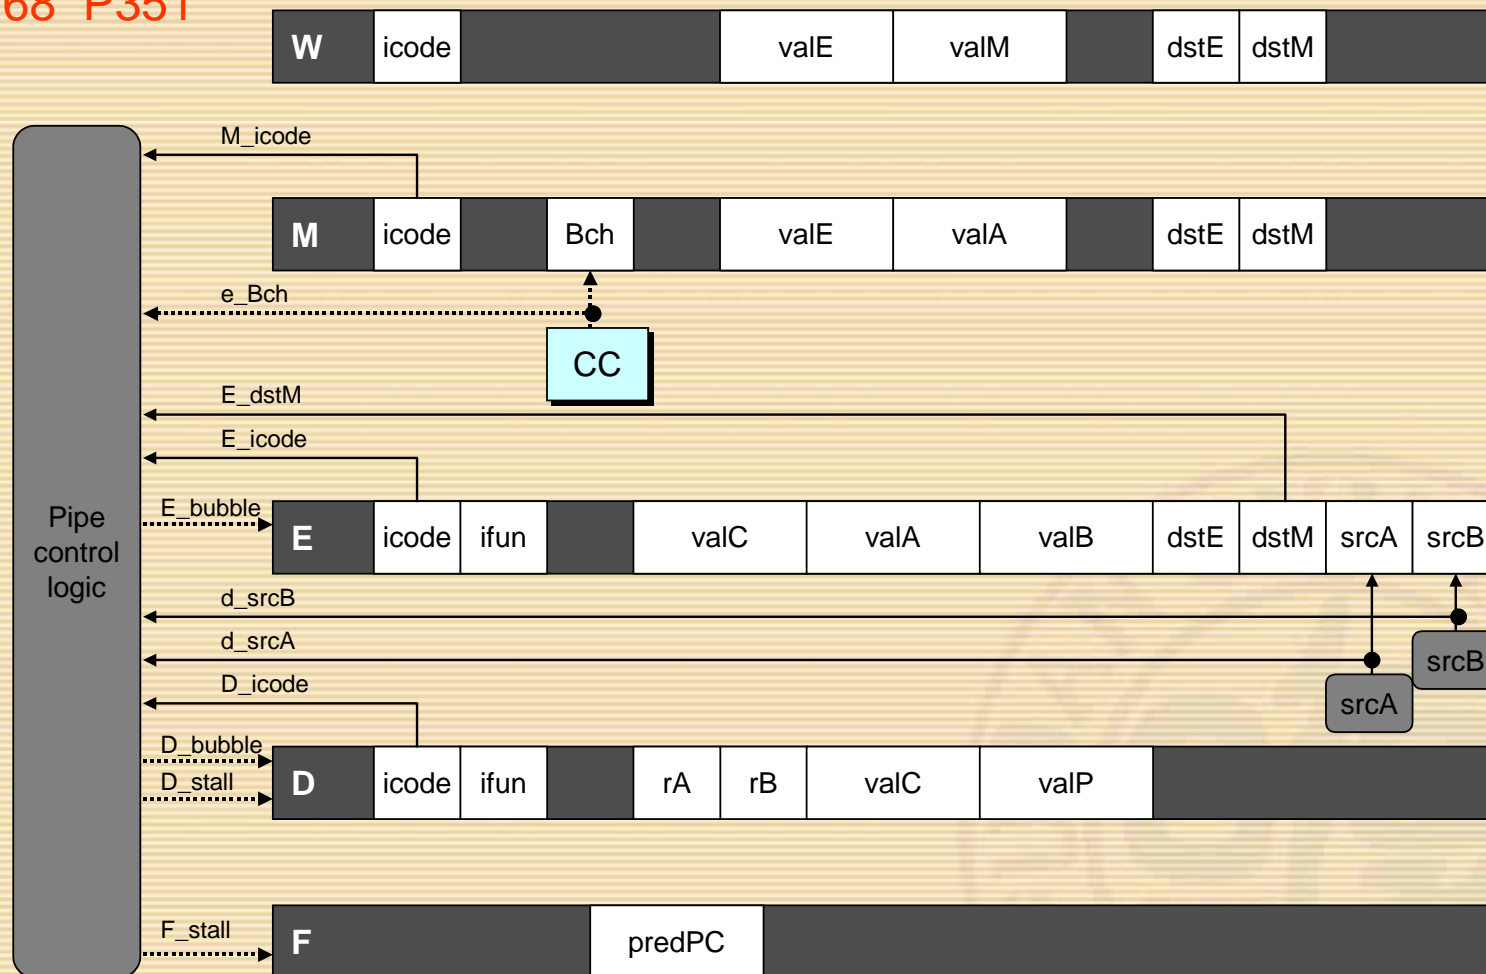
Figure 4.66 P348

- | Condition | F | D | E | M | W |
|---------------------|--------|--------|--------|--------|--------|
| Processing ret | stall | bubble | normal | normal | normal |
| Load/Use Hazard | stall | stall | bubble | normal | normal |
| Mispredicted Branch | normal | bubble | bubble | normal | normal |



Implementing Pipeline Control

Figure 4.68 P351



- Combinational logic generates pipeline control signals
- Action occurs at start of following cycle



Initial Version of Pipeline Control

```
bool F_stall =  
    # Conditions for a load/use hazard  
    E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB } ||  
    # Stalling at fetch while ret passes through pipeline  
    IRET in { D_icode, E_icode, M_icode };  
  
bool D_stall =  
    # Conditions for a load/use hazard  
    E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB };  
  
bool D_bubble =  
    # Mispredicted branch  
    (E_icode == IJXX && !e_Bch) ||  
    # Stalling at fetch while ret passes through pipeline  
    IRET in { D_icode, E_icode, M_icode };  
  
bool E_bubble =  
    # Mispredicted branch  
    (E_icode == IJXX && !e_Bch) ||  
    # Load/use hazard  
    E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB};
```



Control Combinations

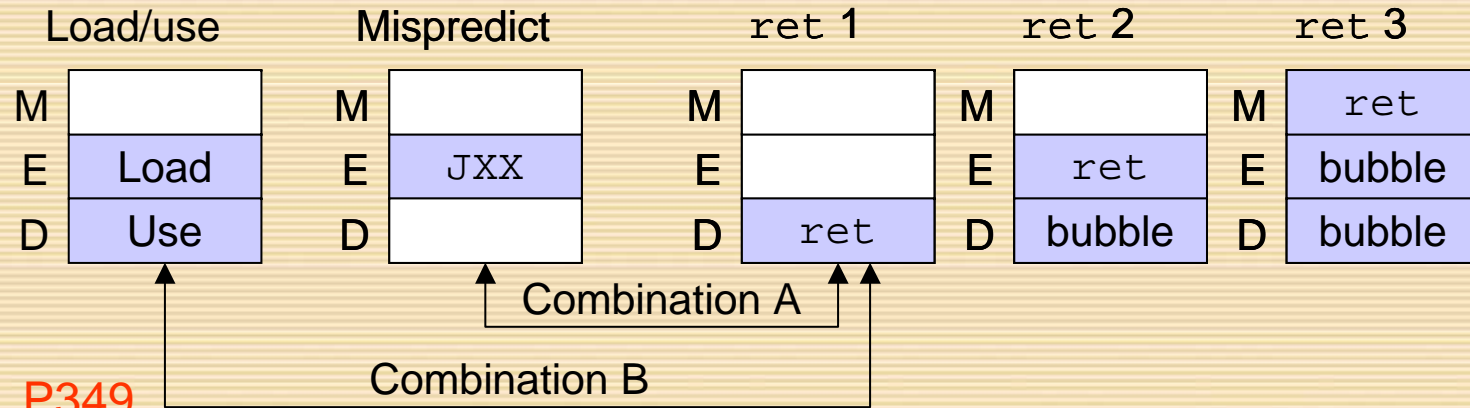
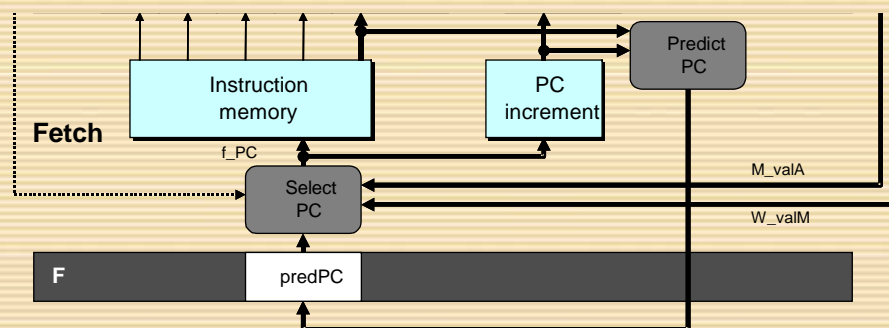
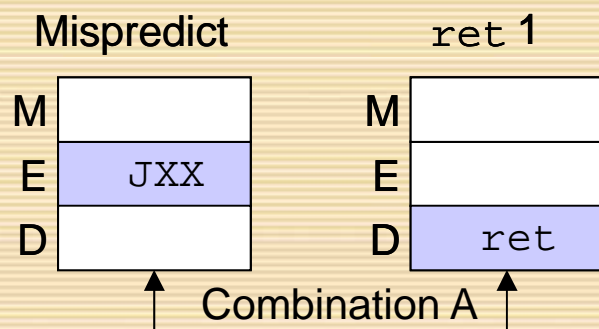


Figure 4.67 P349

- Special cases that can arise on same clock cycle
- Combination A
 - Not-taken branch
 - `ret` instruction at branch target
- Combination B
 - Instruction that reads from memory to `%esp`
 - Followed by `ret` instruction



Control Combination A



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Mispredicted Branch	normal	bubble	bubble	normal	normal
<i>Combination</i>	<i>stall</i>	<i>bubble</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Should be handled as mispredicted branch
- Stalls F pipeline register
- But PC selection logic will be using M_valA anyhow



Control Combination B



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
<i>Combination</i>	<i>stall</i>	<i>bubble + stall</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Would attempt to bubble *and* stall pipeline register D
- Signaled by processor as pipeline error



Handling Control Combination B



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
<i>Combination</i>	<i>stall</i>	<i>stall</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Load/use hazard should get priority
- `ret` instruction should be held in decode stage for additional cycle



Corrected Pipeline Control Logic

```
bool D_bubble =  
    # Mispredicted branch  
    (E_icode == IJXX && !e_Bch) ||  
    # Stalling at fetch while ret passes through pipeline  
    IRET in { D_icode, E_icode, M_icode }  
    # but not condition for a load/use hazard  
    && !(E_icode in { IMRMOVL, IPOPL }  
        && E_dstM in { d_srcA, d_srcB }));
```

Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
<i>Combination</i>	<i>stall</i>	<i>stall</i>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

- Load/use hazard should get priority
- ret instruction should be held in decode stage for additional cycle



Pipeline Summary

- Data Hazards
 - Most handled by forwarding
 - No performance penalty
 - Load/use hazard requires one cycle stall
- Control Hazards
 - Cancel instructions when detect mispredicted branch
 - Two clock cycles wasted
 - Stall fetch stage while `ret` passes through pipeline
 - Three clock cycles wasted
- Control Combinations
 - Must analyze carefully
 - First version had subtle bug
 - Only arises with unusual instruction combination



Performance Metrics

4.5.10

- Clock rate
 - Measured in Megahertz or Gigahertz
 - Function of stage partitioning and circuit design
 - Keep amount of work per stage small
- Rate at which instructions executed
 - CPI: cycles per instruction
 - On average, how many clock cycles does each instruction require?
 - Function of pipeline design and benchmark programs
 - E.g., how frequently are branches mispredicted?



CPI for PIPE

- CPI \approx 1.0
 - Fetch instruction each clock cycle
 - Effectively process new instruction almost every cycle
 - Although each individual instruction has latency of 5 cycles
- CPI $>$ 1.0
 - Sometimes must stall or cancel branches
- Computing CPI
 - C clock cycles
 - I instructions executed to completion
 - B bubbles injected ($C = I + B$)
$$\text{CPI} = C/I = (I+B)/I = 1.0 + B/I$$
 - Factor B/I represents average penalty due to bubbles



CPI for PIPE (Cont.)

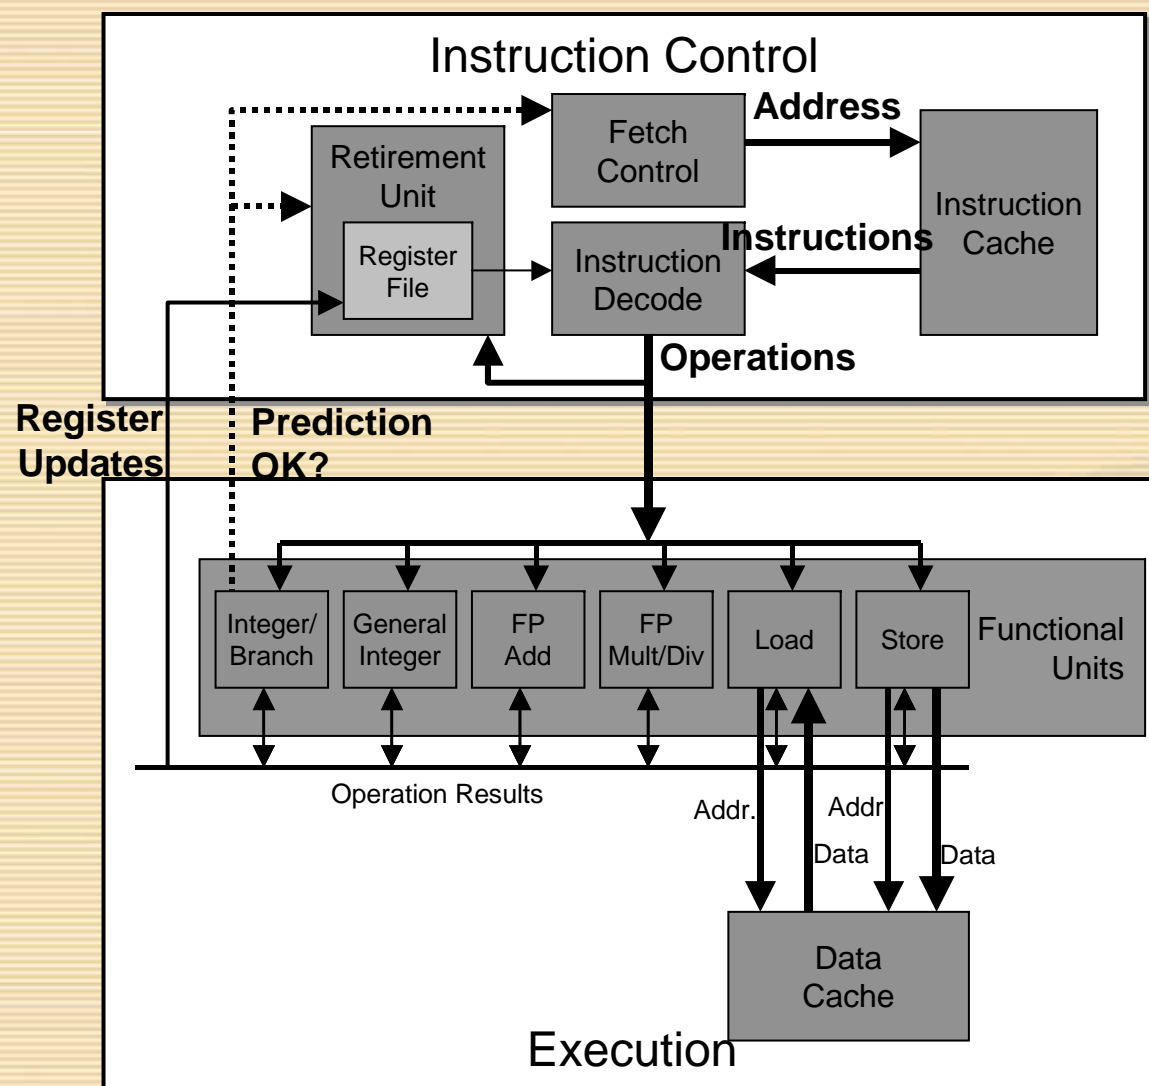
$$B/I = LP + MP + RP$$

Typical Values

- LP: Penalty due to load/use hazard stalling
 - Fraction of instructions that are loads 0.25
 - Fraction of load instructions requiring stall 0.20
 - Number of bubbles injected each time 1
$$\Rightarrow LP = 0.25 * 0.20 * 1 = 0.05$$
- MP: Penalty due to mispredicted branches
 - Fraction of instructions that are cond. jumps 0.20
 - Fraction of cond. jumps mispredicted 0.40
 - Number of bubbles injected each time 2
$$\Rightarrow MP = 0.20 * 0.40 * 2 = 0.16$$
- RP: Penalty due to `ret` instructions
 - Fraction of instructions that are returns 0.02
 - Number of bubbles injected each time 3
$$\Rightarrow RP = 0.02 * 3 = 0.06$$
- Net effect of penalties $0.05 + 0.16 + 0.06 = 0.27$
$$\Rightarrow CPI = 1.27 \quad (\text{Not bad!})$$



Modern CPU Design





Processor Summary

- Design Technique
 - Create uniform framework for all instructions
 - Want to share hardware among instructions
 - Connect standard logic blocks with bits of control logic
- Operation
 - State held in memories and clocked registers
 - Computation done by combinational logic
 - Clocking of registers/memories sufficient to control overall behavior
- Enhancing Performance
 - Pipelining increases throughput and improves resource utilization
 - Must make sure maintains ISA behavior



Code Optimization I: Machine Independent Optimizations



Basic Requirements of Codes

- Robustness
- Efficiency
- Extensibility



Why

- Constant factors
 - Easily see 10:1 performance range depending on how code is written
 - Must optimize at multiple levels:
 - algorithm, data representations, procedures, and loops
- Must understand system to optimize performance
 - How programs are compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality



Optimizing Compilers

- Provide efficient mapping of program to machine
 - register allocation
 - code selection and ordering
 - eliminating minor inefficiencies
- Don't (usually) improve asymptotic efficiency
 - up to programmer to select best overall algorithm
 - big-O savings are (often) more important than constant factors
 - but constant factors also matter
- Have difficulty overcoming “optimization blockers”
 - potential memory aliasing
 - potential procedure side-effects



Limitations of Optimizing Compilers

- Operate Under Fundamental Constraint
 - Must not cause any change in program behavior under any possible condition
 - Often prevents it from making optimizations when would only affect behavior under pathological conditions.
- Behavior that may be obvious to the programmer can be obfuscated by languages and coding styles
 - e.g., data ranges may be more limited than variable types suggest
- Most analysis is performed only within procedures
 - whole-program analysis is too expensive in most cases
- Most analysis is based only on *static* information
 - compiler has difficulty anticipating run-time inputs
- When in doubt, the compiler must be conservative



Machine-Independent Optimizations

- Optimizations that you should do regardless of processor / compiler
- Code Motion
 - Reduce frequency with which computation performed
 - If it will always produce same result
 - Especially moving code out of loop

```
for (i = 0; i < n; i++)  
  for (j = 0; j < n; j++)  
    a[n*i + j] = b[j];
```





Machine-Independent Optimizations

- Optimizations that you should do regardless of processor / compiler
- Code Motion
 - Reduce frequency with which computation performed
 - If it will always produce same result
 - Especially moving code out of loop

```
for (i = 0; i < n; i++)  
    for (j = 0; j < n; j++)  
        a[n*i + j] = b[j];
```



```
for (i = 0; i < n; i++) {  
    int ni = n*i;  
    for (j = 0; j < n; j++)  
        a[ni + j] = b[j];  
}
```



Compiler-Generated Code Motion

- Most compilers do a good job with array code + simple loop structures
- Code Generated by GCC

```
for (i = 0; i < n; i++)  
    for (j = 0; j < n; j++)  
        a[n*i + j] = b[j];
```

```
for (i = 0; i < n; i++) {  
    int ni = n*i;  
    int *p = a+ni;  
    for (j = 0; j < n; j++)  
        *p++ = b[j];  
}
```

```
imull %ebx,%eax      # i*n  
movl 8(%ebp),%edi     # a  
leal (%edi,%eax,4),%edx # p = a+i*n (scaled by 4)  
# Inner Loop  
.L40:  
movl 12(%ebp),%edi    # b  
movl (%edi,%ecx,4),%eax # b+j (scaled by 4)  
movl %eax,(%edx)      # *p = b[j]  
addl $4,%edx          # p++ (scaled by 4)  
incl %ecx             # j++  
jnl .L40              # loop if j<n
```



Reduction in Strength

- Replace costly operation with simpler one
- Shift, add instead of multiply or divide

$16 * x \rightarrow x \ll 4$

- Utility machine dependent
 - Depends on cost of multiply or divide instruction
 - On Pentium II or III, integer multiply only requires 4 CPU cycles
- Recognize sequence of products

```
for (i = 0; i < n; i++)  
  for (j = 0; j < n; j++)  
    a[n*i + j] = b[j];
```





Reduction in Strength

- Replace costly operation with simpler one
- Shift, add instead of multiply or divide

$16 * x \rightarrow x \ll 4$

- Utility machine dependent
 - Depends on cost of multiply or divide instruction
 - On Pentium II or III, integer multiply only requires 4 CPU cycles
- Recognize sequence of products

```
for (i = 0; i < n; i++)  
    for (j = 0; j < n; j++)  
        a[n*i + j] = b[j];
```



```
int ni = 0;  
for (i = 0; i < n; i++) {  
    for (j = 0; j < n; j++)  
        a[ni + j] = b[j];  
    ni += n;  
}
```




Make Use of Registers

- Reading and writing registers is much faster than reading/writing memory
- Limitation
 - Compiler not always able to determine whether variable can be held in register
 - Possibility of *Aliasing*
 - See example later



Machine-Independent Opts. (Cont.)

- Share Common Subexpressions
 - Reuse portions of expressions
 - Compilers often not very sophisticated in exploiting arithmetic properties

```
/* Sum neighbors of i,j */  
up =    val[(i-1)*n + j];  
down =  val[(i+1)*n + j];  
left =  val[i*n    + j-1];  
right = val[i*n    + j+1];  
sum = up + down + left + right;
```

3 multiplications: $i*n$, $(i-1)*n$, $(i+1)*n$


```
leal -1(%edx),%ecx    # i-1  
imull %ebx,%ecx       # (i-1)*n  
leal 1(%edx),%eax     # i+1  
imull %ebx,%eax       # (i+1)*n  
imull %ebx,%edx       # i*n
```

```
int inj = i*n + j;  
up =    val[inj - n];  
down =  val[inj + n];  
left =  val[inj - 1];  
right = val[inj + 1];  
sum = up + down + left + right;
```

1 multiplication: $i*n$



Vector ADT



- Procedures

```
vec_ptr new_vec(int len)
```

- Create vector of specified length

```
int get_vec_element(vec_ptr v, int index, int *dest)
```

- Retrieve vector element, store at *dest
- Return 0 if out of bounds, 1 if successful

```
int *get_vec_start(vec_ptr v)
```

- Return pointer to start of vector data

– Similar to array implementations in Pascal, ML, Java

- E.g., always do bounds checking



Optimization Example

```
void combine1(vec_ptr v, int *dest)
{
    int i;
    *dest = 0;
    for (i = 0; i < vec_length(v); i++) {
        int val;
        get_vec_element(v, i, &val);
        *dest += val;
    }
}
```

- Procedure
 - Compute sum of all elements of vector
 - Store result at destination location




Time Scales

- Absolute Time
 - Typically use nanoseconds
 - 10^{-9} seconds
 - Time scale of computer instructions
- Clock Cycles
 - Most computers controlled by high frequency clock signal
 - Typical Range
 - 100 MHz
 - 10^8 cycles per second
 - Clock period = 10ns
 - 3.2 GHz
 - 3.2×10^9 cycles per second
 - Clock period = 0.3125ns



Cycles Per Element

- Convenient way to express performance of program that operates on vectors or lists
- Length = n
- $T = CPE * n + \text{Overhead}$





Optimization Example

```
void combine1(vec_ptr v, int *dest)
{
    int i;
    *dest = 0;
    for (i = 0; i < vec_length(v); i++) {
        int val;
        get_vec_element(v, i, &val);
        *dest += val;
    }
}
```

- Procedure
 - Compute sum of all elements of integer vector
 - Store result at destination location
 - Vector data structure and operations defined via abstract data type
- Pentium II/III Performance: Clock Cycles / Element
 - 42.06 (Compiled -g) 31.25 (Compiled -O2)



Understanding Loop

```
void combine1-goto(vec_ptr v, int *dest)
{
    int i = 0;
    int val;
    *dest = 0;
    if (i >= vec_length(v))
        goto done;
loop:
    get_vec_element(v, i, &val);
    *dest += val;
    i++;
    if (i < vec_length(v))
        goto loop;
done:
}
```

1 iteration

- Inefficiency
 - Procedure `vec_length` called every iteration
 - Even though result always the same



Move vec_length Call Out of Loop

```
void combine2(vec_ptr v, int *dest)
{
    int i;
    int length = vec_length(v);
    *dest = 0;
    for (i = 0; i < length; i++) {
        int val;
        get_vec_element(v, i, &val);
        *dest += val;
    }
}
```

- Optimization

- Move call to `vec_length` out of inner loop

- Value does not change from one iteration to next
 - Code motion

- CPE: 20.66 (Compiled -O2)

- `vec_length` requires only constant time, but significant overhead



Optimization Blocker: Procedure Calls

- *Why couldn't the compiler move `vec_len` out of the inner loop?*
 - Procedure may have side effects
 - Alters global state each time called
 - Function may not return same value for given arguments
 - Depends on other parts of global state
 - Procedure `lower` could interact with `strlen`
- *Why doesn't compiler look at code for `vec_len`?*
 - Linker may overload with different version
 - Unless declared static
 - Inter-procedural optimization is not used extensively due to cost
- Warning:
 - Compiler treats procedure call as a black box
 - Weak optimizations in and around them



Reduction in Strength

```
void combine3(vec_ptr v, int *dest)
{
    int i;
    int length = vec_length(v);
    int *data = get_vec_start(v);
    *dest = 0;
    for (i = 0; i < length; i++) {
        *dest += data[i];
    }
}
```

- Optimization

- Avoid procedure call to retrieve each vector element
 - Get pointer to start of array before loop
 - Within loop just do pointer reference
 - Not as clean in terms of data abstraction
- CPE: 6.00 (Compiled -O2)
 - Procedure calls are expensive!
 - Bounds checking is expensive



Eliminate Unneeded Memory Refs

```
void combine4(vec_ptr v, int *dest)
{
    int i;
    int length = vec_length(v);
    int *data = get_vec_start(v);
    int sum = 0;
    for (i = 0; i < length; i++)
        sum += data[i];
    *dest = sum;
}
```

- Optimization

- Don't need to store in destination until end
- Local variable sum held in register
- Avoids 1 memory read, 1 memory write per cycle
- CPE: 2.00 (Compiled -O2)
 - Memory references are expensive!



Detecting Unneeded Memory Refs.

Combine3

```
.L18:  
    movl (%ecx,%edx,4),%eax  
    addl %eax,(%edi)  
    incl %edx  
    cmpl %esi,%edx  
    jl .L18
```

Combine4

```
.L24:  
    addl (%eax,%edx,4),%ecx  
  
    incl %edx  
    cmpl %esi,%edx  
    jl .L24
```

- Performance
 - Combine3
 - 5 instructions in 6 clock cycles
 - addl must read and write memory
 - Combine4
 - 4 instructions in 2 clock cycles



Optimization Blocker: Memory Aliasing

- Aliasing
 - Two different memory references specify single location
- Example
 - v: [3, 2, 17]
 - combine3(v, get_vec_start(v)+2)-->?
 - combine4(v, get_vec_start(v)+2)-->?

```
void combine3(vec_ptr v, int *dest)
{
    int i;
    int length = vec_length(v);
    int *data = get_vec_start(v);
    *dest = 0;
    for (i = 0; i < length; i++) {
        *dest += data[i];
    }
}
```

```
void combine4(vec_ptr v, int *dest)
{
    int i;
    int length = vec_length(v);
    int *data = get_vec_start(v);
    int sum = 0;
    for (i = 0; i < length; i++)
        sum += data[i];
    *dest = sum;
}
```



Optimization Blocker: Memory Aliasing

- Aliasing
 - Two different memory references specify single location
- Example
 - `v: [3, 2, 17]`
 - `combine3(v, get_vec_start(v)+2) --> ?`
 - `combine4(v, get_vec_start(v)+2) --> ?`
- Observations
 - Easy to happen in C
 - Since address arithmetic is allowed
 - Direct access to storage structures
 - Get in habit of introducing local variables
 - Accumulating within loops
 - Your way of telling compiler not to check for aliasing



Machine-Independent Opt. Summary

- Code Motion
 - *Compilers are good at this for simple loop/array structures*
 - *Don't do well in presence of procedure calls and memory aliasing*
- Reduction in Strength
 - Shift, add instead of multiply or divide
 - *compilers are (generally) good at this*
 - *Exact trade-offs machine-dependent*
 - Keep data in registers rather than memory
 - *compilers are not good at this, since concerned with aliasing*
- Share Common Subexpressions
 - *compilers have limited algebraic reasoning capabilities*



Important Tools

- Measurement
 - Accurately compute time taken by code
 - Most modern machines have built in cycle counters
 - Using them to get reliable measurements is tricky
 - Profile procedure calling frequencies
 - Unix tool gprof
- Observation
 - Generating assembly code
 - Lets you see what optimizations compiler can make
 - Understand capabilities/limitations of particular compiler



Code Profiling Example

- Task
 - Count word frequencies in document
 - Produce sorted list of words from most frequent to least
- Steps
 - Convert strings to lowercase
 - Apply hash function
 - Read words and insert into hash table
 - Mostly list operations
 - Maintain counter for each unique word
 - Sort results
- Data Set
 - Collected works of Shakespeare
 - 946,596 total words, 26,596 unique
 - Initial implementation: 9.2 seconds

Shakespeare's
most frequent words

29,801	the
27,529	and
21,029	I
20,957	to
18,514	of
15,370	a
14,010	you
12,936	my
11,722	in
11,519	that



Code Profiling

- Augment Executable Program with Timing Functions
 - Computes (approximate) amount of time spent in each function
 - Time computation method
 - Periodically (~ every 10ms) interrupt program
 - Determine what function is currently executing
 - Increment its timer by interval (e.g., 10ms)
 - Also maintains counter for each function indicating number of times called
- Using

```
gcc -O2 -pg prog. -o prog
./prog
```

 - Executes in normal fashion, but also generates file `gmon.out`

```
gprof prog
```

 - Generates profile information based on `gmon.out`




Profiling Results

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
86.60	8.21	8.21	1	8210.00	8210.00	sort_words
5.80	8.76	0.55	946596	0.00	0.00	lower1
4.75	9.21	0.45	946596	0.00	0.00	find_ele_rec
1.27	9.33	0.12	946596	0.00	0.00	h_add

- Call Statistics
 - Number of calls and cumulative time for each function
- Performance Limiter
 - Using inefficient sorting algorithm
 - Single call uses 87% of CPU time




Code Optimizations



- First step: Use more efficient sorting function
- Library function `qsort`



Further Optimizations



- Iter first: Use iterative function to insert elements into linked list
 - Causes code to slow down
- Iter last: Iterative function, places new entry at end of list
 - Tend to place most common words at front of list
- Big table: Increase number of hash buckets
- Better hash: Use more sophisticated hash function
- Linear lower: Move `strlen` out of loop



Profiling Observations

- Benefits
 - Helps identify performance bottlenecks
 - Especially useful when have complex system with many components
- Limitations
 - Only shows performance for data tested
 - E.g., linear lower did not show big gain, since words are short
 - Quadratic inefficiency could remain lurking in code
 - Timing mechanism fairly coarse
 - Only works for programs that run for > 3 seconds



Code Optimization II: Machine Dependent Optimizations



Previous Best Combining Code

```
void combine4(vec_ptr v, int *dest)
{
    int i;
    int length = vec_length(v);
    int *data = get_vec_start(v);
    int sum = 0;
    for (i = 0; i < length; i++)
        sum += data[i];
    *dest = sum;
}
```

- Task
 - Compute sum of all elements in vector
 - Vector represented by C-style abstract data type
 - Achieved CPE of 2.00
 - Cycles per element



General Forms of Combining

```
void abstract_combine4(vec_ptr v, data_t *dest)
{
    int i;
    int length = vec_length(v);
    data_t *data = get_vec_start(v);
    data_t t = IDENT;
    for (i = 0; i < length; i++)
        t = t OP data[i];
    *dest = t;
}
```

- Data Types

- Use different declarations for data_t
- int
- float
- double

- Operations

- Use different definitions of OP and IDENT
- + / 0
- * / 1



Machine Independent Opt. Results

- Optimizations
 - Reduce function calls and memory references within loop

Method	Integer		Floating Point	
	+	*	+	*
Abstract -g	42.06	41.86	41.44	160.00
Abstract -O2	31.25	33.25	31.25	143.00
Move vec_length	20.66	21.25	21.15	135.00
data access	6.00	9.00	8.00	117.00
Accum. in temp	2.00	4.00	3.00	5.00

- Performance Anomaly
 - Computing FP product of all elements exceptionally slow.
 - Very large speedup when accumulate in temporary
 - Caused by quirk of IA32 floating point
 - Memory uses 64-bit format, register use 80
 - Benchmark data caused overflow of 64 bits, but not 80



Pointer Code

```
void combine4p(vec_ptr v, int *dest)
{
    int length = vec_length(v);
    int *data = get_vec_start(v);
    int *dend = data+length;
    int sum = 0;
    while (data < dend) {
        sum += *data;
        data++;
    }
    *dest = sum;
}
```

- Optimization
 - Use pointers rather than array references
 - CPE: 3.00 (Compiled -O2)
 - We're not making progress here!

Warning: Some compilers do better job optimizing array code



Pointer vs. Array Code Inner Loops

- Array Code

```
.L24:                # Loop:
    addl (%eax,%edx,4),%ecx # sum += data[i]
    incl %edx              # i++
    cmpl %esi,%edx         # i:length
    jl  .L24              # if < goto Loop
```

- Pointer Code

```
.L30:                # Loop:
    addl (%eax),%ecx      # sum += *data
    addl $4,%eax          # data ++
    cmpl %edx,%eax        # data:dend
    jb  .L30             # if < goto Loop
```

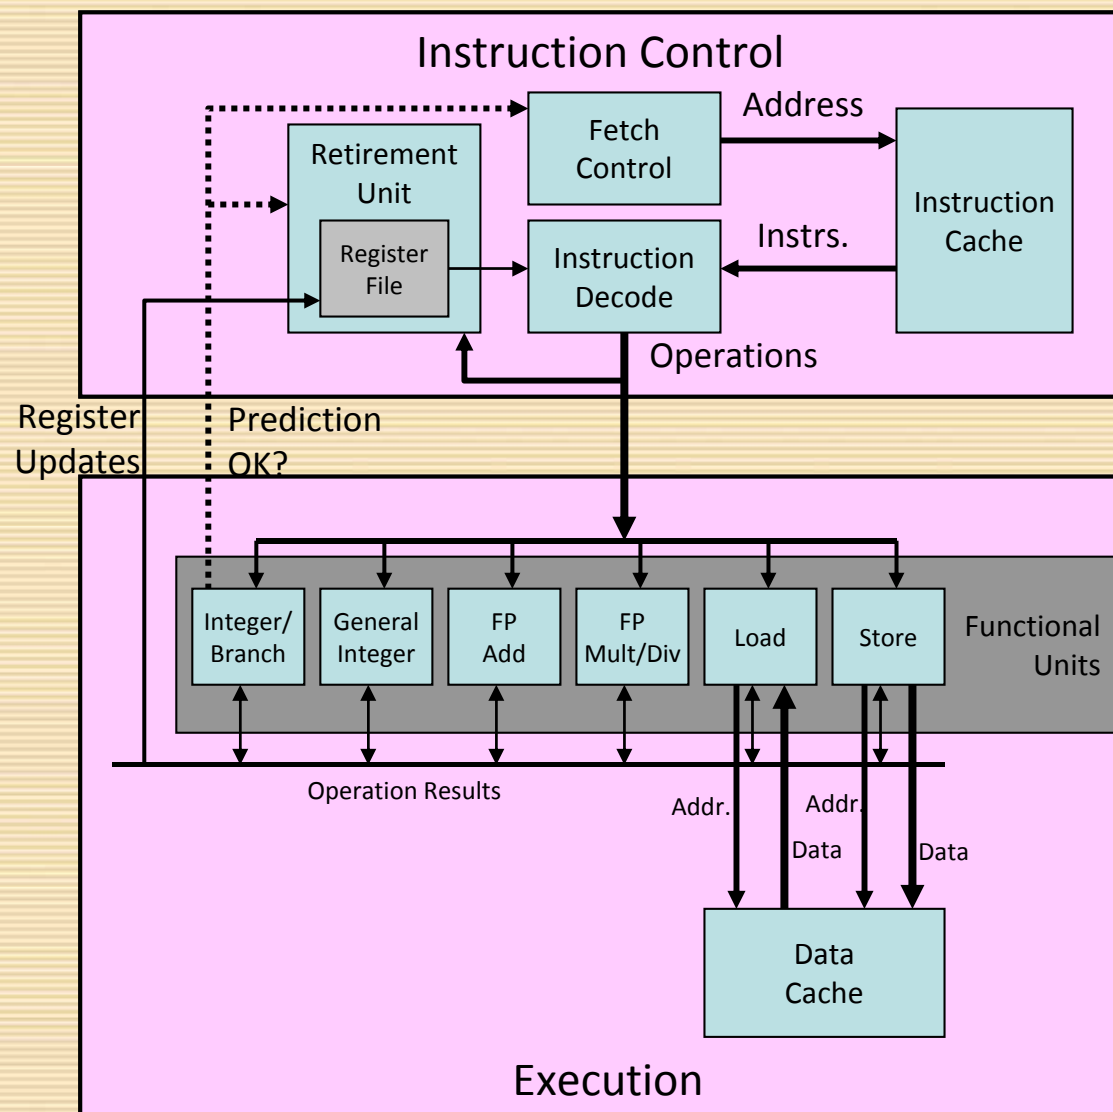
- Performance

- Array Code: 4 instructions in 2 clock cycles
- Pointer Code: Almost same 4 instructions in 3 clock cycles



Modern CPU Design

Page 396





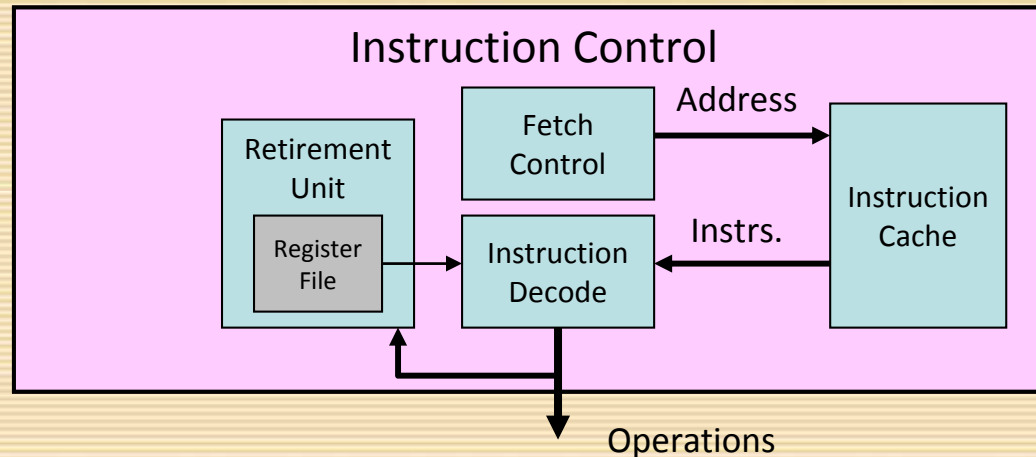
CPU Capabilities of Pentium III

- Multiple Instructions Can Execute in Parallel
 - 1 load
 - 1 store
 - 2 integer (one may be branch)
 - 1 FP Addition
 - 1 FP Multiplication or Division
- Some Instructions Take > 1 Cycle, but Can be Pipelined

– Instruction	Latency	Cycles/Issue
– Load / Store	3	1
– Integer Multiply	4	1
– Integer Divide	36	36
– Double/Single FP Multiply	5	2
– Double/Single FP Add	3	1
– Double/Single FP Divide	38	38



Instruction Control



- Grabs Instruction Bytes From Memory
 - Based on current PC + predicted targets for predicted branches
 - Hardware dynamically guesses whether branches taken/not taken and (possibly) branch target
- Translates Instructions Into *Operations*
 - Primitive steps required to perform instruction
 - Typical instruction requires 1–3 operations
- Converts Register References Into *Tags*
 - Abstract identifier linking destination of one operation with sources of later operations



Translation Example

- Version of Combine4
 - Integer data, multiply operation

```
.L24:                                # Loop:
    imull (%eax,%edx,4),%ecx         # t *= data[i]
    incl %edx                        # i++
    cmpl %esi,%edx                   # i:length
    jl .L24                          # if < goto Loop
```

- Translation of First Iteration

```
.L24:
    imull (%eax,%edx,4),%ecx

    incl %edx
    cmpl %esi,%edx
    jl .L24
```

```
load (%eax,%edx.0,4) → t.1
imull t.1, %ecx.0     → %ecx.1
incl %edx.0           → %edx.1
cmpl %esi, %edx.1     → cc.1
jl-taken cc.1
```



Translation Example #1

```
imull (%eax,%edx,4),%ecx
```

```
load (%eax,%edx.0,4) → t.1  
imull t.1, %ecx.0 → %ecx.1
```

– Split into two operations

- `load` reads from memory to generate temporary result `t.1`
- Multiply operation just operates on registers

– Operands

- Register `%eax` does not change in loop. Values will be retrieved from register file during decoding
- Register `%ecx` changes on every iteration. Uniquely identify different versions as `%ecx.0`, `%ecx.1`, `%ecx.2`, ...
 - Register *renaming*
 - Values passed directly from producer to consumers



Translation Example #2

```
incl %edx
```

```
incl %edx.0 → %edx.1
```

- Register `%edx` changes on each iteration.
Rename as `%edx.0`, `%edx.1`, `%edx.2`, ...



Translation Example #3

```
cmpl %esi,%edx
```

```
cmpl %esi, %edx.1 → cc.1
```

- Condition codes are treated similar to registers
- Assign tag to define connection between producer and consumer



Translation Example #4


jl .L24

jl-taken cc.1

- Instruction control unit determines destination of jump
- Predicts whether will be taken and target
- Starts fetching instruction at predicted destination
- Execution unit simply checks whether or not prediction was OK
- If not, it signals instruction control
 - Instruction control then “invalidates” any operations generated from misfetched instructions
 - Begins fetching and decoding instructions at correct target



Visualizing Operations




```
load (%eax,%edx,4) ➔ t.1
imull t.1, %ecx.0 ➔ %ecx.1
incl %edx.0 ➔ %edx.1
cml %esi, %edx.1 ➔ cc.1
jl-taken cc.1
```

- Operations
 - Vertical position denotes time at which executed
 - Cannot begin operation until operands available
 - Height denotes latency
- Operands
 - Arcs shown only for operands that are passed within execution unit



Visualizing Operations (cont.)



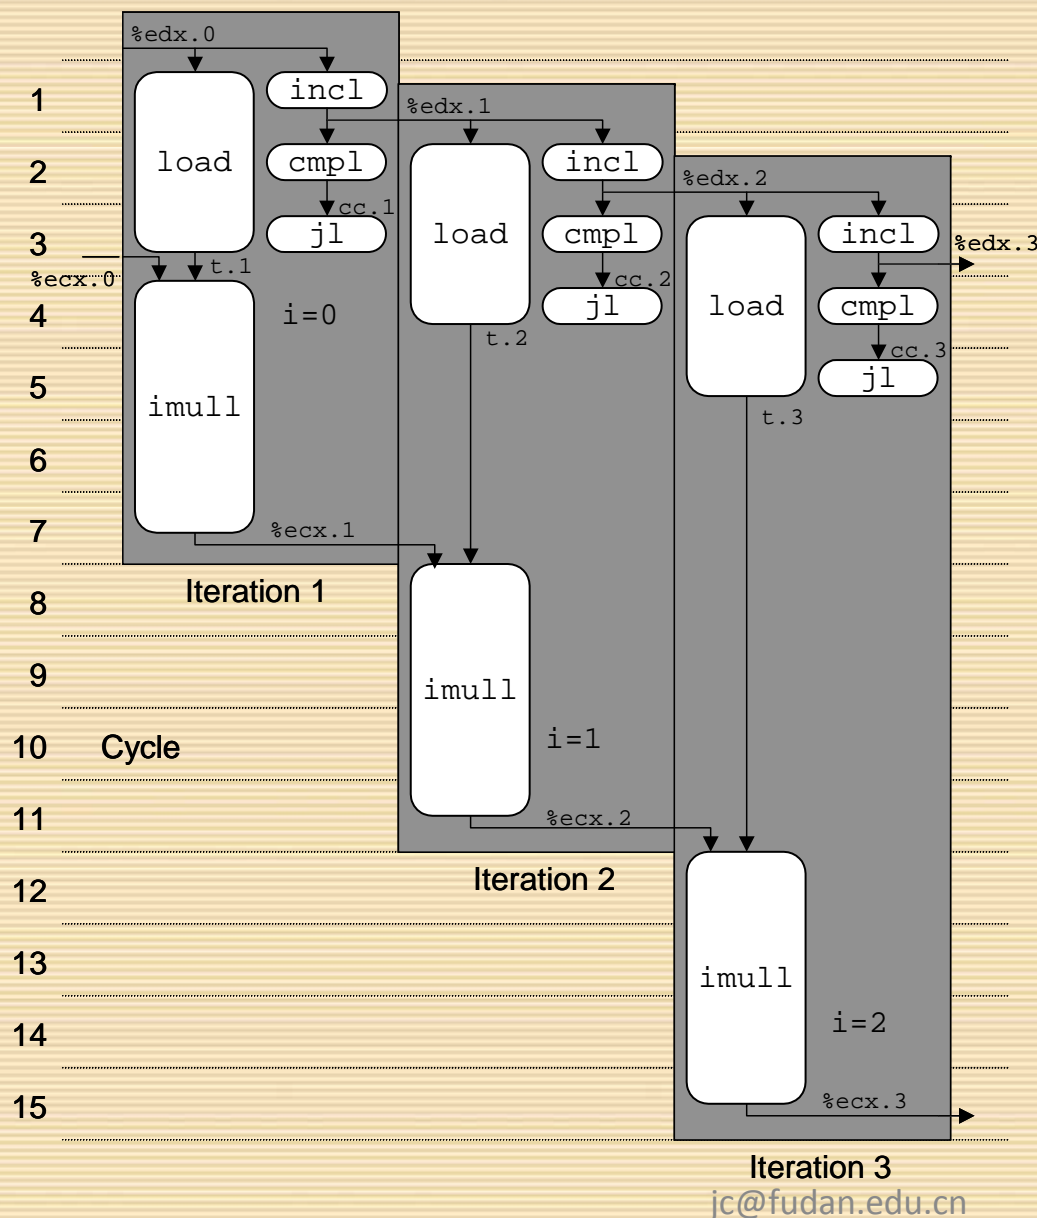
```
load (%eax,%edx,4) → t.1  
iaddl t.1, %ecx.0 → %ecx.1  
incl %edx.0 → %edx.1  
cml %esi, %edx.1 → cc.1  
jl-taken cc.1
```

- Operations
 - Same as before, except that add has latency of 1



3 Iterations of Combining Product

Page 404

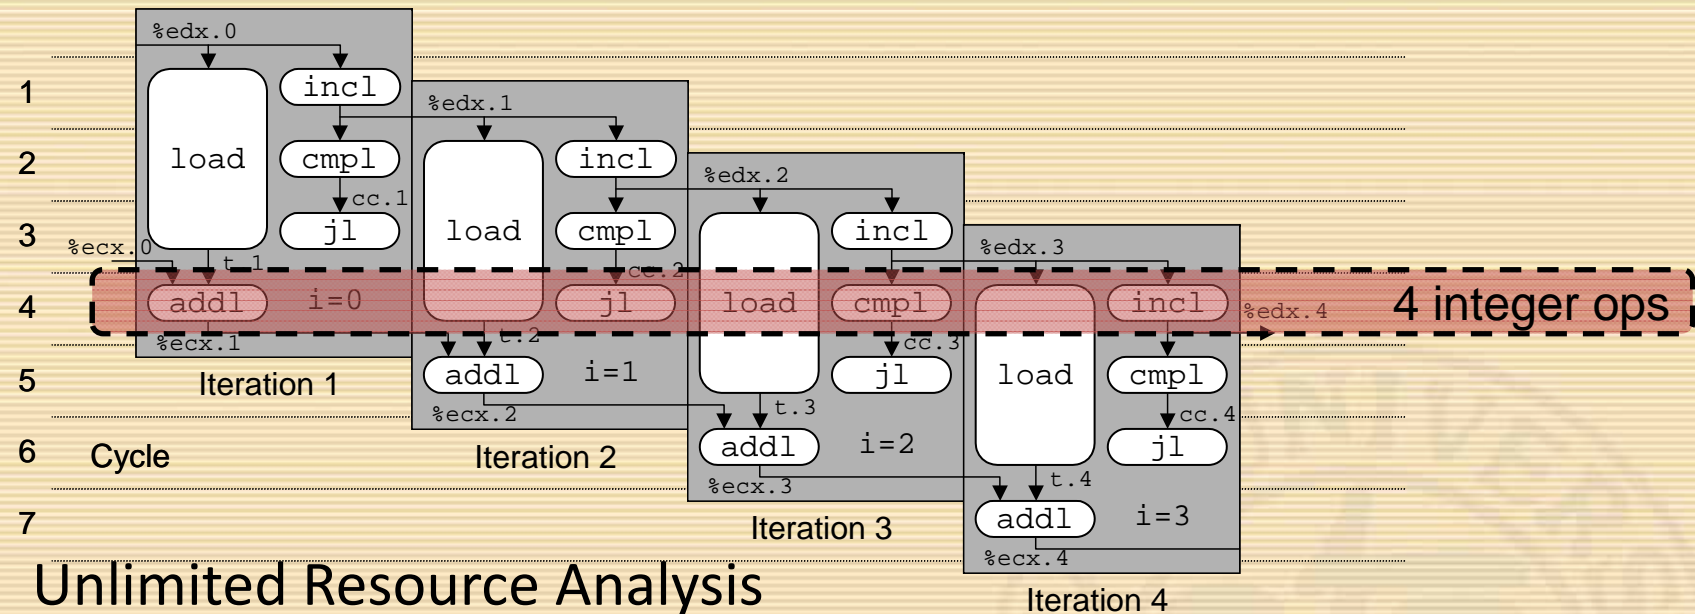


- Unlimited Resource Analysis
 - Assume operation can start as soon as operands available
 - Operations for multiple iterations overlap in time
- Performance
 - Limiting factor becomes latency of integer multiplier
 - Gives CPE of 4.0



4 Iterations of Combining Sum

Page 405

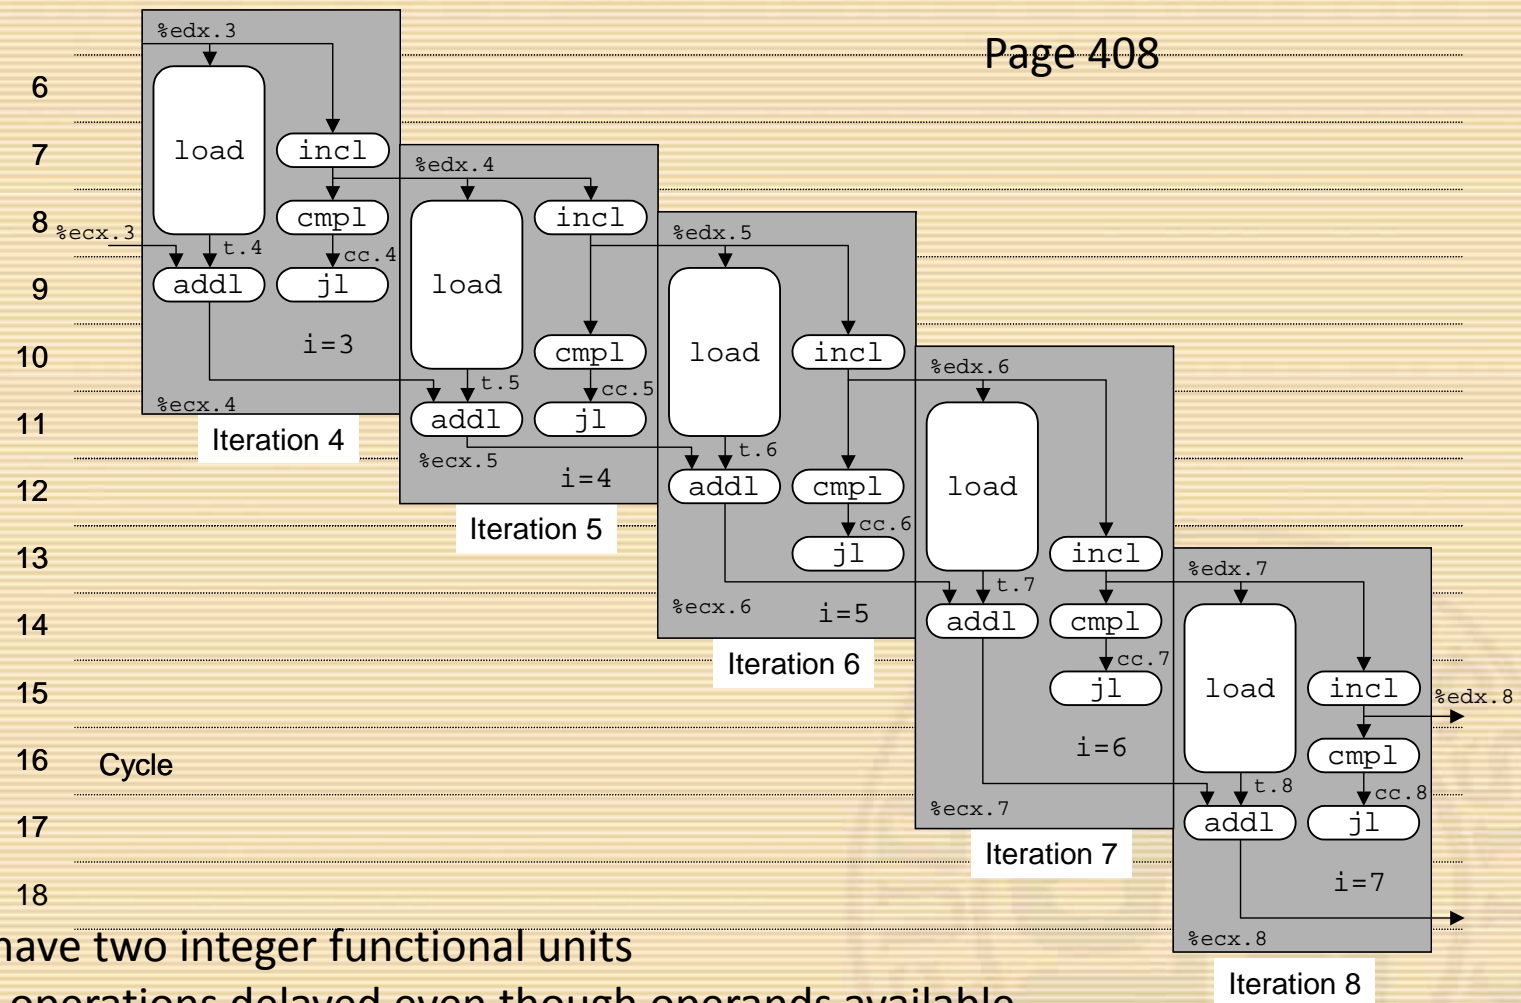


- Unlimited Resource Analysis
- Performance
 - Can begin a new iteration on each clock cycle
 - Should give CPE of 1.0
 - Would require executing 4 integer operations in parallel



Combining Sum: Resource Constraints

Page 408



- Only have two integer functional units
- Some operations delayed even though operands available
- Set priority based on program order
- **Performance**
 - Sustain CPE of 2.0



Loop Unrolling

```
void combine5(vec_ptr v, int *dest)
{
    int length = vec_length(v);
    int limit = length-2;
    int *data = get_vec_start(v);
    int sum = 0;
    int i;
    /* Combine 3 elements at a time */
    for (i = 0; i < limit; i+=3) {
        sum += data[i] + data[i+2]
              + data[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        sum += data[i];
    }
    *dest = sum;
}
```

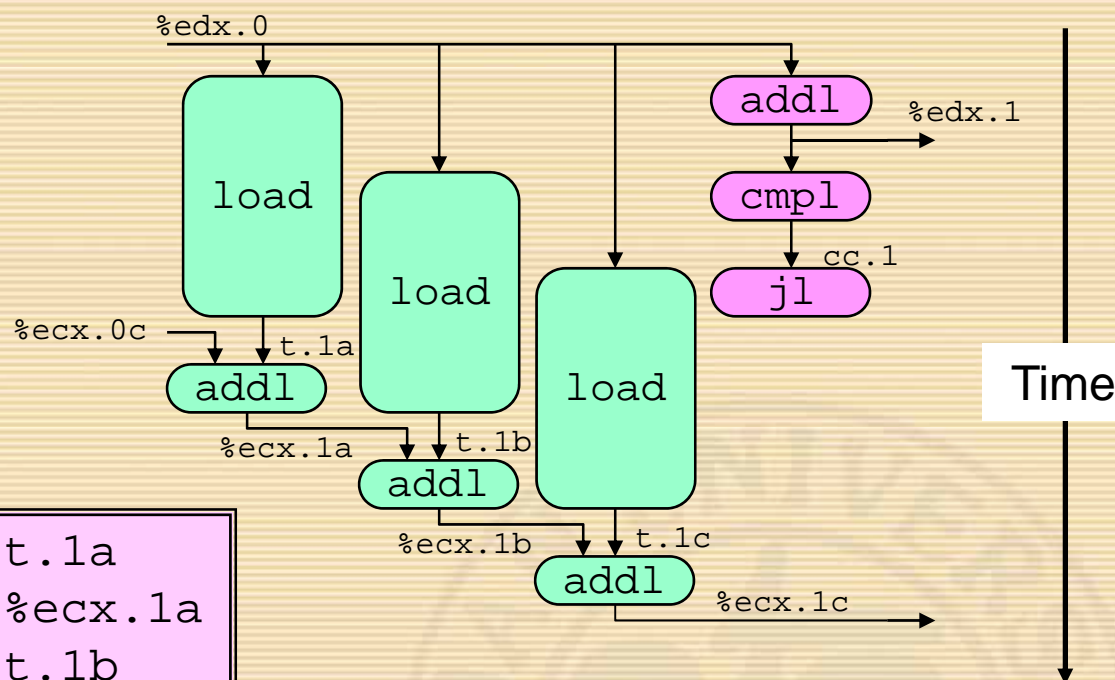
- Optimization

- Combine multiple iterations into single loop body
- Amortizes loop overhead across multiple iterations
- Finish extras at end
- Measured CPE = 1.33



Visualizing Unrolled Loop

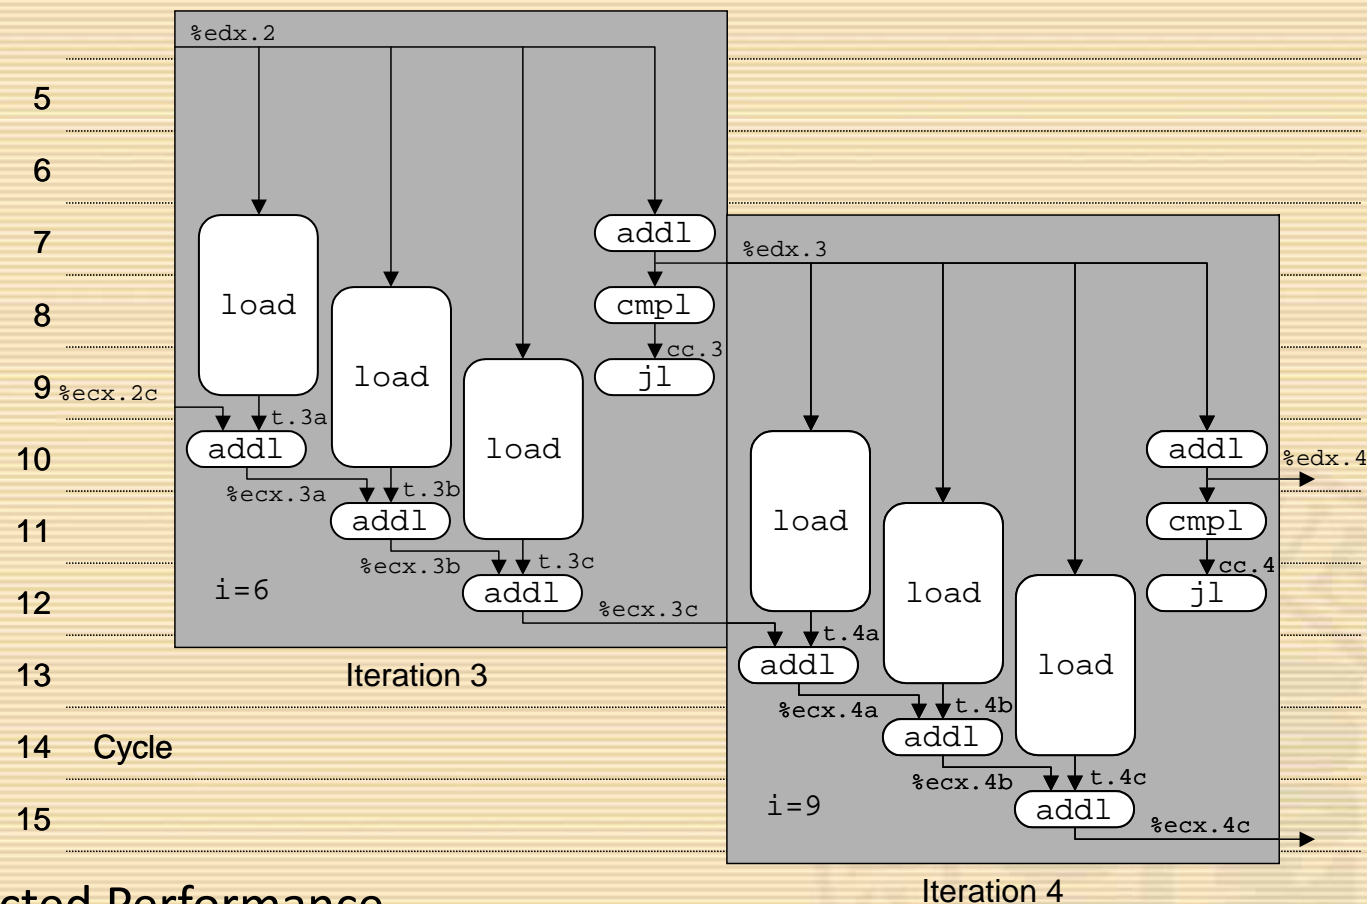
- Loads can pipeline, since don't have dependencies
- Only one set of loop control operations



```
load (%eax,%edx.0,4)    ➔ t.1a
iaddl t.1a, %ecx.0c      ➔ %ecx.1a
load 4(%eax,%edx.0,4)   ➔ t.1b
iaddl t.1b, %ecx.1a     ➔ %ecx.1b
load 8(%eax,%edx.0,4)   ➔ t.1c
iaddl t.1c, %ecx.1b     ➔ %ecx.1c
iaddl $3,%edx.0         ➔ %edx.1
cmpl %esi, %edx.1       ➔ cc.1
jl-taken cc.1
```



Executing with Loop Unrolling




- Predicted Performance
 - Can complete iteration in 3 cycles
 - Should give CPE of 1.0
- Measured Performance
 - CPE of 1.33, One iteration every 4 cycles



Effect of Unrolling

Unrolling Degree		1	2	3	4	8	16
Integer	Sum	2.00	1.50	1.33	1.50	1.25	1.06
Integer	Product	4.00					
FP	Sum	3.00					
FP	Product	5.00					

- Only helps integer sum for our examples
 - Other cases constrained by functional unit latencies
- Effect is nonlinear with degree of unrolling
 - Many subtle effects determine exact scheduling of operations




Serial Computation

- Computation

$$(((((((((((1 * x_0) * x_1) * x_2) * x_3) * x_4) * x_5) * x_6) * x_7) * x_8) * x_9) * x_{10}) * x_{11})$$

- Performance

- N elements, D cycles/operation
- N*D cycles





Parallel Loop Unrolling

```
void combine6(vec_ptr v, int *dest)
{
    int length = vec_length(v);
    int limit = length-1;
    int *data = get_vec_start(v);
    int x0 = 1;
    int x1 = 1;
    int i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x0 *= data[i];
        x1 *= data[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x0 *= data[i];
    }
    *dest = x0 * x1;
}
```

- Code Version
 - Integer product
- Optimization
 - Accumulate in two different products
 - Can be performed simultaneously
 - Combine at end
- Performance
 - CPE = 2.0
 - 2X performance



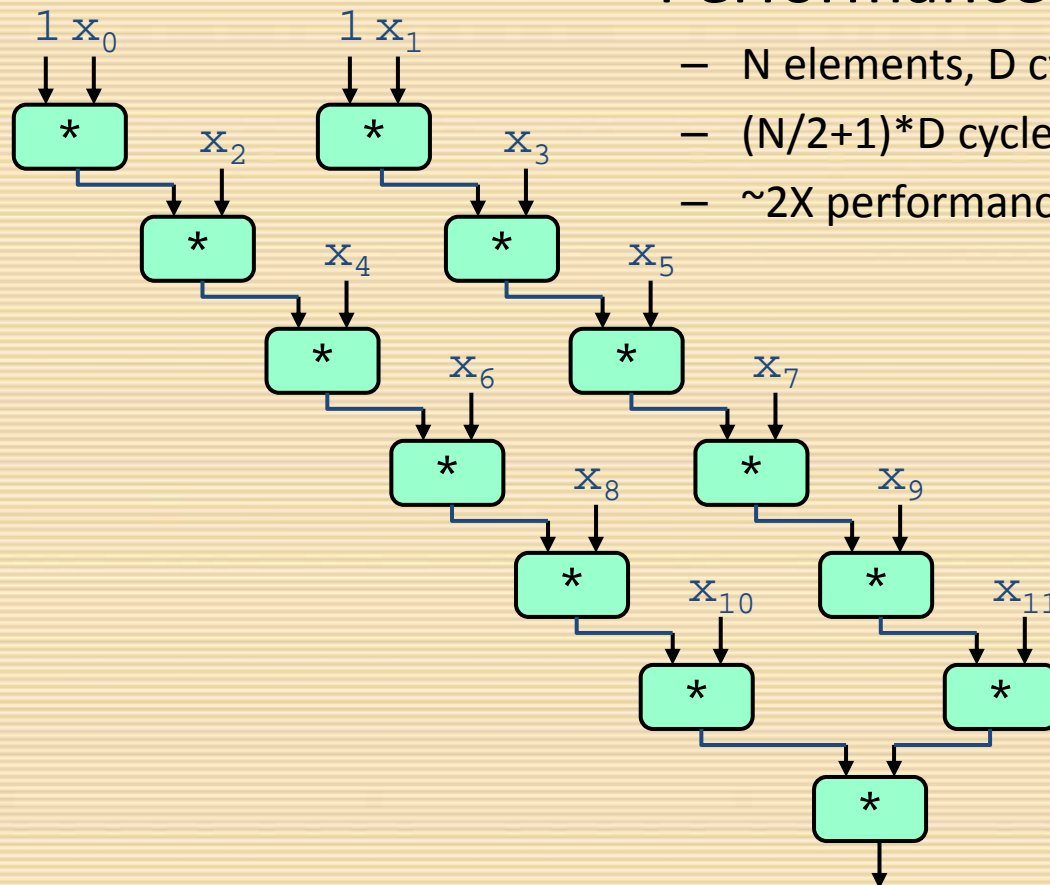
Dual Product Computation

- Computation

$$\begin{aligned} & ((((((1 * x_0) * x_2) * x_4) * x_6) * x_8) * x_{10}) * \\ & ((((((1 * x_1) * x_3) * x_5) * x_7) * x_9) * x_{11})) \end{aligned}$$

- Performance

- N elements, D cycles/operation
- $(N/2+1)*D$ cycles
- ~2X performance improvement





Requirements for Parallel Computation

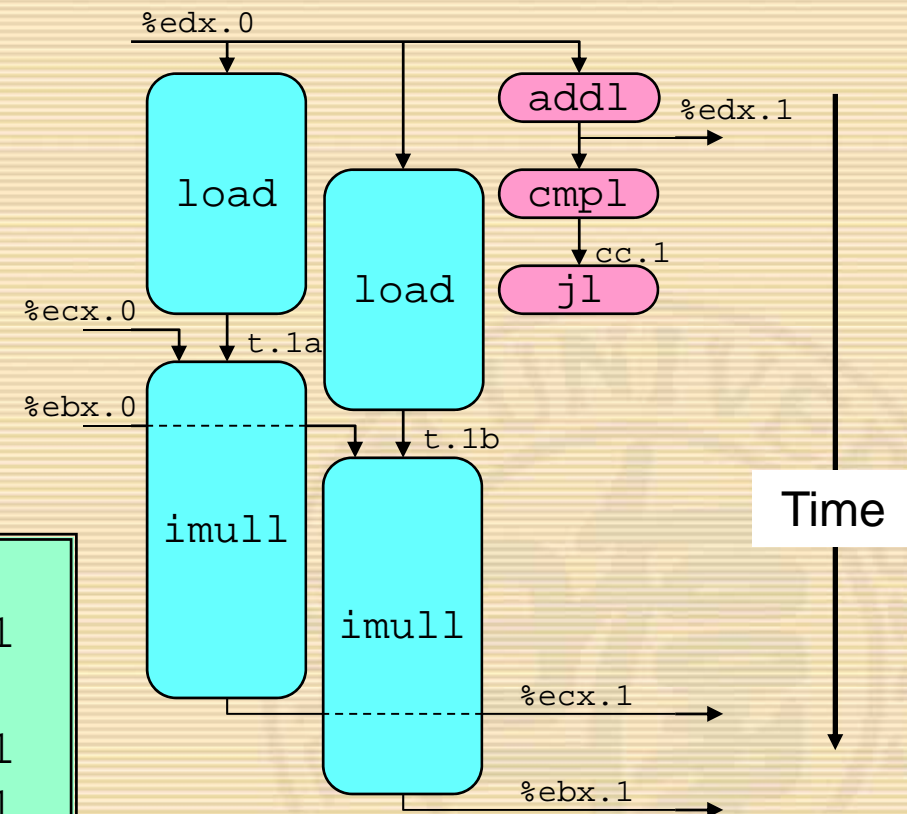
- Mathematical
 - Combining operation must be associative & commutative
 - OK for integer multiplication
 - Not strictly true for floating point
 - OK for most applications
- Hardware
 - Pipelined functional units
 - Ability to dynamically extract parallelism from code



Visualizing Parallel Loop

- Two multiplies within loop no longer have data dependency
- Allows them to pipeline

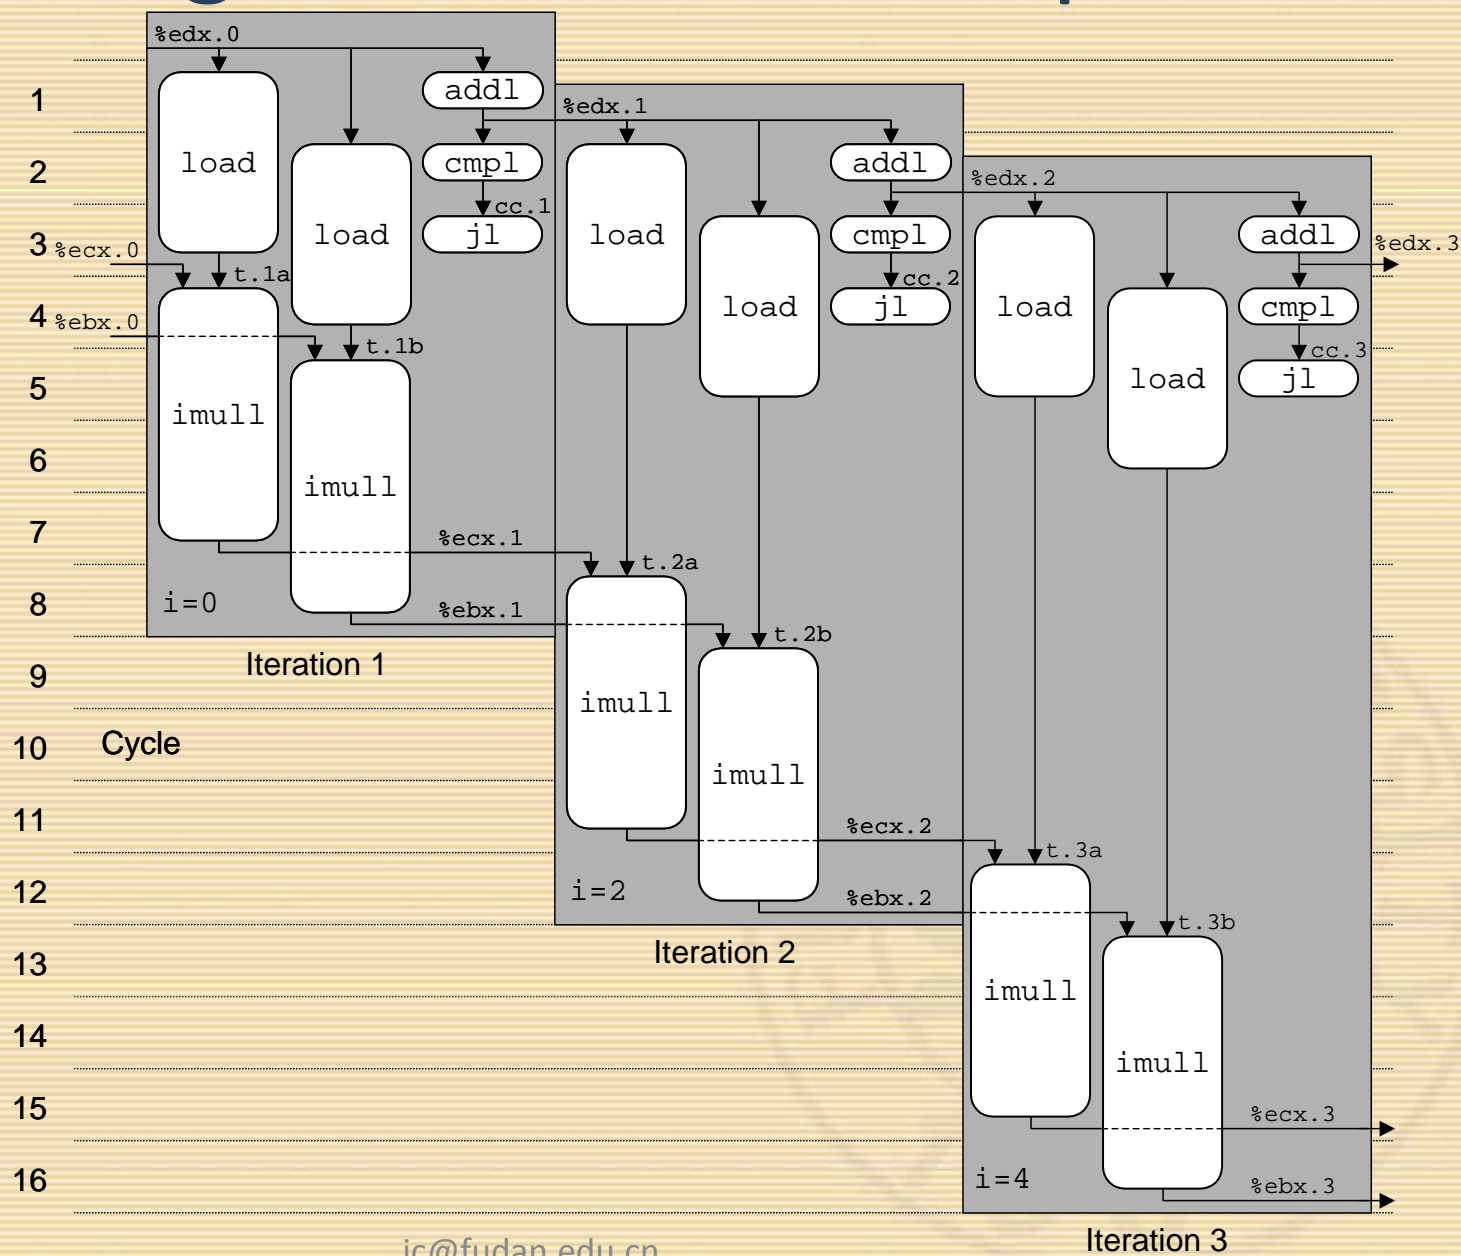
```
load (%eax,%edx.0,4)    → t.1a
imull t.1a, %ecx.0       → %ecx.1
load 4(%eax,%edx.0,4)   → t.1b
imull t.1b, %ebx.0       → %ebx.1
iaddl $2,%edx.0         → %edx.1
cmpl %esi, %edx.1       → cc.1
jl-taken cc.1
```





Executing with Parallel Loop

- Predicted Performance
 - Can keep 4-cycle multiplier busy performing two simultaneous multiplications
 - Gives CPE of 2.0





Optimization Results for Combining

Method	Integer		Floating Point	
	+	*	+	*
Abstract -g	42.06	41.86	41.44	160.00
Abstract -O2	31.25	33.25	31.25	143.00
Move vec_length	20.66	21.25	21.15	135.00
data access	6.00	9.00	8.00	117.00
Accum. in temp	2.00	4.00	3.00	5.00
Pointer	3.00	4.00	3.00	5.00
Unroll 4	1.50	4.00	3.00	5.00
Unroll 16	1.06	4.00	3.00	5.00
2 X 2	1.50	2.00	2.00	2.50
4 X 4	1.50	2.00	1.50	2.50
8 X 4	1.25	1.25	1.50	2.00
Theoretical Opt.	1.00	1.00	1.00	2.00
<i>Worst : Best</i>	39.7	33.5	27.6	80.0



Parallel Unrolling: Method #2

```
void combine6aa(vec_ptr v, int *dest)
{
    int length = vec_length(v);
    int limit = length-1;
    int *data = get_vec_start(v);
    int x = 1;
    int i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x *= (data[i] * data[i+1]);
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x *= data[i];
    }
    *dest = x;
}
```

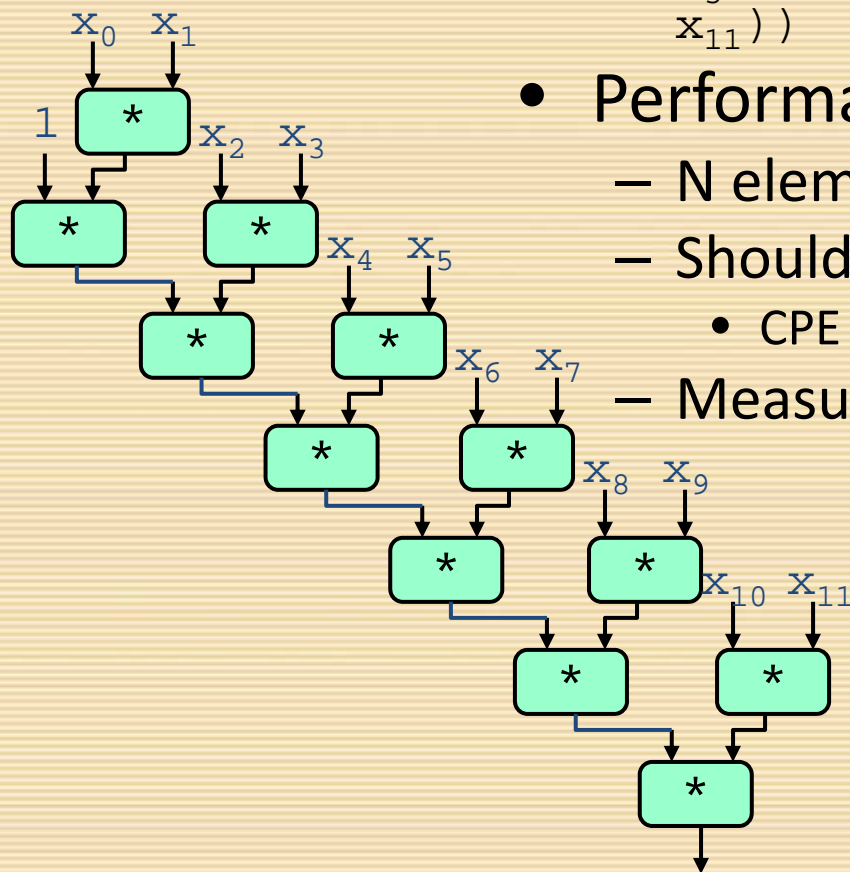
- Code Version
 - Integer product
- Optimization
 - Multiply pairs of elements together
 - And then update product
 - “Tree height reduction”
- Performance
 - CPE = 2.5



Method #2 Computation

- Computation

$$(((((((1 * (x_0 * x_1)) * (x_2 * x_3)) * (x_4 * x_5)) * (x_6 * x_7)) * (x_8 * x_9)) * (x_{10} * x_{11}))$$



- Performance

- N elements, D cycles/operation
- Should be $(N/2+1)*D$ cycles
 - CPE = 2.0
- Measured CPE worse

Unrolling	CPE (measured)	CPE (theoretical)
2	2.50	2.00
3	1.67	1.33
4	1.50	1.00
6	1.78	1.00




Understanding Parallelism

```
/* Combine 2 elements at a time */  
for (i = 0; i < limit; i+=2) {  
    x = (x * data[i]) * data[i+1];  
}
```

- CPE = 4.00
- All multiplies performed in sequence

```
/* Combine 2 elements at a time */  
for (i = 0; i < limit; i+=2) {  
    x = x * (data[i] * data[i+1]);  
}
```

- CPE = 2.50
- Multiplies overlap





Limitations of Parallel Execution

- Need Lots of Registers (Register Spilling)
 - To hold sums/products
 - Only 6 usable integer registers
 - Also needed for pointers, loop conditions
 - 8 FP registers
 - When not enough registers, must spill temporaries onto stack
 - Wipes out any performance gains
 - Not helped by renaming
 - Cannot reference more operands than instruction set allows
 - Major drawback of IA32 instruction set



Register Spilling Example

- Example
 - 8 X 8 integer product
 - 7 local variables share 1 register
 - See that are storing locals on stack
 - E.g., at $-8(\%ebp)$

.L165:

```
imull (%eax),%ecx
movl -4(%ebp),%edi
imull 4(%eax),%edi
movl %edi,-4(%ebp)
movl -8(%ebp),%edi
imull 8(%eax),%edi
movl %edi,-8(%ebp)
movl -12(%ebp),%edi
imull 12(%eax),%edi
movl %edi,-12(%ebp)
movl -16(%ebp),%edi
imull 16(%eax),%edi
movl %edi,-16(%ebp)
```

...

```
addl $32,%eax
addl $8,%edx
cmpl -32(%ebp),%edx
jl .L165
```



Summary: Results for Pentium III

Method	Integer		Floating Point	
	+	*	+	*
Abstract -g	42.06	41.86	41.44	160.00
Abstract -O2	31.25	33.25	31.25	143.00
Move vec_length	20.66	21.25	21.15	135.00
data access	6.00	9.00	8.00	117.00
Accum. in temp	2.00	4.00	3.00	5.00
Unroll 4	1.50	4.00	3.00	5.00
Unroll 16	1.06	4.00	3.00	5.00
4 X 2	1.50	2.00	1.50	2.50
8 X 4	1.25	1.25	1.50	2.00
8 X 8	1.88	1.88	1.75	2.00
Worst : Best	39.7	33.5	27.6	80.0

—Biggest gain doing basic optimizations



Results for Alpha Processor

Method	Integer		Floating Point	
	+	*	+	*
Abstract -g	40.14	47.14	52.07	53.71
Abstract -O2	25.08	36.05	37.37	32.02
Move vec_length	19.19	32.18	28.73	32.73
data access	6.26	12.52	13.26	13.01
Accum. in temp	1.76	9.01	8.08	8.01
Unroll 4	1.51	9.01	6.32	6.32
Unroll 16	1.25	9.01	6.33	6.22
4 X 2	1.19	4.69	4.44	4.45
8 X 4	1.15	4.12	2.34	2.01
8 X 8	1.11	4.24	2.36	2.08
Worst : Best	36.2	11.4	22.3	26.7

- Overall trends very similar to those for Pentium III
- Even though very different architecture and compiler



Results for Pentium 4

Method	Integer		Floating Point	
	+	*	+	*
Abstract -g	35.25	35.34	35.85	38.00
Abstract -O2	26.52	30.26	31.55	32.00
Move vec_length	18.00	25.71	23.36	24.25
data access	3.39	31.56	27.50	28.35
Accum. in temp	2.00	14.00	5.00	7.00
Unroll 4	1.01	14.00	5.00	7.00
Unroll 16	1.00	14.00	5.00	7.00
4 X 2	1.02	7.00	2.63	3.50
8 X 4	1.01	3.98	1.82	2.00
8 X 8	1.63	4.50	2.42	2.31
Worst : Best	35.2	8.9	19.7	19.0

- Higher latencies (int * = 14, fp + = 5.0, fp * = 7.0)
 - Clock runs at 2.0 GHz
 - Not an improvement over 1.0 GHz P3 for integer *
- Avoids FP multiplication anomaly



What About Branches?

- Challenge
 - Instruction Control Unit must work well ahead of Exec. Unit
 - To generate enough operations to keep EU busy

```
80489f3: movl    $0x1,%ecx
80489f8: xorl    %edx,%edx
80489fa: cmpl    %esi,%edx
80489fc: jnl     8048a25
80489fe: movl    %esi,%esi
8048a00: imull   (%eax,%edx,4),%ecx
```

} Executing

} Fetching &
Decoding

- When encountered with conditional branch, cannot reliably determine where to continue fetching



Avoiding Branches

- On Modern Processor, Branches are Very Expensive
 - Unless prediction can be reliable
 - When possible, best to avoid altogether
- Example
 - Compute maximum of two values
 - 14 cycles when prediction correct, 29 cycles when incorrect
 - Overall: 20 cycles(?)

```
int max(int x, int y)
{
    return (x < y) ? y : x;
}
```

```
movl 12(%ebp),%edx    # Get y
movl 8(%ebp),%eax     # rval=x
cmpl %edx,%eax        # rval:y
jge L11               # skip when >=
movl %edx,%eax        # rval=y
L11:
```



Avoiding Branches with Bit Tricks

- In style of Lab #1
- Use mask rather than conditionals

```
int bmax(int x, int y)
{
    int mask = -(x>y);
    return (mask & x) | (~mask & y);
}
```

- Compiler still uses conditional
 - 16 cycles when predict correctly, 32 otherwise

```
xorl %edx,%edx      # mask = 0
movl 8(%ebp),%eax
movl 12(%ebp),%ecx
cmpl %ecx,%eax
jle L13             # skip if x<=y
movl $-1,%edx       # mask = -1
L13:
```



Avoiding Branches with Bit Tricks

– Force compiler to generate desired code

```
int bvmx(int x, int y)
{
    volatile int t = (x>y);
    int mask = -t;
    return (mask & x) |
           (~mask & y);
}
```

```
movl 8(%ebp),%ecx    # Get x
movl 12(%ebp),%edx   # Get y
cmpl %edx,%ecx       # x:y
setg %al             # (x>y)
movzbl %al,%eax      # Zero extend
movl %eax,-4(%ebp)   # Save as t
movl -4(%ebp),%eax   # Retrieve t
```

– `volatile` declaration forces value to be written to memory

- Compiler must therefore generate code to compute `t`
- Simplest way is `setg/movzbl` combination

– Not very elegant!

- A hack to get control over compiler

– 22 clock cycles on all data

- Better than misprediction



Conditional Move

- Added with P6 microarchitecture (PentiumPro onward)
- `cmovXXl %edx, %eax`
 - If condition `XX` holds, copy `%edx` to `%eax`
 - Doesn't involve any branching
 - Handled as operation within Execution Unit

```
movl 8(%ebp),%edx    # Get x
movl 12(%ebp),%eax   # rval=y
cmpl %edx, %eax      # rval:x
cmovll %edx,%eax     # If <, rval=x
```

- Current version of GCC won't use this instruction
 - Thinks it's compiling for a 386
- Performance
 - 14 cycles on all data



Machine-Dependent Opt. Summary

- Pointer Code
 - Look carefully at generated code to see whether helpful
- Loop Unrolling
 - Some compilers do this automatically
 - Generally not as clever as what can achieve by hand
- Exposing Instruction-Level Parallelism
 - Very machine dependent
- Warning:
 - Benefits depend heavily on particular machine
 - Best if performed by compiler
 - But GCC on IA32/Linux is not very good
 - Do only for performance-critical parts of code



Amdahl's Law P443

$$\begin{aligned}T_{\text{new}} &= (1-\alpha)T_{\text{old}} + (\alpha T_{\text{old}})/k \\ &= T_{\text{old}}[(1-\alpha) + \alpha/k]\end{aligned}$$

$$S = T_{\text{old}} / T_{\text{new}} = 1/[(1-\alpha) + \alpha/k]$$

$$S_{\infty} = 1/(1-\alpha)$$



Role of Programmer

How should I write my programs, given that I have a good, optimizing compiler?

- Don't: Smash Code into Oblivion
 - Hard to read, maintain, & assure correctness
- Do:
 - Select best algorithm
 - Write code that's readable & maintainable
 - Procedures, recursion, without built-in constant limits
 - Even though these factors can slow down code
 - Eliminate optimization blockers
 - Allows compiler to do its job
- Focus on Inner Loops
 - Do detailed optimizations where code will be executed repeatedly
 - Will get most performance gain here



Chap. 5 Summary

- Machine Independent Optimization
 - Code Motion
 - Simpler Operation, Register / Memory, Sub-expression
 - Loop
 - Repeatedly Called Function
 - Repeatedly Referenced Memory Location (Aliasing)
- Machine Dependant Optimization
 - Instructions -> Operations, Latency, Issue Time
 - Pointer / Array, Loop Unrolling, Parallel
 - Load, Store



The Memory Hierarchy



Random-Access Memory (RAM)

- Key features
 - **RAM** is packaged as a chip.
 - Basic storage unit is a **cell** (one bit per cell).
 - Multiple RAM chips form a memory.
- Static RAM (**SRAM**)
 - Each cell stores bit with a six-transistor circuit.
 - Retains value indefinitely, as long as it is kept powered.
 - Relatively insensitive to disturbances such as electrical noise.
 - Faster and more expensive than DRAM.
- Dynamic RAM (**DRAM**)
 - Each cell stores bit with a capacitor and transistor.
 - Value must be refreshed every 10-100 ms.
 - Sensitive to disturbances.
 - Slower and cheaper than SRAM.




SRAM vs DRAM Summary

	Tran. per bit	Access time	Persist?	Sensitive?	Cost	Applications
SRAM	6	1X	Yes	No	100x	cache memories
DRAM	1	10X	No	Yes	1X	Main memories, frame buffers



Conventional DRAM Organization

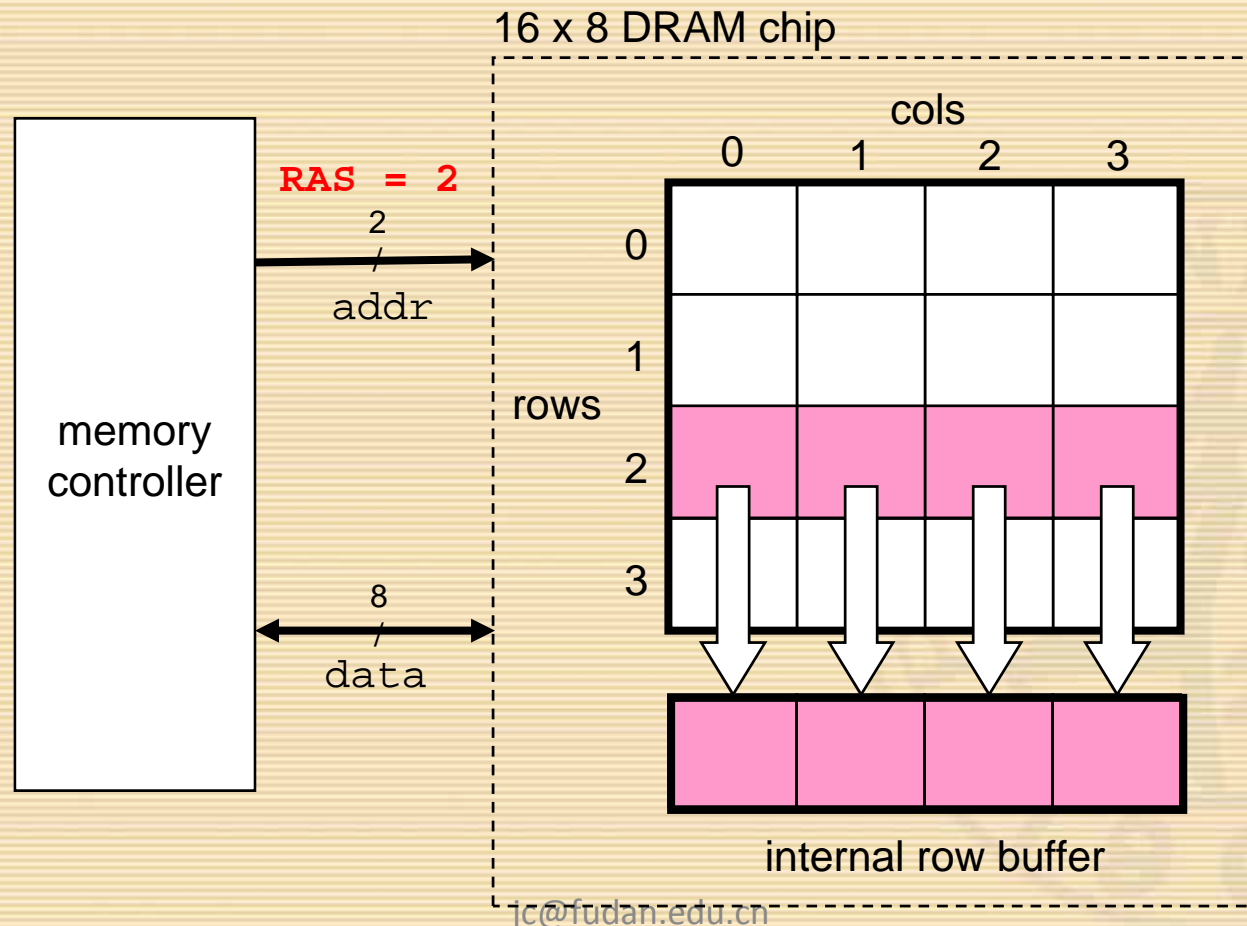
- $d \times w$ DRAM:
 - dw total bits organized as d **supercells** of size w bits





Reading DRAM Supercell (2,1)

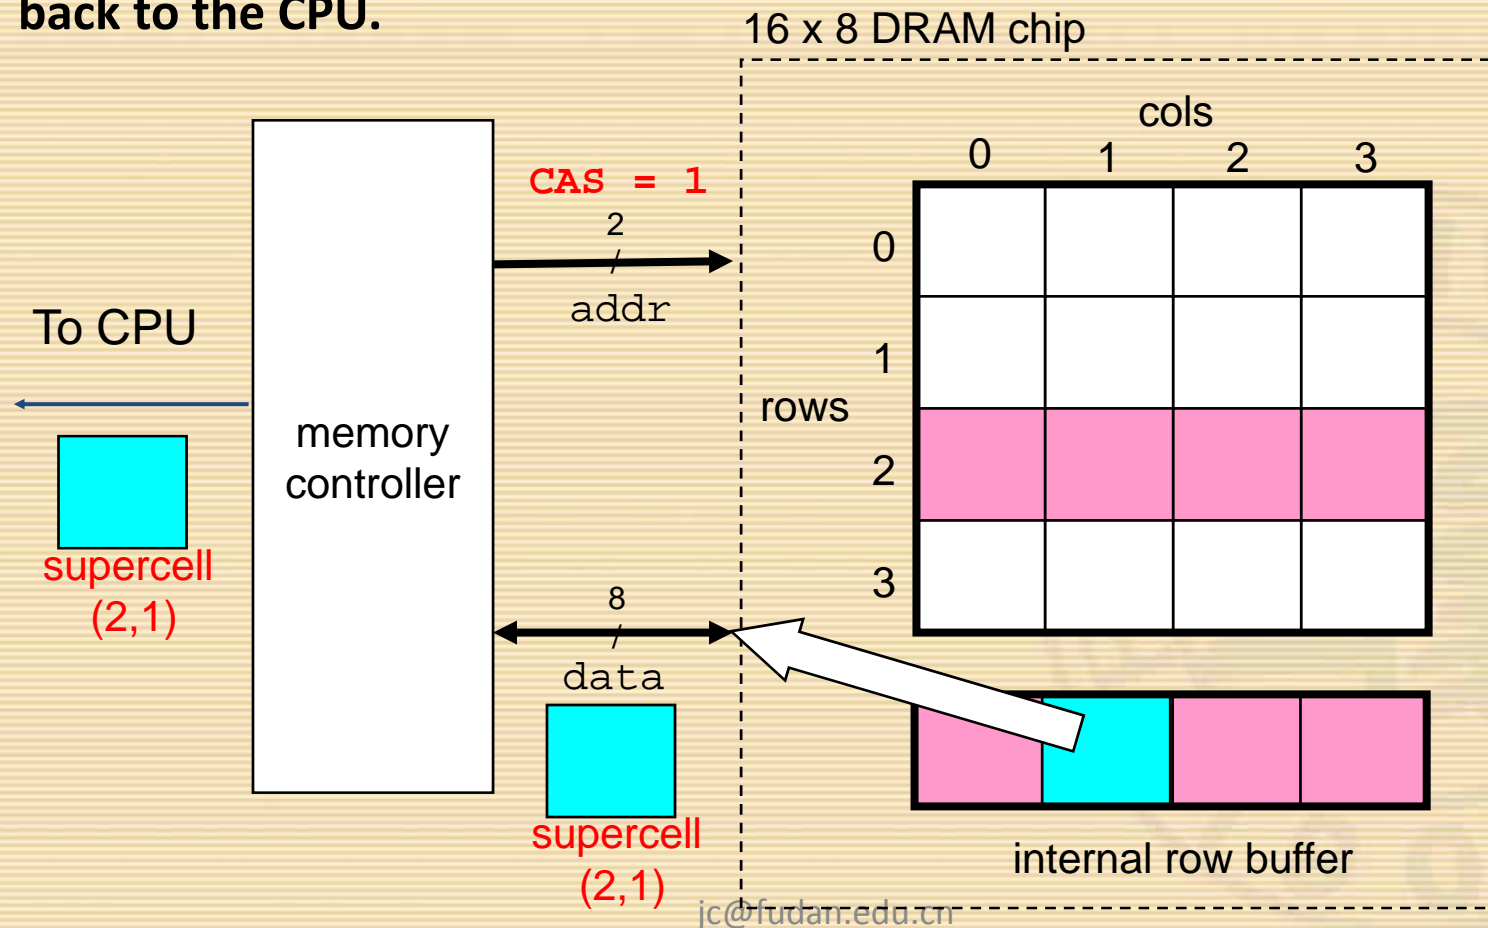
- Step 1(a): Row access strobe (**RAS**) selects row 2.
- Step 1(b): Row 2 copied from DRAM array to row buffer.





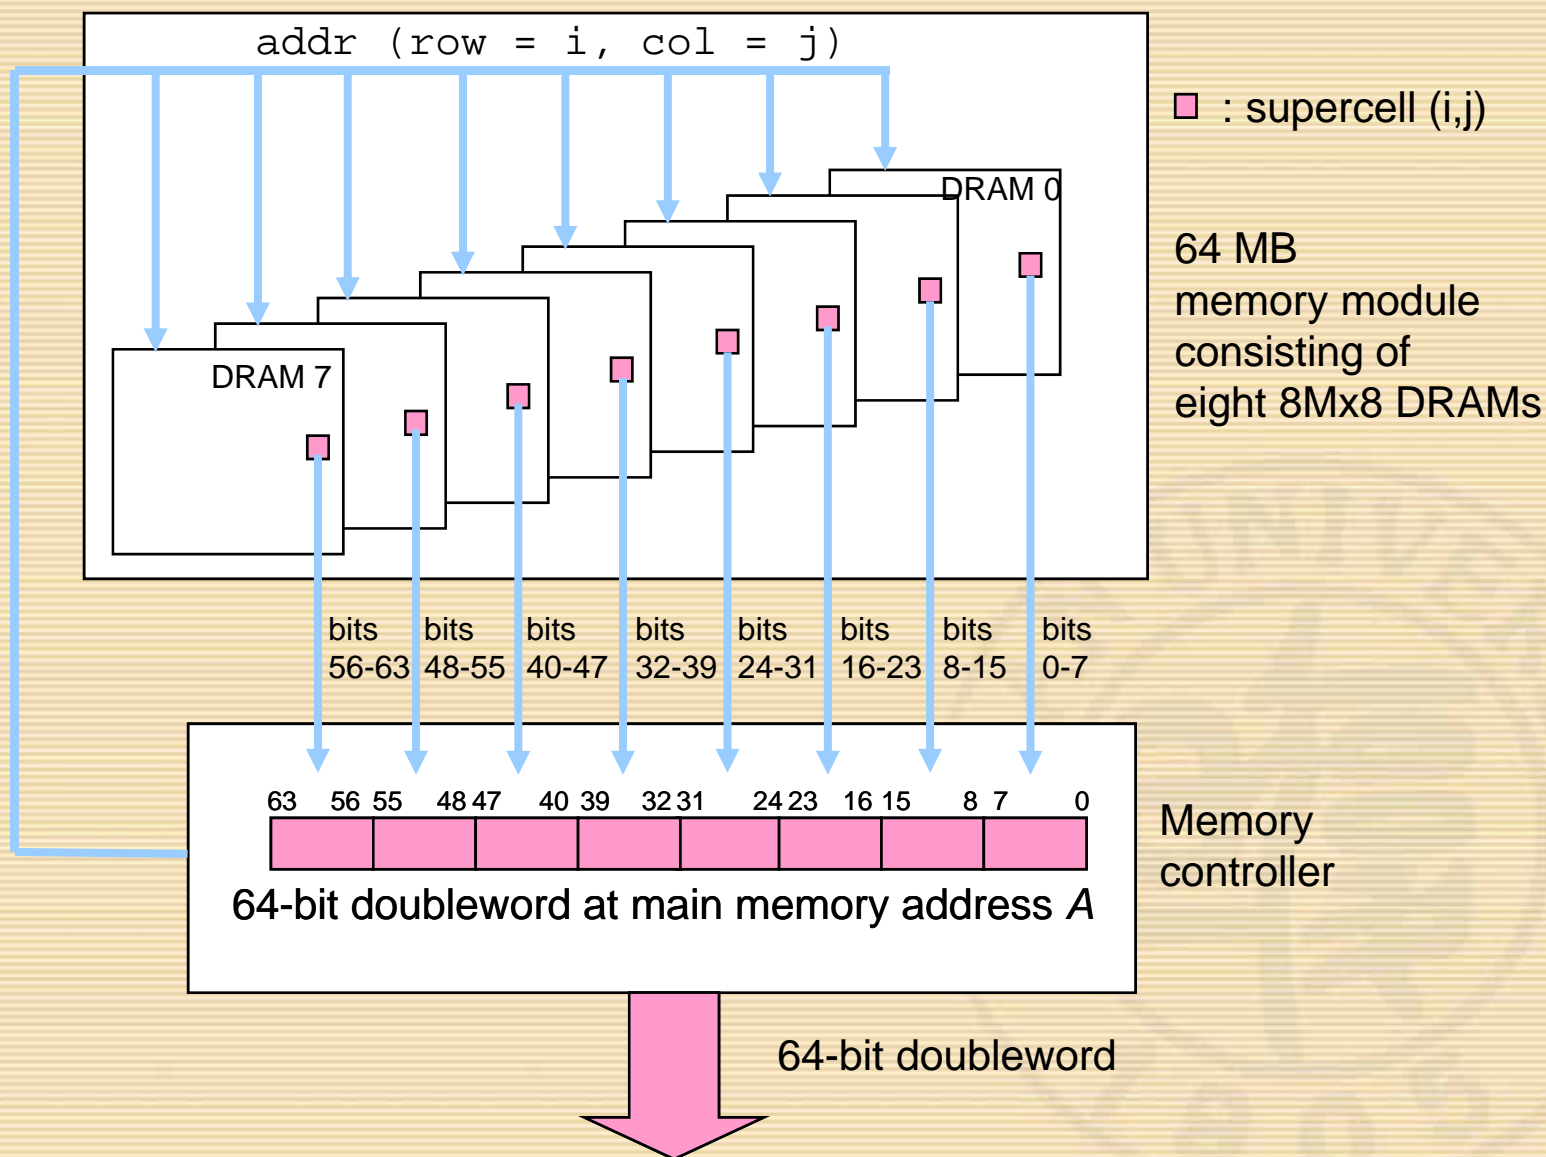
Reading DRAM Supercell (2,1)

- Step 2(a): Column access strobe (**CAS**) selects column 1.
- Step 2(b): Supercell (2,1) copied from buffer to data lines, and eventually back to the CPU.





Memory Modules





Enhanced DRAMs

- All enhanced DRAMs are built around the conventional DRAM core.
 - Fast page mode DRAM (**FPM DRAM**)
 - Access contents of row with [RAS, CAS, CAS, CAS, CAS] instead of [(RAS,CAS), (RAS,CAS), (RAS,CAS), (RAS,CAS)].
 - Extended data out DRAM (**EDO DRAM**)
 - Enhanced FPM DRAM with more closely spaced CAS signals.
 - Synchronous DRAM (**SDRAM**)
 - Driven with rising clock edge instead of asynchronous control signals.
 - Double data-rate synchronous DRAM (**DDR SDRAM**)
 - Enhancement of SDRAM that uses both clock edges as control signals.
 - Video RAM (**VRAM**)
 - Like FPM DRAM, but output is produced by shifting row buffer
 - Dual ported (allows concurrent reads and writes)



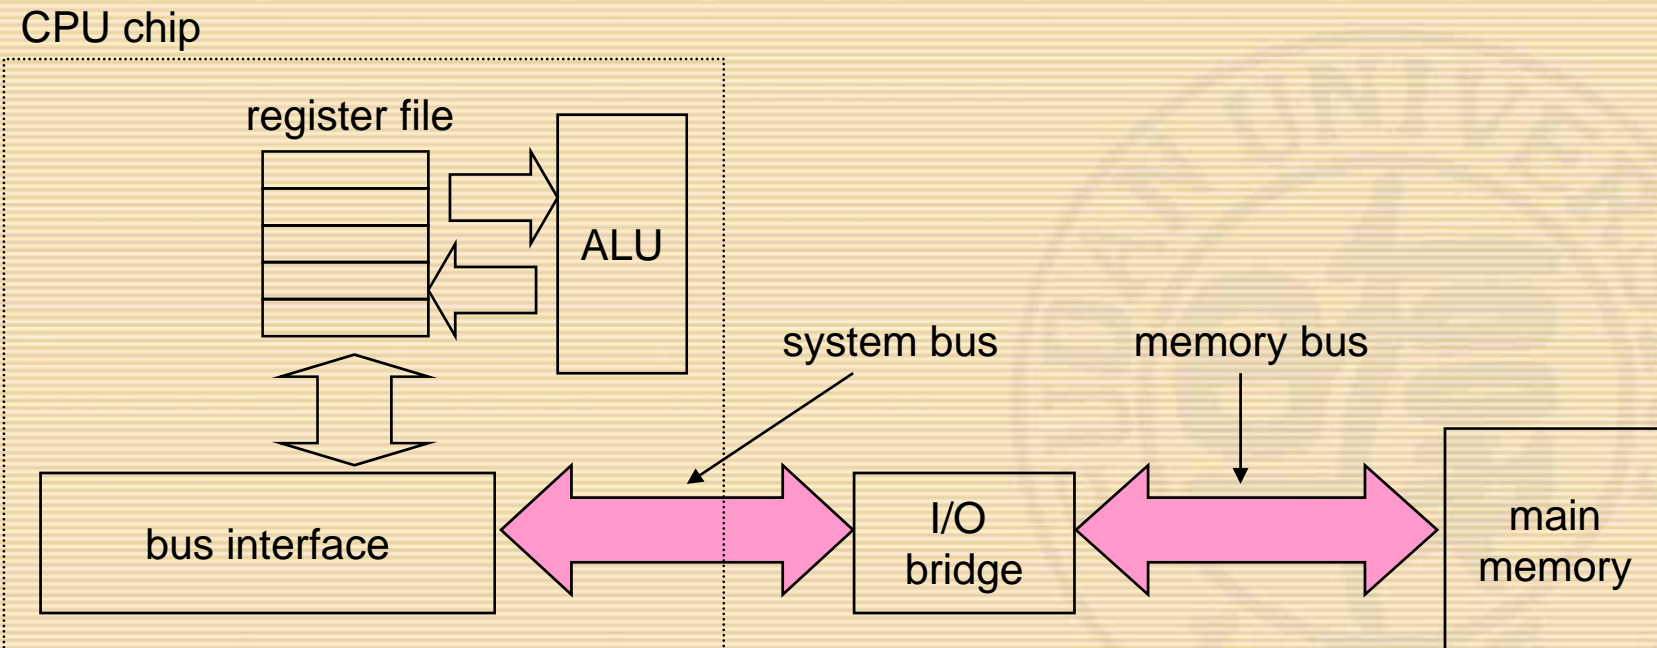
Nonvolatile Memories

- DRAM and SRAM are volatile memories
 - Lose information if powered off.
- Nonvolatile memories retain value even if powered off.
 - Generic name is read-only memory (**ROM**).
 - Misleading because some ROMs can be read and modified.
- Types of ROMs
 - Programmable ROM (**PROM**)
 - Erasable programmable ROM (**EPROM**)
 - Electrically erasable PROM (**EEPROM**)
 - Flash memory
- **Firmware**
 - Program stored in a ROM
 - Boot time code, BIOS (basic input/output system)
 - graphics cards, disk controllers.



Typical Bus Structure Connecting CPU and Memory

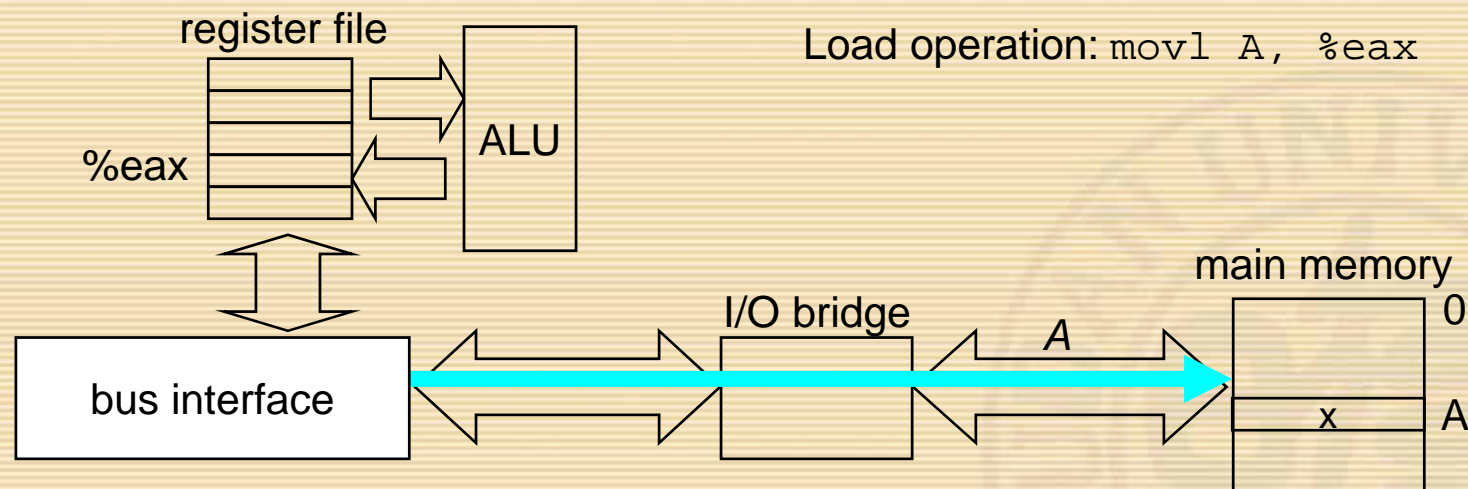
- A **bus** is a collection of parallel wires that carry address, data, and control signals.
- Buses are typically shared by multiple devices.





Memory Read Transaction (1)

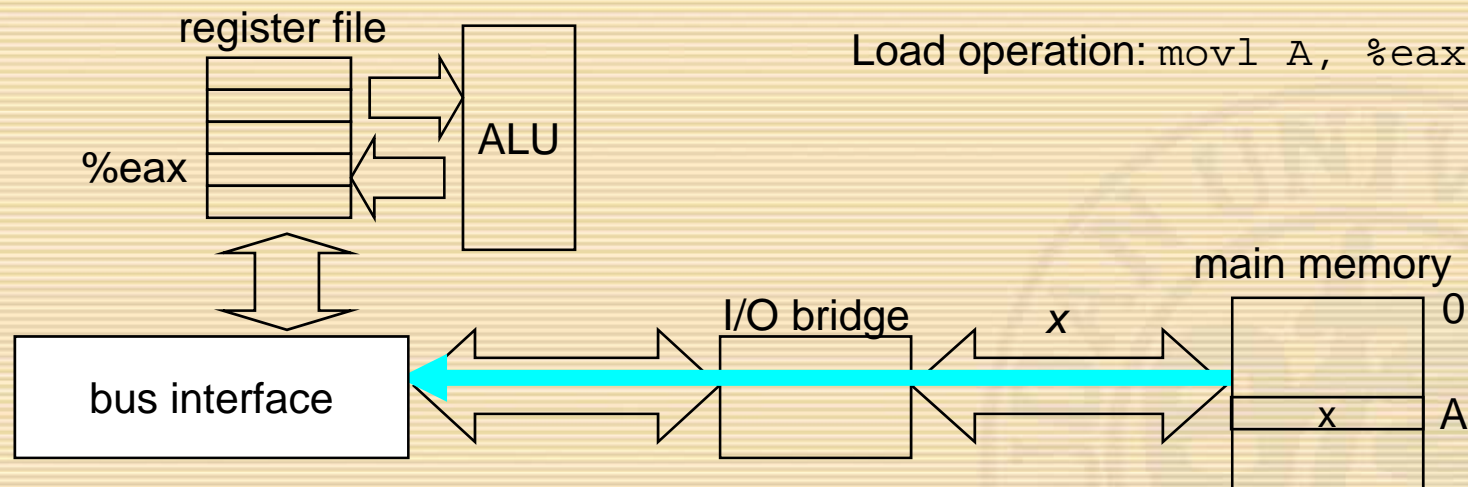
- CPU places address A on the memory bus.





Memory Read Transaction (2)


- Main memory reads *A* from the memory bus, retrieves word *x*, and places it on the bus.





Memory Read Transaction (3)

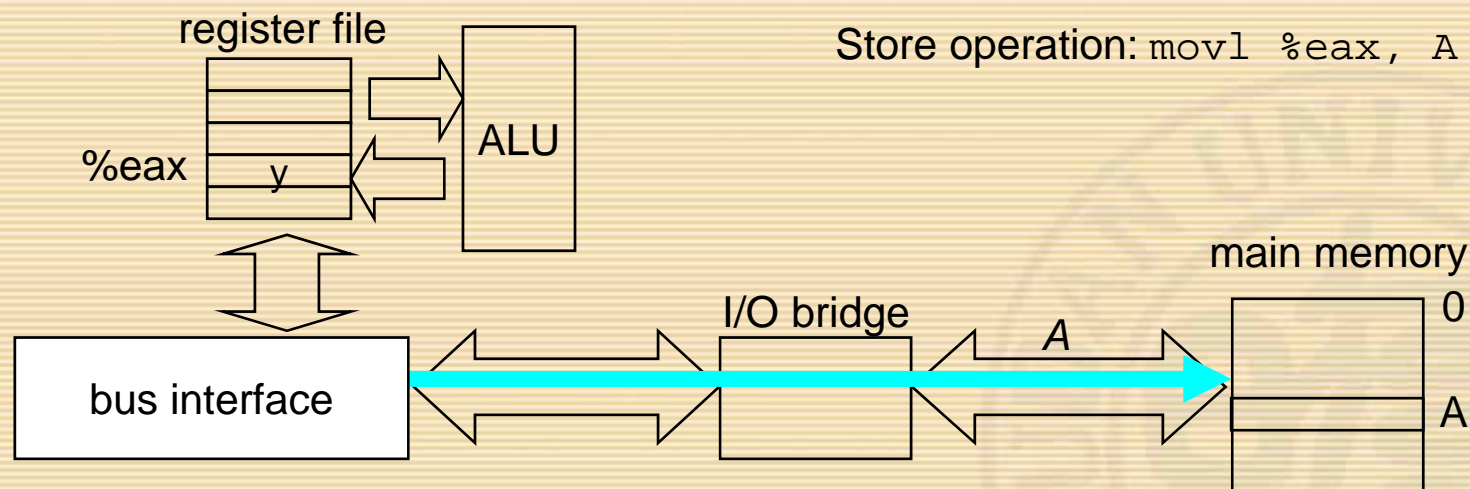
- CPU read word x from the bus and copies it into register `%eax`.





Memory Write Transaction (1)


- CPU places address A on bus. Main memory reads it and waits for the corresponding data word to arrive.





Memory Write Transaction (2)

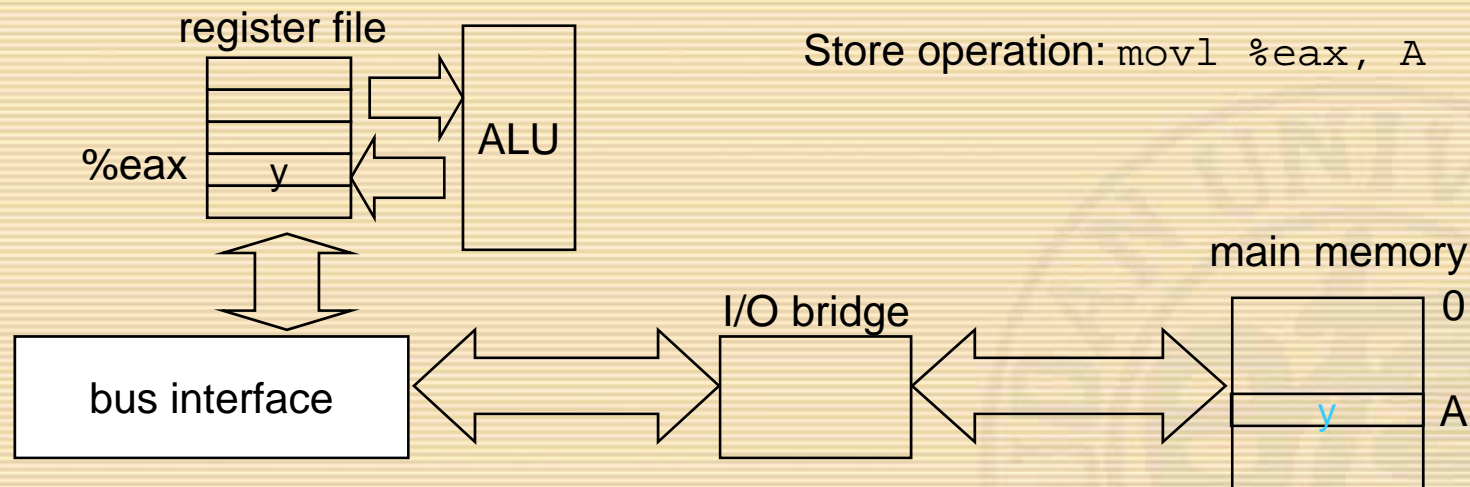
- CPU places data word y on the bus.





Memory Write Transaction (3)


- Main memory read data word y from the bus and stores it at address A .





Disk Geometry


- Disks consist of **platters**, each with two **surfaces**.
- Each surface consists of concentric rings called **tracks**.
- Each track consists of **sectors** separated by **gaps**.





Disk Geometry (Multiple-Platter View)

- Aligned tracks form a cylinder.





Disk Capacity

- **Capacity**: maximum number of bits that can be stored.
 - Vendors express capacity in units of gigabytes (GB), where $1 \text{ GB} = 10^9 \text{ B}$.
- Capacity is determined by these technology factors:
 - **Recording density** (bits/in): number of bits that can be squeezed into a 1 inch segment of a track.
 - **Track density** (tracks/in): number of tracks that can be squeezed into a 1 inch radial segment.
 - **Areal density** (bits/in²): product of recording and track density.
- Modern disks partition tracks into disjoint subsets called **recording zones**
 - Each track in a zone has the same number of sectors, determined by the circumference of innermost track.
 - Each zone has a different number of sectors/track




Computing Disk Capacity

- Capacity = (# bytes/sector) x (avg. # sectors/track) x
- (# tracks/surface) x (# surfaces/platter) x
- (# platters/disk)
- Example:
 - 512 bytes/sector
 - 300 sectors/track (on average)
 - 20,000 tracks/surface
 - 2 surfaces/platter
 - 5 platters/disk
- Capacity = $512 \times 300 \times 20000 \times 2 \times 5B$
- = 30,720,000,000
- = 30.72 GB



Disk Operation (Single-Platter View)

The disk surface spins at a fixed rotational rate




The read/write *head* is attached to the end of the *arm* and flies over the disk surface on a thin cushion of air.

By moving radially, the arm can position the read/write head over any track.



Disk Operation (Multi-Platter View)





Disk Access Time

- Average time to access some target sector approximated by :
 - $T_{\text{access}} = T_{\text{avg seek}} + T_{\text{avg rotation}} + T_{\text{avg transfer}}$
- **Seek time** ($T_{\text{avg seek}}$)
 - Time to position heads over cylinder containing target sector.
 - Typical $T_{\text{avg seek}} = 9 \text{ ms}$
- **Rotational latency** ($T_{\text{avg rotation}}$)
 - Time waiting for first bit of target sector to pass under r/w head.
 - $T_{\text{avg rotation}} = 1/2 \times 1/\text{RPMs} \times 60 \text{ sec}/1 \text{ min}$
- **Transfer time** ($T_{\text{avg transfer}}$)
 - Time to read the bits in the target sector.
 - $T_{\text{avg transfer}} = 1/\text{RPM} \times 1/(\text{avg \# sectors/track}) \times 60 \text{ secs}/1 \text{ min.}$



Disk Access Time Example

- Given:
 - Rotational rate = 7,200 RPM
 - Average seek time = 9 ms.
 - Avg # sectors/track = 400.
- Derived:
 - $T_{\text{avg rotation}} = 1/2 \times (60 \text{ secs}/7200 \text{ RPM}) \times 1000 \text{ ms/sec} = 4 \text{ ms}.$
 - $T_{\text{avg transfer}} = 60/7200 \text{ RPM} \times 1/400 \text{ secs/track} \times 1000 \text{ ms/sec} = 0.02 \text{ ms}$
 - $T_{\text{access}} = 9 \text{ ms} + 4 \text{ ms} + 0.02 \text{ ms}$
- Important points:
 - Access time dominated by seek time and rotational latency.
 - First bit in a sector is the most expensive, the rest are free.
 - SRAM access time is about 4 ns/doubleword, DRAM about 60 ns
 - Disk is about 40,000 times slower than SRAM,
 - 2,500 times slower than DRAM.

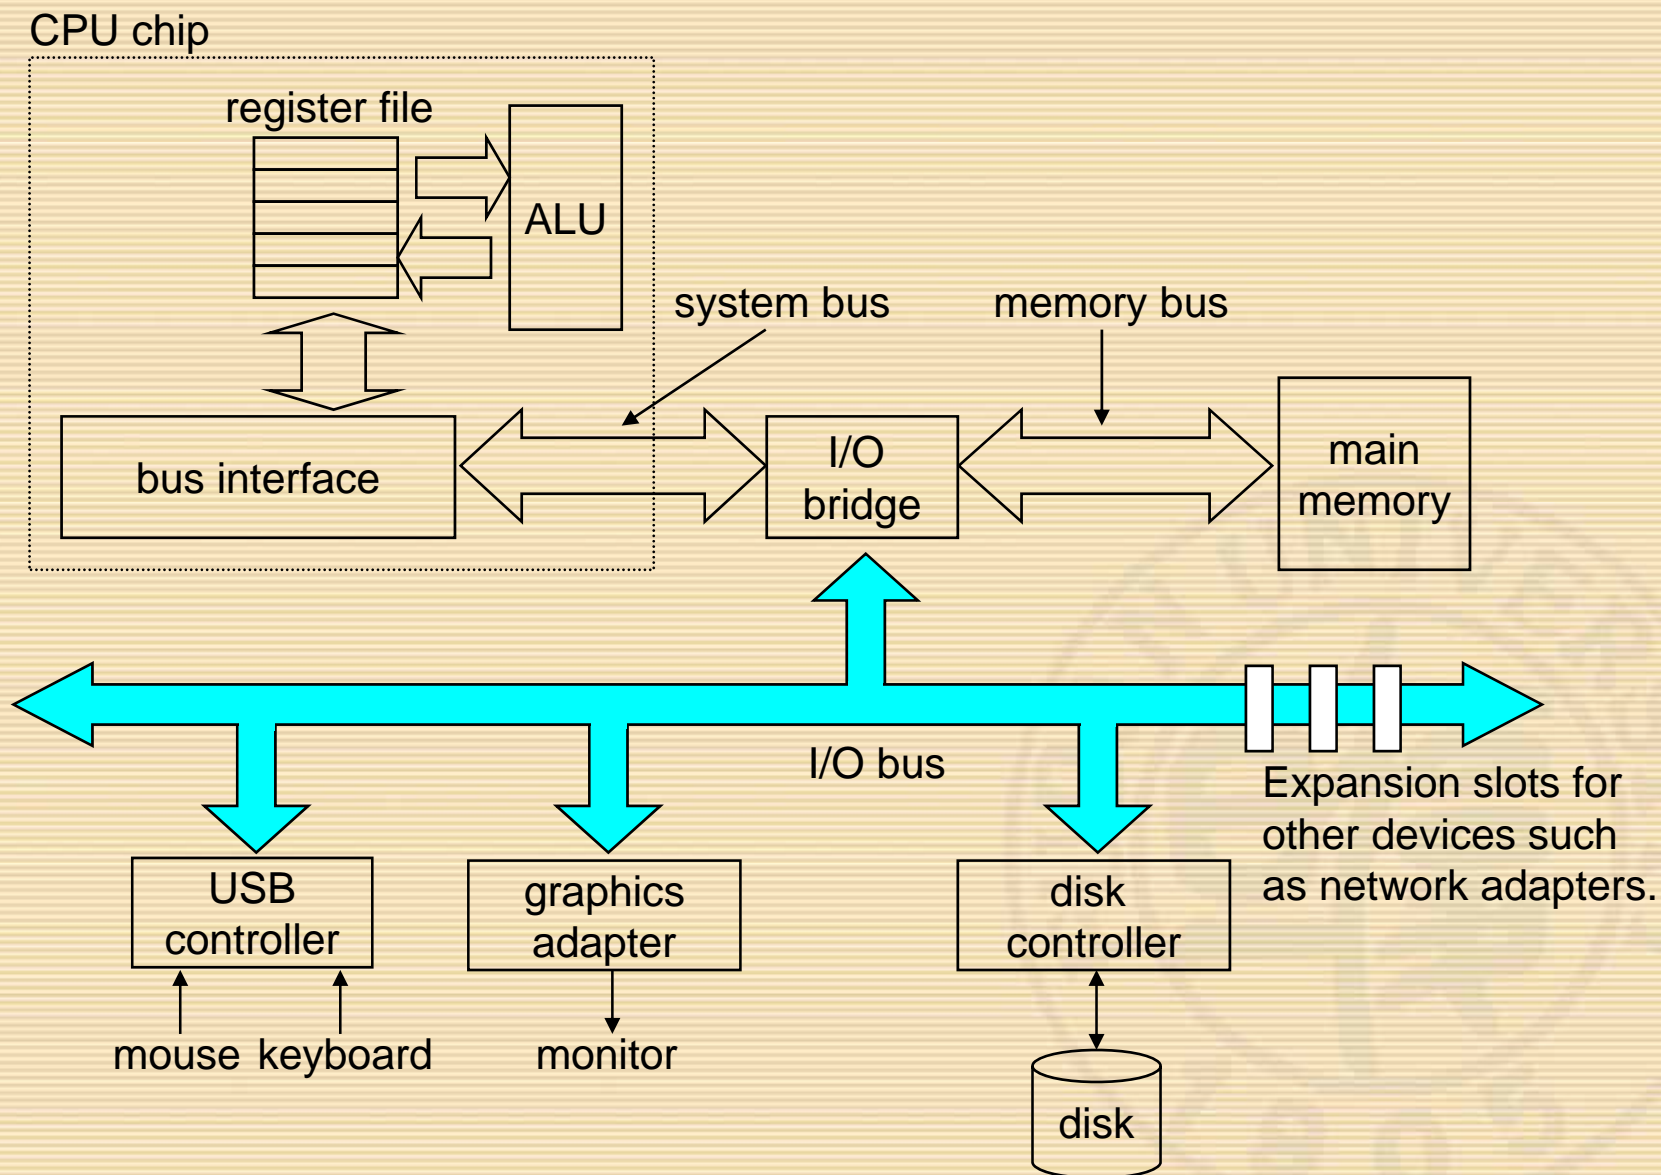


Logical Disk Blocks

- Modern disks present a simpler abstract view of the complex sector geometry:
 - The set of available sectors is modeled as a sequence of b-sized **logical blocks** (0, 1, 2, ...)
- Mapping between logical blocks and actual (physical) sectors
 - Maintained by hardware/firmware device called disk controller.
 - Converts requests for logical blocks into (surface, track, sector) triples.
- Allows controller to set aside spare cylinders for each zone. P471
 - Accounts for the difference in “**formatted capacity**” and “**maximum capacity**”.

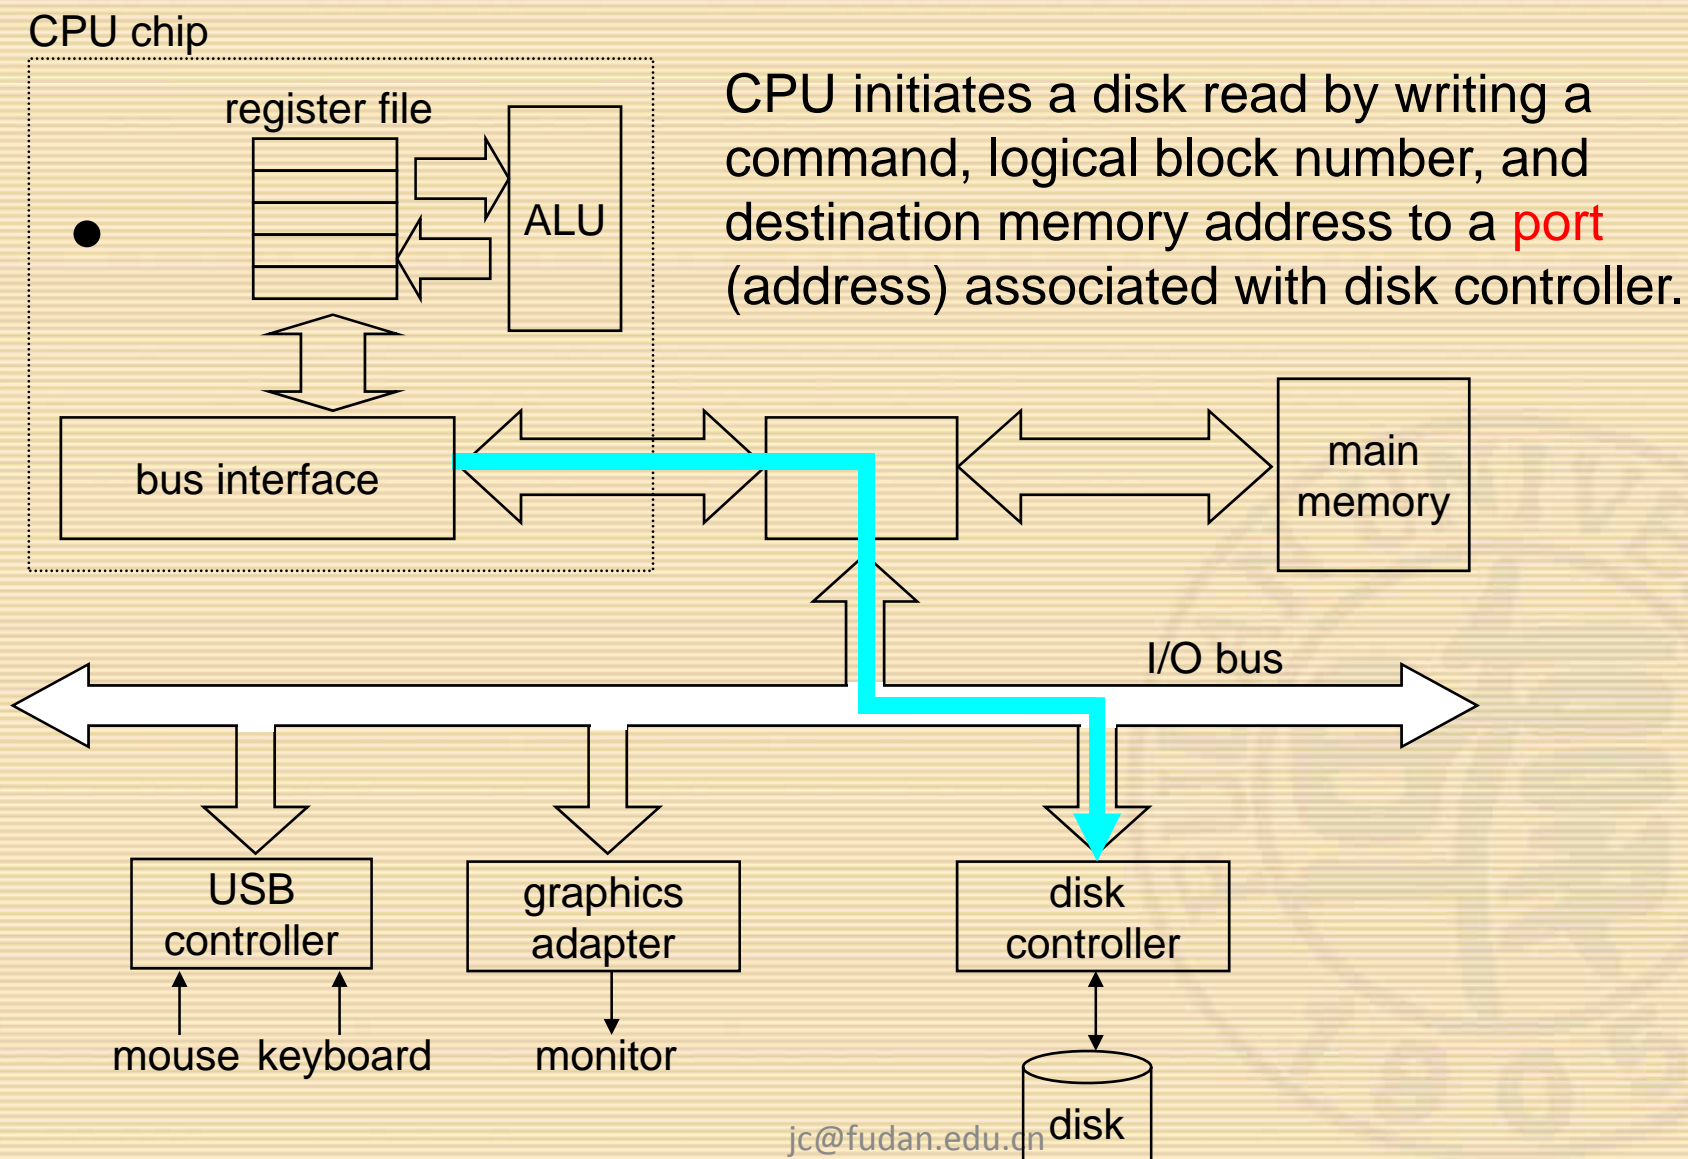


I/O Bus



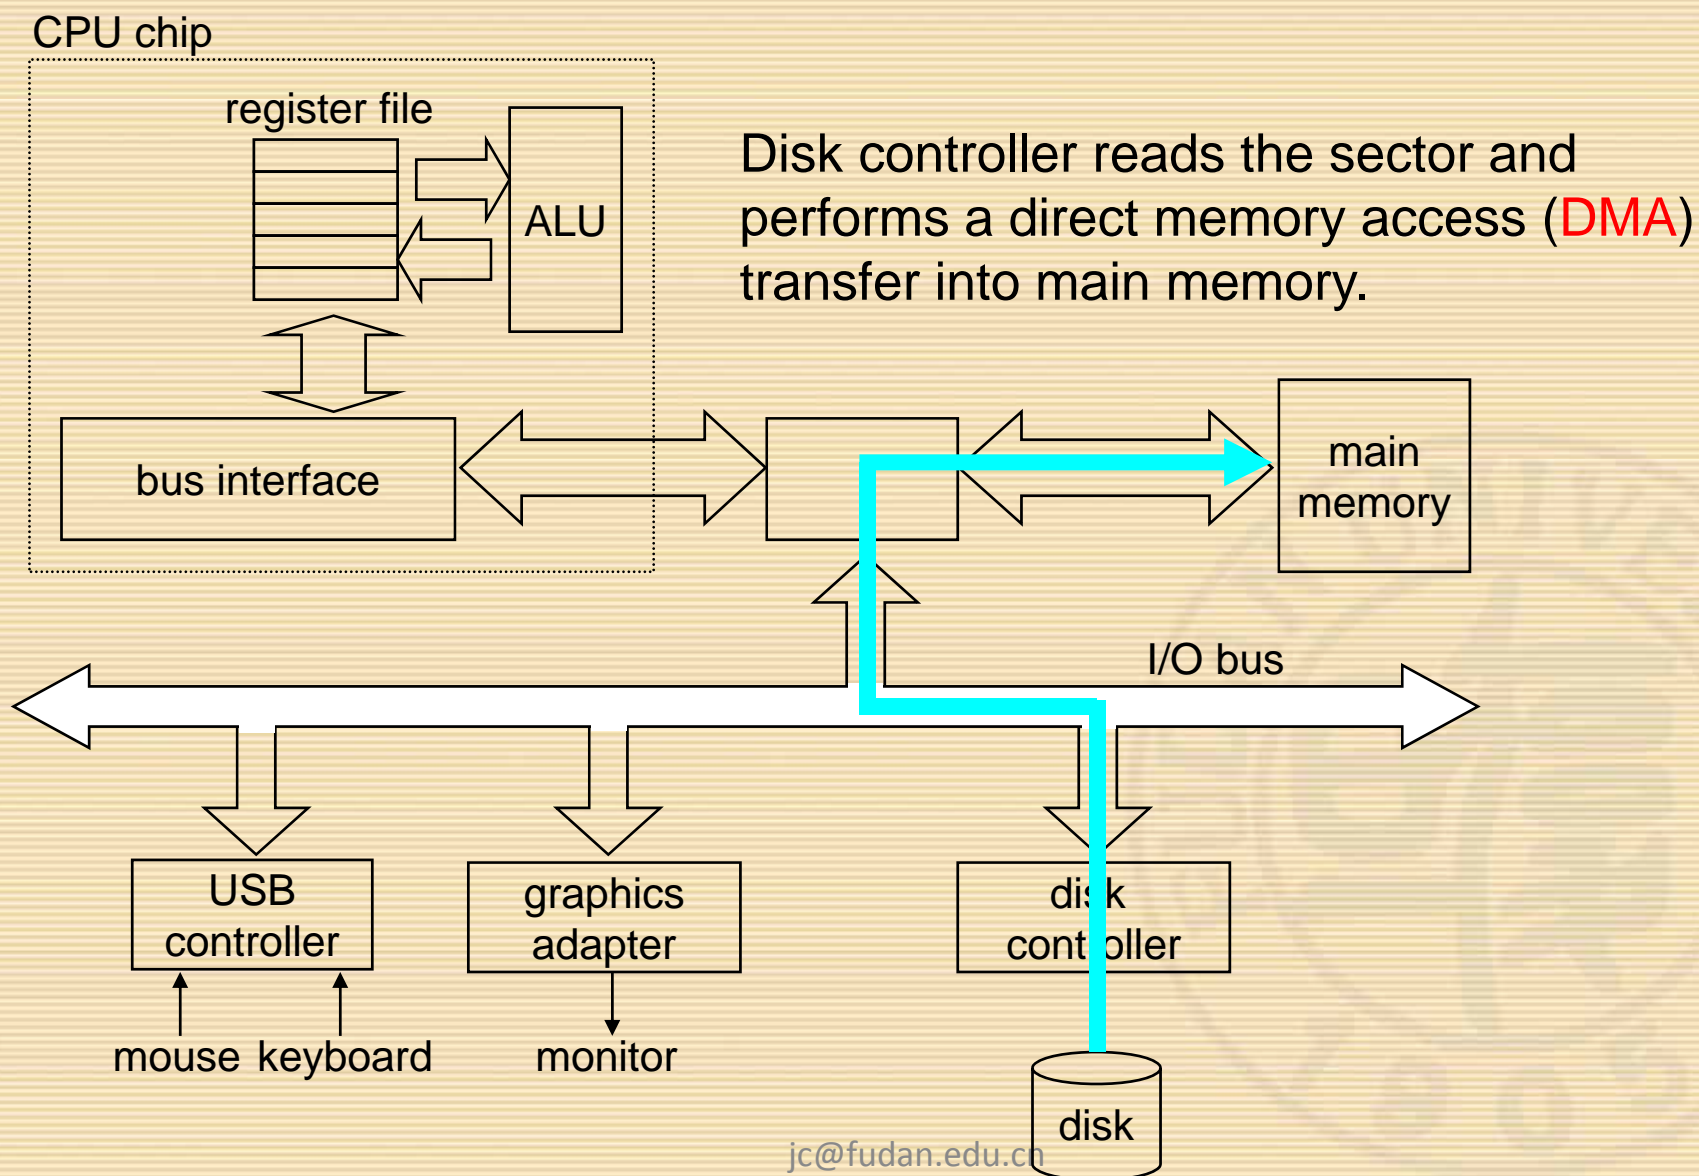


Reading a Disk Sector (1)



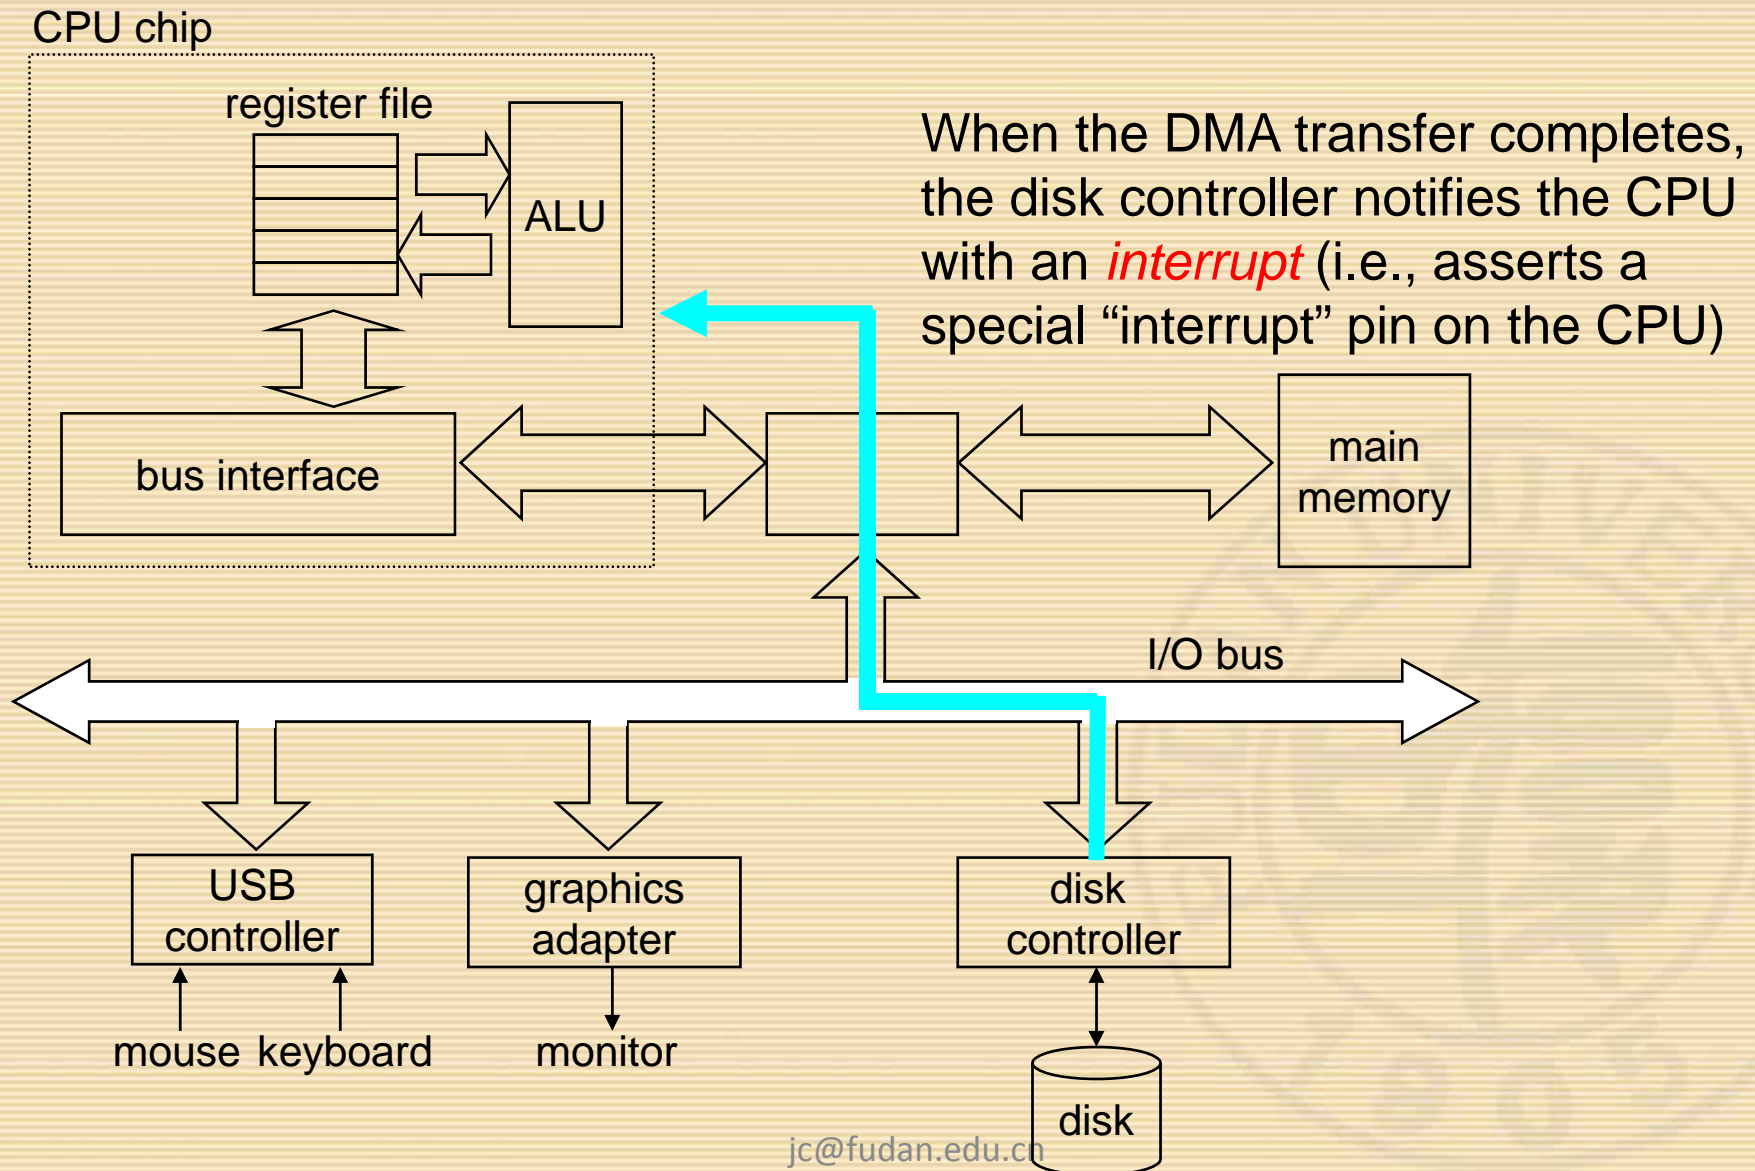


Reading a Disk Sector (2)





Reading a Disk Sector (3)





Storage Trends

SRAM

metric	1980	1985	1990	1995	2000	2000:1980
\$/MB	19,200	2,900	320	256	100	190
access (ns)	300	150	35	15	2	100

DRAM

metric	1980	1985	1990	1995	2000	2000:1980
\$/MB	8,000	880	100	30	1	8,000
access (ns)	375	200	100	70	60	6
typical size(MB)	0.064	0.256	4	16	64	1,000

Disk

metric	1980	1985	1990	1995	2000	2000:1980
\$/MB	500	100	8	0.30	0.05	10,000
access (ms)	87	75	28	10	8	11
typical size(MB)	1	10	160	1,000	9,000	9,000

(Culled from back issues of Byte and PC Magazine)




CPU Clock Rates

	1980	1985	1990	1995	2000	2000:1980
processor 8080	286	386	Pent	P-III		
clock rate(MHz)	1	6	20	150	750	750
cycle time(ns)	1,000	166	50	6	1.6	750



The CPU-Memory Gap

- The increasing gap between DRAM, disk, and CPU speeds.





Locality

- Principle of Locality:
 - Programs tend to reuse data and instructions near those they have used recently, or that were recently referenced themselves.
 - **Temporal locality**: Recently referenced items are likely to be referenced in the near future.
 - **Spatial locality**: Items with nearby addresses tend to be referenced close together in time.

Locality Example:

- Data
 - Reference array elements in succession (stride-1 reference pattern): **Spatial locality**
 - Reference `sum` each iteration: **Temporal locality**
- Instructions
 - Reference instructions in sequence: **Spatial locality**
 - Cycle through loop repeatedly: **Temporal locality**

```
sum = 0;
for (i = 0; i < n; i++)
    sum += a[i];
return sum;
```



Locality Example

- **Claim:** Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.
- **Question:** Does this function have good locality?

```
int sumarrayrows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```



Locality Example

- **Question:** Does this function have good locality?

```
int sumarraycols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum
}
```



Locality Example

- **Question:** Can you permute the loops so that the function scans the 3-d array `a[]` with a stride-1 reference pattern (and thus has good spatial locality)?

```
int sumarray3d(int a[N][N][N])
{
    int i, j, k, sum = 0;

    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < N; k++)
                sum += a[k][i][j];
    return sum;
}
```

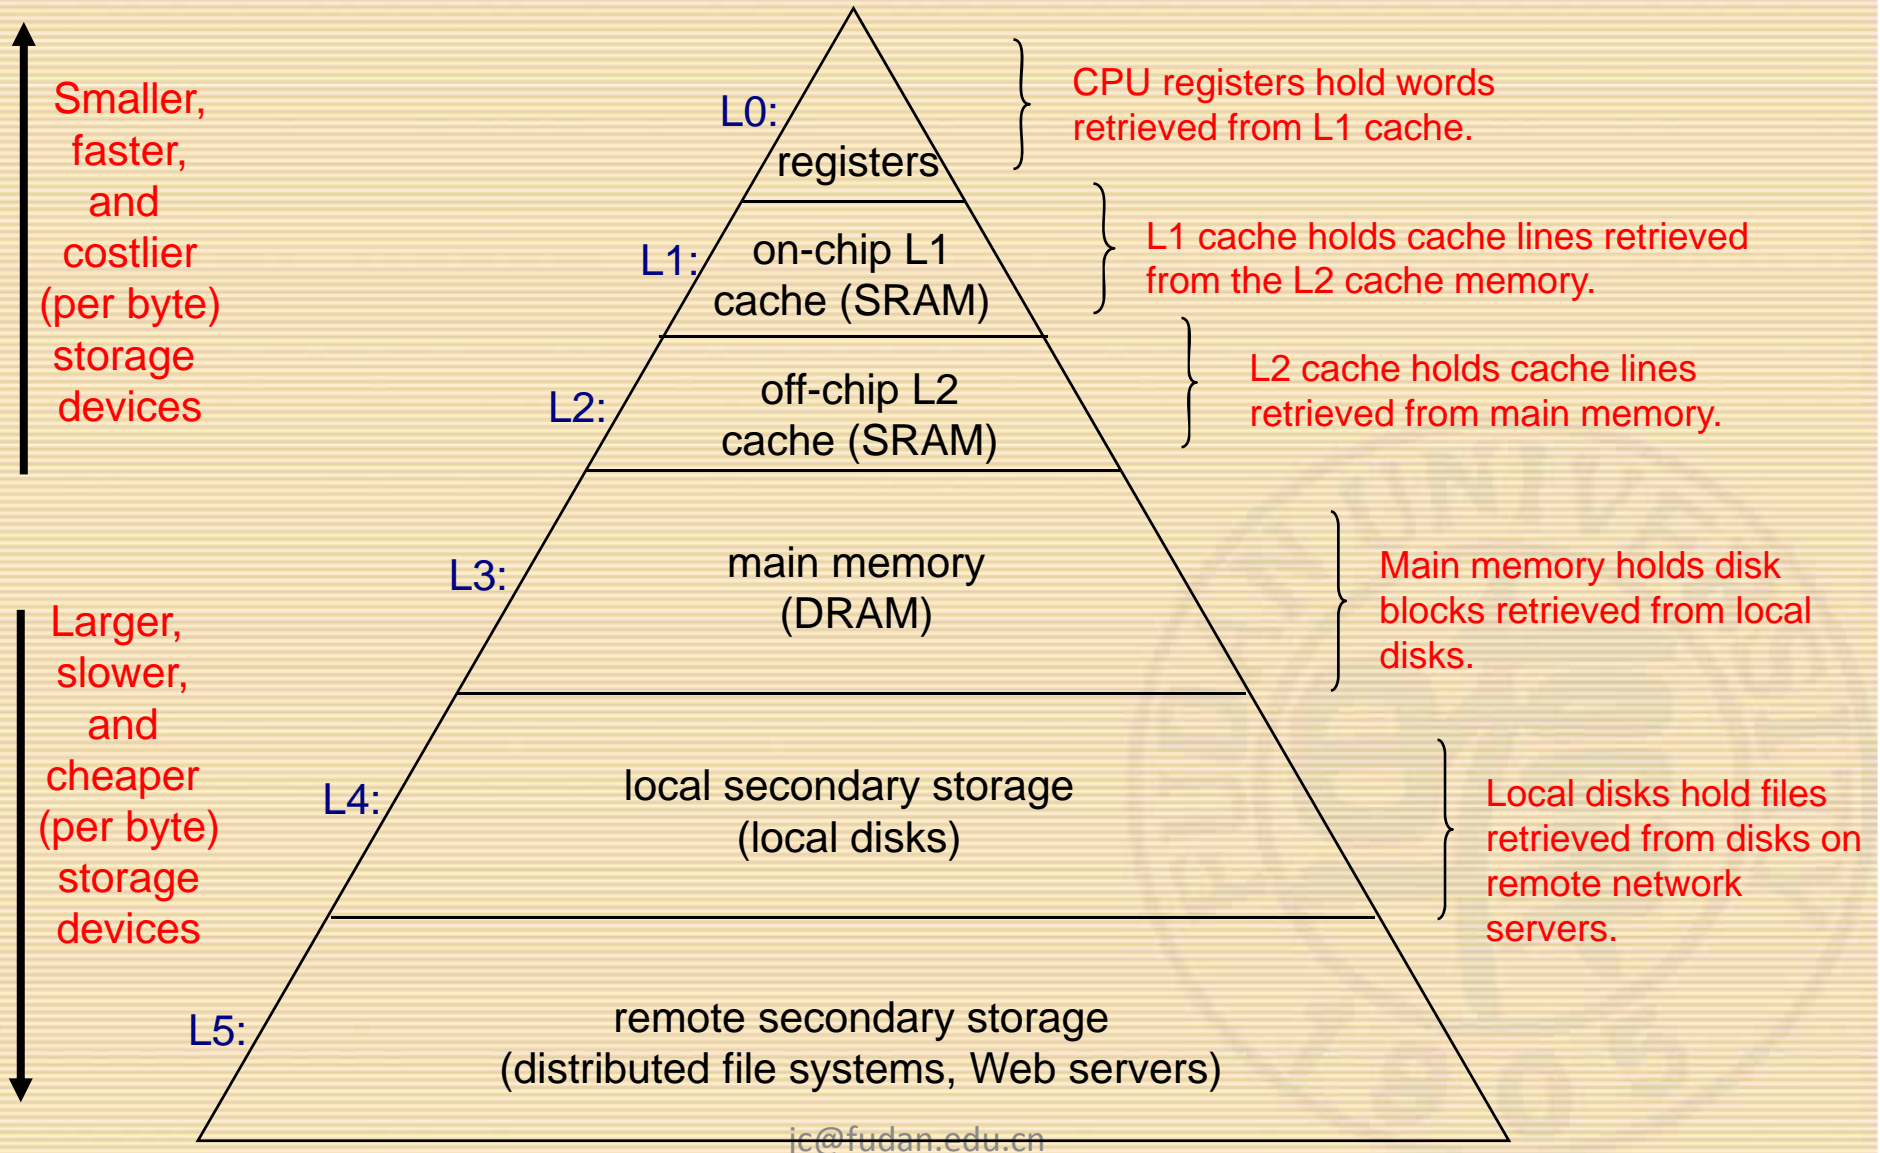


Memory Hierarchies

- Some fundamental and enduring properties of hardware and software:
 - Fast storage technologies cost more per byte and have less capacity.
 - The gap between CPU and main memory speed is widening.
 - Well-written programs tend to exhibit good locality.
- These fundamental properties complement each other beautifully.
- They suggest an approach for organizing memory and storage systems known as a **memory hierarchy**.



An Example Memory Hierarchy






Caches


- **Cache:** A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.
- Fundamental idea of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$.
- Why do memory hierarchies work?
 - Programs tend to access the data at level k more often than they access the data at level $k+1$.
 - Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per bit.
 - **Net effect:** A large pool of memory that costs as much as the cheap storage near the bottom, but that serves data to programs at the rate of the fast storage near the top.



Caching in a Memory Hierarchy




Smaller, faster, more expensive device at level k caches a subset of the blocks from level k+1



Data is copied between levels in block-sized transfer units

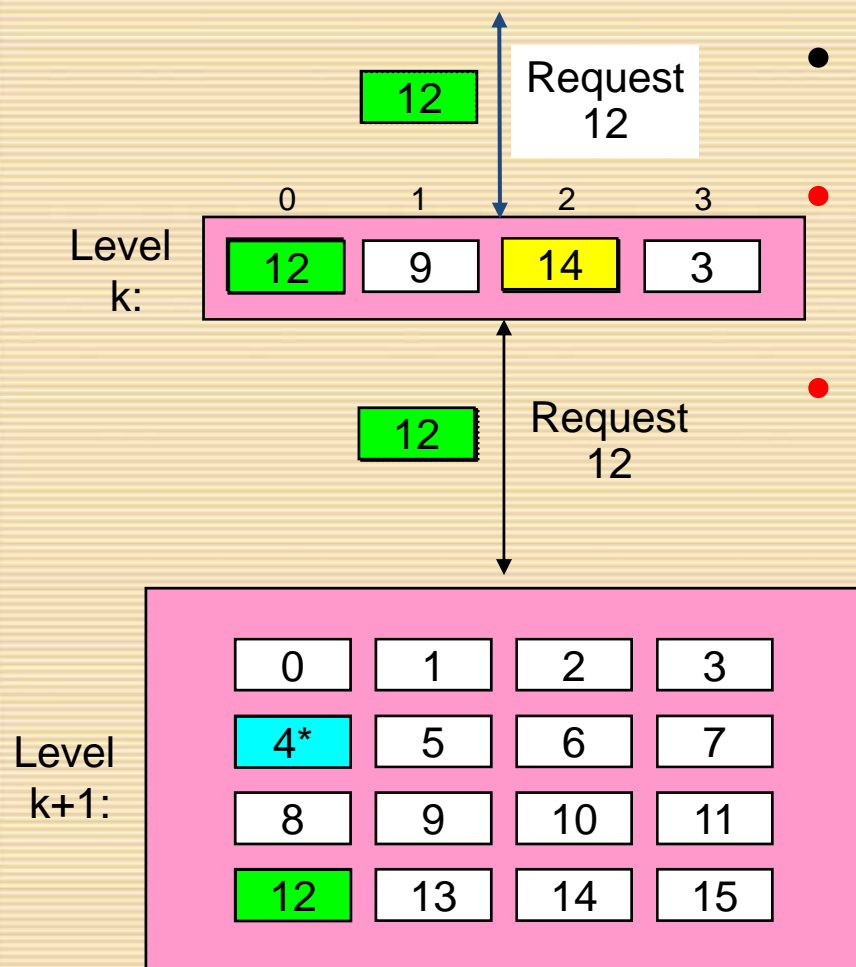
Level k+1:



Larger, slower, cheaper storage device at level k+1 is partitioned into blocks.



General Caching Concepts



- Program needs object d, which is stored in some block b.
- **Cache hit**
 - Program finds b in the cache at level k. E.g., block 14.
- **Cache miss**
 - b is not at level k, so level k cache must fetch it from level k+1. E.g., block 12.
 - If level k cache is full, then some current block must be replaced (evicted). Which one is the “victim”?
 - **Placement policy:** where can the new block go? E.g., $b \bmod 4$
 - **Replacement policy:** which block should be evicted? E.g., LRU



General Caching Concepts

- Types of cache misses:
 - Cold (compulsary) miss
 - Cold misses occur because the cache is empty.
 - Conflict miss
 - Most caches limit blocks at level $k+1$ to a small subset (sometimes a singleton) of the block positions at level k .
 - E.g. Block i at level $k+1$ must be placed in block $(i \bmod 4)$ at level k .
 - Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block.
 - E.g. Referencing blocks 0, 8, 0, 8, 0, 8, ... would miss every time.
 - Capacity miss
 - Occurs when the set of active cache blocks (working set) is larger than the cache.



Examples of Caching in the Hierarchy

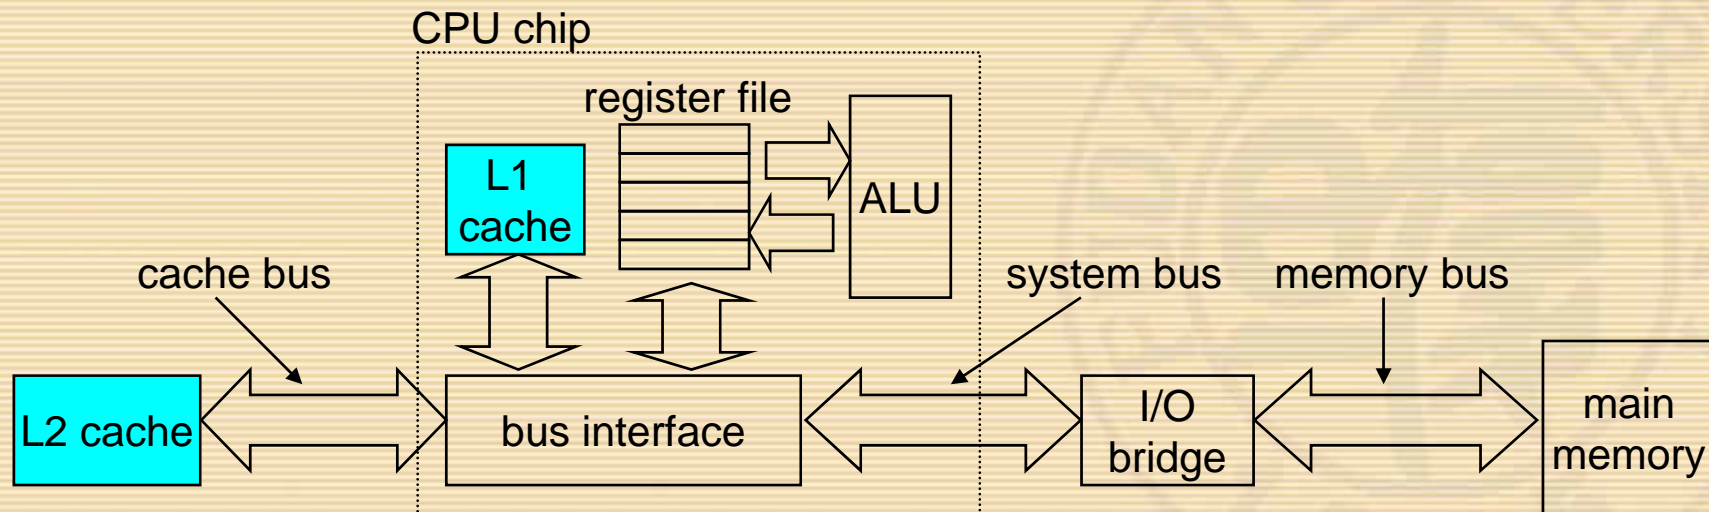
Cache Type	What Cached	Where Cached	Latency (cycles)	Managed By
Registers	4-byte word	CPU registers	0	Compiler
TLB	Address translations	On-Chip TLB	0	Hardware
L1 cache	32-byte block	On-Chip L1	1	Hardware
L2 cache	32-byte block	Off-Chip L2	10	Hardware
Virtual Memory	4-KB page	Main memory	100	Hardware+ OS
Buffer cache	Parts of files	Main memory	100	OS
Network buffer cache	Parts of files	Local disk	10,000,000	AFS/NFS client
Browser cache	Web pages	Local disk	10,000,000	Web browser
Web cache	Web pages	Remote server disks	1,000,000,000	Web proxy server

TLB: Translation Lookaside Buffer, maps virtual addresses to physical ones



Cache Memories

- Cache memories are small, fast SRAM-based memories managed automatically in hardware.
 - Hold frequently accessed blocks of main memory
- CPU looks first for data in L1, then in L2, then in main memory.
- Typical bus structure:

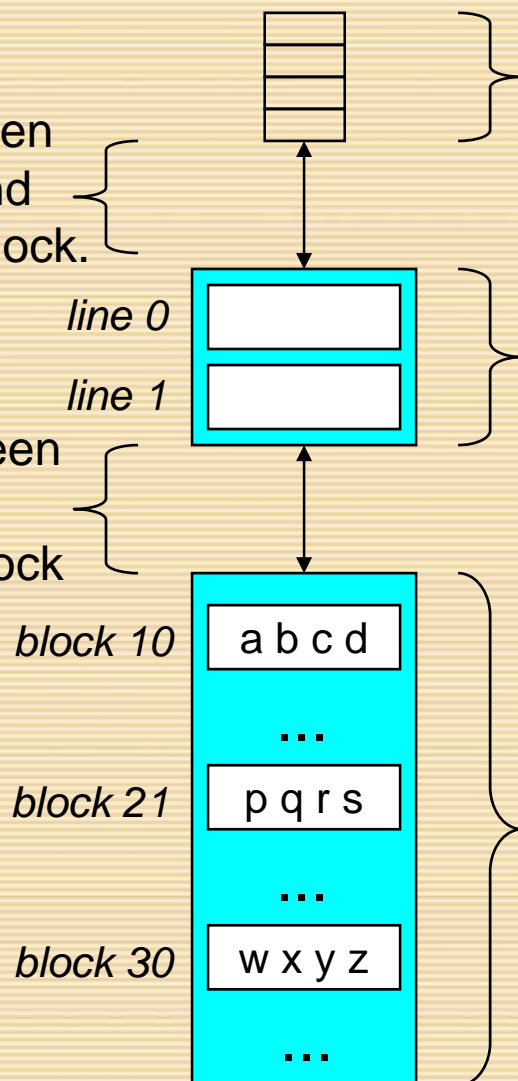




Inserting an L1 Cache Between the CPU and Main Memory

The transfer unit between the CPU register file and the cache is a 4-byte block.

The transfer unit between the cache and main memory is a 4-word block (16 bytes).



The tiny, very fast CPU register file has room for four 4-byte words.

The small fast L1 cache has room for two 4-word blocks.

The big slow main memory has room for many 4-word blocks.



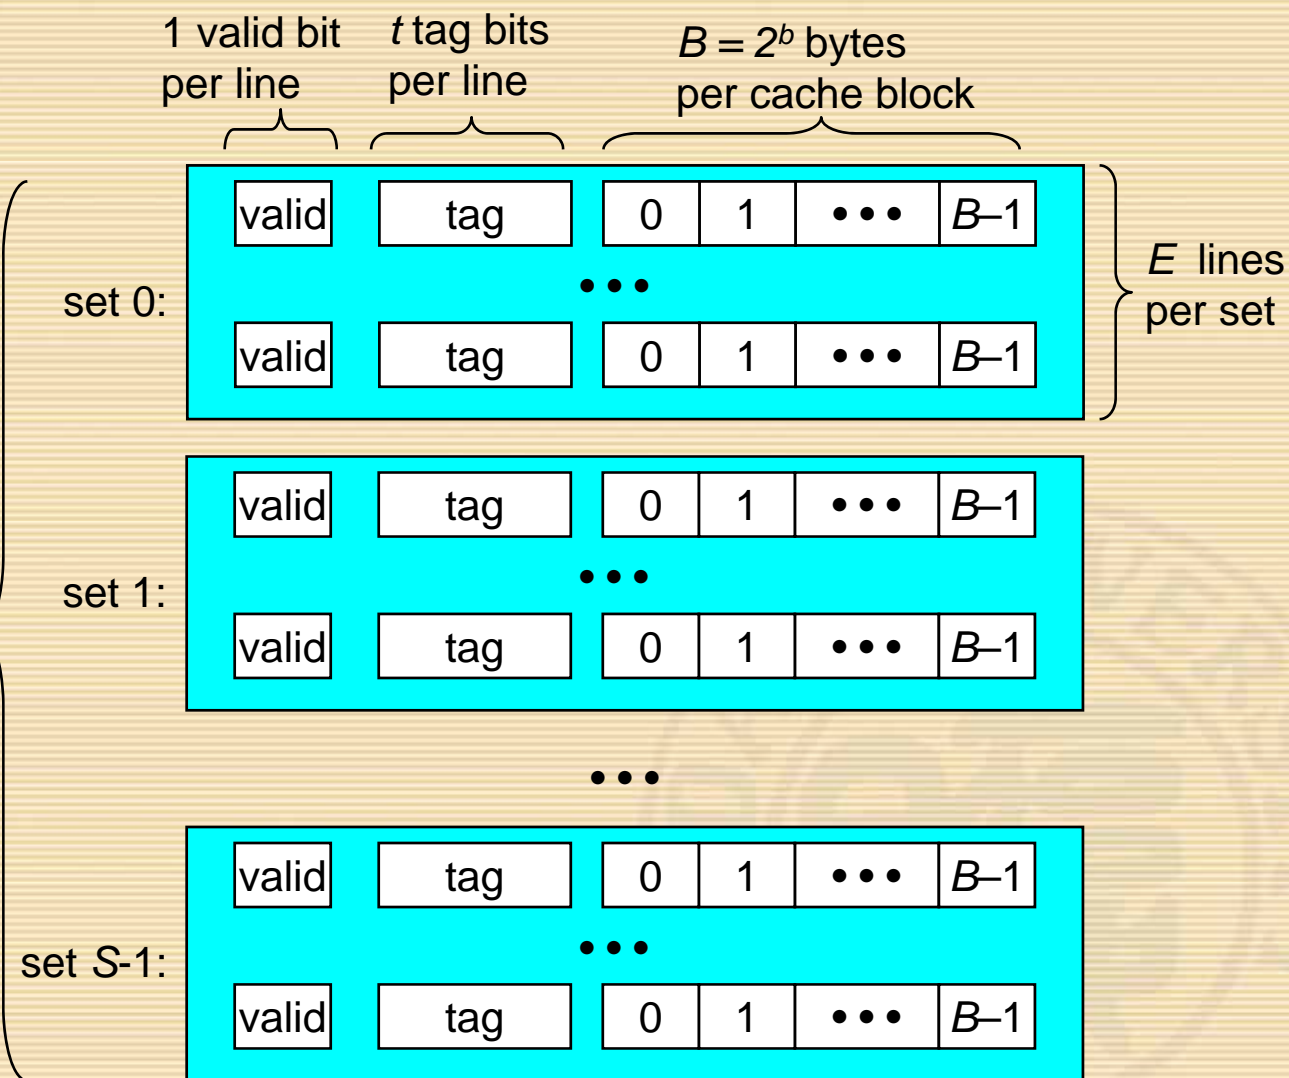
General Org of a Cache Memory

Cache is an array of sets.

Each set contains one or more lines.

Each line holds a block of data.

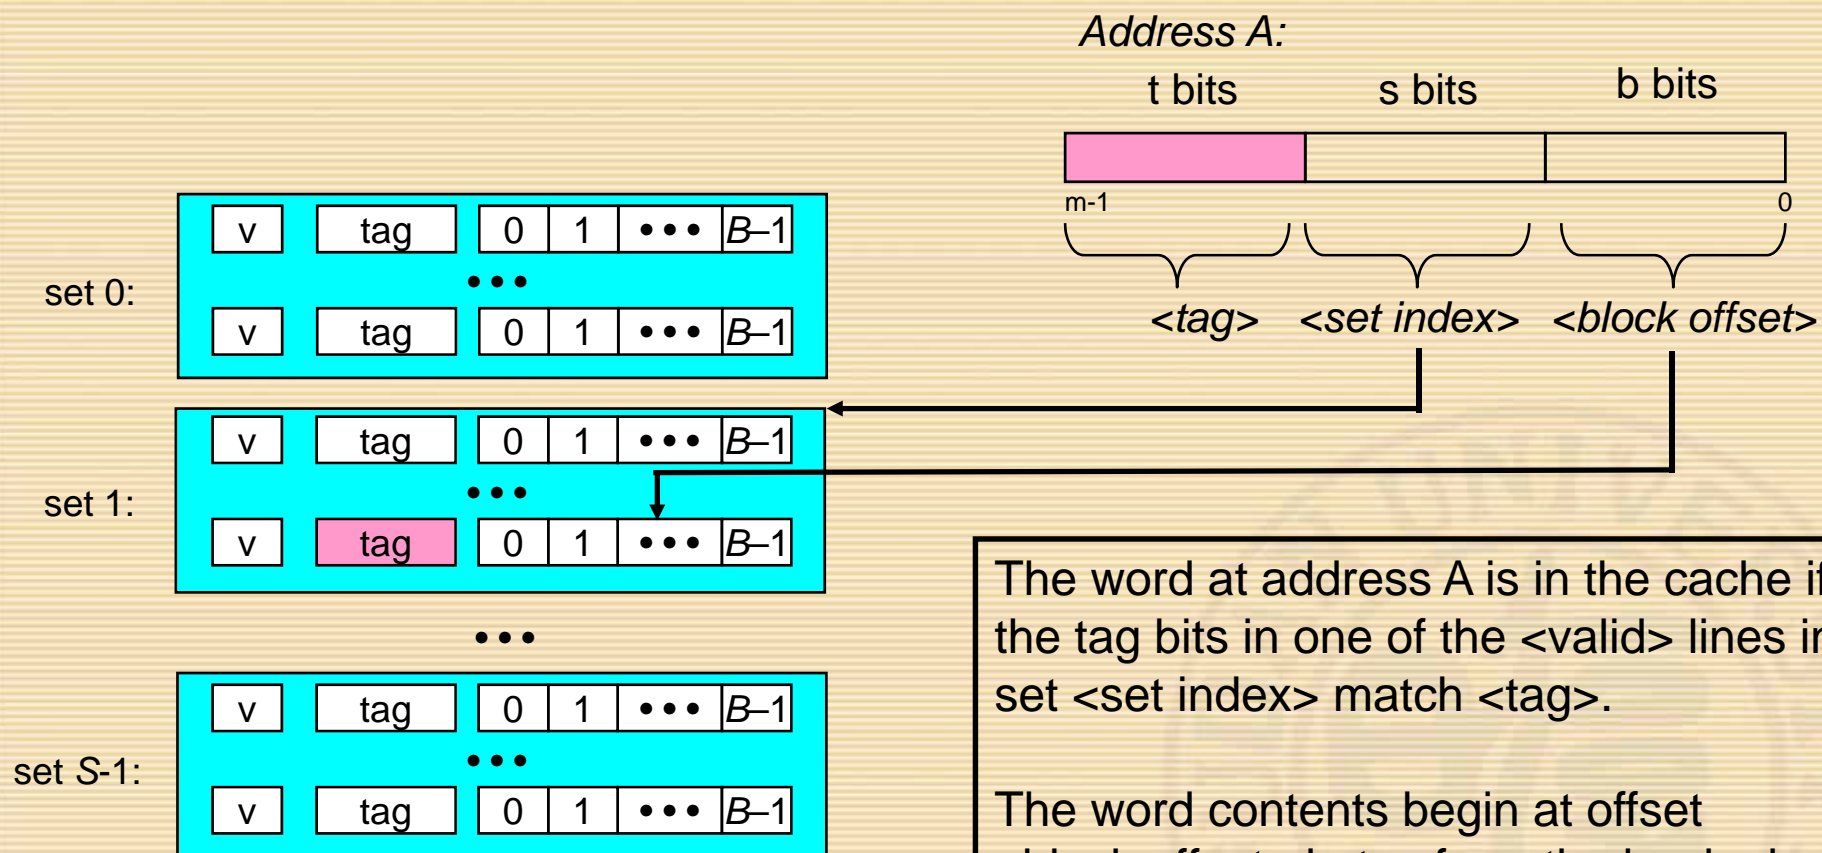
$$S = 2^s \text{ sets}$$



$$\text{Cache size: } C = B \times E \times S \text{ data bytes}$$



Addressing Caches




The word at address A is in the cache if the tag bits in one of the <valid> lines in set <set index> match <tag>.

The word contents begin at offset <block offset> bytes from the beginning of the block.



Direct-Mapped Cache

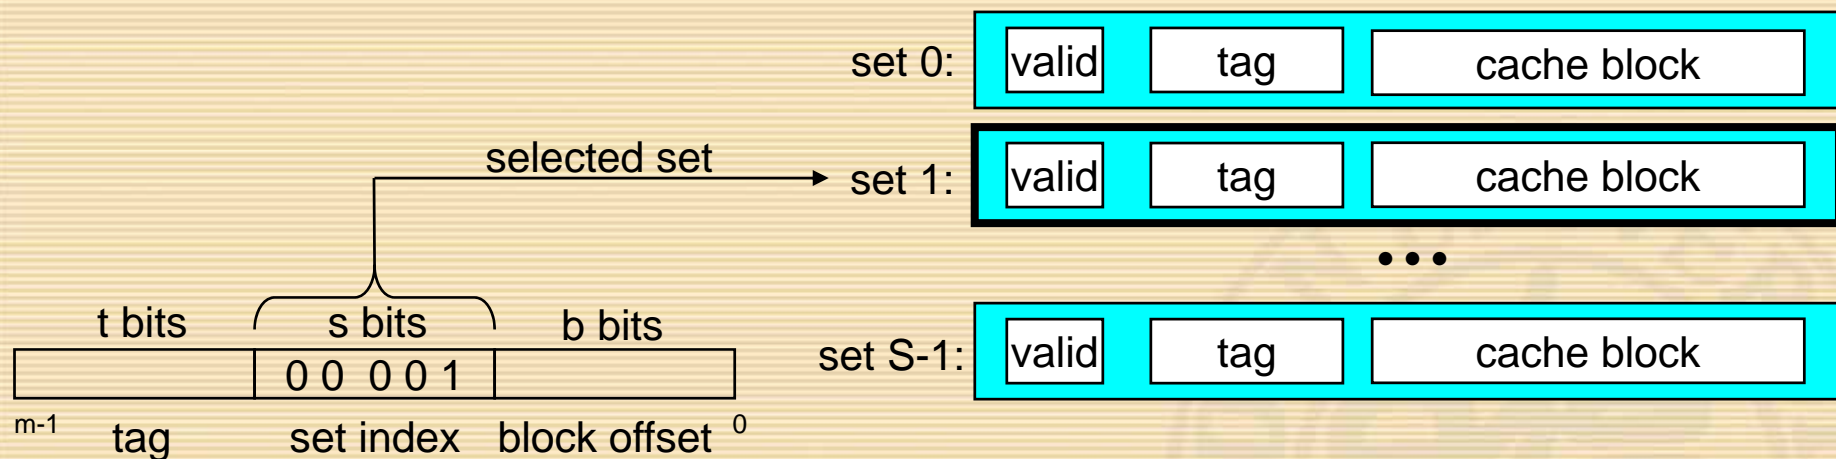
- Simplest kind of cache
- Characterized by exactly one line per set.





Accessing Direct-Mapped Caches


- Set selection
 - Use the set index bits to determine the set of interest.





Accessing Direct-Mapped Caches


- Line matching and word selection
 - Line matching**: Find a valid line in the selected set with a matching tag
 - Word selection**: Then extract the word





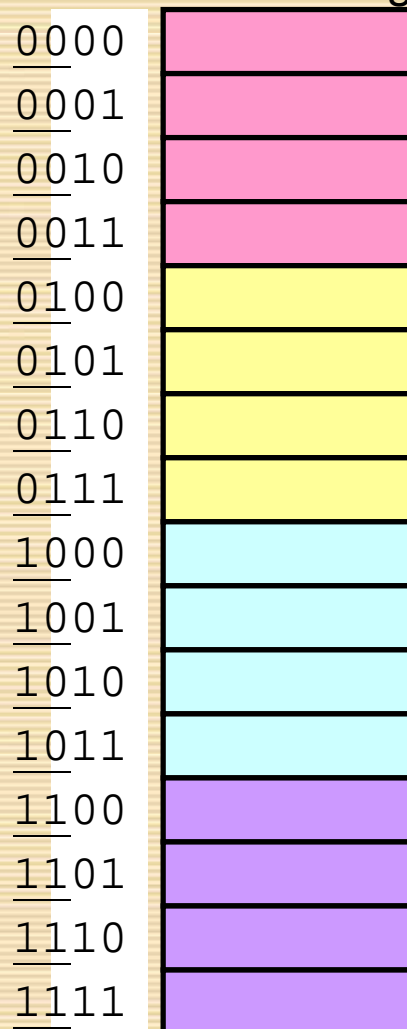
Why Use Middle Bits as Set Index?

4-line Cache




- High-Order Bit Indexing
 - Adjacent memory lines would map to same cache entry
 - Poor use of spatial locality
- Middle-Order Bit Indexing
 - Consecutive memory lines map to different cache lines
 - Can hold C-byte region of address space in cache at one time

High-Order
Bit Indexing



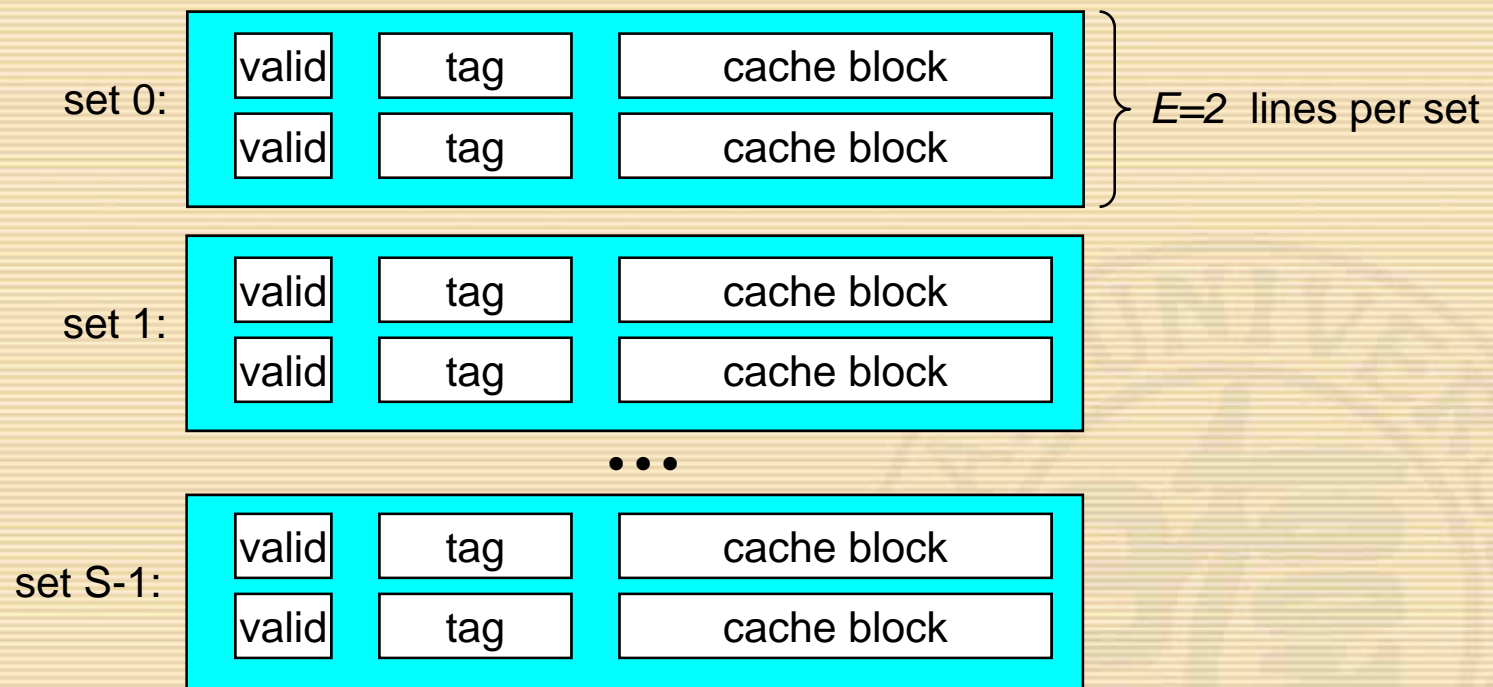
Middle-Order
Bit Indexing





Set Associative Caches

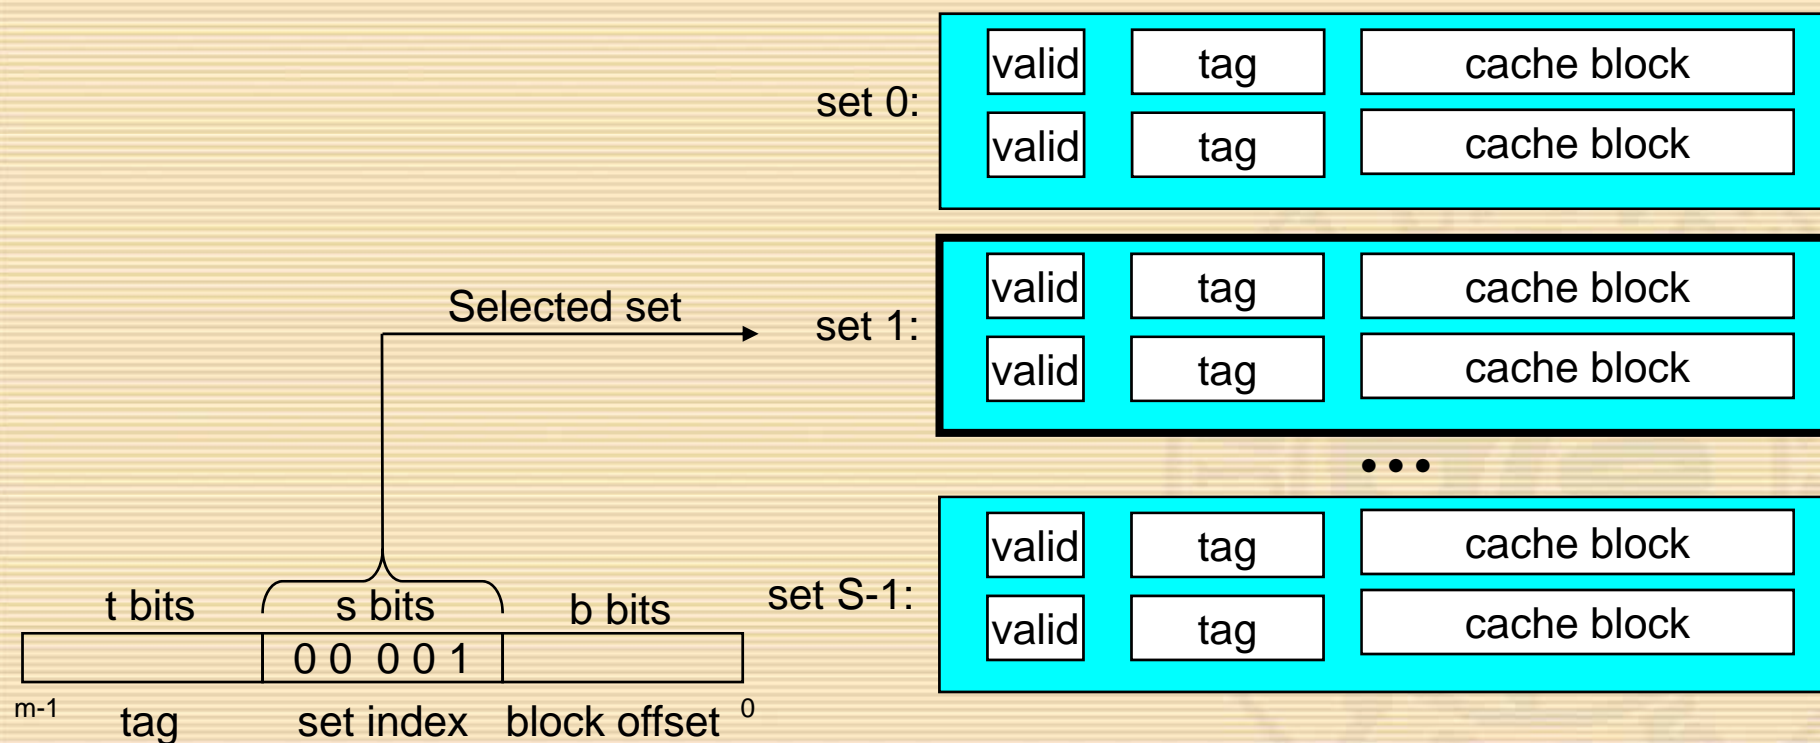
- Characterized by more than one line per set





Accessing Set Associative Caches

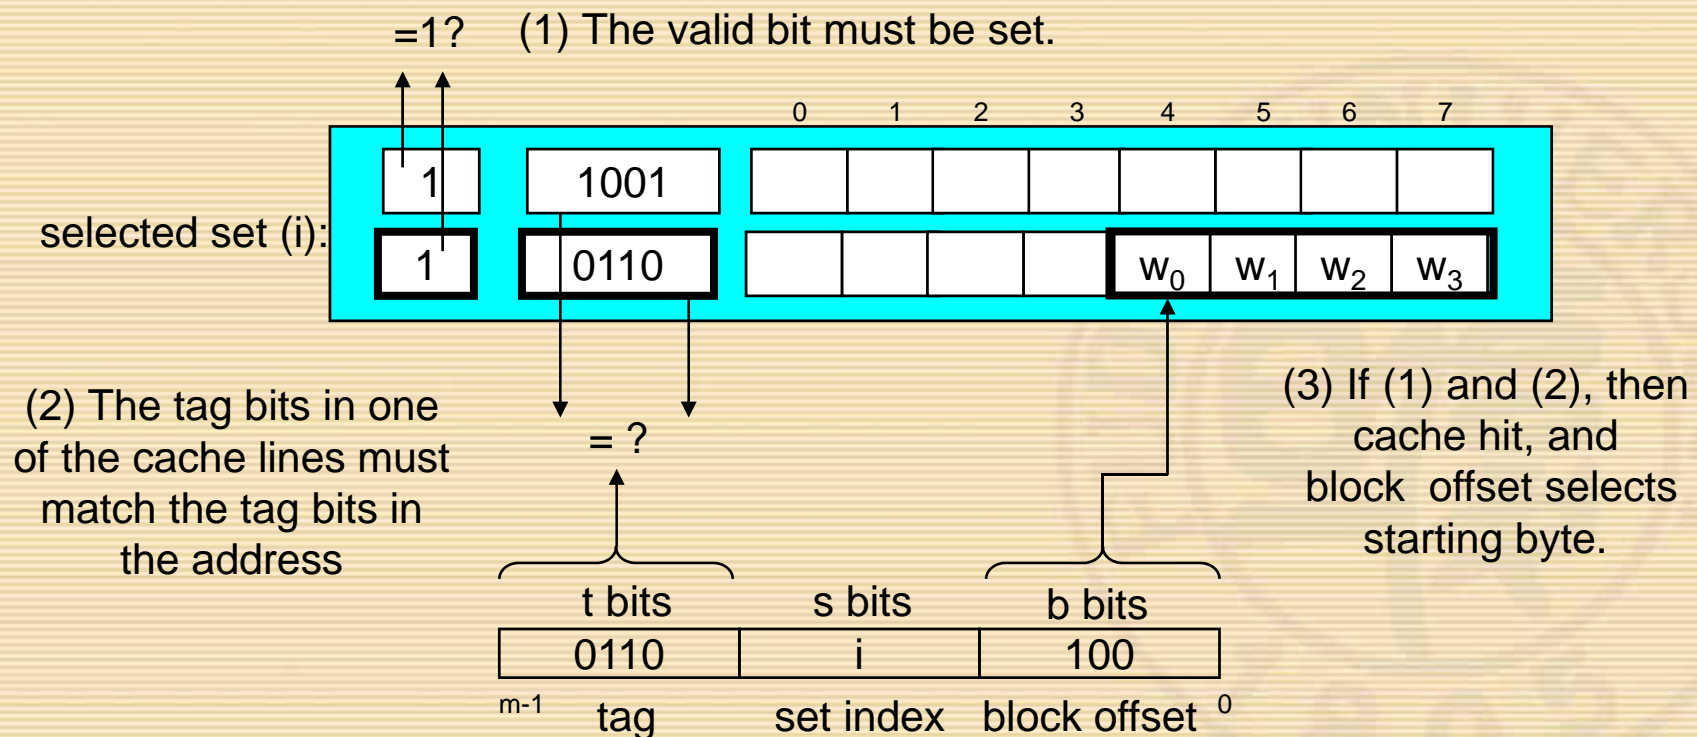
- Set selection
 - identical to direct-mapped cache





Accessing Set Associative Caches

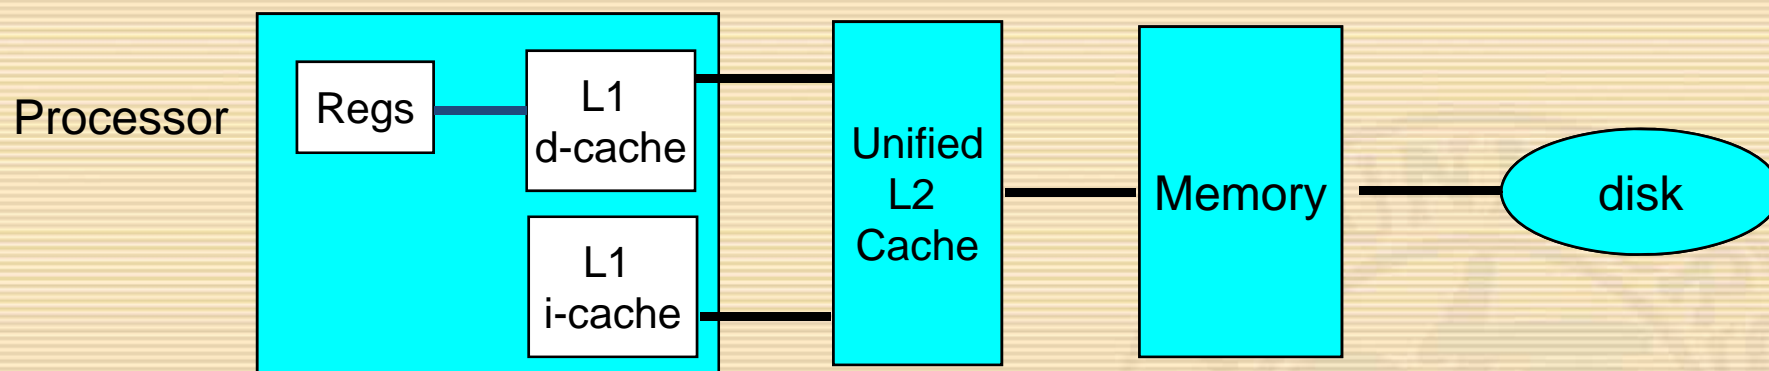
- Line matching and word selection
 - must compare the tag in each valid line in the selected set.





Multi-Level Caches

- Options: separate **data** and **instruction caches**, or a **unified cache**



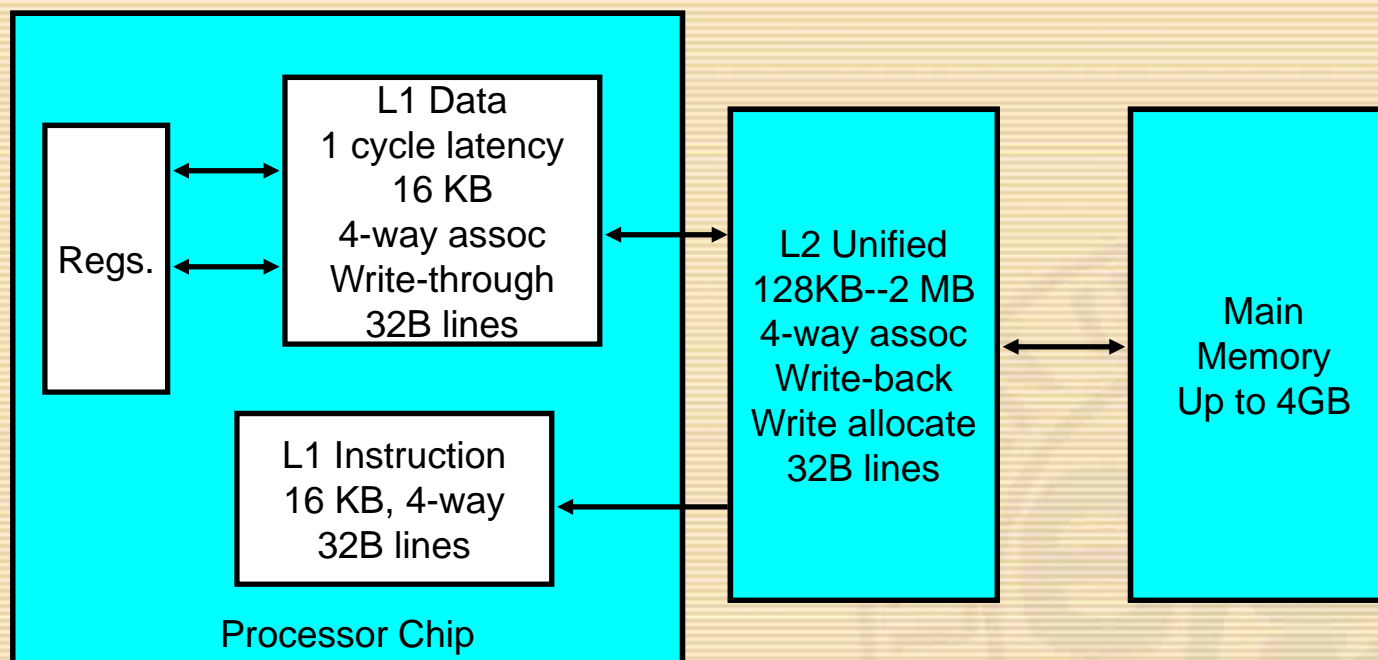
size:	200 B	8-64 KB	1-4MB SRAM	128 MB DRAM	30 GB
speed:	3 ns	3 ns	6 ns	60 ns	8 ms
\$/Mbyte:			\$100/MB	\$1.50/MB	\$0.05/MB
line size:	8 B	32 B	32 B	8 KB	

larger, slower, cheaper





Intel Pentium Cache Hierarchy





Cache Performance Metrics

- Miss Rate
 - Fraction of memory references not found in cache (misses/references)
 - Typical numbers:
 - 3-10% for L1
 - can be quite small (e.g., < 1%) for L2, depending on size, etc.
- Hit Time
 - Time to deliver a line in the cache to the processor (includes time to determine whether the line is in the cache)
 - Typical numbers:
 - 1 clock cycle for L1
 - 3-8 clock cycles for L2
- Miss Penalty
 - Additional time required because of a miss
 - Typically 25-100 cycles for main memory



Writing Cache Friendly Code

- Repeated references to variables are good (temporal locality)
- Stride-1 reference patterns are good (spatial locality)
- Examples:
 - cold cache, 4-byte words, 4-word cache blocks

```
int sumarrayrows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = $1/4 = 25\%$

```
int sumarraycols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = 100%

ScienceMark Membench				
	AMD Opteron 265	Intel Pentium D 830	Intel Petium M T2600	Dual Xeon 3.0GHz
Memory Bandwidth	4817.58 MB/s	4429.8 MB/s	3444.99 MB/s	4073.71 MB/s
L1 Cache Latency				
32 Bytes Stride	3 cycles/1.67ns	4 cycles/1.33ns	3 cycles/1.38ns	3 cycles/1.00ns
L2 Cache Latency				
4 Bytes Stride	3 cycles/1.67 ns	6 cycles/2.00 ns	3 cycles/1.38 ns	6 cycles/2.00 ns
16 Bytes Stride	5 cycles/2.79 ns	13 cycles/4.33 ns	5 cycles/2.31 ns	13 cycles/4.33 ns
64 Bytes Stride	17 cycles/9.47 ns	29 cycles/9.67 ns	14 cycles/6.46 ns	27 cycles/9.00 ns
256 Bytes Stride	12 cycles/6.69 ns	28 cycles/9.33 ns	14 cycles/6.46 ns	26 cycles/8.67 ns
512 Bytes Stride	13 cycles/7.24 ns	26 cycles/8.67 ns	14 cycles/6.46 ns	25 cycles/8.33 ns
Memory Latency				
4 Bytes Stride	3 cycles/1.67 ns	7 cycles/2.33 ns	4 cycles/1.85 ns	6 cycles/ 2.00 ns
16 Bytes Stride	12 cycles/6.69 ns	15 cycles/5.00 ns	13 cycles/6.00 ns	15 cycles/5.00 ns
64 Bytes Stride	48 cycles/26.75 ns	43 cycles/14.33 ns	53 cycles/24.46 ns	49 cycles/16.33 ns
256 Bytes Stride	103 cycles/57.40 ns	270 cycles/90.00 ns	202 cycles/93.22 ns	376 cycles/125.33 ns
512 Bytes Stride	106 cycles/59.07 ns	284 cycles/94.66 ns	205 cycles/94.61 ns	395 cycles/131.66 ns



Matrix Multiplication Example

- Major Cache Effects to Consider

- Total cache size

- Exploit temporal locality and keep the working set small (e.g., by using blocking)

- Block size

- Exploit spatial locality

- Description:

- Multiply $N \times N$ matrices

- $O(N^3)$ total operations

- Accesses

- N reads per source element
- N values summed per destination
 - but may be able to hold in register


```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

Variable sum held in register



Miss Rate Analysis for Matrix Multiply

- Assume:
 - Line size = 32B (big enough for 4 64-bit words)
 - Matrix dimension (N) is very large
 - Approximate $1/N$ as 0.0
 - Cache is not even big enough to hold multiple rows
- Analysis Method:
 - Look at access pattern of inner loop





Layout of C Arrays in Memory (review)


- C arrays allocated in row-major order
 - each row in contiguous memory locations
- Stepping through columns in one row:
 - `for (i = 0; i < N; i++)`
 `sum += a[0][i];`
 - accesses successive elements
 - if block size (B) > 4 bytes, exploit spatial locality
 - compulsory miss rate = 4 bytes / B
- Stepping through rows in one column:
 - `for (i = 0; i < n; i++)`
 `sum += a[i][0];`
 - accesses distant elements
 - no spatial locality!
 - compulsory miss rate = 1 (i.e. 100%)



Matrix Multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

Inner loop:



- Misses per Inner Loop Iteration:


A	B	C
0.25	1.0	0.0



Matrix Multiplication (jik)

```
/* jik */
for (j=0; j<n; j++) {
  for (i=0; i<n; i++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

Inner loop:



- Misses per Inner Loop Iteration:

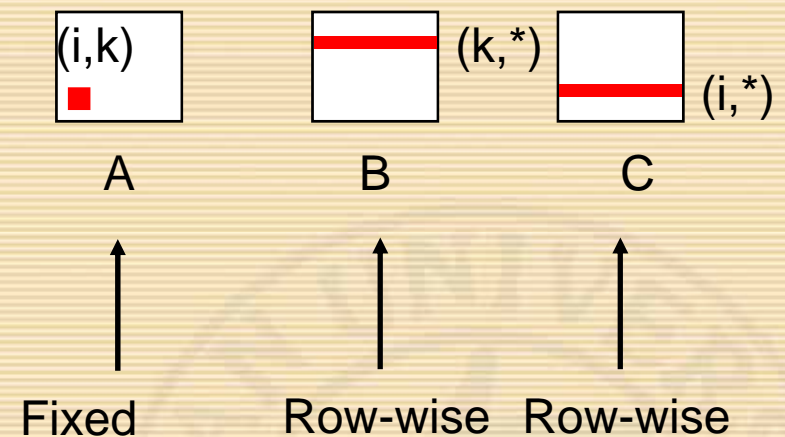
A	B	C
0.25	1.0	0.0



Matrix Multiplication (kij)

```
/* kij */
for (k=0; k<n; k++) {
  for (i=0; i<n; i++) {
    r = a[i][k];
    for (j=0; j<n; j++)
      c[i][j] += r * b[k][j];
  }
}
```

Inner loop:



- Misses per Inner Loop Iteration:


A	B	C
0.0	0.25	0.25



Matrix Multiplication (ikj)

```
/* ikj */
for (i=0; i<n; i++) {
  for (k=0; k<n; k++) {
    r = a[i][k];
    for (j=0; j<n; j++)
      c[i][j] += r * b[k][j];
  }
}
```

Inner loop:



- Misses per Inner Loop Iteration:


A	B	C
0.0	0.25	0.25



Matrix Multiplication (jki)


```
/* jki */
for (j=0; j<n; j++) {
  for (k=0; k<n; k++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```

Inner loop:



A

Column -
wise



B

Fixed



C

Column-
wise

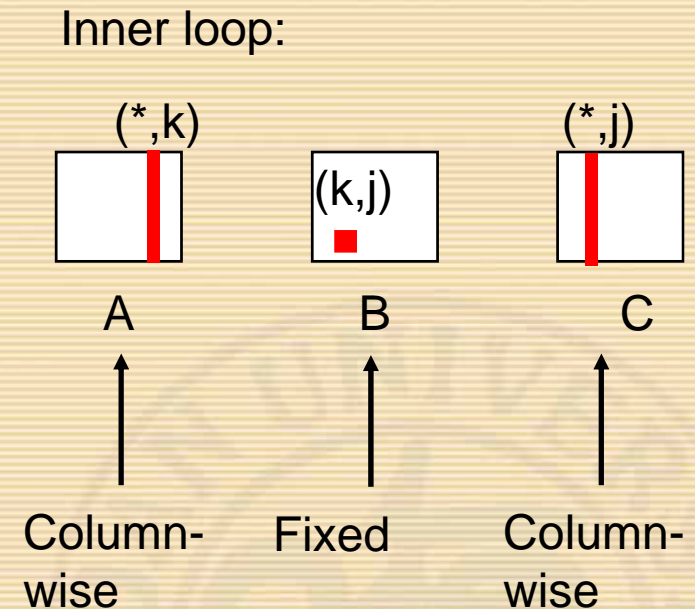
- Misses per Inner Loop Iteration:

A	B	C
1.0	0.0	1.0



Matrix Multiplication (kji)

```
/* kji */
for (k=0; k<n; k++) {
  for (j=0; j<n; j++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```



- Misses per Inner Loop Iteration:

A	B	C
1.0	0.0	1.0



Summary of Matrix Multiplication

ijk (& jik):

- 2 loads, 0 stores
- misses/iter = 1.25

```
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

kij (& ikj):

- 2 loads, 1 store
- misses/iter = 0.5

```
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

jki (& kji):


- 2 loads, 1 store
- misses/iter = 2.0

```
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```



Pentium Matrix Multiply Performance

- Miss rates are helpful but not perfect predictors.
 - Code scheduling matters, too.





Improving Temporal Locality by Blocking

- Example: Blocked matrix multiplication
 - “block” (in this context) does not mean “cache block”.
 - Instead, it mean a sub-block within the matrix.
 - Example: $N = 8$; sub-block size = 4

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix}$$

Key idea: Sub-blocks (i.e., A_{xy}) can be treated just like scalars.

$$C_{11} = A_{11}B_{11} + A_{12}B_{21} \quad C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21} \quad C_{22} = A_{21}B_{12} + A_{22}B_{22}$$



Blocked Matrix Multiply (bijk)

```
for (jj=0; jj<n; jj+=bsize) {  
    for (i=0; i<n; i++)  
        for (j=jj; j < min(jj+bsize,n); j++)  
            c[i][j] = 0.0;  
    for (kk=0; kk<n; kk+=bsize) {  
        for (i=0; i<n; i++) {  
            for (j=jj; j < min(jj+bsize,n); j++) {  
                sum = 0.0  
                for (k=kk; k < min(kk+bsize,n); k++) {  
                    sum += a[i][k] * b[k][j];  
                }  
                c[i][j] += sum;  
            }  
        }  
    }  
}
```




Blocked Matrix Multiply Analysis

- Innermost loop pair multiplies a $1 \times bsize$ sliver of A by a $bsize \times bsize$ block of B and accumulates into $1 \times bsize$ sliver of C
- Loop over i steps through n row slivers of A & C , using same B

```
for (i=0; i<n; i++) {  
    for (j=jj; j < min(jj+bsize,n); j++) {  
        sum = 0.0  
        for (k=kk; k < min(kk+bsize,n); k++) {  
            sum += a[i][k] * b[k][j];  
        }  
        c[i][j] += sum;  
    }  
}
```


Innermost
Loop Pair





Blocked Matrix Multiply Performance

- Blocking (bijk and bikj) improves performance by a factor of two over unblocked versions (ijk and jik)
 - relatively insensitive to array size.





Concluding Observations

- Programmer can optimize for cache performance
 - How data structures are organized
 - How data are accessed
 - Nested loop structure
 - Blocking is a general technique
- All systems favor “cache friendly code”
 - Getting absolute optimum performance is very platform specific
 - Cache sizes, line sizes, associativities, etc.
 - Can get most of the advantage with generic code
 - Keep working set reasonably small (temporal locality)
 - Use small strides (spatial locality)



Virtual Memory




Motivations for Virtual Memory

- Use Physical DRAM as a Cache for the Disk
 - Address space of a process can exceed physical memory size
 - Sum of address spaces of multiple processes can exceed physical memory
- Simplify Memory Management
 - Multiple processes resident in main memory.
 - Each process with its own address space
 - Only “active” code and data is actually in memory
 - Allocate more memory to process as needed.
- Provide Protection
 - One process can’t interfere with another.
 - because they operate in different address spaces.
 - User process cannot access privileged information
 - different sections of address spaces have different permissions.



Locating an Object in Cache

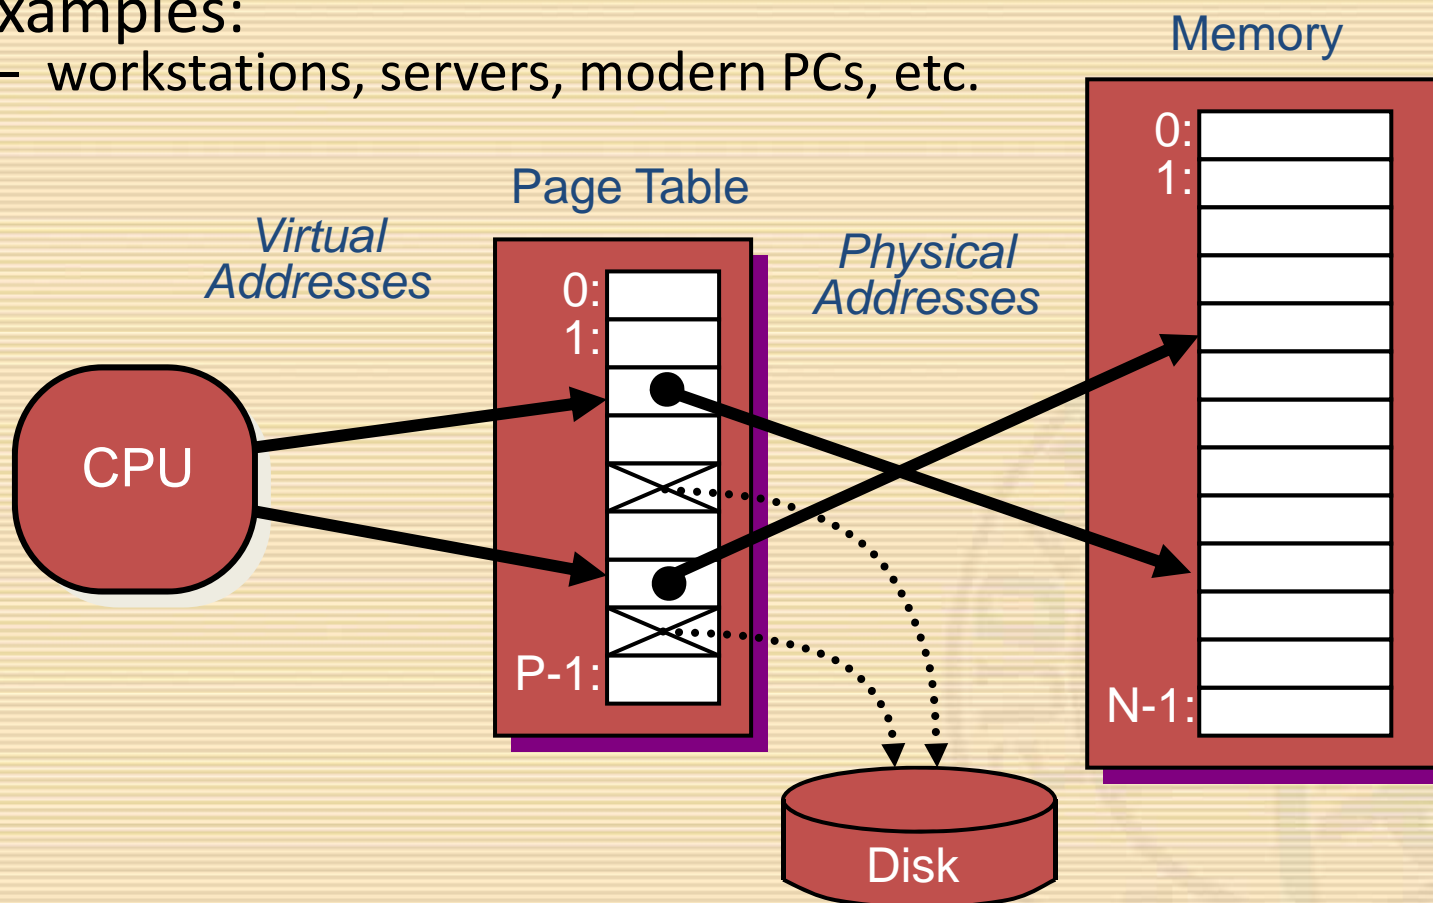
- DRAM Cache
 - Each allocated page of virtual memory has entry in *page table*
 - Mapping from virtual pages to physical pages
 - From uncached form to cached form
 - Page table entry even if page not in memory
 - Specifies disk address
 - Only way to indicate where to find page
 - OS retrieves information





A System with Virtual Memory

- Examples:
 - workstations, servers, modern PCs, etc.

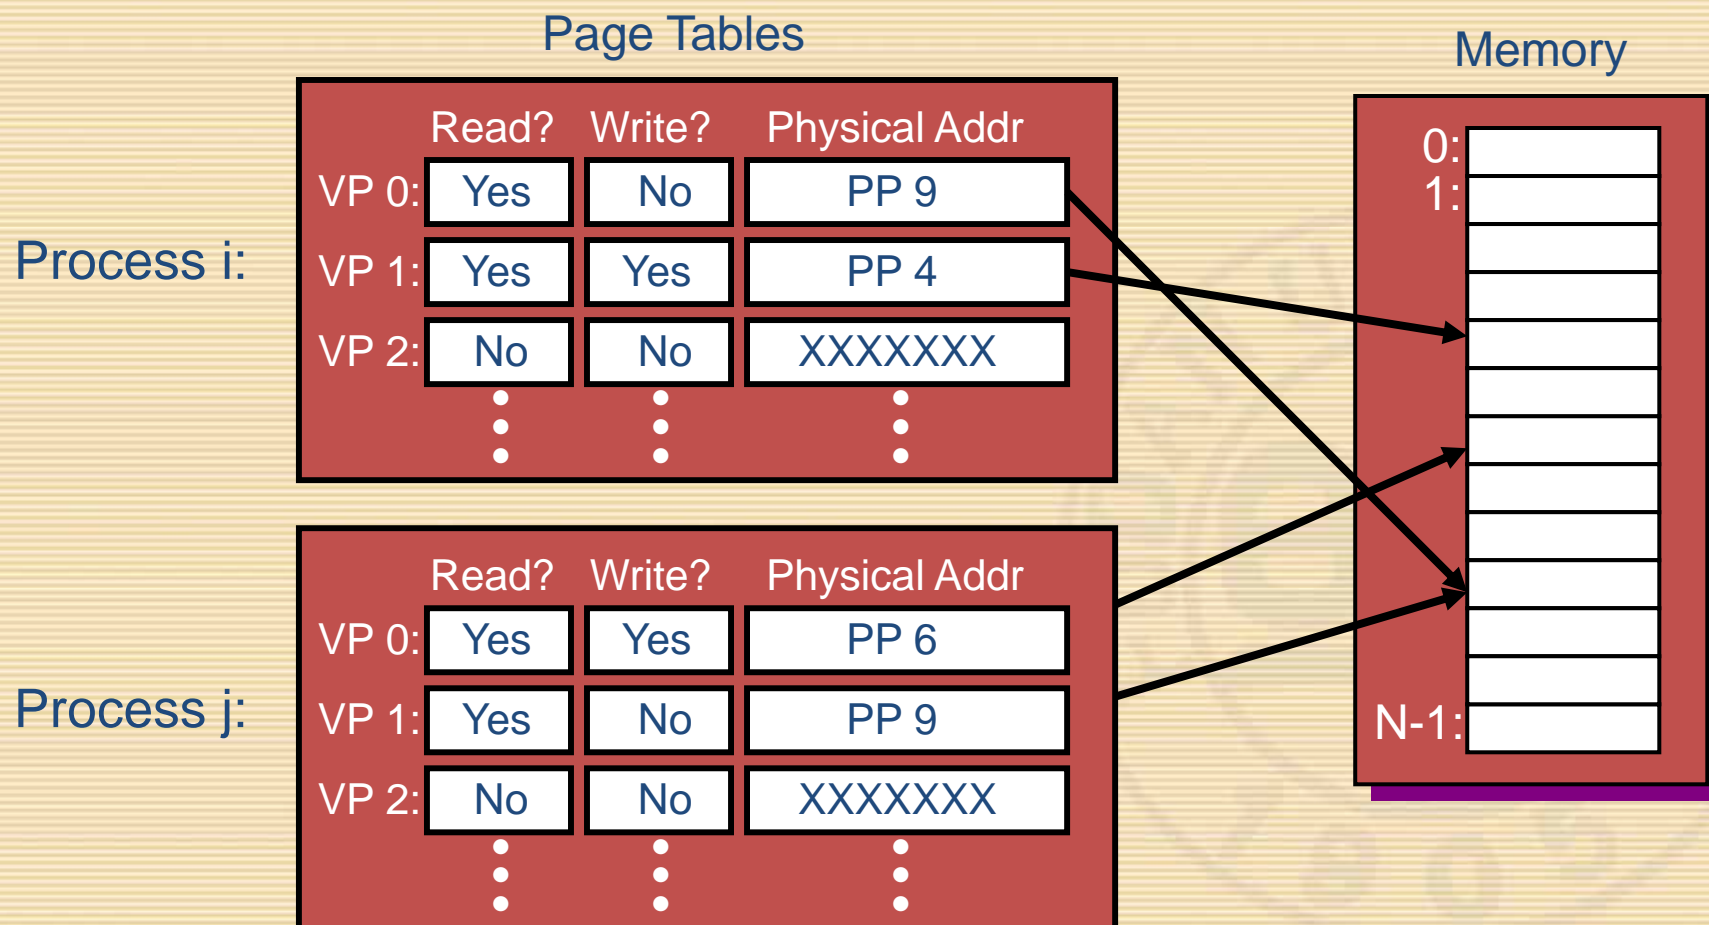


- Address Translation: Hardware converts virtual addresses to physical addresses via OS-managed lookup table (page table)



Protection

- Page table entry contains access rights information
 - hardware enforces this protection (trap into OS if violation occurs)






VM Address Translation

- Virtual Address Space
 - $V = \{0, 1, \dots, N-1\}$
- Physical Address Space
 - $P = \{0, 1, \dots, M-1\}$
 - $M < N$
- Address Translation
 - $\text{MAP}: V \rightarrow P \cup \{\emptyset\}$
 - For virtual address a :
 - $\text{MAP}(a) = a'$ if data at virtual address a at physical address a' in P
 - $\text{MAP}(a) = \emptyset$ if data at virtual address a not in physical memory
 - Either invalid or stored on disk



VM Address Translation

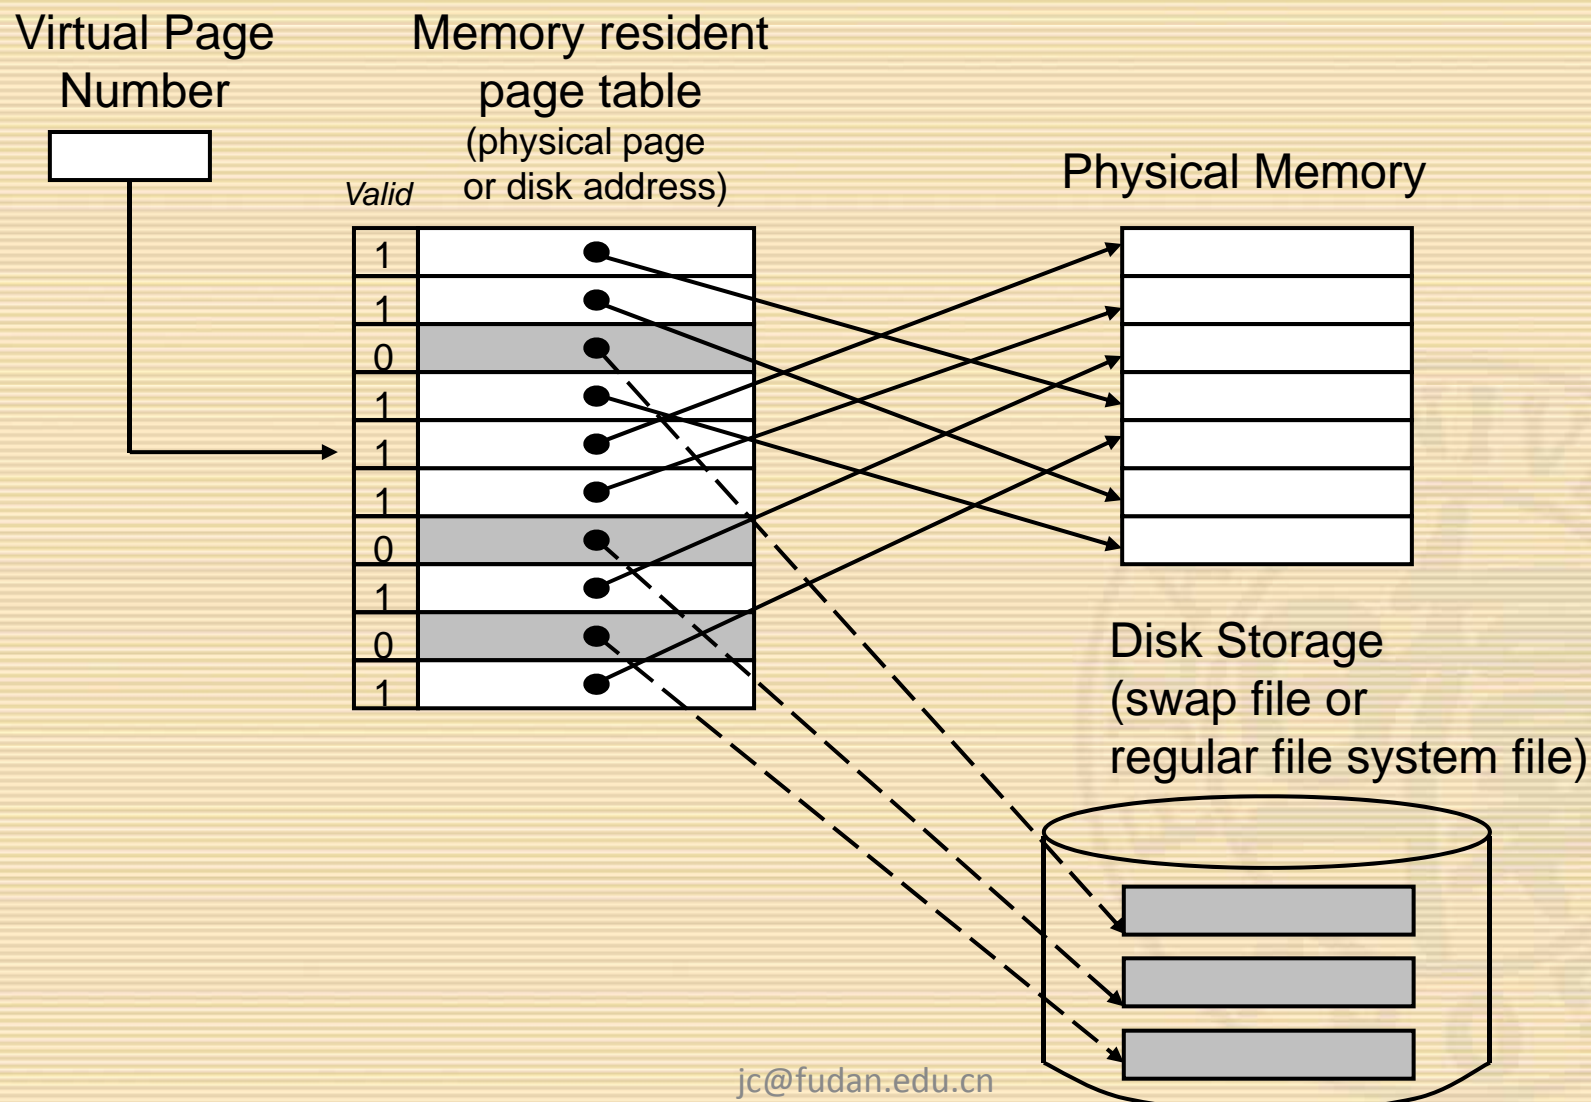
- Parameters
 - $P = 2^p$ = page size (bytes).
 - $N = 2^n$ = Virtual address limit
 - $M = 2^m$ = Physical address limit



Page offset bits don't change as a result of translation

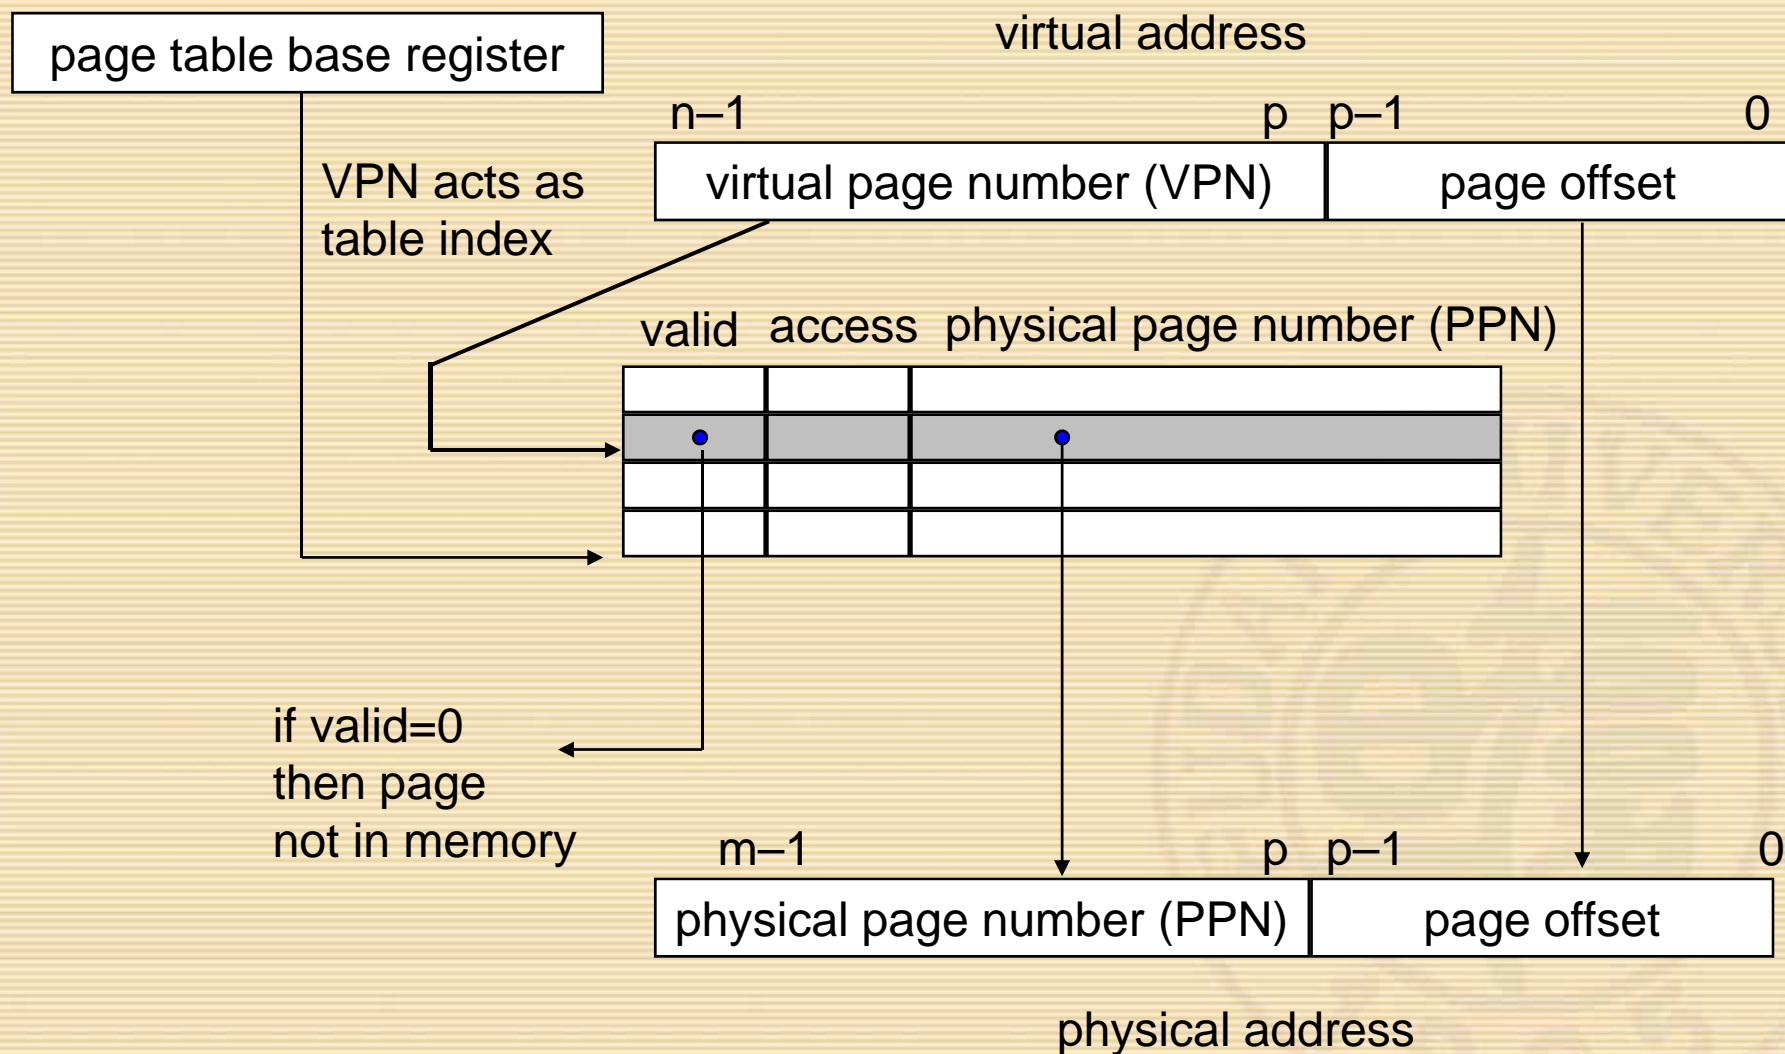


Page Tables





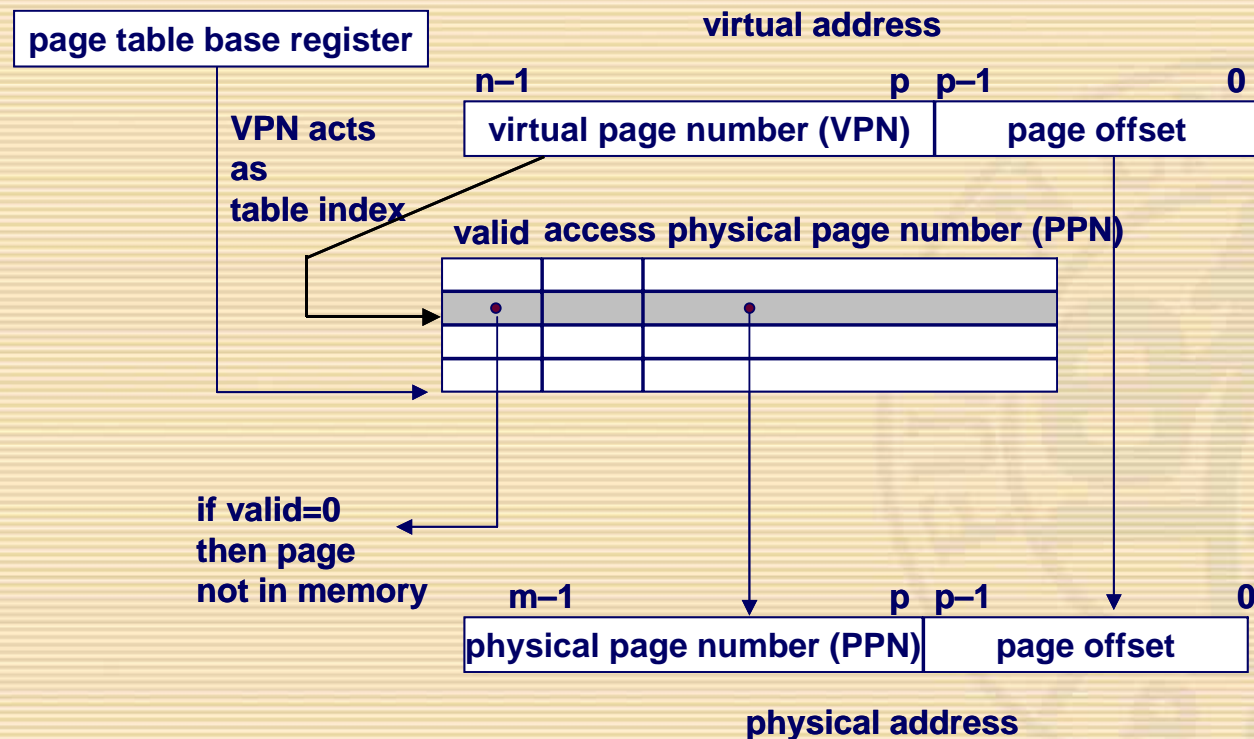
Address Translation via Page Table





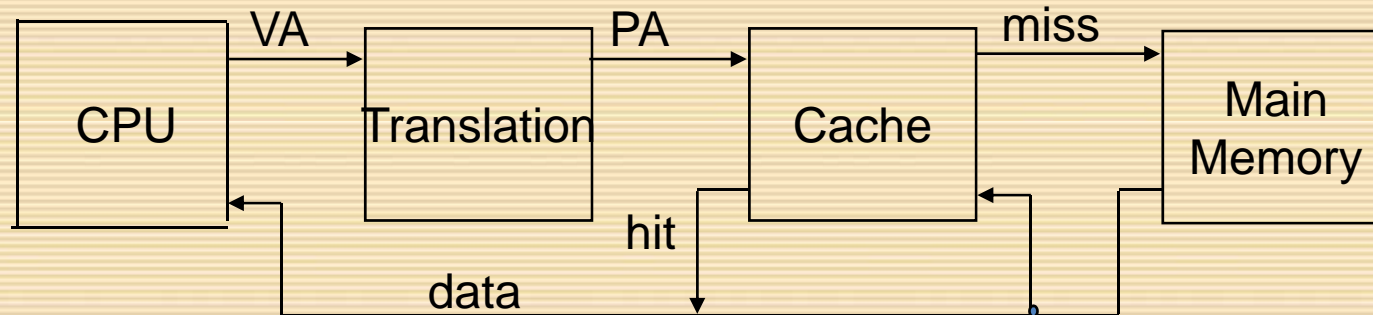
Page Table Operation

- Translation
 - Separate (set of) page table(s) per process
 - VPN forms index into page table (points to a page table entry)





Integrating VM and Cache

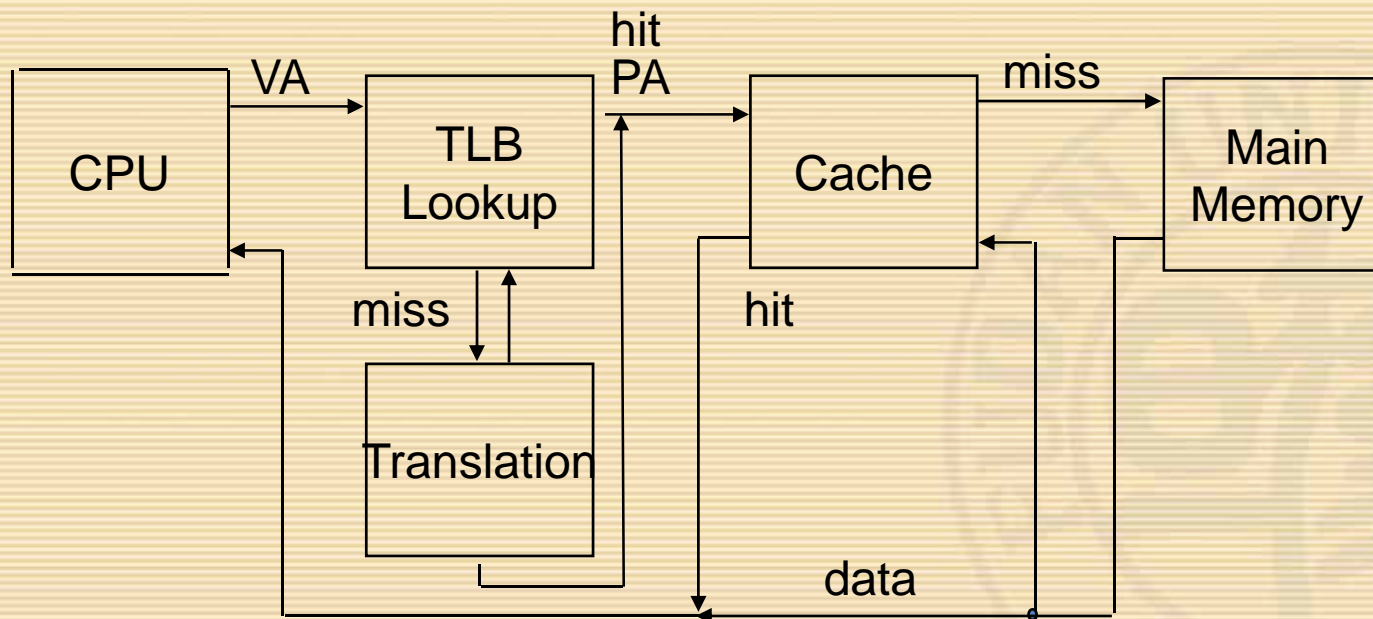


- Most Caches “Physically Addressed”
 - Accessed by physical addresses
 - Allows multiple processes to have blocks in cache at same time
 - Allows multiple processes to share pages
 - Cache doesn’t need to be concerned with protection issues
 - Access rights checked as part of address translation
- Perform Address Translation Before Cache Lookup
 - But this could involve a memory access itself (of the PTE)
 - Of course, page table entries can also become cached



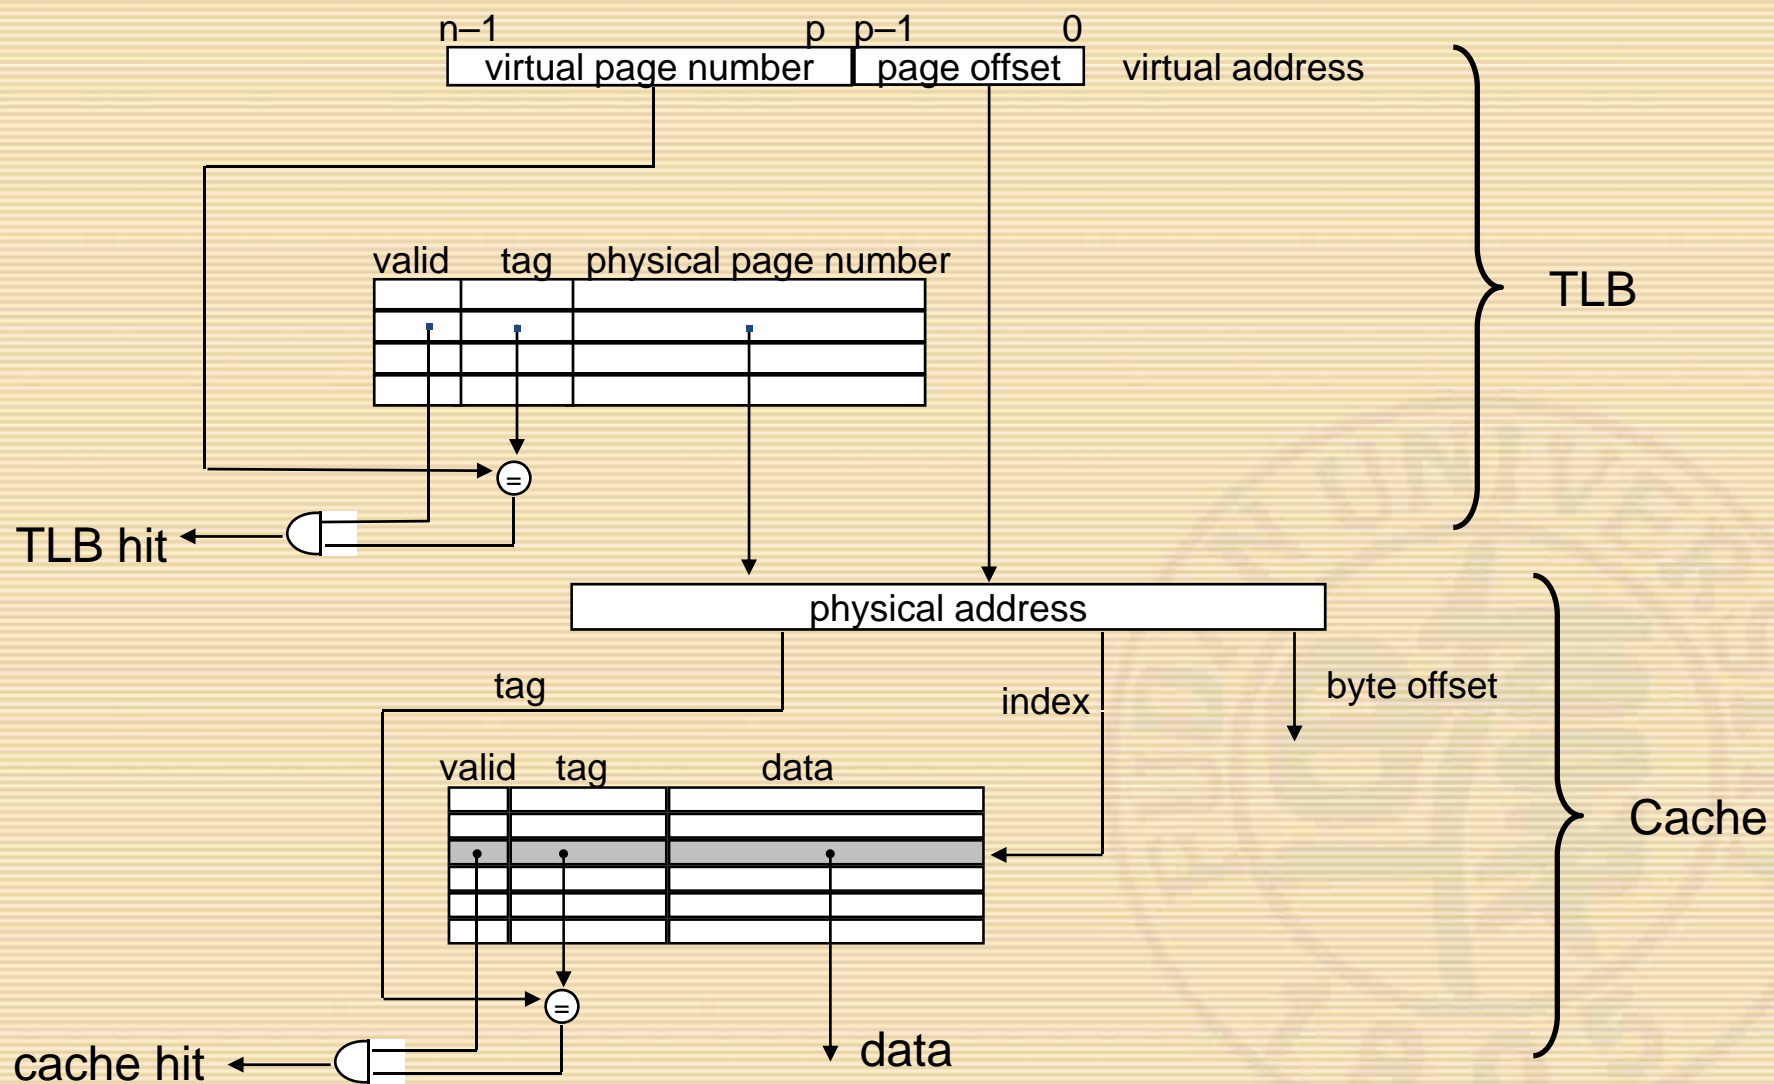
Speeding up Translation with a TLB

- “Translation Lookaside Buffer” (TLB)
 - Small hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages






Address Translation with a TLB





Simple Memory System Example

- Addressing
 - 14-bit virtual addresses
 - 12-bit physical address
 - Page size = 64 bytes





Simple Memory System Page Table


– Only show first 16 entries

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1



Simple Memory System TLB

- TLB
 - 16 entries
 - 4-way associative




Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0



Simple Memory System Cache

- Cache
 - 16 lines
 - 4-byte line size
 - Direct mapped




Idx	Tag	Valid	B0	B1	B2	B3	Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11	8	24	1	3A	00	51	89
1	15	0	–	–	–	–	9	2D	0	–	–	–	–
2	1B	1	00	02	04	08	A	2D	1	93	15	DA	3B
3	36	0	–	–	–	–	B	0B	0	–	–	–	–
4	32	1	43	6D	8F	09	C	12	0	–	–	–	–
5	0D	1	36	72	F0	1D	D	16	1	04	96	34	15
6	31	0	–	–	–	–	E	13	1	83	77	1B	D3
7	16	1	11	C2	DF	03	F	14	0	–	–	–	–

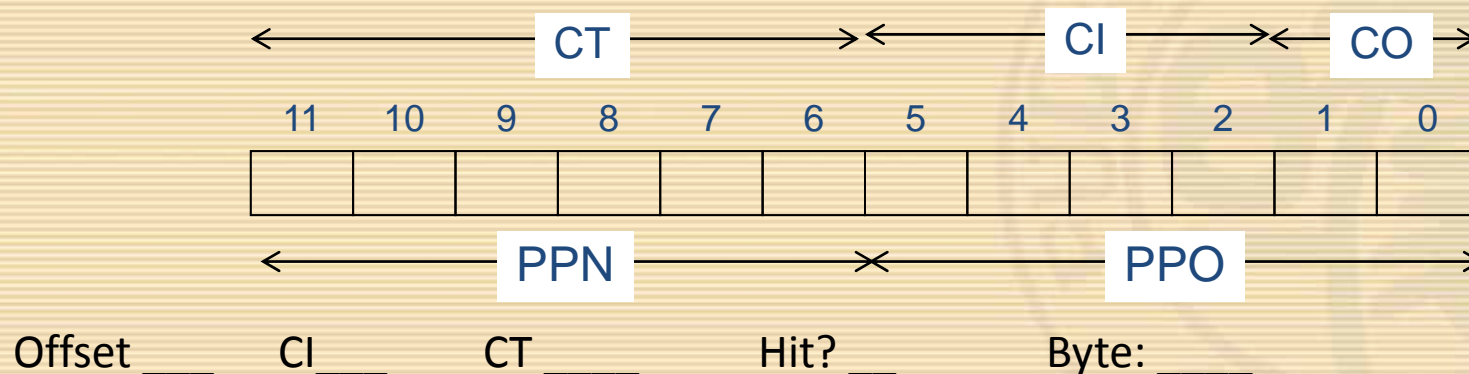


Address Translation Example #1

- Virtual Address 0x03D4



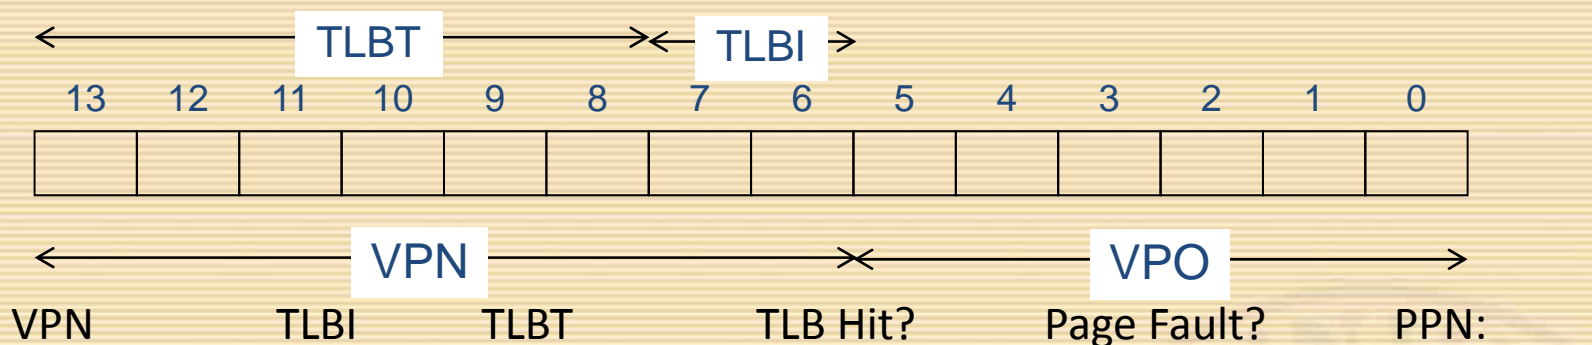
- Physical Address





Address Translation Example #2

- Virtual Address 0x0B8F



- Physical Address

